

CS 380 - GPU and GPGPU Programming

Lecture 4: GPU Architecture 2

Markus Hadwiger, KAUST

Reading Assignment #2 (until Sep 14)



Read (required):

- Orange book (GLSL), chapter 4 (*The OpenGL Programmable Pipeline*)
- Brief GLSL overview:

https://en.wikipedia.org/wiki/OpenGL_Shading_Language

- GPU Gems 2 book, chapter 30 (*The GeForce 6 Series GPU Architecture*)
available online:

http://download.nvidia.com/developer/GPU_Gems_2/GPU_Gems2_ch30.pdf

Programming Assignments: Schedule (tentative)



Assignment #1:

- Querying the GPU (OpenGL/GLSL and CUDA) due Sep 7

Assignment #2:

- Phong shading and procedural texturing (GLSL) due Sep 21

Assignment #3:

- Image Processing with GLSL due Oct 5

Assignment #4:

- Image Processing with CUDA
- Convolutional layers with CUDA due Oct 26

Assignment #5:

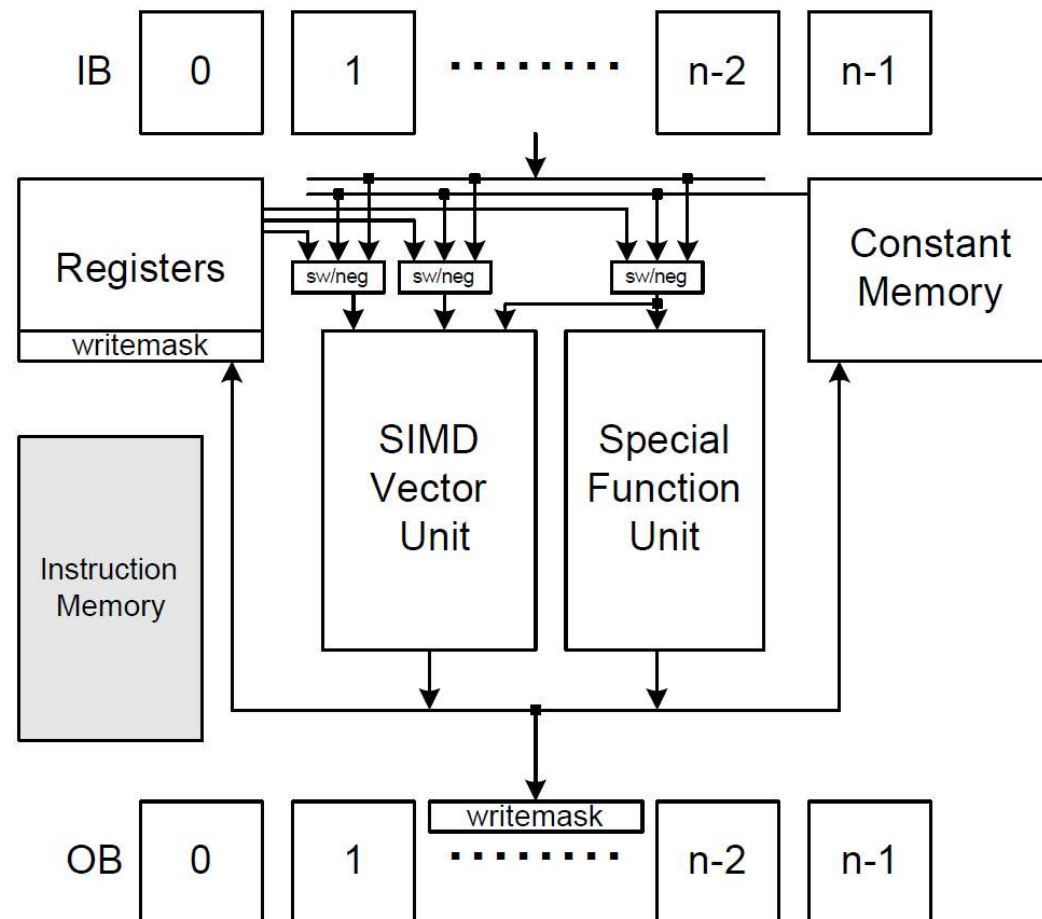
- Linear Algebra (CUDA) due Nov 16

Legacy Vertex Shading Unit (1)



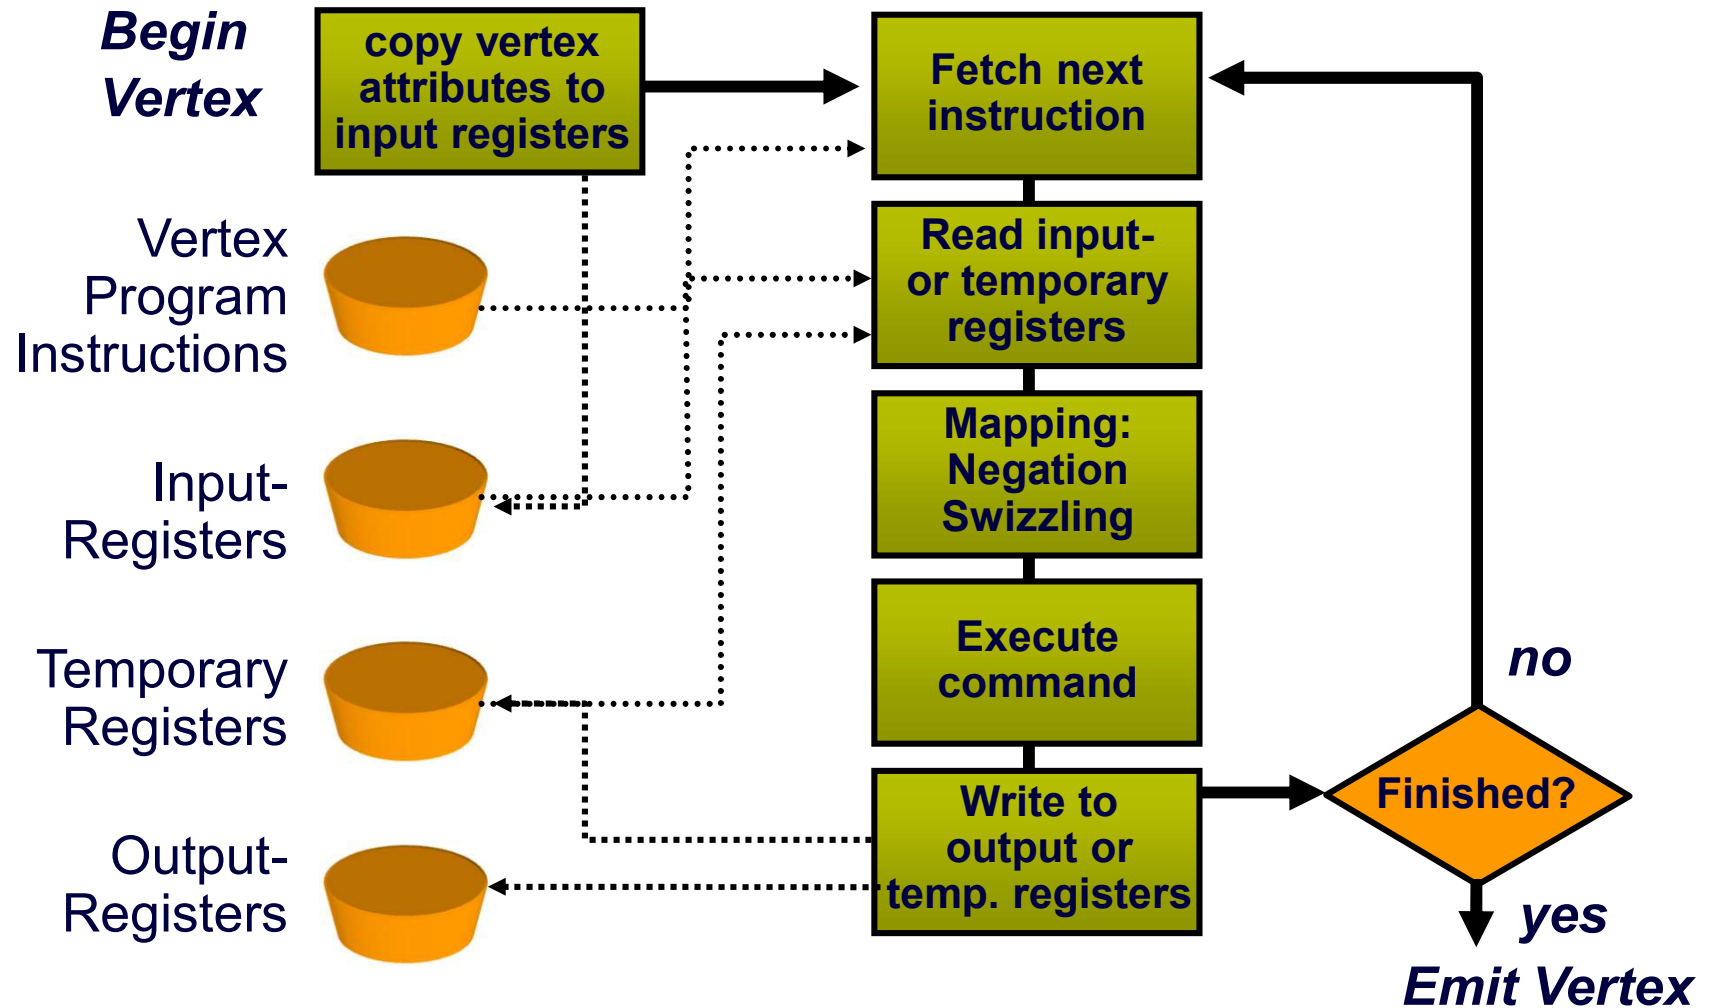
Geforce 3 (NV20), 2001

- floating point
4-vector
vertex engine
- still very
instructive for
understanding
GPUs in general



Lindholm et al., A User-Programmable Vertex Engine, SIGGRAPH 2001

Vertex Processor



Legacy Vertex Shading Unit (2)



Input
attributes

Vertex Attribute Register	Conventional Per-vertex Parameter	Conventional Per-vertex Parameter Command	Conventional Component Mapping
0	Vertex position	glVertex	<i>x,y,z,w</i>
1	Vertex weights	glVertexWeightEXT	<i>w,0,0,1</i>
2	Normal	glNormal	
3	Primary color	glColor	<i>r,g,b,a</i>
4	Secondary color	glSecondaryColorEXT	<i>r,g,b,1</i>
5	Fog coordinate	glFogCoordEXT	<i>f,0,0,1</i>
6	-	-	-
7	-	-	-
8	Texture coord 0	glMultiTexCoordARB (GL_TEXTURE0...)	<i>s,t,r,q</i>
9	Texture coord 1	glMultiTexCoordARB (GL_TEXTURE1...)	<i>s,t,r,q</i>
10	Texture coord 2	glMultiTexCoordARB (GL_TEXTURE2...)	<i>s,t,r,q</i>
11	Texture coord 3	glMultiTexCoordARB (GL_TEXTURE3...)	<i>s,t,r,q</i>
12	Texture coord 4	glMultiTexCoordARB (GL_TEXTURE4...)	<i>s,t,r,q</i>
13	Texture coord 5	glMultiTexCoordARB (GL_TEXTURE5...)	<i>s,t,r,q</i>
14	Texture coord 6	glMultiTexCoordARB (GL_TEXTURE6...)	<i>s,t,r,q</i>
15	Texture coord 7	glMultiTexCoordARB (GL_TEXTURE7...)	<i>s,t,r,q</i>

Code
examples

```

DP4 o[HPOS].x, c[0], v[OPOS];
MUL R1, R0.zxyw, R2.yzxw ;
MAD R1, R0.yzxw, R2.zxyw, -R1;
    
```

swizzling!

Legacy Vertex Shading Unit (3)



Vector instruction set, very few instructions; no branching yet!

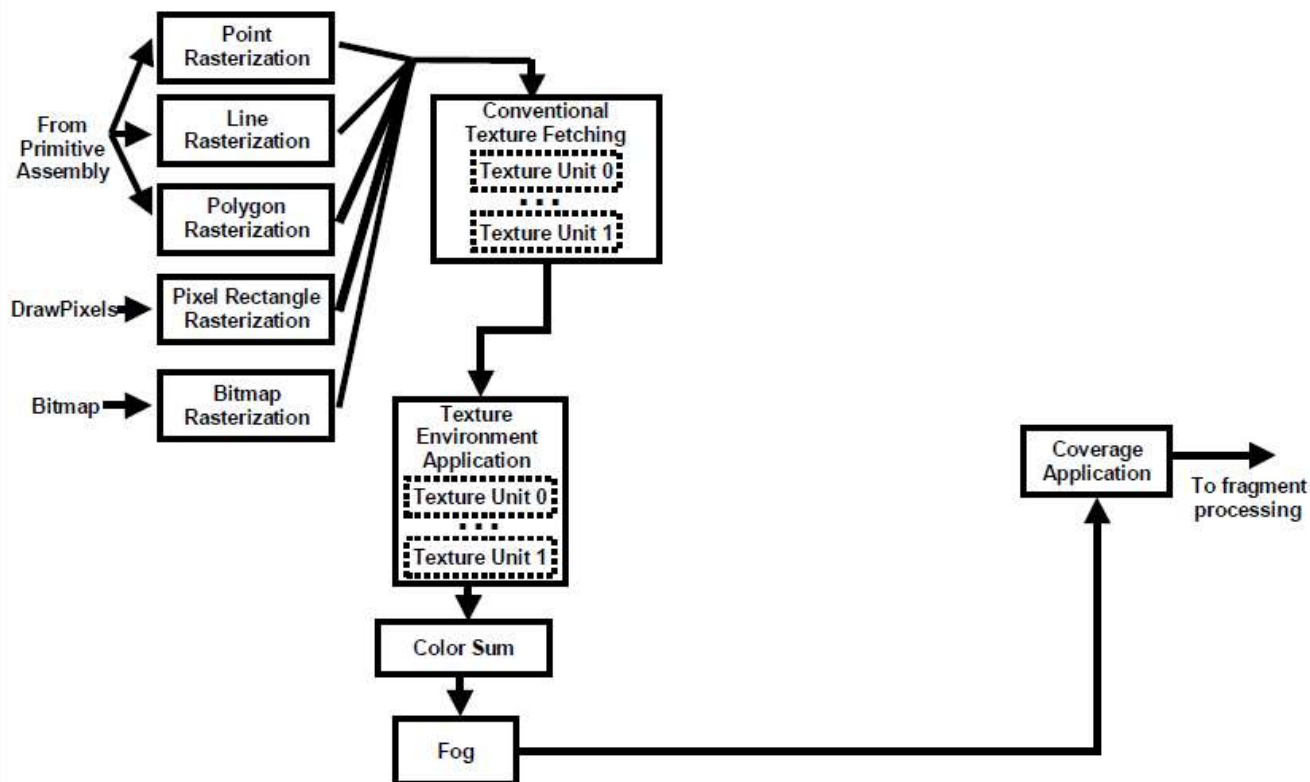
OpCode	Full Name	Description
MOV	Move	vector -> vector
MUL	Multiply	vector -> vector
ADD	Add	vector -> vector
MAD	Multiply and add	vector -> vector
DST	Distance	vector -> vector
MIN	Minimum	vector -> vector
MAX	Maximum	vector -> vector
SLT	Set on less than	vector -> vector
SGE	Set on greater or equal	vector -> vector
RCP	Reciprocal	scalar-> replicated scalar
RSQ	Reciprocal square root	scalar-> replicated scalar
DP3	3 term dot product	vector-> replicated scalar
DP4	4 term dot product	vector-> replicated scalar
LOG	Log base 2	miscellaneous
EXP	Exp base 2	miscellaneous
LIT	Phong lighting	miscellaneous
ARL	Address register load	miscellaneous

Fast Forward to Programm. Fragment Shading



Core OpenGL Fragment Texturing & Coloring

< 1999



NVIDIA Proprietary


Courtesy Mark Kilgard

Fast Forward to Programm. Fragment Shading



NV10 OpenGL Fragment Texturing & Coloring

GeForce 256,
1999



NVIDIA Proprietary

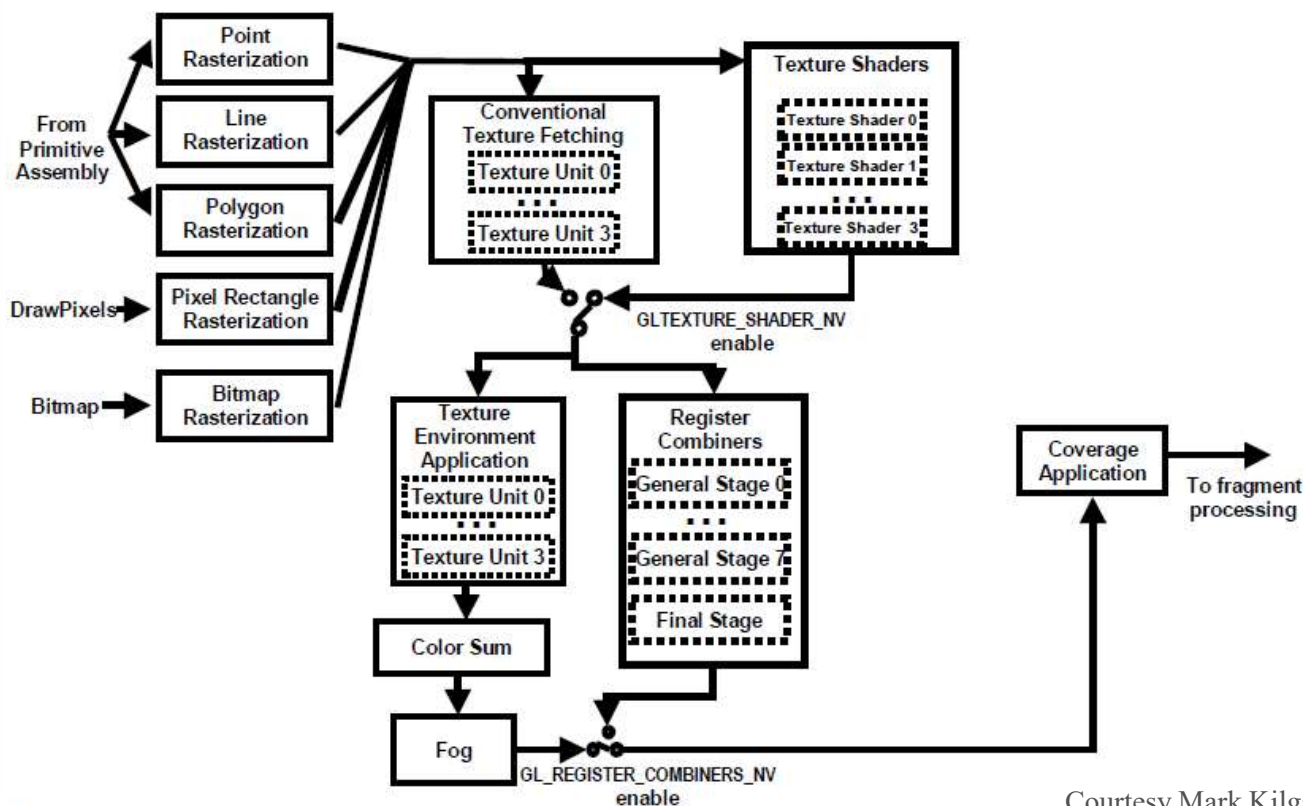
Courtesy Mark Kilgard

Fast Forward to Programm. Fragment Shading



NV20 OpenGL Fragment Texturing & Coloring

GeForce 3,
2001



NVIDIA Proprietary

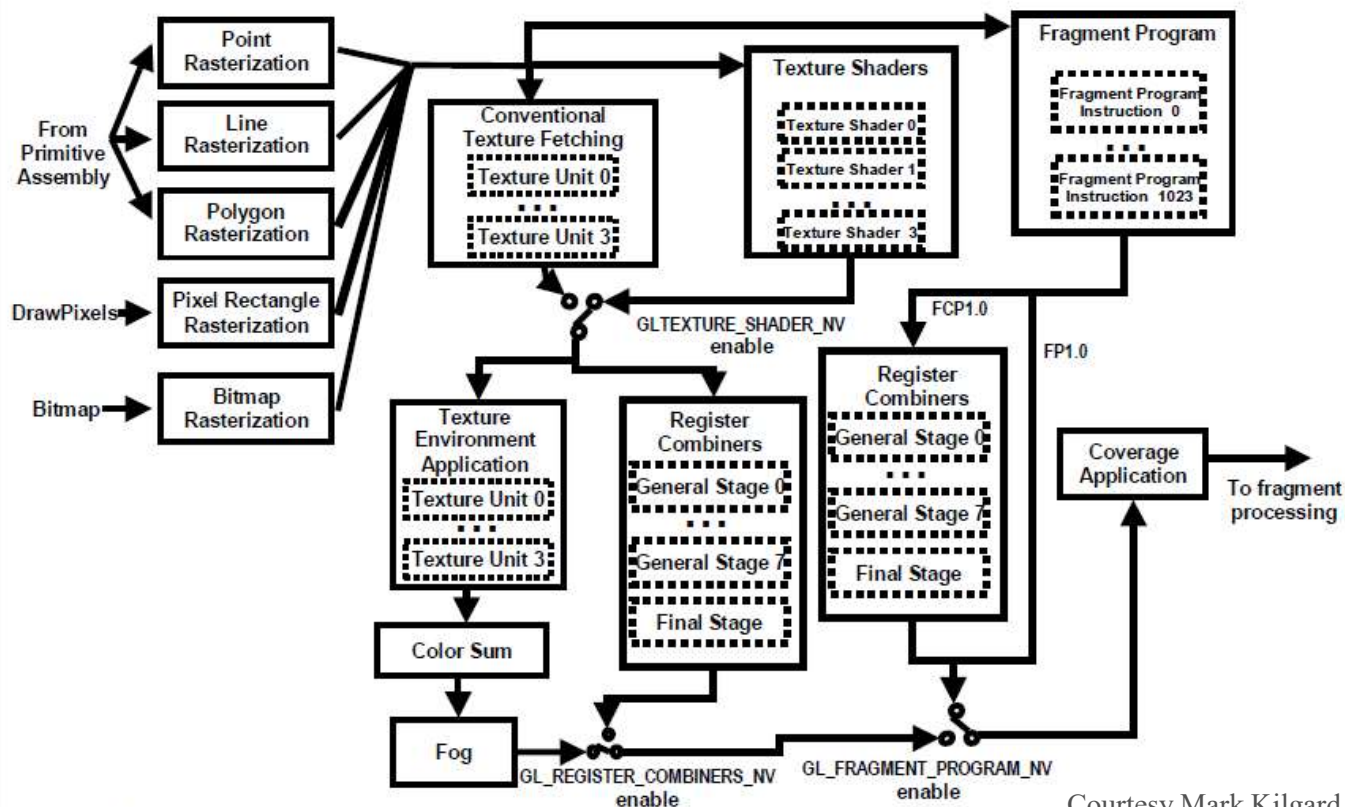
Courtesy Mark Kilgard

Fast Forward to Programm. Fragment Shading



NV30 OpenGL Fragment Texturing & Coloring

GeForce FX (5),
2003



NVIDIA Proprietary

Courtesy Mark Kilgard

GPU Structure Before Unified Shaders

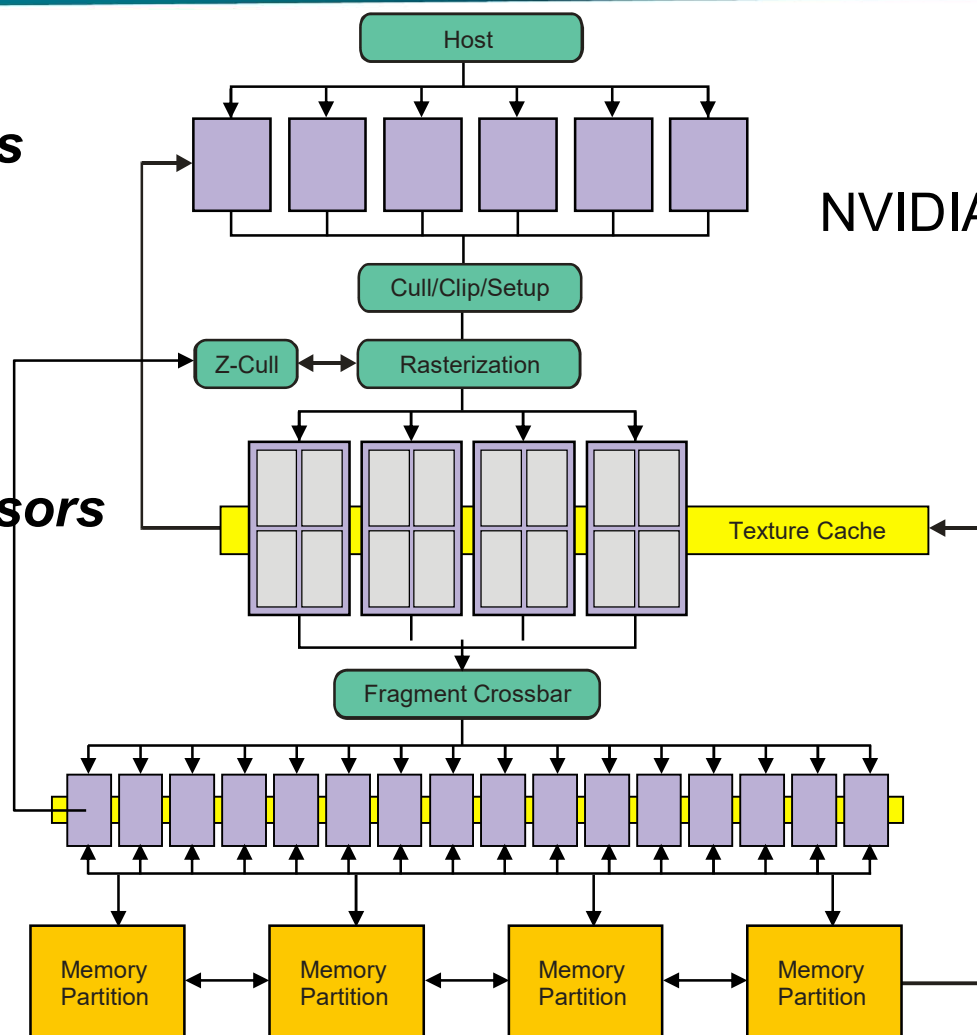


Vertex Processors

Example
NVIDIA GeForce 6/7,
2004, 2005

Fragment Processors

Memory Access
Z-Compare and
Blending



Legacy Fragment Shading Unit (1)



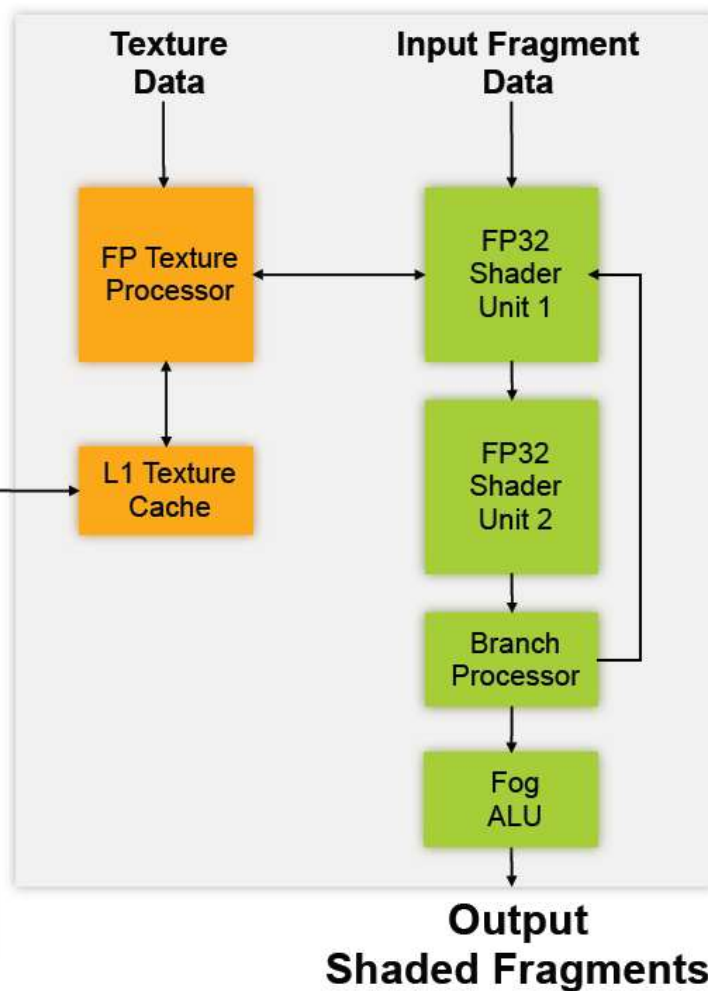
GeForce 6 (NV40), 2004

- dynamic branching

Texture Filter
Bi / Tri / Aniso
1 texture @ full speed
4-tap filter @ full speed
16:1 Aniso w/ Trilinear (128-tap)
FP16 Texture Filtering

L2 Texture
Cache

SIMD Architecture
Dual Issue / Co-Issue
FP32 Computation
Shader Model 3.0



Shader Unit 1
4 FP Ops / pixel
Dual/Co-Issue
Texture Address Calc
Free fp16 normalize
+ mini ALU

Shader Unit 2
4 FP Ops / pixel
Dual/Co-Issue
+ mini ALU

Legacy Fragment Shading Unit (2)



Example code

```
!!ARBfp1.0

ATTRIB unit_tc = fragment.texcoord[ 0 ];
PARAM  mvp_inv[] = { state.matrix.mvp.inverse };
PARAM  constants = {0, 0.999, 1, 2};

TEMP pos_win, temp;

TEX pos_win.z, unit_tc, texture[ 1 ], 2D;

ADD pos_win.w, constants.y, -pos_win.z;
KIL pos_win.w;

MOV result.color.w, pos_win.z;

MOV pos_win.xyw, unit_tc;
MAD pos_win.xyz, pos_win, constants.a, -constants.b;

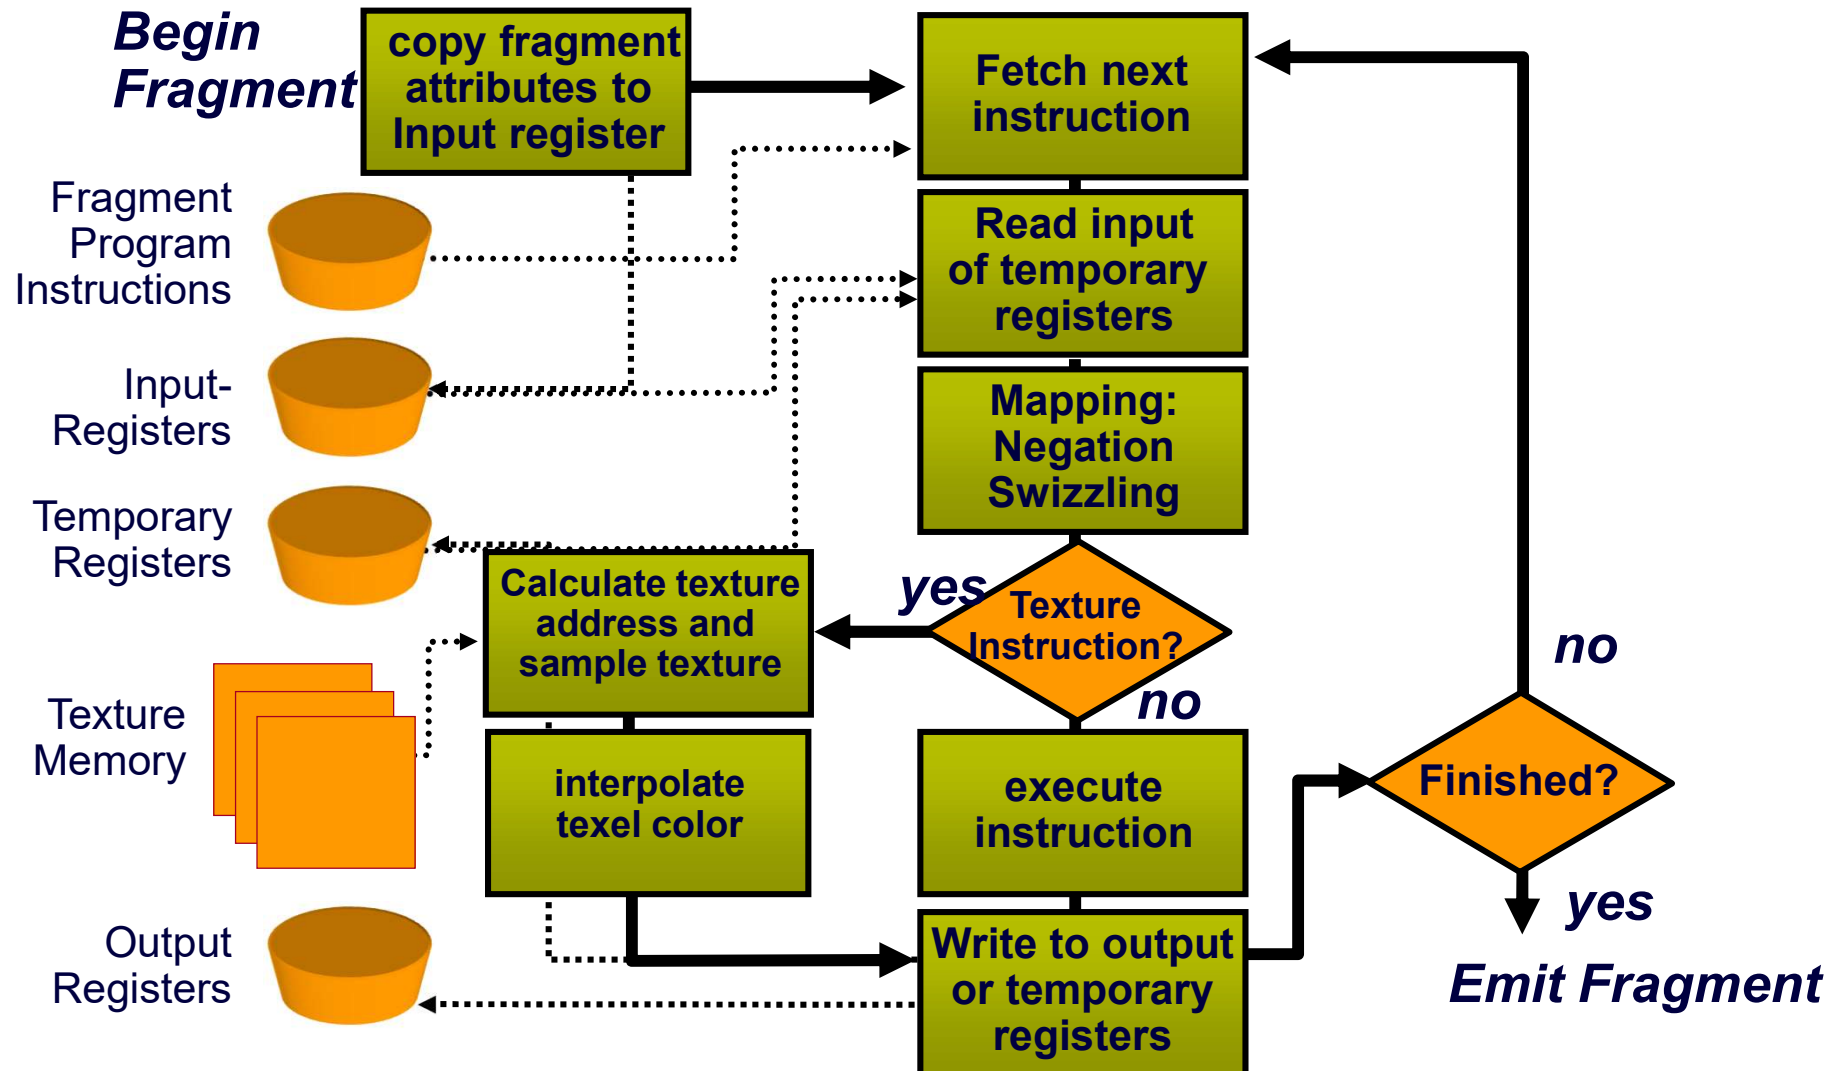
DP4 temp.w, mvp_inv[ 3 ], pos_win;
RCP temp.w, temp.w;

MUL pos_win, pos_win, temp.w;

DP4 result.color.x, mvp_inv[ 0 ], pos_win;
DP4 result.color.y, mvp_inv[ 1 ], pos_win;
DP4 result.color.z, mvp_inv[ 2 ], pos_win;

END
```


Fragment Processor



A diffuse reflectance shader

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

Independent, but no explicit parallelism

Compile shader

1 unshaded fragment input record



```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp ( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```



```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul   r3, v0, cb0[0]  
madd  r3, v1, cb0[1], r3  
madd  r3, v2, cb0[2], r3  
clmp  r3, r3, l(0.0), l(1.0)  
mul   o0, r0, r3  
mul   o1, r1, r3  
mul   o2, r2, r3  
mov   o3, l(1.0)
```



1 shaded fragment output record



Per-Pixel(Fragment) Lighting




Simulating smooth surfaces by calculating illumination for each fragment

Example: specular highlights (Phong illumination/shading)

Phong shading:
per-fragment evaluation

Gouraud shading:
linear interpolation from vertices



Per-Pixel Phong Lighting (Cg)



```
void main(float4 position : TEXCOORD0,  
          float3 normal   : TEXCOORD1,  
  
          out float4 oColor : COLOR,  
  
          uniform float3 ambientCol,  
          uniform float3 lightCol,  
          uniform float3 lightPos,  
          uniform float3 eyePos,  
          uniform float3 Ka,  
          uniform float3 Kd,  
          uniform float3 Ks,  
          uniform float  shiny)  
{
```

Per-Pixel Phong Lighting (Cg)



```
float3 P = position.xyz;
float3 N = normal;
float3 V = normalize(eyePosition - P);
float3 H = normalize(L + V);

float3 ambient = Ka * ambientCol;

float3 L      = normalize(lightPos - P);
float  diffLight = max(dot(L, N), 0);
float3 diffuse  = Kd * lightCol * diffLight;

float  specLight = pow(max(dot(H, N), 0), shiny);
float3 specular  = Ks * lightCol * specLight;

oColor.xyz = ambient + diffuse + specular;
oColor.w = 1;
}
```

Thank you.