

CS 380 - GPU and GPGPU Programming

Lecture 7: GPU Architecture 5

Markus Hadwiger, KAUST

Reading Assignment #4 (until Sep 28)



Read (required):

- Programming Massively Parallel Processors book, Chapter 2 (*History of GPU Computing*)
- GLSL book, Chapter 7 (OpenGL Shading Language API)
- OpenGL Shading Language 4.6 specification: Chapter 2
<http://www.opengl.org/registry/doc/GLSLangSpec.4.60.pdf>
- Download OpenGL 4.6 specification
<http://www.opengl.org/registry/doc/glspec46.core.pdf>

Read (optional):

- OpenGL 4 Shading Language Cookbook, Chapter 3
- http://www.opengl.org/wiki/History_of_OpenGL

See more at: <http://www.opengl.org/documentation/specs/>

Quiz #1: Sep 28



Organization

- First 30 min of lecture
- No material (book, notes, ...) allowed

Content of questions

- Lectures (both actual lectures and slides)
- Reading assignments
- Programming assignments (algorithms, methods)
- Solve short practical examples

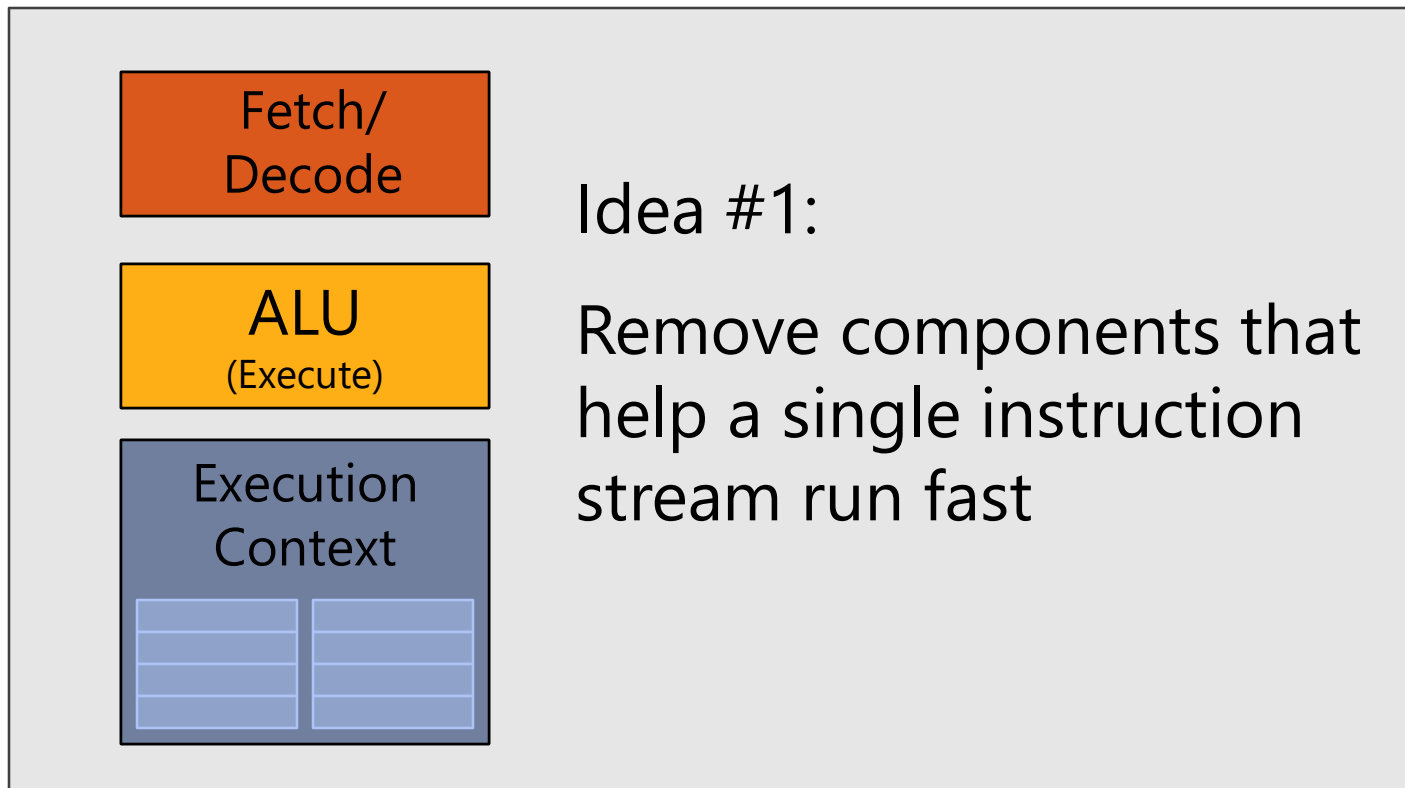
Where this is going...



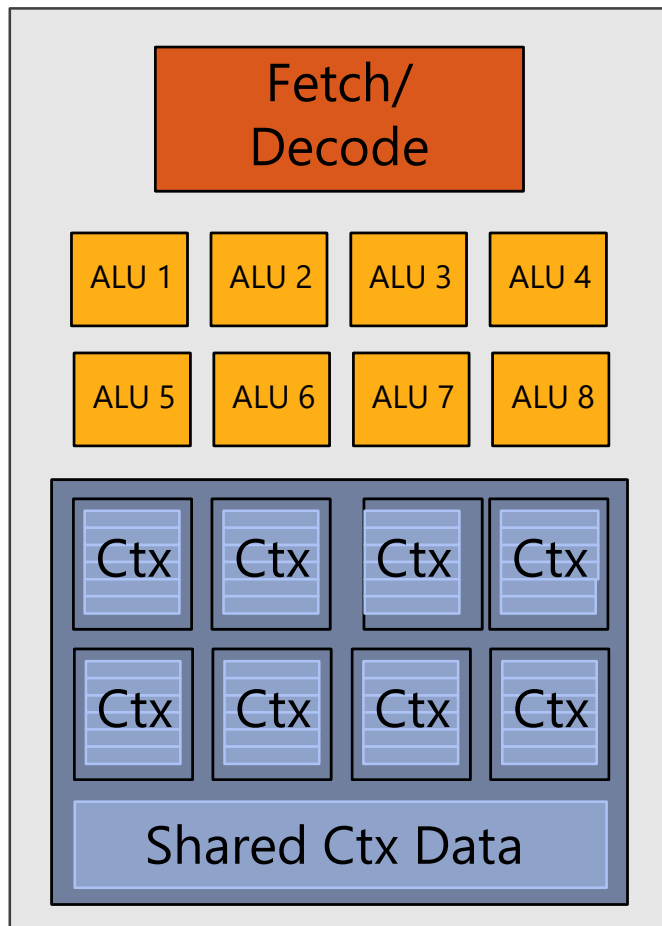
Summary: three key ideas for high-throughput execution

- 1. Use many “slimmed down cores,” run them in parallel**
- 2. Pack cores full of ALUs (by sharing instruction stream overhead across groups of fragments)**
 - Option 1: Explicit SIMD vector instructions**
 - Option 2: Implicit sharing managed by hardware**
- 3. Avoid latency stalls by interleaving execution of many groups of fragments**
 - When one group stalls, work on another group**

Idea #1: Slim down



Idea #2: Add ALUs



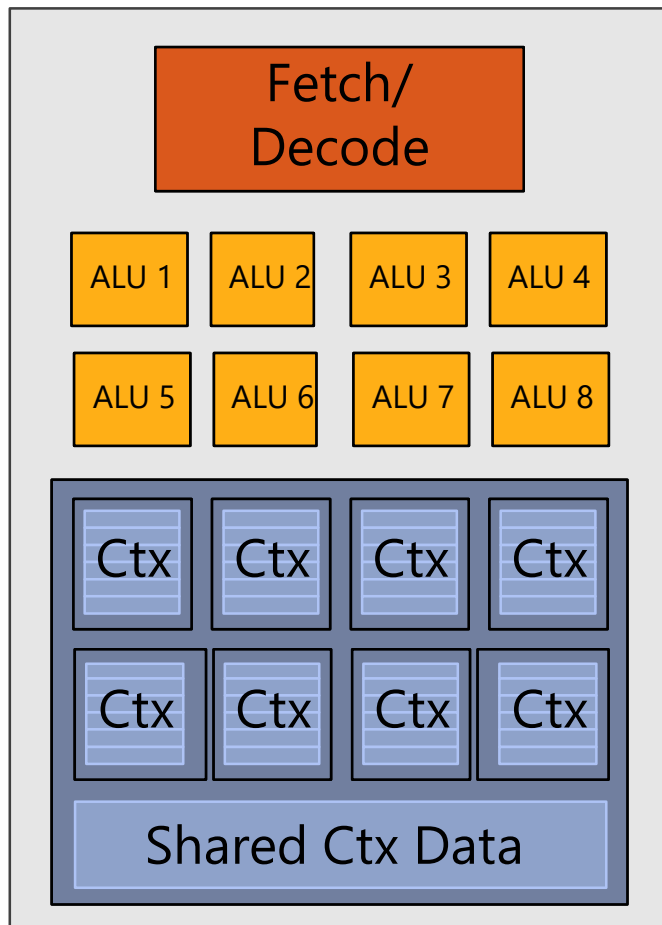
Idea #2:

Amortize cost/complexity of managing an instruction stream across many ALUs

SIMD processing

(or SIMT, SPMD)

Modifying the shader





```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul  r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, l(0.0), l(1.0)  
mul  o0, r0, r3  
mul  o1, r1, r3  
mul  o2, r2, r3  
mov  o3, l(1.0)
```

Original compiled shader:



Processes one fragment
using scalar ops on scalar registers

128 fragments in parallel



16 cores = 128 ALUs
= 16 simultaneous instruction streams


128 [vertices / fragments primitives CUDA threads OpenCL work items compute shader threads] in parallel



Clarification


SIMD processing does not imply SIMD instructions

- Option 1: Explicit vector instructions
 - Intel/AMD x86 SSE, Intel Larrabee
- Option 2: Scalar instructions, implicit HW vectorization
 - HW determines instruction stream sharing across ALUs (amount of sharing hidden from software)
 - NVIDIA GeForce (“SIMT” warps), AMD Radeon architectures



In practice: 16 to 64 fragments share an instruction stream

But what about branches?




<unconditional
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional
shader code>

But what about branches?

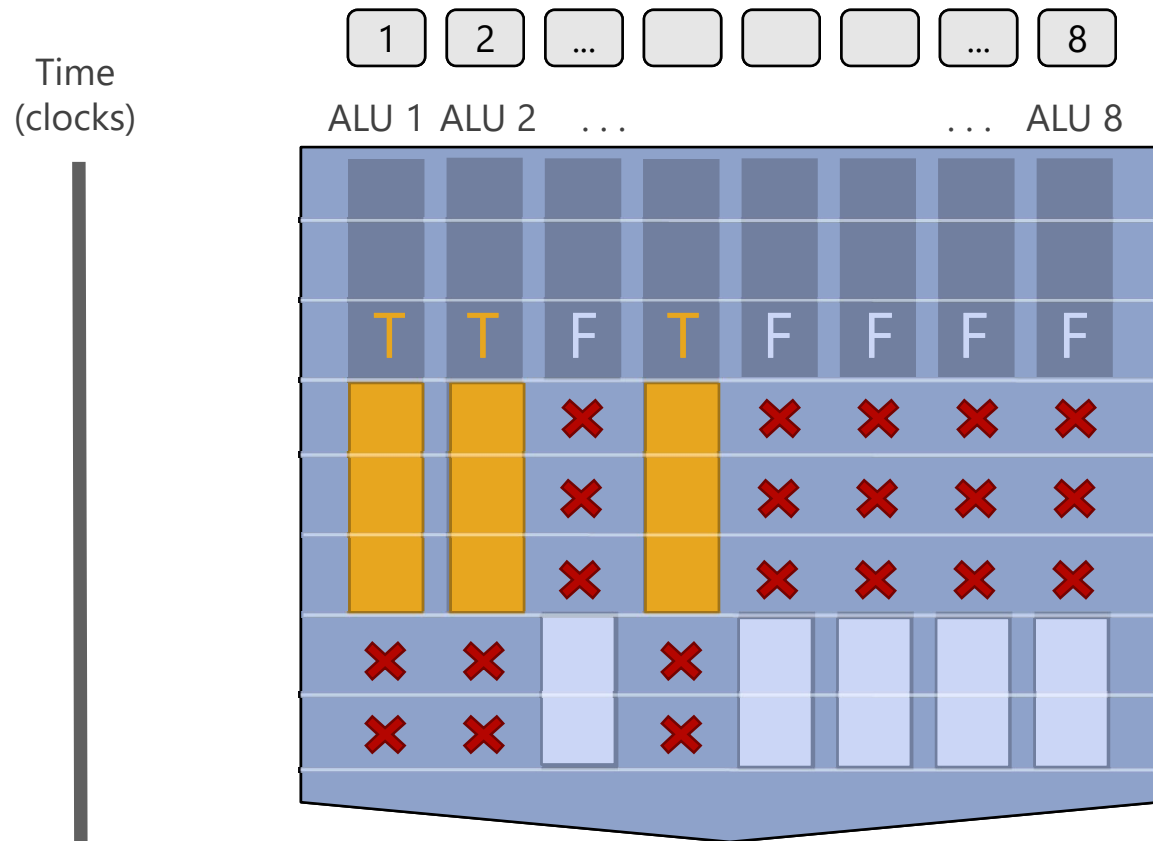


<unconditional
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional
shader code>

But what about branches?



Not all ALUs do useful work!
Worst case: 1/8
performance

<unconditional
shader code>

```
if (x > 0) {
```

```
    y = pow(x, exp);
```

```
    y *= Ks;
```

```
    refl = y + Ka;
```

```
} else {
```


```
    x = 0;
```

```
    refl = Ka;
```

```
}
```

<resume unconditional
shader code>

But what about branches?



<unconditional
shader code>

```
if (x > 0) {
```

```
    y = pow(x, exp);
```

```
    y *= Ks;
```

```
    refl = y + Ka;
```

```
} else {
```


```
    x = 0;
```

```
    refl = Ka;
```

```
}
```

<resume unconditional
shader code>

Idea #3: Store multiple contexts



Stalls!

Stalls occur when a core cannot run the next instruction because of a dependency on a previous operation.

Texture access latency = 100's to 1000's of cycles

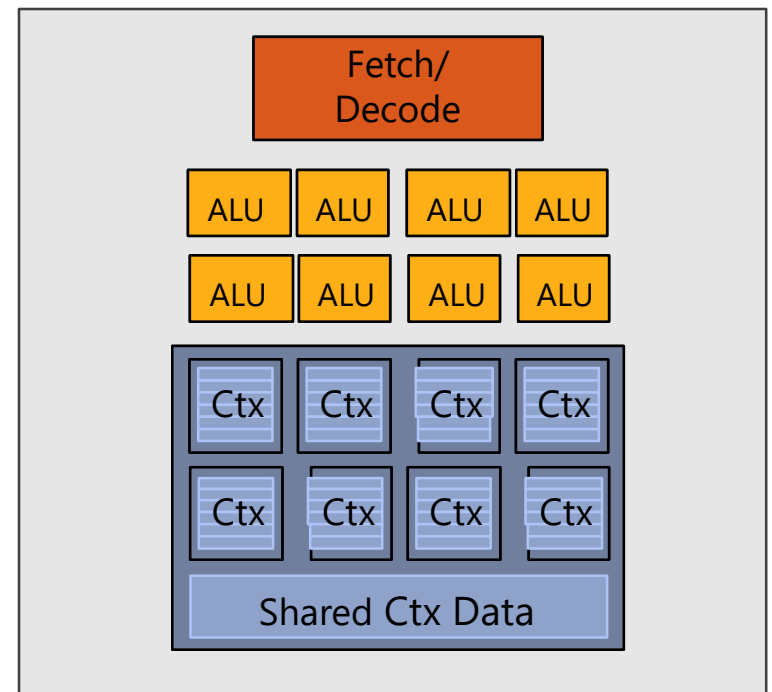

We've removed the fancy caches and logic that helps avoid stalls.

But we have **LOTS** of independent fragments.

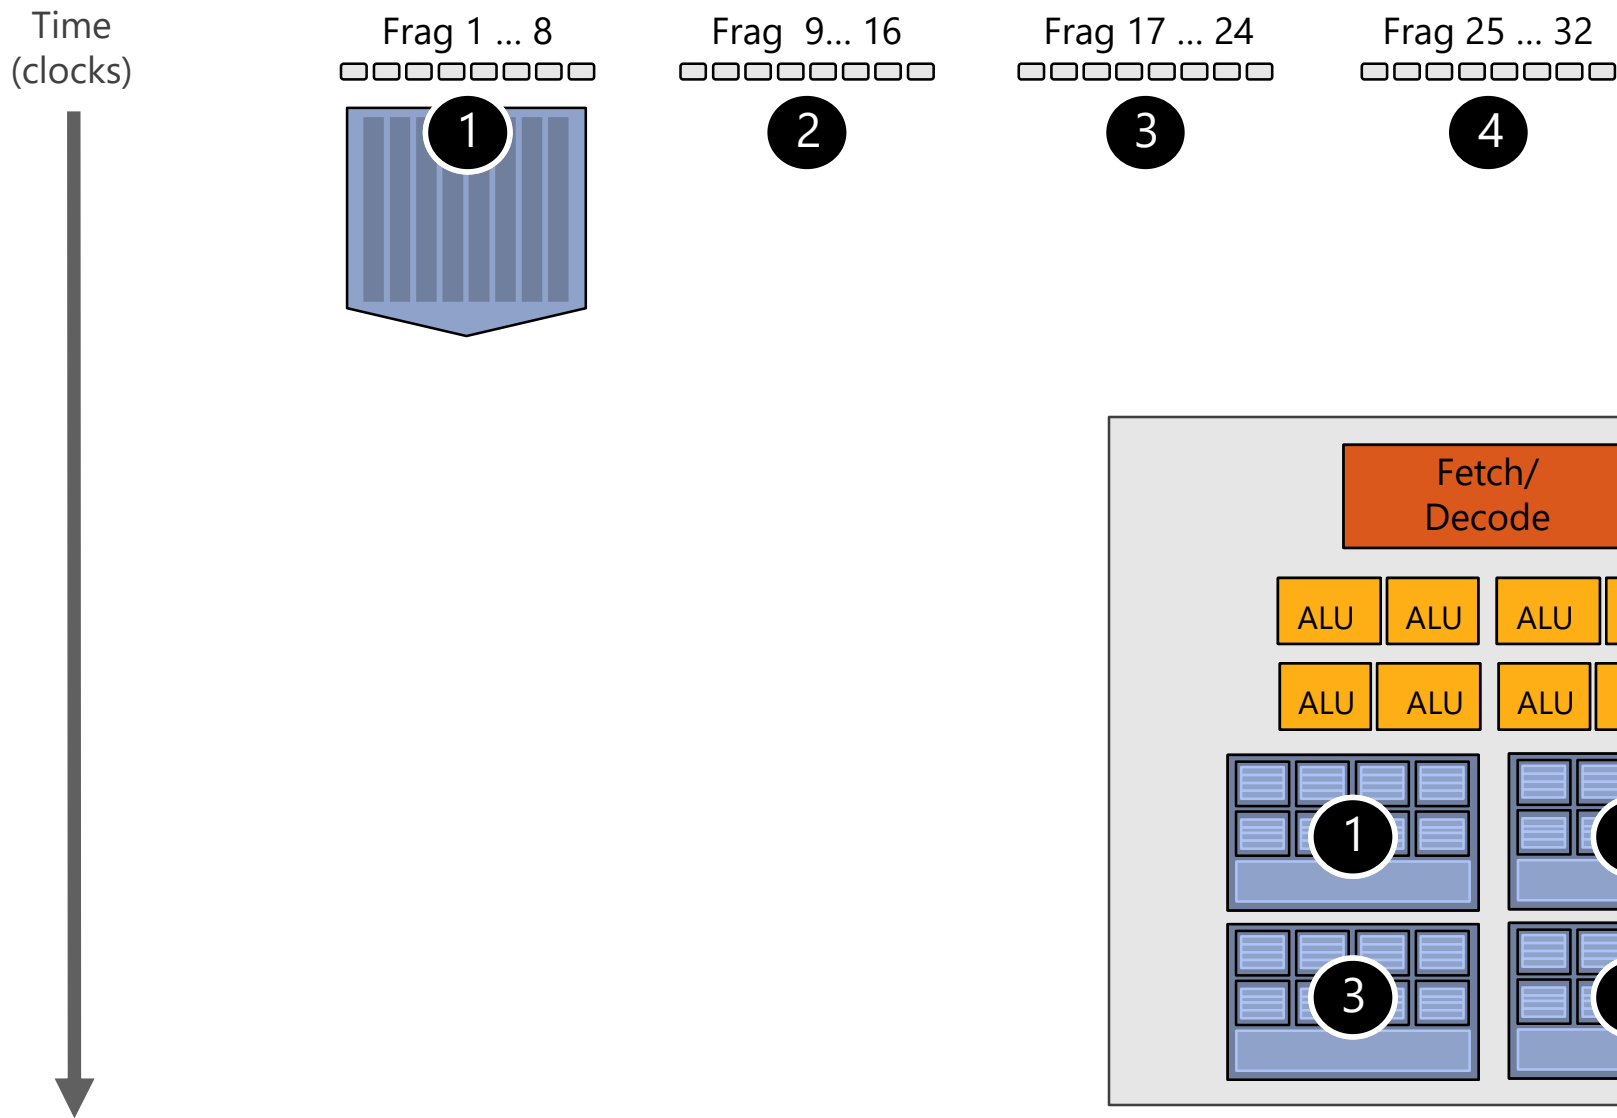
Idea #3:

Interleave processing of many fragments on a single core
to avoid stalls caused by high latency operations.


Hiding shader stalls




Hiding shader stalls




Hiding shader stalls



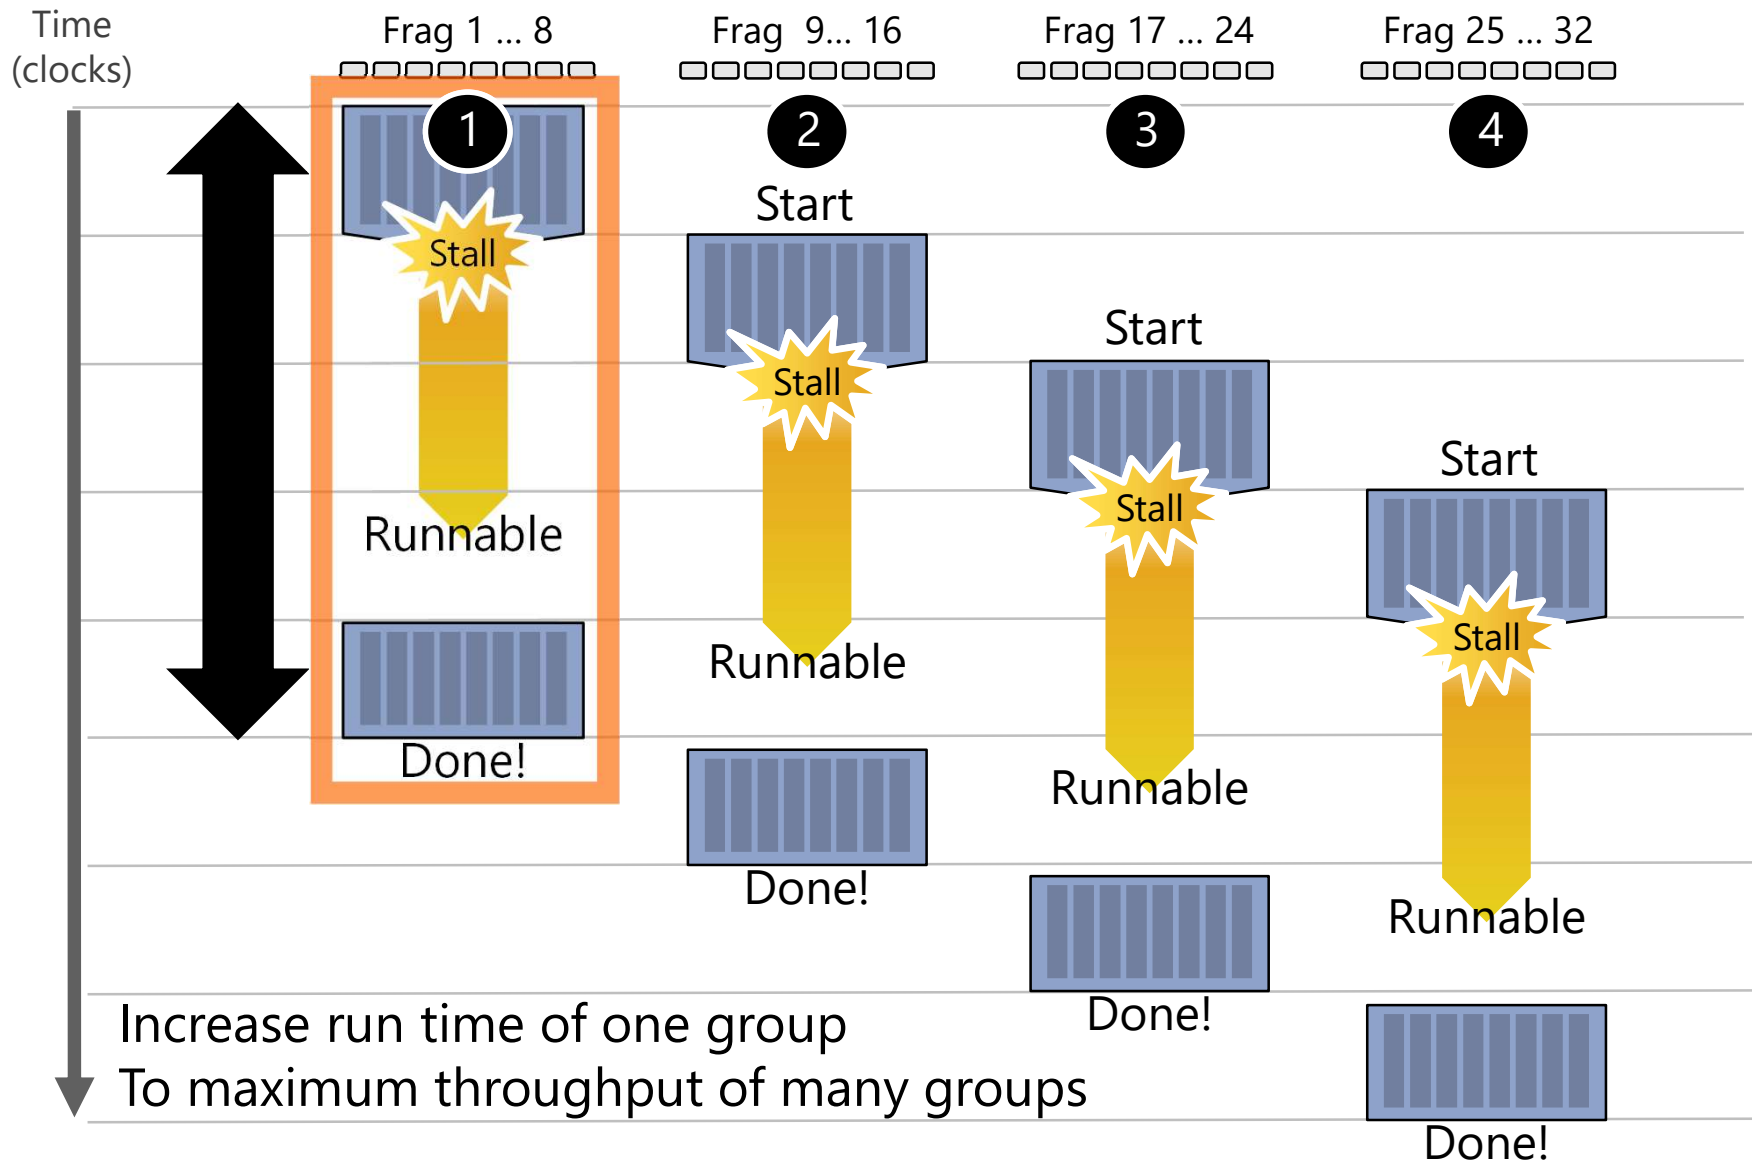
Hiding shader stalls




Hiding shader stalls



Throughput!

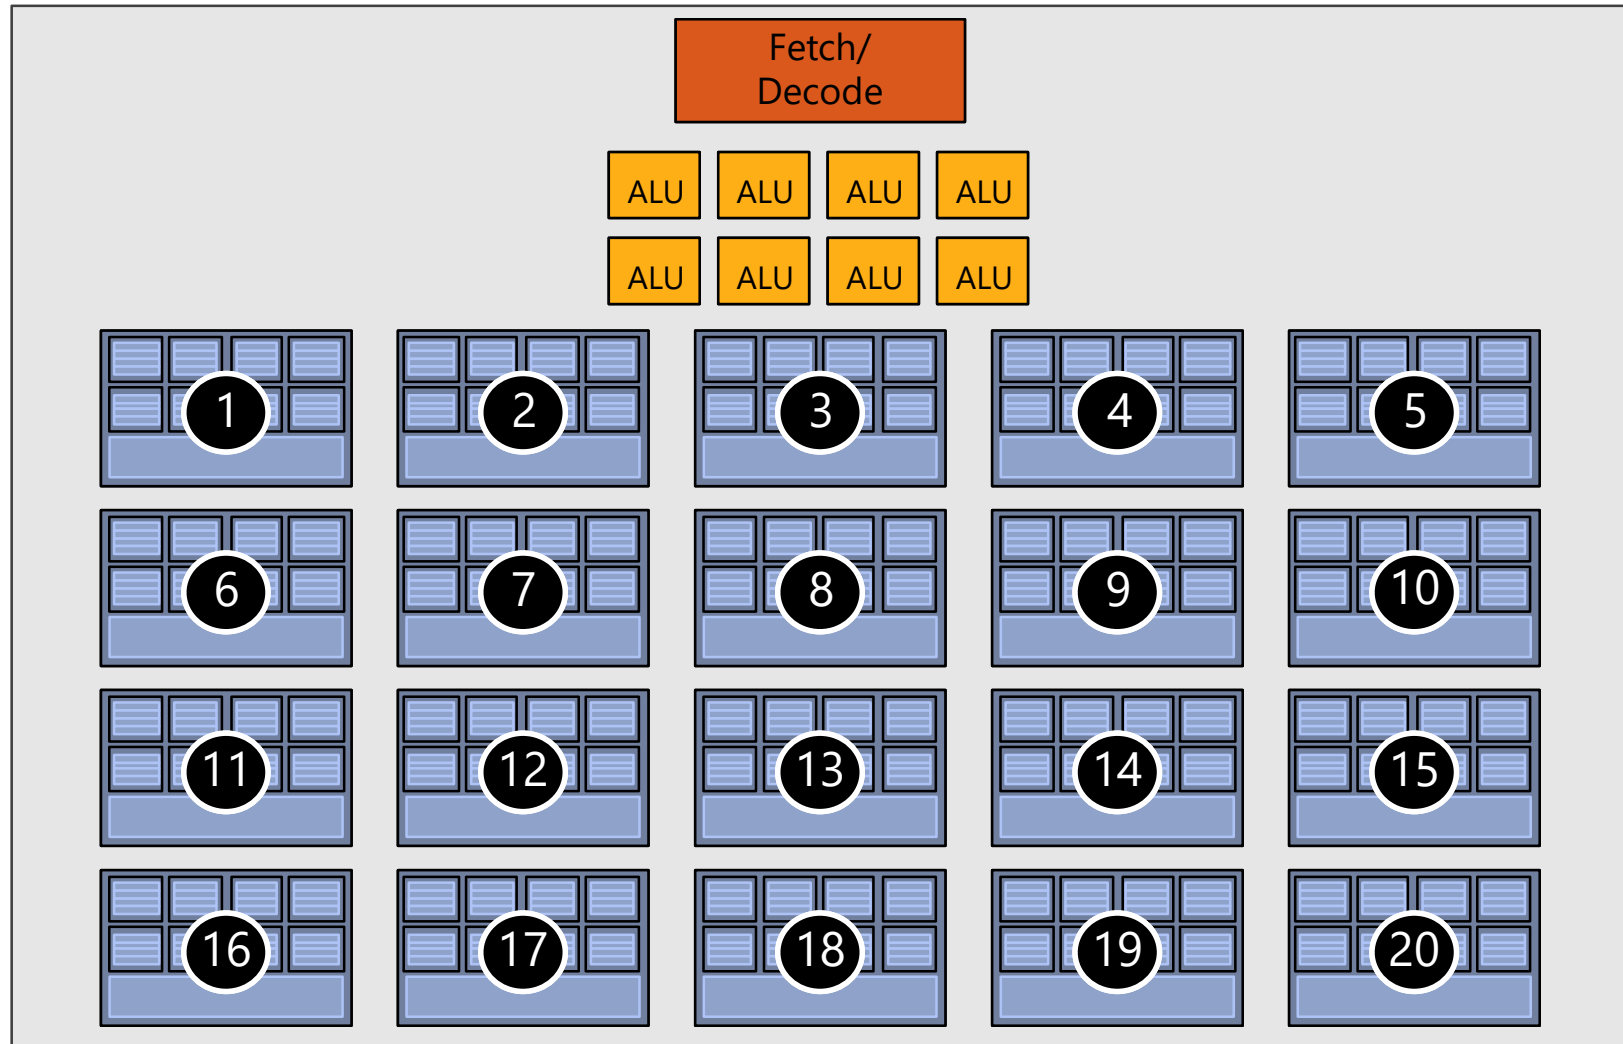


Storing contexts

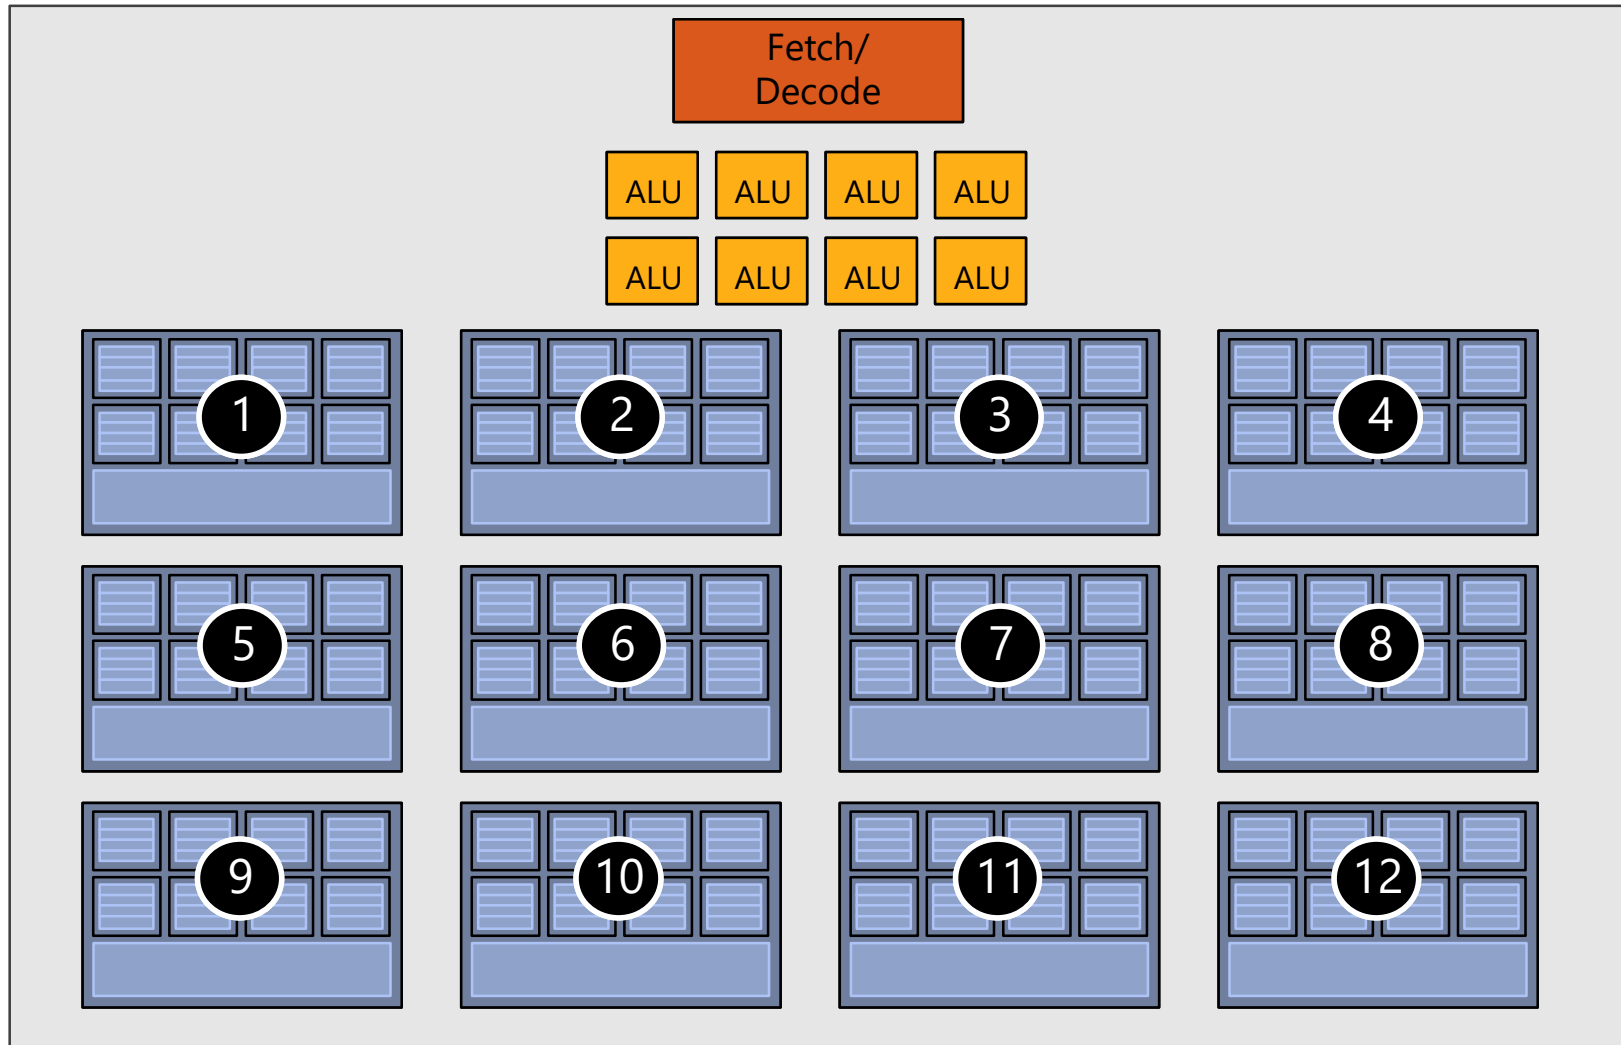


Twenty small contexts

(maximal latency hiding ability)




Twelve medium contexts



Four large contexts

(low latency hiding ability)



My chip!

16 cores


8 mul-add ALUs per core
(128 total)

16 simultaneous
instruction streams


64 concurrent (but interleaved)
instruction streams

512 concurrent fragments

= 256 GFLOPs (@ 1GHz)



My “enthusiast” chip!



32 cores, 16 ALUs per core (512 total) = 1 TFLOP (@ 1 GHz)



Summary: three key ideas for high-throughput execution

- 1. Use many “slimmed down cores,” run them in parallel**
- 2. Pack cores full of ALUs (by sharing instruction stream overhead across groups of fragments)**
 - Option 1: Explicit SIMD vector instructions**
 - Option 2: Implicit sharing managed by hardware**
- 3. Avoid latency stalls by interleaving execution of many groups of fragments**
 - When one group stalls, work on another group**



End of material for Quiz #1 !

Thank you.