

CS 380 - GPU and GPGPU Programming

Lecture 25: Parallel Scan Bank Conflicts; Shuffle Instructions

Markus Hadwiger, KAUST

Reading Assignment #14 (until Dec 7)



Read (required):

- Warp Shuffle Functions
 - CUDA Programming Guide 11.1, Appendix B.21
- CUDA Cooperative Groups (Volta + Turing)
 - <https://devblogs.nvidia.com/cooperative-groups/>
 - CUDA Programming Guide 11.1, Appendix C
- Programming Tensor Cores
 - <https://devblogs.nvidia.com/programming-tensor-cores-cuda-9/>
 - CUDA Programming Guide 11.1, Appendix B.23

Read (optional):

- CUDA Warp-Level Primitives
 - <https://developer.nvidia.com/blog/using-cuda-warp-level-primitives/>
- Warp-aggregated atomics
 - <https://devblogs.nvidia.com/cuda-pro-tip-optimized-filtering-warp-aggregated-atomics/>

Quiz #4: Dec 9



Organization

- First 30 min of lecture
- No material (book, notes, ...) allowed

Content of questions

- Lectures (both actual lectures and slides)
- Reading assignments
- Programming assignments (algorithms, methods)
- Solve short practical examples

Semester Project Presentation Event(s)




Sunday/Monday, Dec 13/14

tbd


$O(n \log n)$ Scan

Courtesy John Owens



- Step efficient ($\log n$ steps)
- Not work efficient ($n \log n$ work)
- Requires barriers at each step (WAR dependencies)


A First-Attempt Parallel Scan Algorithm




1. Read input from device memory to shared memory. Set first element to zero and shift others right by one.
2. Iterate $\log(n)$ times: Threads *stride* to *n*: Add pairs of elements *stride* elements apart. Double *stride* at each iteration. (note must double buffer shared mem arrays)
3. Write output to device memory.

$O(n)$ Scan [Blelloch]


Courtesy John Owens



- Work efficient ($O(n)$ work)
- Bank conflicts, and lots of ‘em




Build the Sum Tree



Iterate $\log(n)$ times. Each thread adds value $stride / 2$ elements away to its own value.

Note that this algorithm operates in-place: no need for double buffering

Build Scan From Partial Sums




Done! We now have a completed scan that we can write out to device memory.

Total steps: $2 * \log(n)$.

Total work: $2 * (n-1)$ adds = $O(n)$ **Work Efficient!**

Build Scan From Partial Sums



Done! We now have a completed scan that we can write out to device memory.

Total steps: $2 * \log(n)$.

Total work: $< 2 * (n-1)$ adds = $O(n)$ **Work Efficient!**

Bank Conflicts in Scan

- Non-power-of-two -

Initial Bank Conflicts on Load


- **Each thread loads two shared mem data elements**
- **Tempting to interleave the loads**

```
temp[2*thid]      = g_idata[2*thid];  
temp[2*thid+1]    = g_idata[2*thid+1];
```
- **Threads:(0,1,2,...,8,9,10,...)→banks:(0,2,4,...,0,2,4,...)**
- **Better to load one element from each half of the array**

```
temp[thid]          = g_idata[thid];  
temp[thid + (n/2)]  = g_idata[thid + (n/2)];
```

Bank Conflicts in the tree algorithm

- When we build the sums, each thread reads two shared memory locations and writes one:
- Th(0,8) access bank 0




First iteration: 2 threads access each of 8 banks.

Each corresponds to a single thread.

Like-colored arrows represent simultaneous memory accesses

Bank Conflicts in the tree algorithm

- When we build the sums, each thread reads two shared memory locations and writes one:
- Th(1,9) access bank 2, etc.




First iteration: 2 threads access each of 8 banks.

Each corresponds to a single thread.

Like-colored arrows represent simultaneous memory accesses

Bank Conflicts in the tree algorithm

- **2nd iteration: even worse!**
 - 4-way bank conflicts; for example:
 $\text{Th}(0,4,8,12)$ access bank 1, $\text{Th}(1,5,9,13)$ access Bank 5, etc.



Each corresponds to a single thread.

Like-colored arrows represent simultaneous memory accesses

Scan Bank Conflicts (1)

- A full binary tree with 64 leaf nodes:

Scale (s)	Thread addresses																																				
1	0	2	4	6	8	10	12	14	16	18	20	22	24	26	28	30	32	34	36	38	40	42	44	46	48	50	52	54	56	58	60	62					
2	0	4	8	12	16	20	24	28	32	36	40	44	48	52	56	60																					
4	0	8	16	24	32	40	48	56																													
8	0	16	32	48																																	
16	0	32																																			
32	0																																				
Conflicts	Banks																																				
2-way	0	2	4	6	8	10	12	14	0	2	4	6	8	10	12	14	0	2	4	6	8	10	12	14	0	2	4	6	8	10	12	14					
4-way	0	4	8	12	0	4	8	12	0	4	8	12	0	4	8	12																					
4-way	0	8	0	8	0	8	0	8																													
4-way	0	0	0	0																																	
2-way	0	0																																			
None	0																																				

- Multiple 2-and 4-way bank conflicts
- Shared memory cost for whole tree
 - 1 32-thread warp = 6 cycles per thread w/o conflicts
 - Counting 2 shared mem reads and one write ($s[a] += s[b]$)
 - $6 * (2+4+4+4+2+1) = 102$ cycles
 - 36 cycles if there were no bank conflicts ($6 * 6$)

Scan Bank Conflicts (2)


- **It's much worse with bigger trees!**
- **A full binary tree with 128 leaf nodes**
 - Only the last 6 iterations shown (root and 5 levels below)

Scale (s)	Thread addresses																												
2	0	4	8	12	16	20	24	28	32	36	40	44	48	52	56	60	64	68	72	76	80	84	88	92	96				
4	0	8	16	24	32	40	48	56	64	72	80	88	96	104	112	120													
8	0	16	32	48	64	80	96	112																					
16	0	32	64	96																									
32	0	64																											
64	0																												
Conflicts	Banks																												
4-way	0	4	8	12	0	4	8	12	0	4	8	12	0	4	8	12	0	4	8	12	0	4	8	12	0	4	8	10	
8-way	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0	8	0
8-way	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4-way	0	0	0	0																									
2-way	0	0																											
None	0																												

- **Cost for whole tree:**
 - $12*2 + 6*(4+8+8+4+2+1) = 186$ cycles
 - 48 cycles if there were no bank conflicts! $12*1 + (6*6)$

Bank Conflicts in the tree algorithm

- We can use padding to prevent bank conflicts
 - Just add a word of padding every 16 words:
- No more conflicts! **32 for full warps!**



Now, within a 16-thread half-warp, all threads access different banks.

32-thread full warp!

(Note that only arrows with the same color happen simultaneously.)

Use Padding to Reduce Conflicts

- **This is a simple modification to the last exercise**
- **After you compute a shared mem address like this:**

```
Address =      stride * thid;
```

- **Add padding like this:**

```
Address += (Address >> 4); // divide by NUM_BANKS
```

- **This removes most bank conflicts**
 - Not all, in the case of deep trees

Fixing Scan Bank Conflicts

- Insert padding every NUM_BANKS elements

```
const int LOG_NUM_BANKS = 4; // 16 banks
int tid = threadIdx.x;
int s = 1;
// Traversal from leaves up to root
for (d = n>>1; d > 0; d >>= 1)
{
    if (thid <= d)
    {
        int a = s*(2*tid); int b = s*(2*tid+1)
        a += (a >> LOG_NUM_BANKS); // insert pad word
        b += (b >> LOG_NUM_BANKS); // insert pad word
        shared[a] += shared[b];
    }
}
```

Fixing Scan Bank Conflicts

- A full binary tree with 64 leaf nodes

Leaf Nodes	Scale (s)	Thread addresses
64	1	0 2 4 6 8 10 12 14 17 19 21 23 25 27 29 31 34 36 38 40 42 44 46 48 51 53 55 57 59 61 63
	2	0 4 8 12 17 21 25 29 34 38 42 46 51 55 59 63
	4	0 8 17 25 34 42 51 59
	8	0 17 34 51
	16	0 34
	32	0
Conflicts Banks		
None	0 2 4 6 8 10 12 14 1 3 5 7 9 11 13 15 2 4 6 8 10 12 14 0 3 5 7 9 11 13 15	
None	0 4 8 12 1 5 9 13 2 6 10 14 3 7 11 15	
None	0 8 1 9 2 10 3 11	
None	0 1 2 3	
None	0 2	
None	0	

- No more bank conflicts!
 - However, there are ~8 cycles overhead for addressing
 - For each $s[a] += s[b]$ (8 cycles/iter. * 6 iter. = 48 extra cycles)
 - So just barely worth the overhead on a small tree
 - 84 cycles vs. 102 with conflicts vs. 36 optimal

Fixing Scan Bank Conflicts

- A full binary tree with 128 leaf nodes
 - Only the last 6 iterations shown (root and 5 levels below)

Scale (s)	Thread addresses															
2	0	4	8	12	17	21	25	29	34	38	42	46	51	55	59	63
4	0	8	17	25	34	42	51	59	68	76	85	93	102	110	119	127
8	0	17	34	51	68	85	102	119								
16	0	34	68	102												
32	0	68														
64	0															

Conflicts	Banks
None	0 4 8 12 1 5 9 13 2 6 10 14 3 7 11 15 4 8 12 0 5 9 13 1 6 10 14 2 7 11 15 3
None	0 8 1 9 2 10 3 11 4 12 5 13 6 14 7 15
None	0 1 2 3 4 5 6 7
None	0 2 4 6
None	0 4
None	0

- No more bank conflicts!
 - Significant performance win:
 - 106 cycles vs. 186 with bank conflicts vs. 48 optimal

Fixing Scan Bank Conflicts

- A full binary tree with 512 leaf nodes
 - Only the last 6 iterations shown (root and 5 levels below)

Scale (s)	Thread addresses																																		
8	0	17	34	51	68	85	102	119	136	153	170	187	204	221	238	255	272	289	306	323	340	357	374	391	408	425	442	459	476	493	510	527			
16	0	34	68	102	136	170	204	238	272	306	340	374	408	442	476	510																			
32	0	68	136	204	272	340	408	476																											
64	0	136	272	408																															
128	0	272																																	
256	0																																		

Conflicts	Banks																																						
None	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15							
2-way	0	2	4	6	8	10	12	14	0	2	4	6	8	10	12	14																							
2-way	0	4	8	12	0	4	8	12																															
2-way	0	8	0	8																																			
2-way	0	0																																					
None	0																																						

- Wait, we still have bank conflicts
 - Method is not foolproof, but still much improved
 - 304 cycles vs. 570 with bank conflicts vs. 120 optimal
- But it does not pay off to optimize for the rest. Address calculations are getting too expensive

Summary

- **Parallel Programming requires careful planning**
 - of the branching behavior
 - of the memory access patterns
 - of the work efficiency
- **Vector Reduction**
 - branch efficient
 - bank efficient
- **Scan Algorithm**
 - based in Balanced Tree principle:
bottom up, top down traversal

GPU TECHNOLOGY
CONFERENCE

Shuffle: Tips and Tricks

Julien Demouth, NVIDIA

Glossary

Safer with cooperative thread groups!


- Warp
 - ~~Implicitly synchronized~~ group of threads (32 on current HW)
- Warp ID (`warpid`)
 - Identifier of the warp in a block: `threadIdx.x / 32`
- Lane ID (`laneid`)
 - Coordinate of the thread in a warp: `threadIdx.x % 32`
 - Special register (available from PTX): `%laneid`

Shuffle (SHFL)

- Instruction to exchange data in a warp
- Threads can “read” other threads’ registers
- No shared memory is needed
- It is available starting from SM 3.0

Variants

- 4 variants (idx, up, down, bfly):



Now: Use `_sync` variants / shuffle in cooperative thread groups!

Instruction (PTX)

Optional dst. predicate Lane/offset/mask
shfl.mode.b32 d[|p], a, b, c;
 ↑
 Dst. register Src. register Bound

Now: Use _sync variants / shuffle in cooperative thread groups!

Implement SHFL for 64b Numbers

```
__device__ __inline__ double shfl(double x, int lane)
{
    // Split the double number into 2 32b registers.
    int lo, hi;
    asm volatile( "mov.b32 {%0,%1}, %2;" : "=r"(lo), "=r"(hi) : "d"(x));

    // Shuffle the two 32b registers.
    lo = __shfl(lo, lane);
    hi = __shfl(hi, lane);


    // Recreate the 64b number.
    asm volatile( "mov.b64 %0, {%1,%2};" : "=d(x)" : "r"(lo), "r"(hi));

    return x;
}
```


- Generic SHFL: <https://github.com/BryanCatanzaro/generics>

Performance Experiment

- One element per thread



- Each thread takes its right neighbor



Performance Experiment

- We run the following test on a K20

```
T x = input[tidx];
for(int i = 0 ; i < 4096 ; ++i)
    x = get_right_neighbor(x);
output[tidx] = x;
```

- We launch 26 blocks of 1024 threads
 - On K20, we have 13 SMs
 - We need 2048 threads per SM to have 100% of occupancy
- We time different variants of that kernel

Performance Experiment

- Shared memory (SMEM)

```
smem[threadIdx.x] = smem[32*warpid + ((laneid+1) % 32)];  
__syncthreads();
```



- Shuffle (SHFL)

```
x = __shfl(x, (laneid+1) % 32);
```



- Shared memory without __syncthreads + volatile (*unsafe*)

```
__shared__ volatile T *smem = ...;  
smem[threadIdx.x] = smem[32*warpid + ((laneid+1) % 32)];
```

Performance Experiment (fp32)



Performance Experiment (fp64)



Performance Experiment

- Always faster than shared memory
- Much safer than using no `__syncthreads` (and volatile)
 - And never slower
- Does not require shared memory
 - Useful when occupancy is limited by SMEM usage

Broadcast

Now: Use cooperative thread groups!

- All threads read from a single lane

```
x = __shfl(x, 0); // All the threads read x from laneid 0.
```

- More complex example

```
// All threads evaluate a predicate.  
int predicate = ...;  
  
// All threads vote.  
unsigned vote = __ballot(predicate);  
  
// All threads get x from the "last" lane which evaluated the predicate to true.  
if(vote)  
    x = __shfl(x, __bfind(vote));  
  
// __bind(unsigned i): Find the most significant bit in a 32/64 number (PTX).  
__bfind(&b, i) { asm volatile("bfind.u32 %0, %1;" : "=r"(b) : "r"(i)); }
```

Reduce

■ Code

```
// Threads want to reduce the value in x.  
  
float x = ...;  
  
#pragma unroll  
for(int mask = WARP_SIZE / 2 ; mask > 0 ; mask >>= 1)  
    x += __shfl_xor(x, mask);  
  
// The x variable of Laneid 0 contains the reduction.
```

■ Performance

- Launch 26 blocks of 1024 threads
- Run the reduction 4096 times

Execution Time fp32 (ms)



SMEM per Block fp32 (KB)




Scan

■ Code

```
#pragma unroll
for( int offset = 1 ; offset < 32 ; offset <= 1 )
{
    float y = __shfl_up(x, offset);
    if(laneid() >= offset)
        x += y;
}
```

■ Performance

- Launch 26 blocks of 1024 threads
- Run the reduction 4096 times




Scan

- Use the predicate from SHFL


```
#pragma unroll
for( int offset = 1 ; offset < 32 ; offset <= 1 )
{
    asm volatile( "{"
        "    .reg .f32 r0;" 
        "    .reg .pred p;" 
        "    shfl.up.b32 r0|p, %0, %1, 0x0;" 
        "    @p add.f32 r0, r0, %0;" 
        "    mov.f32 %0, r0;" 
    "}" : "+f"(x) : "r"(offset));
}
```

- Use CUB:
<https://nvlabs.github.com/cub>




Bitonic Sort

x: [11 3 8 5 10 15 9 7 12 4 2 0 14 13 6 1 ...]



Bitonic Sort




Bitonic Sort

```
int swap(int x, int mask, int dir)
{
    int y = __shfl_xor(x, mask);
    return x < y == dir ? y : x;
}


x = swap(x, 0x01, bfe(laneid, 1) ^ bfe(laneid, 0)); // 2
x = swap(x, 0x02, bfe(laneid, 2) ^ bfe(laneid, 1)); // 4
x = swap(x, 0x01, bfe(laneid, 2) ^ bfe(laneid, 0));
x = swap(x, 0x04, bfe(laneid, 3) ^ bfe(laneid, 2)); // 8
x = swap(x, 0x02, bfe(laneid, 3) ^ bfe(laneid, 1));
x = swap(x, 0x01, bfe(laneid, 3) ^ bfe(laneid, 0));
x = swap(x, 0x08, bfe(laneid, 4) ^ bfe(laneid, 3)); // 16
x = swap(x, 0x04, bfe(laneid, 4) ^ bfe(laneid, 2));
x = swap(x, 0x02, bfe(laneid, 4) ^ bfe(laneid, 1));
x = swap(x, 0x01, bfe(laneid, 4) ^ bfe(laneid, 0));
x = swap(x, 0x10, bfe(laneid, 4)); // 32
x = swap(x, 0x08, bfe(laneid, 3));
x = swap(x, 0x04, bfe(laneid, 2));
x = swap(x, 0x02, bfe(laneid, 1));
x = swap(x, 0x01, bfe(laneid, 0));

// int bfe(int i, int k): Extract k-th bit from i
// PTX: bfe dst, src, start, len (see p.81, ptx_isa_3.1)
```

Execution Time int32 (ms)




SMEM per Block (KB)



Transpose

- When threads load or store arrays of structures, transposes enable fully coalesced memory operations
- e.g. when loading, have the warp perform coalesced loads, then transpose to send the data to the appropriate thread



Transpose

- You can use SMEM to implement this transpose, or you can use SHFL
- Code:
<http://github.com/bryancatanzaro/trove>
- Performance
 - Launch 104 blocks of 256 threads
 - Run the transpose 4096 times

Execution Time 7*int 32





SMEM per Block (KB)



Array of Structures Access via Transpose

- Transpose speeds access to arrays of structures
- High-level interface: `coalesced_ptr<T>`
 - Just dereference like any pointer
 - Up to 6x faster than direct compiler generated access



Conclusion

- SHFL is available for SM \geq SM 3.0
- It is always faster than “safe” shared memory
- It is never slower than “unsafe” shared memory
- It can be used in many different algorithms

Thank you.

- Hendrik Lensch, Robert Strzodka
- NVIDIA