

# CS 380 - GPU and GPGPU Programming

## Lecture 17: GPU Texturing, Pt. 4; Virtual Texturing, Virtual Geometry

Markus Hadwiger, KAUST



# Reading Assignment #10 (until Nov 8)

Read (required):

- **Brook for GPUs: Stream Computing on Graphics Hardware**

Ian Buck et al., SIGGRAPH 2004

<http://graphics.stanford.edu/papers/brookgpu/>

Read (optional):

- **The Imagine Stream Processor**

Ujval Kapasi et al.; IEEE ICCD 2002

<http://cva.stanford.edu/publications/2002/imagine-overview-iccd/>

- **Merrimac: Supercomputing with Streams**

Bill Dally et al.; SC 2003

<https://dl.acm.org/citation.cfm?doid=1048935.1050187>



# Quiz #2: Nov 10

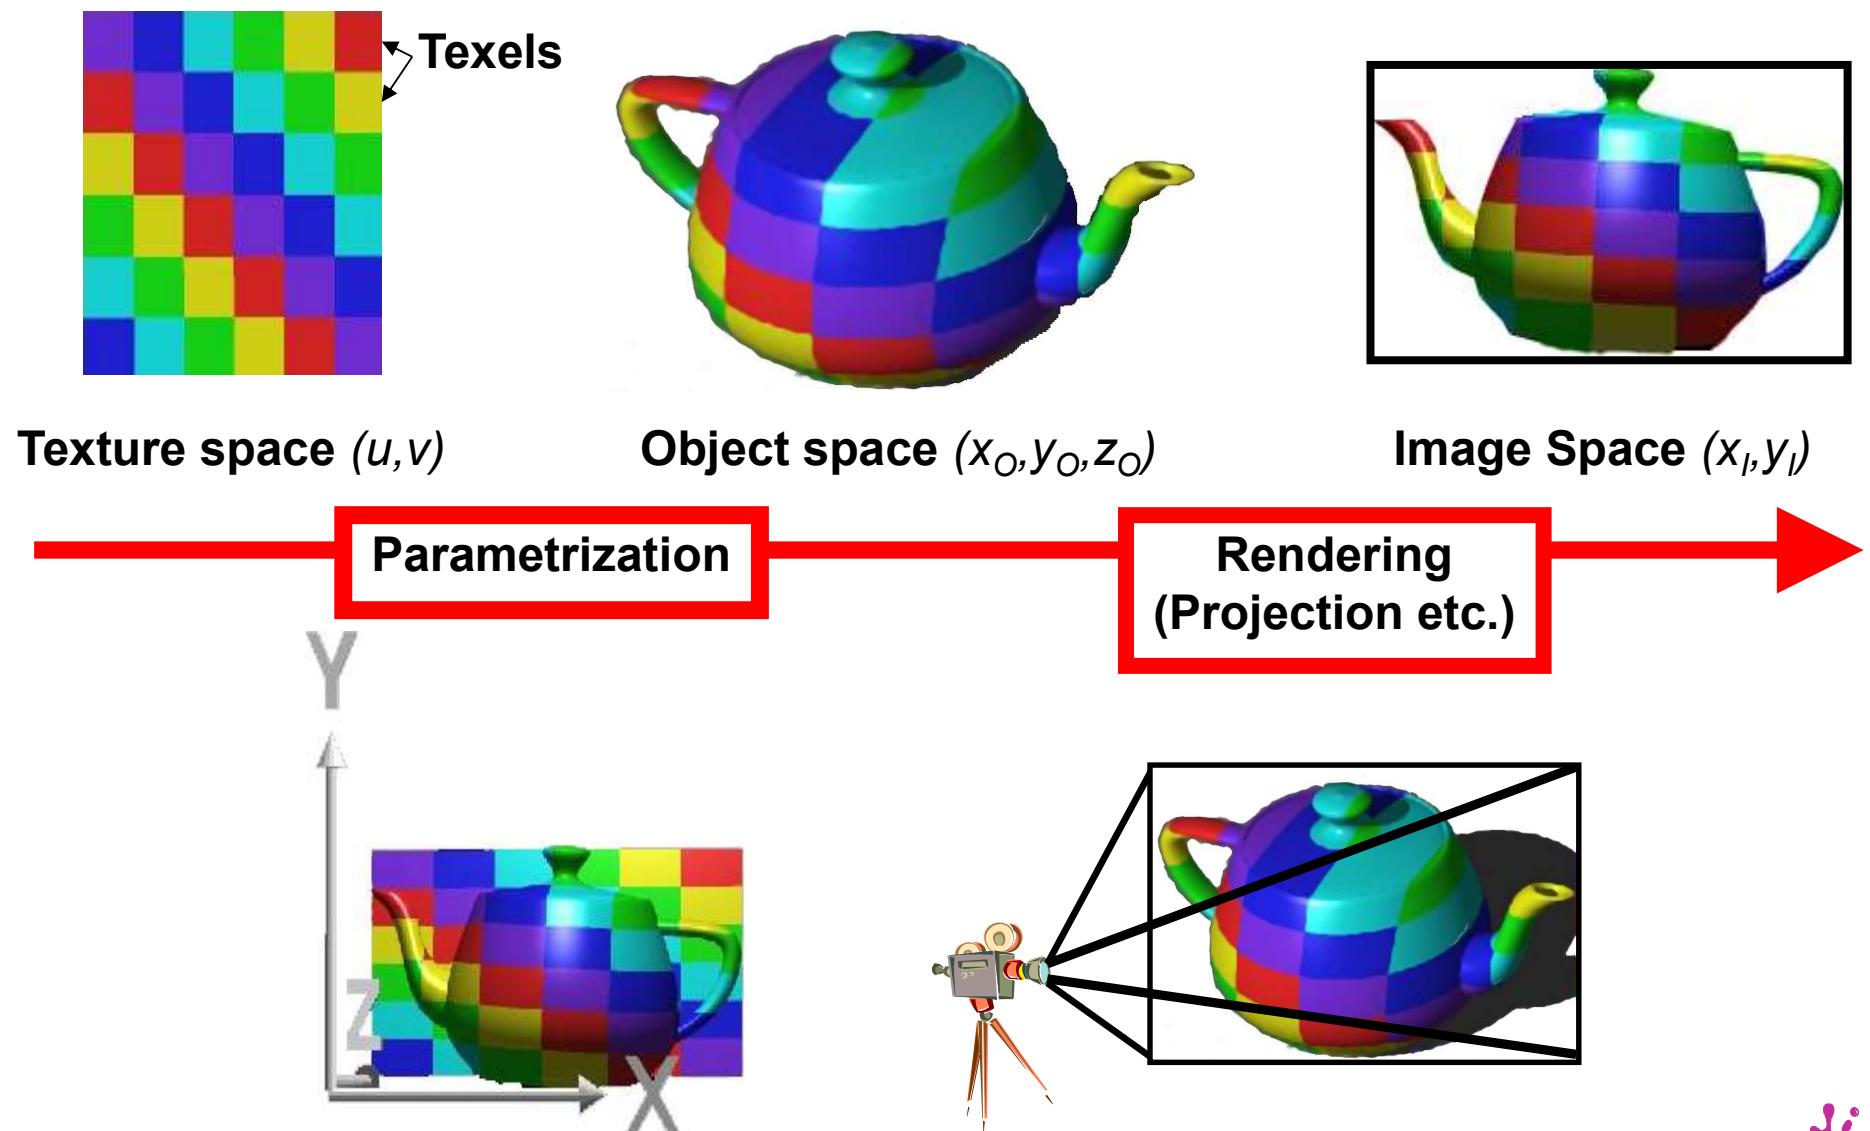
## Organization

- First 30 min of lecture
- No material (book, notes, ...) allowed

## Content of questions

- Lectures (both actual lectures and slides)
- Reading assignments
- Programming assignments (algorithms, methods)
- Solve short practical examples

# Texturing: General Approach





# Interpolation Type + Purpose #2: **Interpolation of Samples in Texture Space**

*(Multi-Linear Interpolation)*

# Magnification (Bi-linear Filtering Example)



Original image



Nearest neighbor



Bi-linear filtering

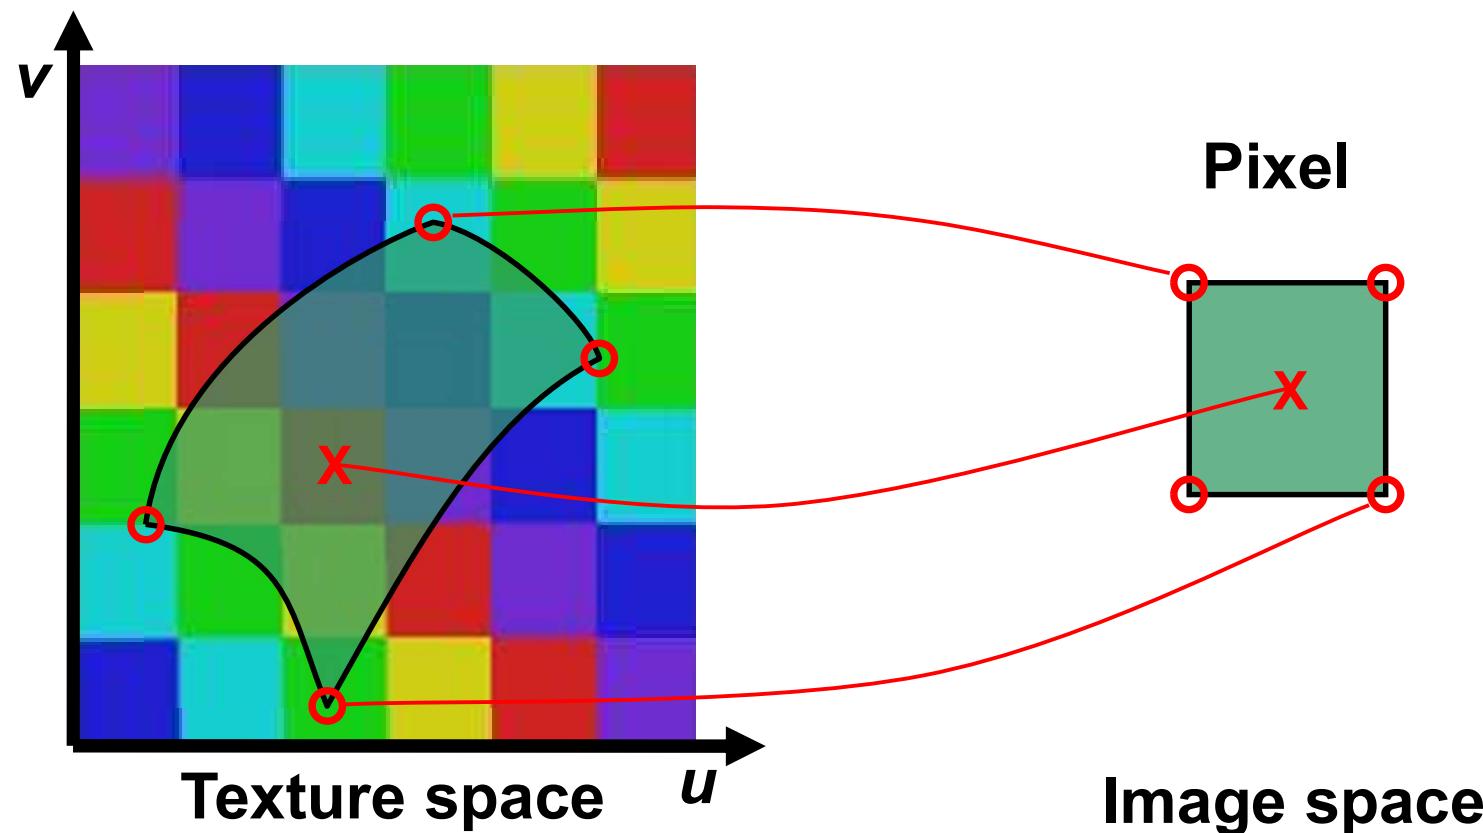




# Texture Minification

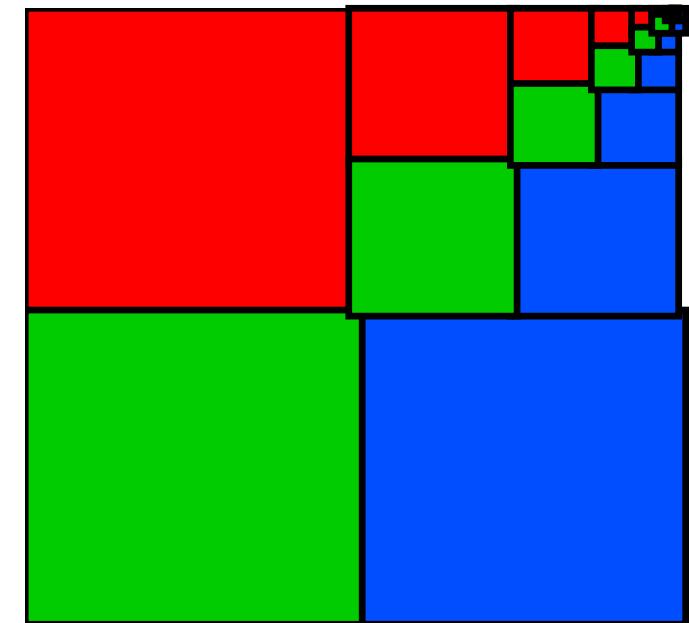
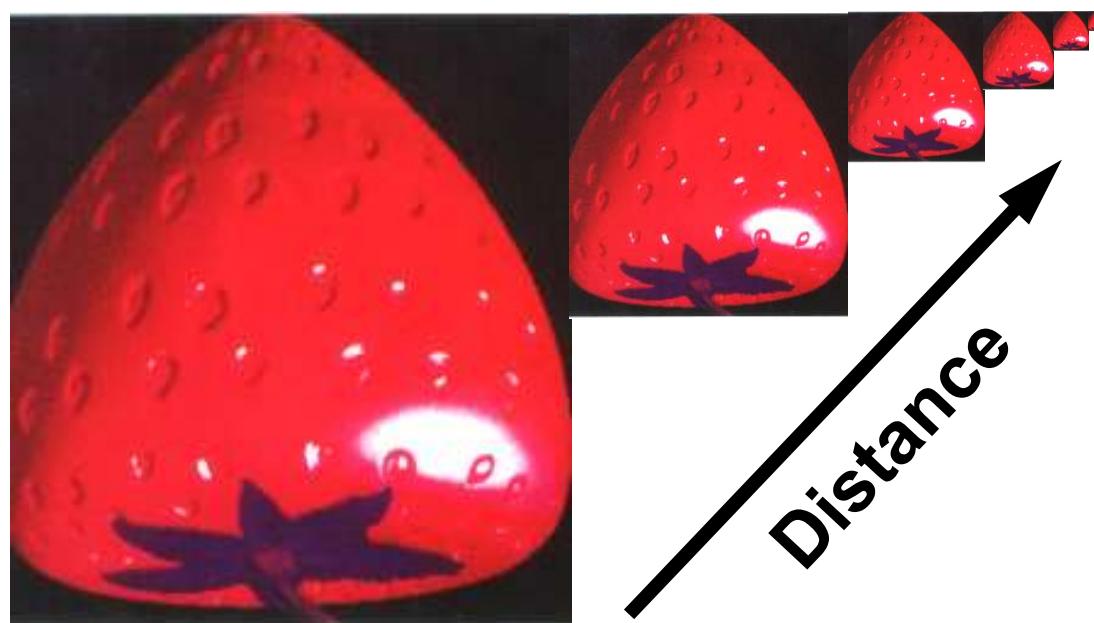
# Texture Anti-Aliasing: Minification

- A good pixel value is the weighted mean of the pixel area projected into texture space



# Texture Anti-Aliasing: MIP Mapping

- MIP Mapping (“Multum In Parvo”)
  - Texture size is reduced by factors of 2 (*downsampling* = "many things in a small place")
  - Simple (4 pixel average) and memory efficient
  - Last image is only ONE texel



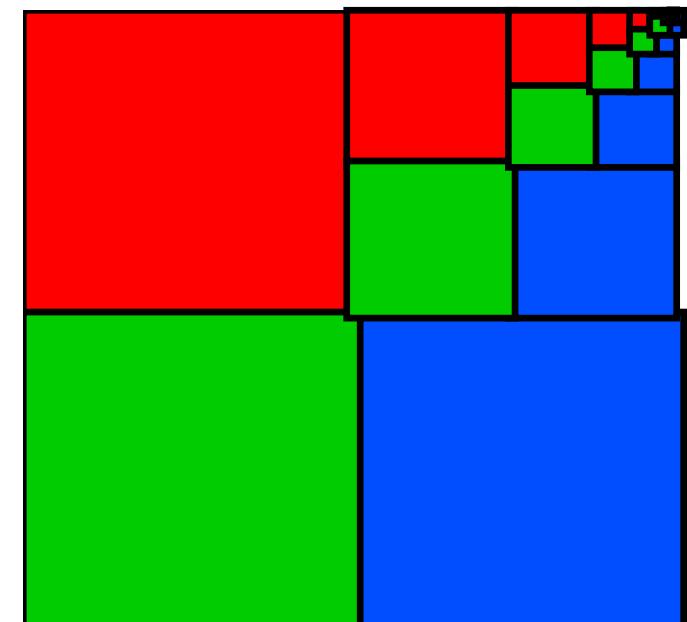
# Texture Anti-Aliasing: MIP Mapping

- MIP Mapping (“Multum In Parvo”)
  - Texture size is reduced by factors of 2 (*downsampling* = "many things in a small place")
  - Simple (4 pixel average) and memory efficient
  - Last image is only ONE texel

geometric series:

$$a + ar + ar^2 + ar^3 + \dots + ar^{n-1} =$$

$$= \sum_{k=0}^{n-1} ar^k = a \left( \frac{1 - r^n}{1 - r} \right)$$

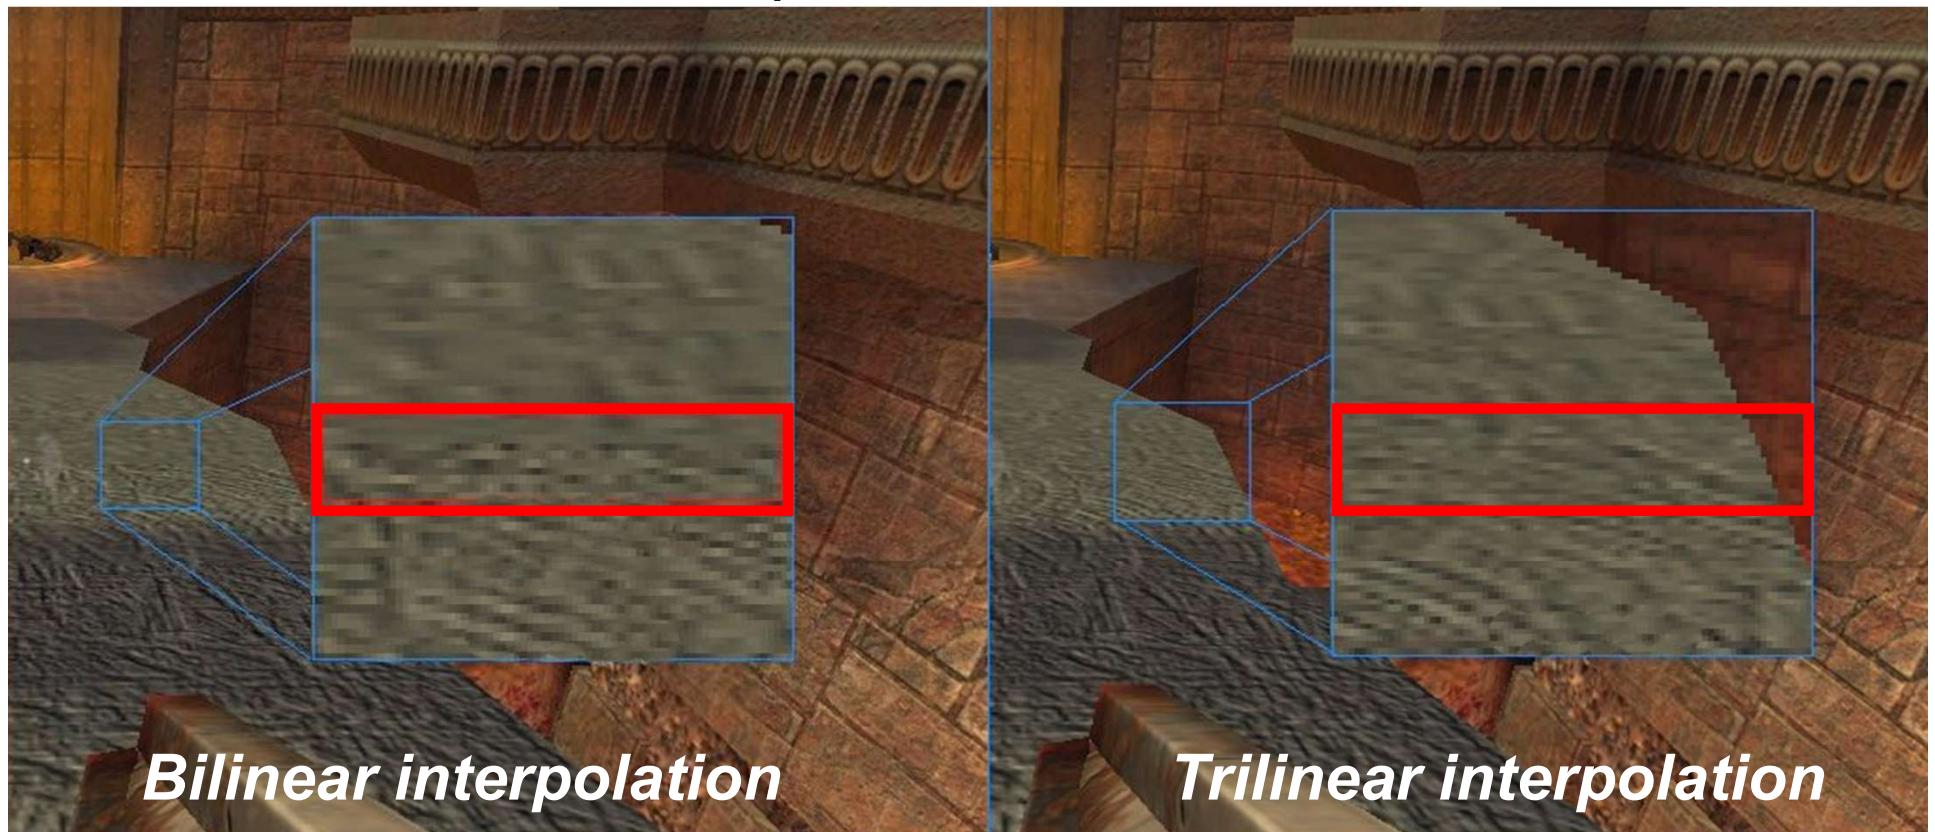
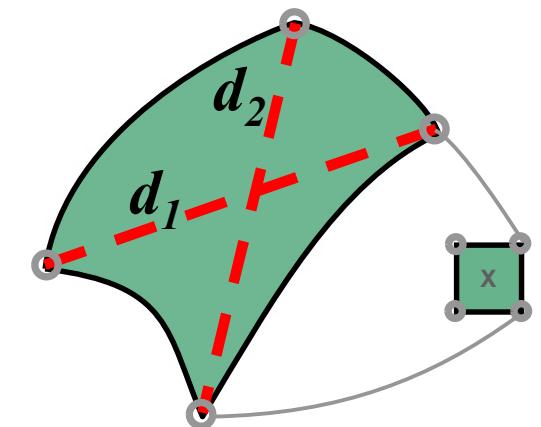


# Texture Anti-Aliasing: MIP Mapping

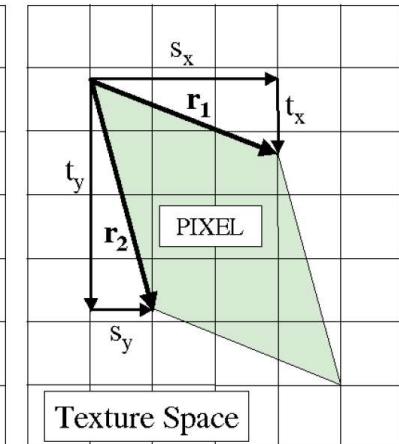
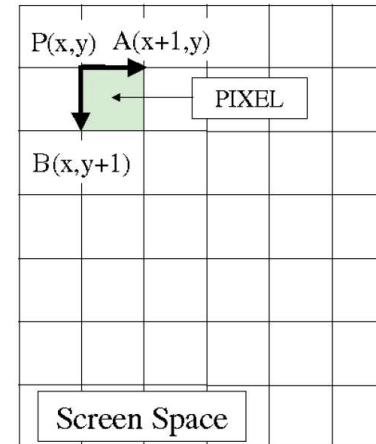
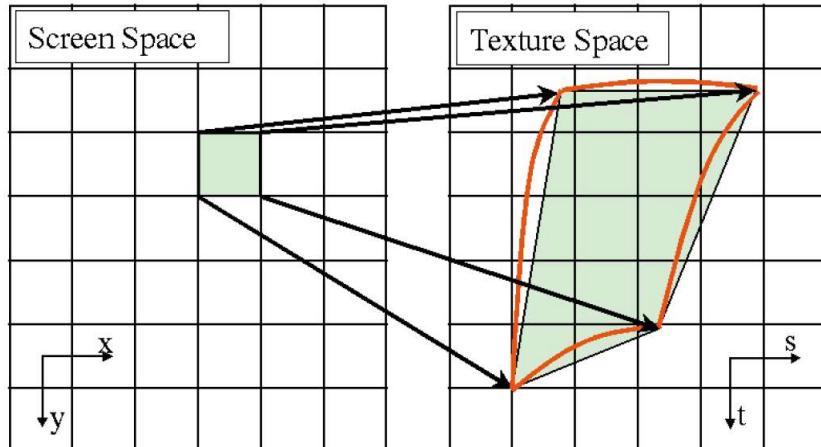
- MIP Mapping Algorithm

- $D := ld(\max(d_1, d_2))$       "Mip Map level"
- $T_0 := \text{value from texture } D_0 = \text{trunc}(D)$

- Use *bilinear interpolation*



# MIP-Map Level Computation



- Use the partial derivatives of texture coordinates with respect to screen space coordinates
- This is the Jacobian matrix
- Area of parallelogram is the absolute value of the Jacobian determinant (the *Jacobian*)

$$\begin{pmatrix} \partial u / \partial x & \partial u / \partial y \\ \partial v / \partial x & \partial v / \partial y \end{pmatrix} = \begin{pmatrix} s_x & s_y \\ t_x & t_y \end{pmatrix}$$



# MIP-Map Level Computation (OpenGL)

- OpenGL 4.6 core specification, pp. 251-264  
(3D tex coords!)

$$\lambda_{base}(x, y) = \log_2[\rho(x, y)]$$

$$\rho = \max \left\{ \sqrt{\left(\frac{\partial u}{\partial x}\right)^2 + \left(\frac{\partial v}{\partial x}\right)^2 + \left(\frac{\partial w}{\partial x}\right)^2}, \sqrt{\left(\frac{\partial u}{\partial y}\right)^2 + \left(\frac{\partial v}{\partial y}\right)^2 + \left(\frac{\partial w}{\partial y}\right)^2} \right\}$$

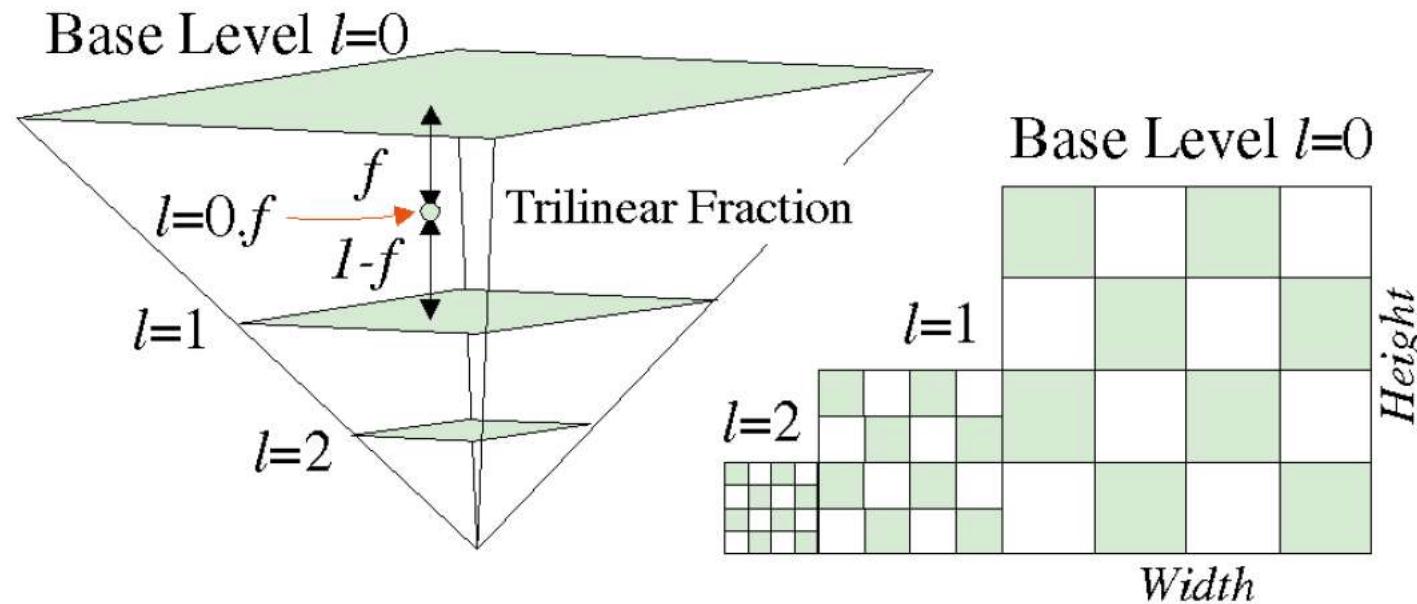
Does not use area of parallelogram but greater hypotenuse [Heckbert, 1983]

- Approximation without square-roots

$$m_u = \max \left\{ \left| \frac{\partial u}{\partial x} \right|, \left| \frac{\partial u}{\partial y} \right| \right\} \quad m_v = \max \left\{ \left| \frac{\partial v}{\partial x} \right|, \left| \frac{\partial v}{\partial y} \right| \right\} \quad m_w = \max \left\{ \left| \frac{\partial w}{\partial x} \right|, \left| \frac{\partial w}{\partial y} \right| \right\}$$

$$\max\{m_u, m_v, m_w\} \leq f(x, y) \leq m_u + m_v + m_w$$

# MIP-Map Level Interpolation

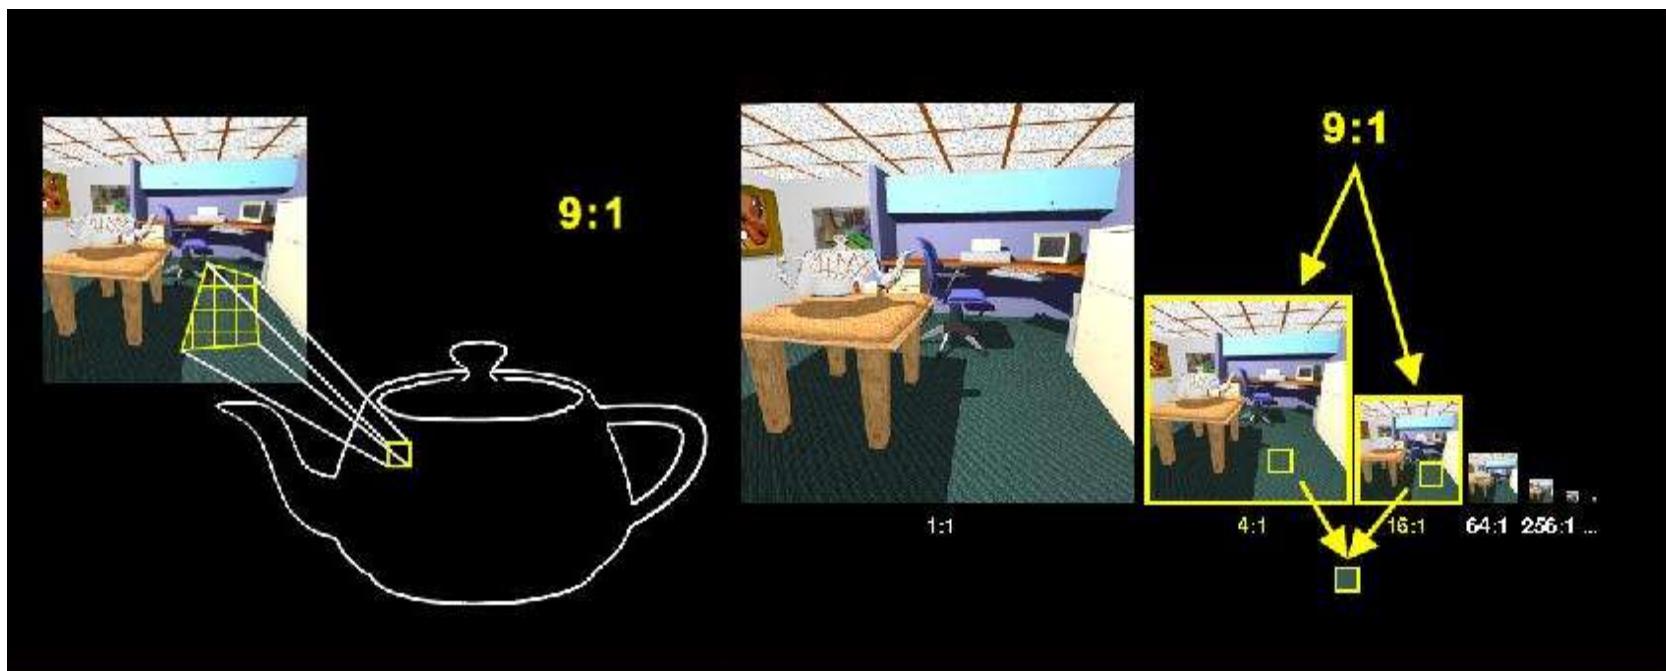


- Level of detail value is **fractional!**
- Use fractional part to blend (lin.) between two adjacent mipmap levels

# Texture Anti-Aliasing: MIP Mapping

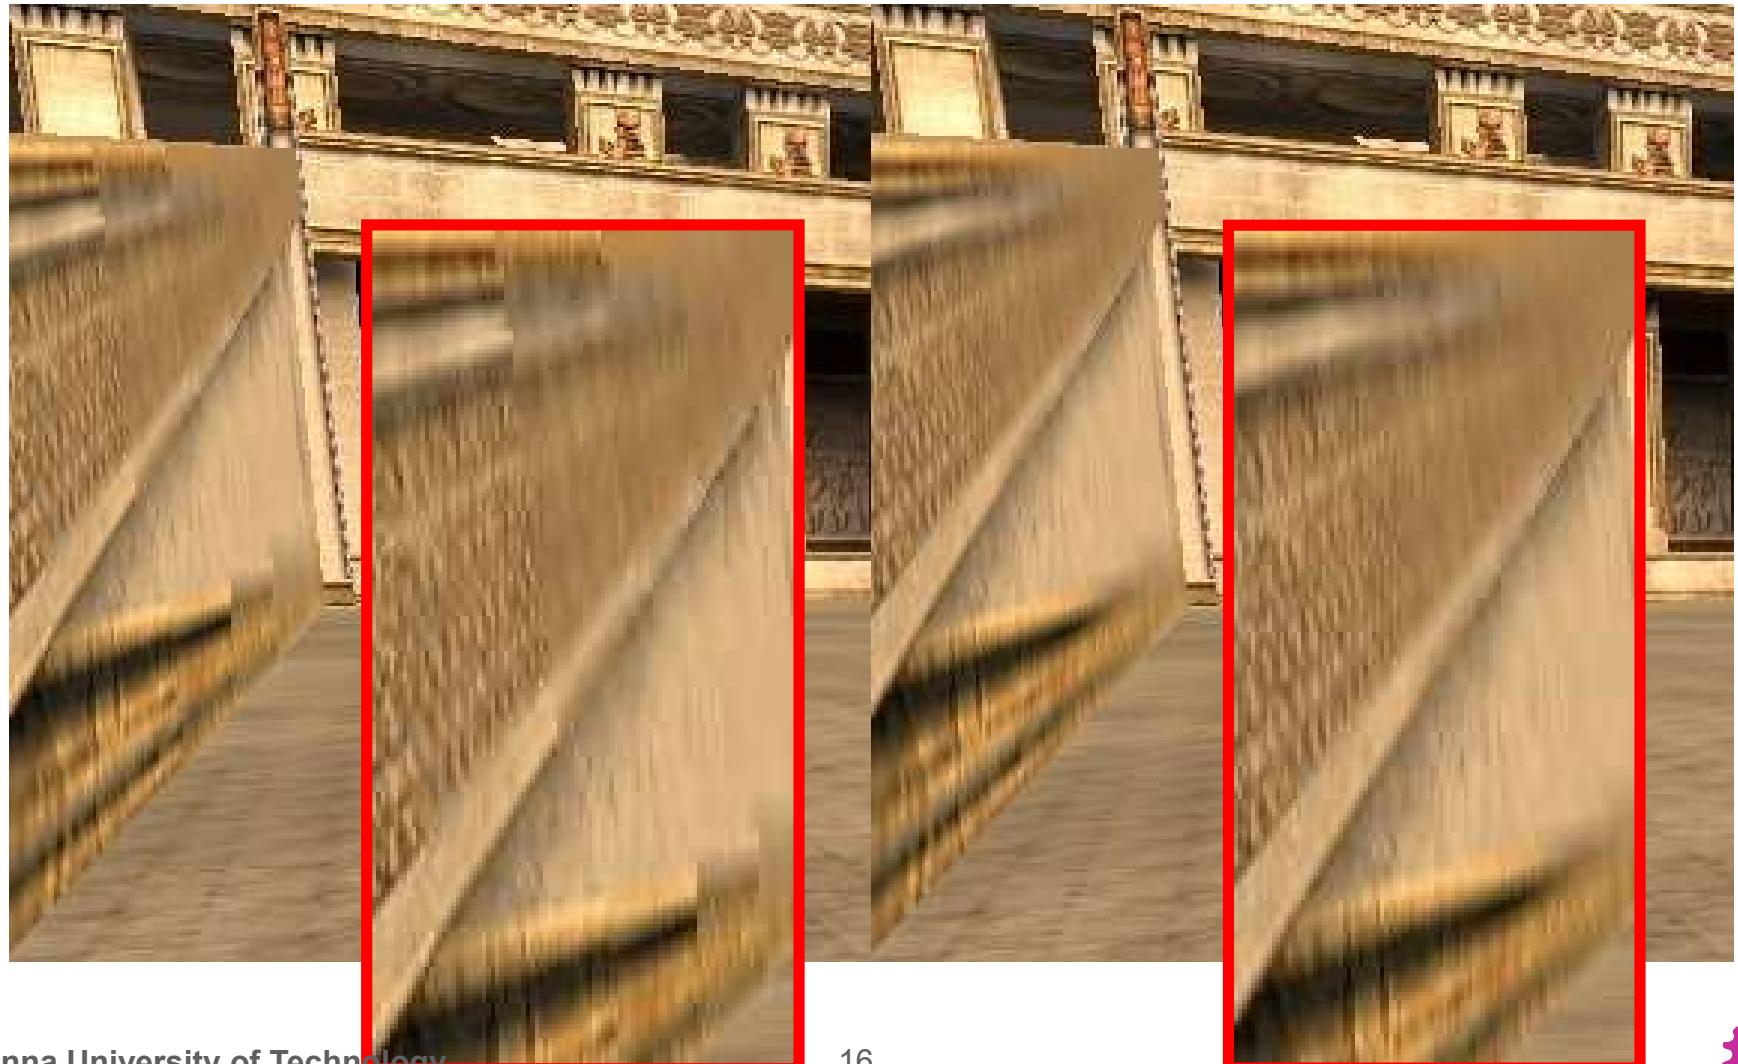
## ■ Trilinear interpolation:

- $T_1 :=$  value from texture  $D_1 = D_0 + 1$  (bilin.interpolation)
- Pixel value :=  $(D_1 - D) \cdot T_0 + (D - D_0) \cdot T_1$ 
  - Linear interpolation between successive MIP Maps
- Avoids "Mip banding" (but doubles texture lookups)



# Texture Anti-Aliasing: MIP Mapping

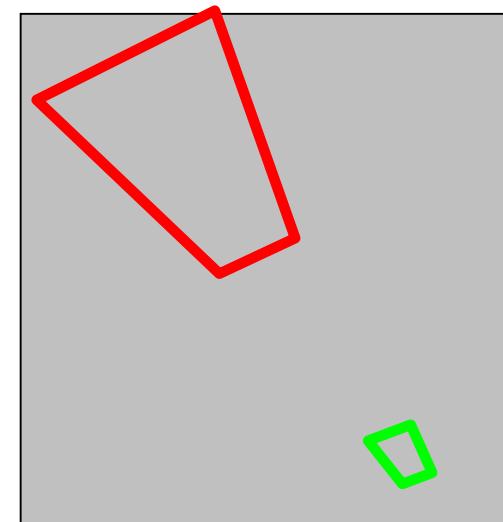
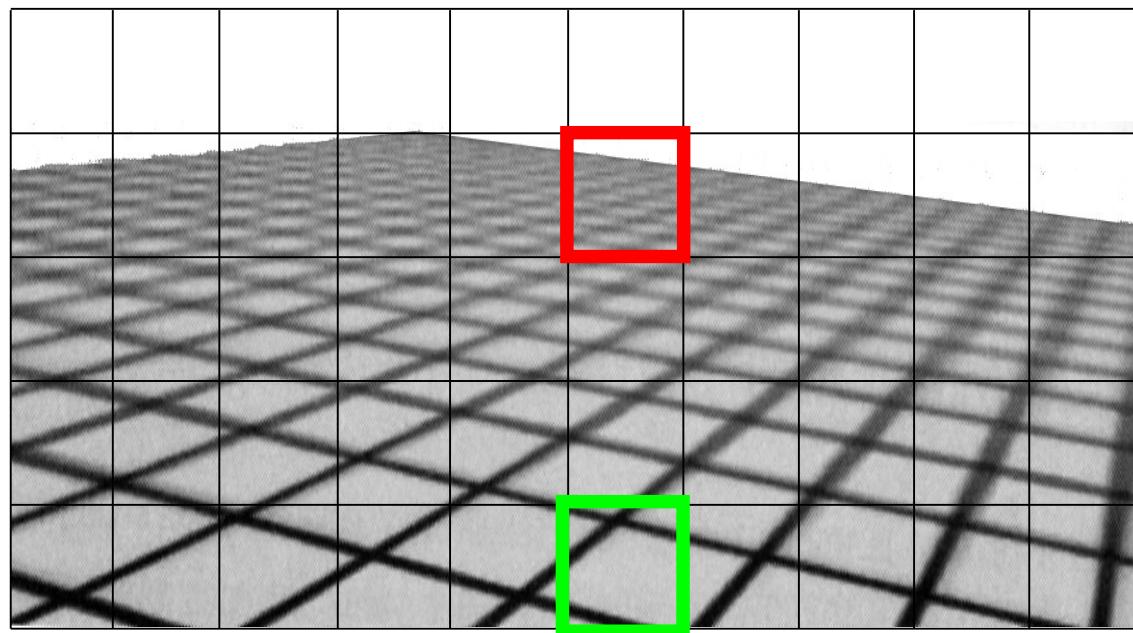
- Other example for bilinear vs. trilinear filtering



# Anti-Aliasing: Anisotropic Filtering

## ■ Anisotropic filtering

- View-dependent filter kernel
- Implementation: *summed area table*, "*RIP Mapping*", *footprint assembly*, *elliptical weighted average* (EWA)

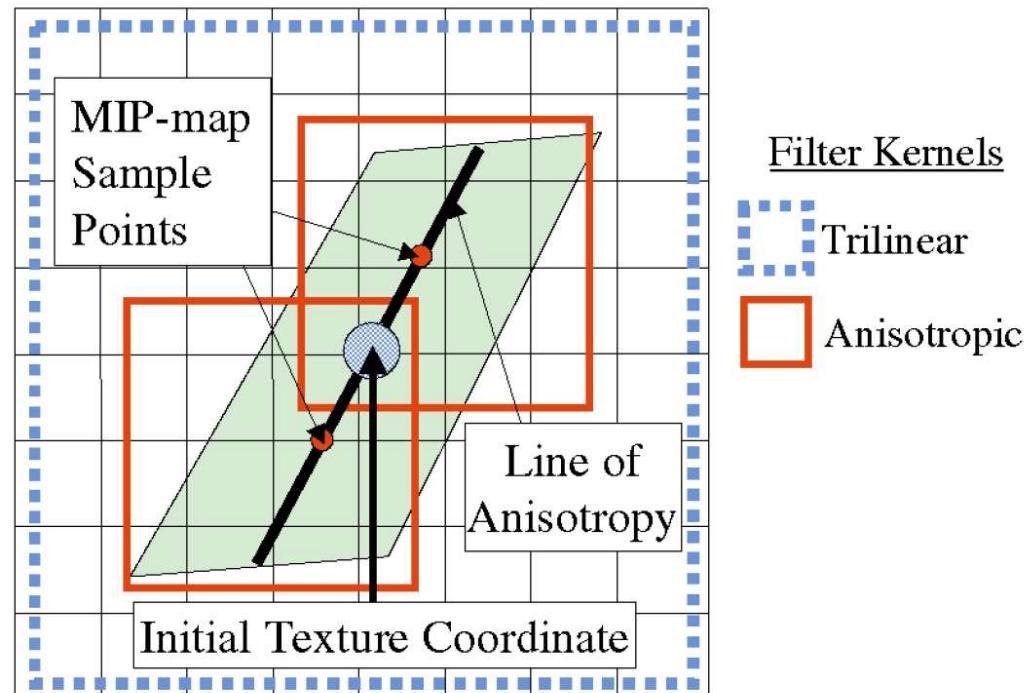


**Texture space**



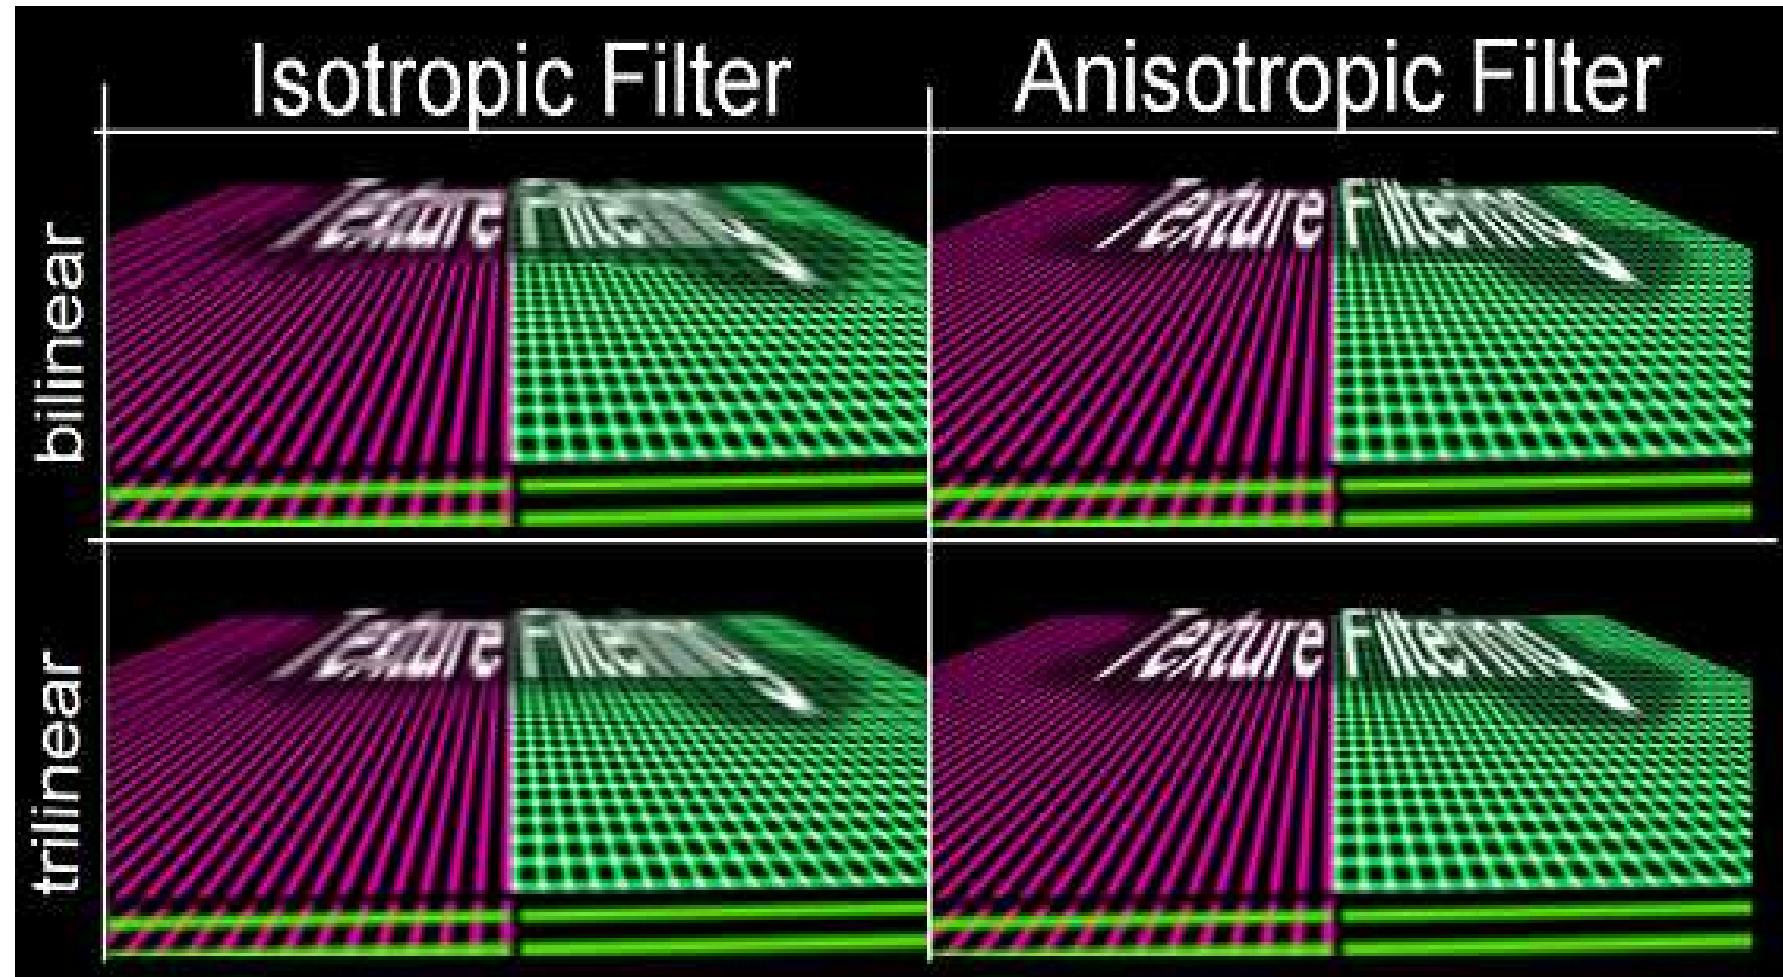


# Anisotropic Filtering: Footprint Assembly



# Anti-Aliasing: Anisotropic Filtering

## ■ Example



- Basically, everything done in hardware
- gluBuild2DMipmaps () generates MIPmaps
- Set parameters in glTexParameter()
  - GL\_TEXTURE\_MAG\_FILTER: GL\_NEAREST, GL\_LINEAR, ...
  - GL\_TEXTURE\_MIN\_FILTER: GL\_LINEAR\_MIPMAP\_NEAREST
- Anisotropic filtering is an extension:
  - GL\_EXT\_texture\_filter\_anisotropic
  - Number of samples can be varied (4x,8x,16x)
    - Vendor specific support and extensions





# Virtual Textures

# Virtual Texturing



## Example #1:

### **ARB Sparse Textures (originally: AMD Partially Resident Textures)**

ARB\_sparse\_texture / ARB\_sparse\_texture2

[https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB\\_sparse\\_texture.txt](https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB_sparse_texture.txt)

[https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB\\_sparse\\_texture2.txt](https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB_sparse_texture2.txt)

Hardware Virtual Texturing, Graham Sellers,  
from SIGGRAPH 2013 course “Rendering Massive Virtual Worlds”

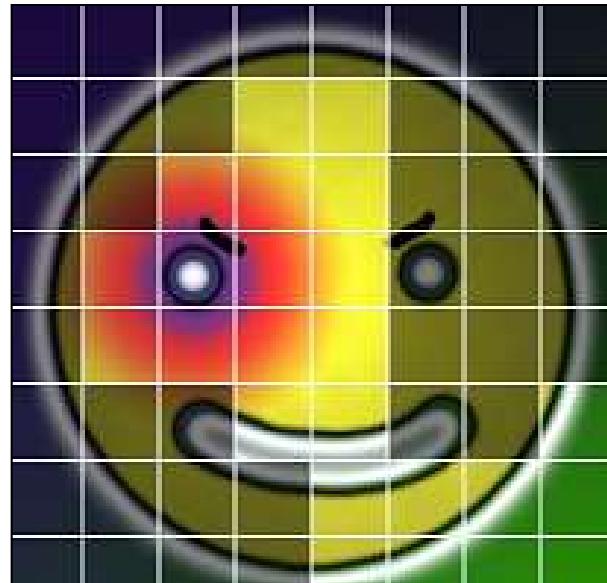
[https://cesiumjs.org/hosted-apps/massiveworlds/downloads/Graham/Hardware\\_Virtual\\_Textures.pptx](https://cesiumjs.org/hosted-apps/massiveworlds/downloads/Graham/Hardware_Virtual_Textures.pptx)

# Virtual Texturing

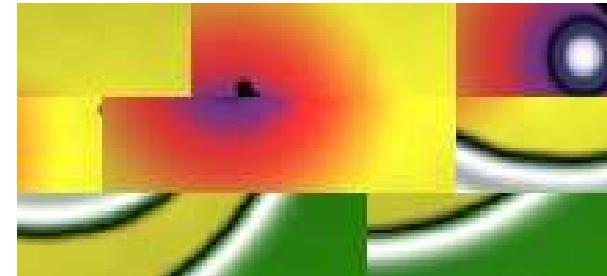


Divide texture up into tiles

- Commit only *used* tiles to memory
- Store data in separate physical texture



Virtual Texture

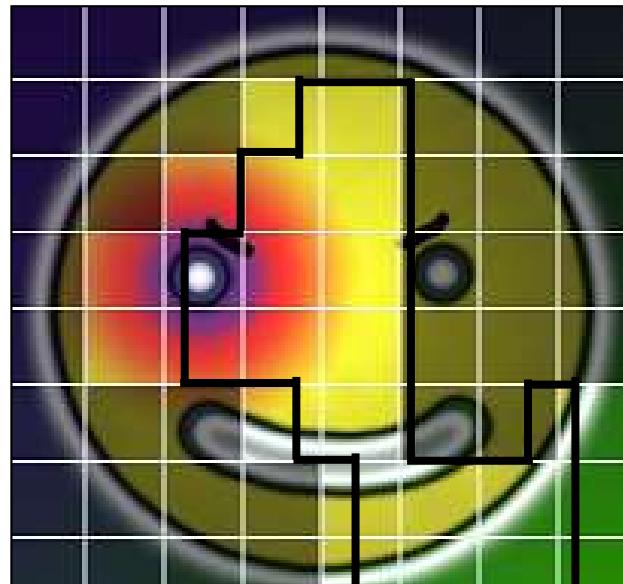


Physical Texture

# Virtual Texturing



Memory requirements set by number of resident tiles, not texture dimensions



RGBA8, 1024x1024, 64 tiles

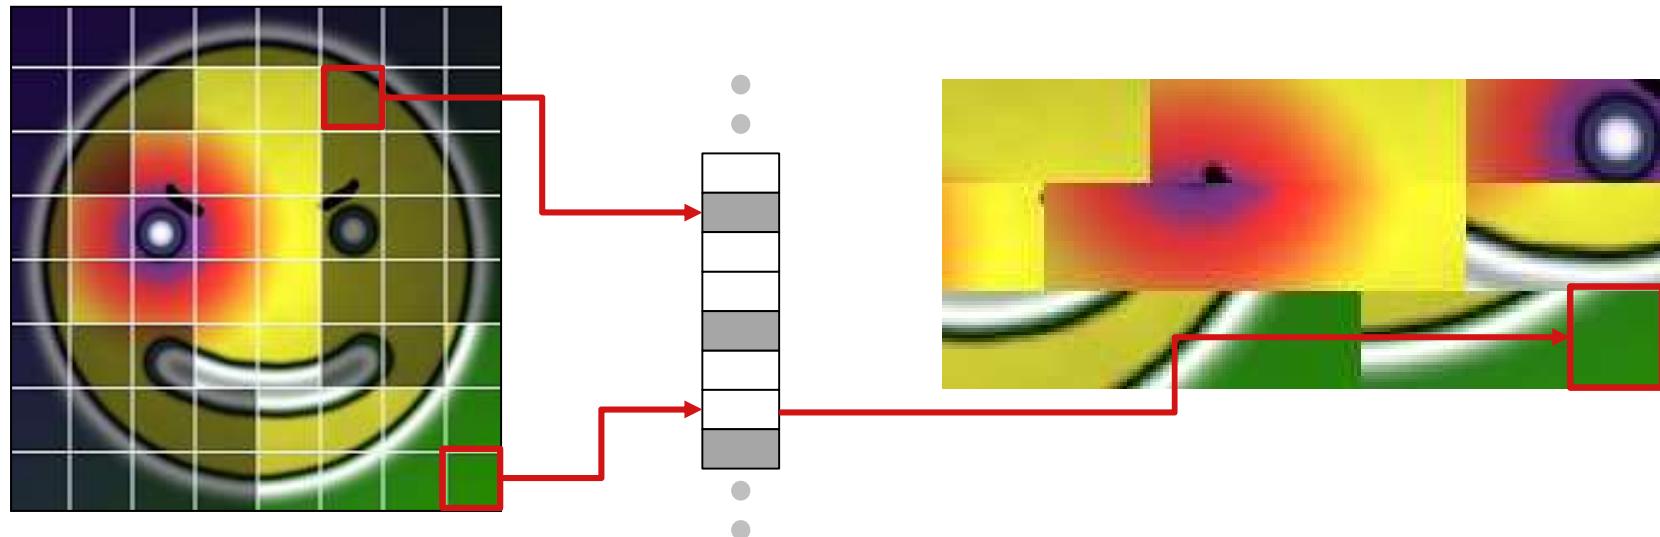
	Virtual	Physical
Memory	4096 kB	1536 kB

# Virtual Texturing



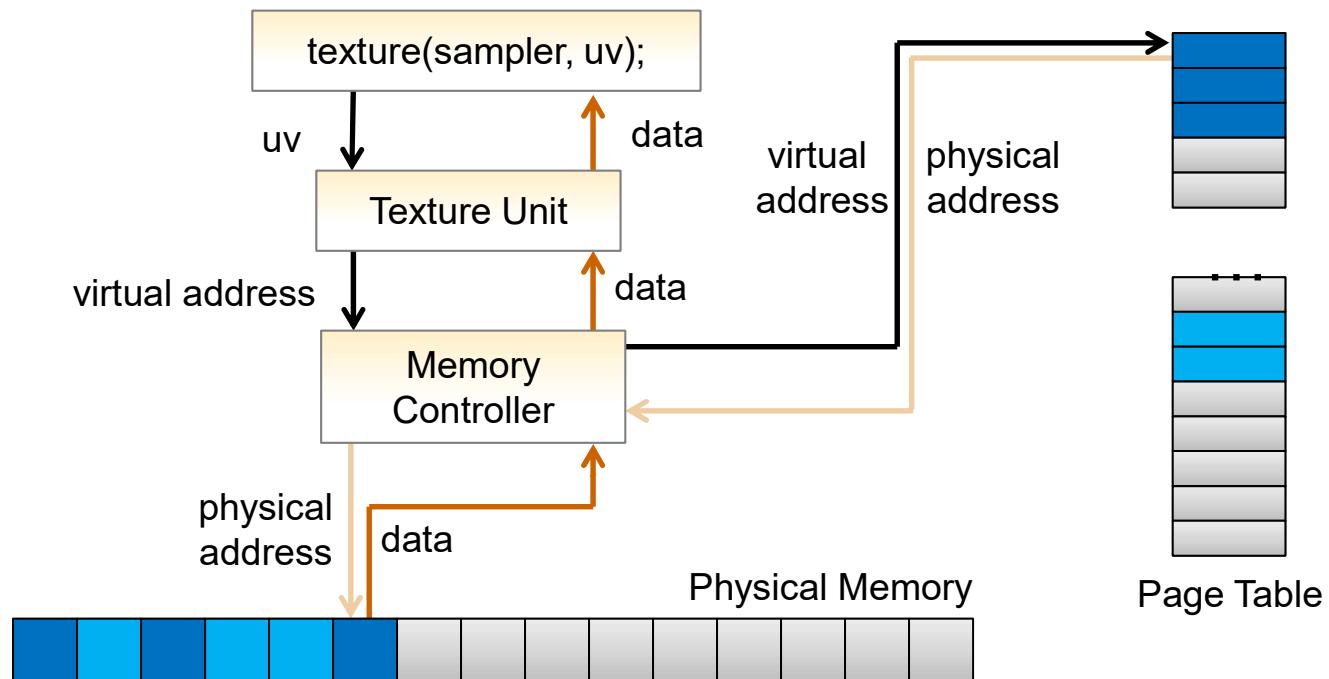
Use indirection table to map virtual to physical

- This is also known as a *page table*





# GPU Virtual Memory



# Summary (Shader vs. Full Hardware Support)



	SVTs	HVTs
Address translation	Shader code	HW page table
Filtering	HW + shader code	HW only
# of texture fetches	2, dependent	1
Supported formats	The ones implemented	All supported by HW
Supported texture types	The ones implemented	All supported by HW

# Virtual Texturing



## Example #2:

### **Adaptive Shadow Maps (ASM)**

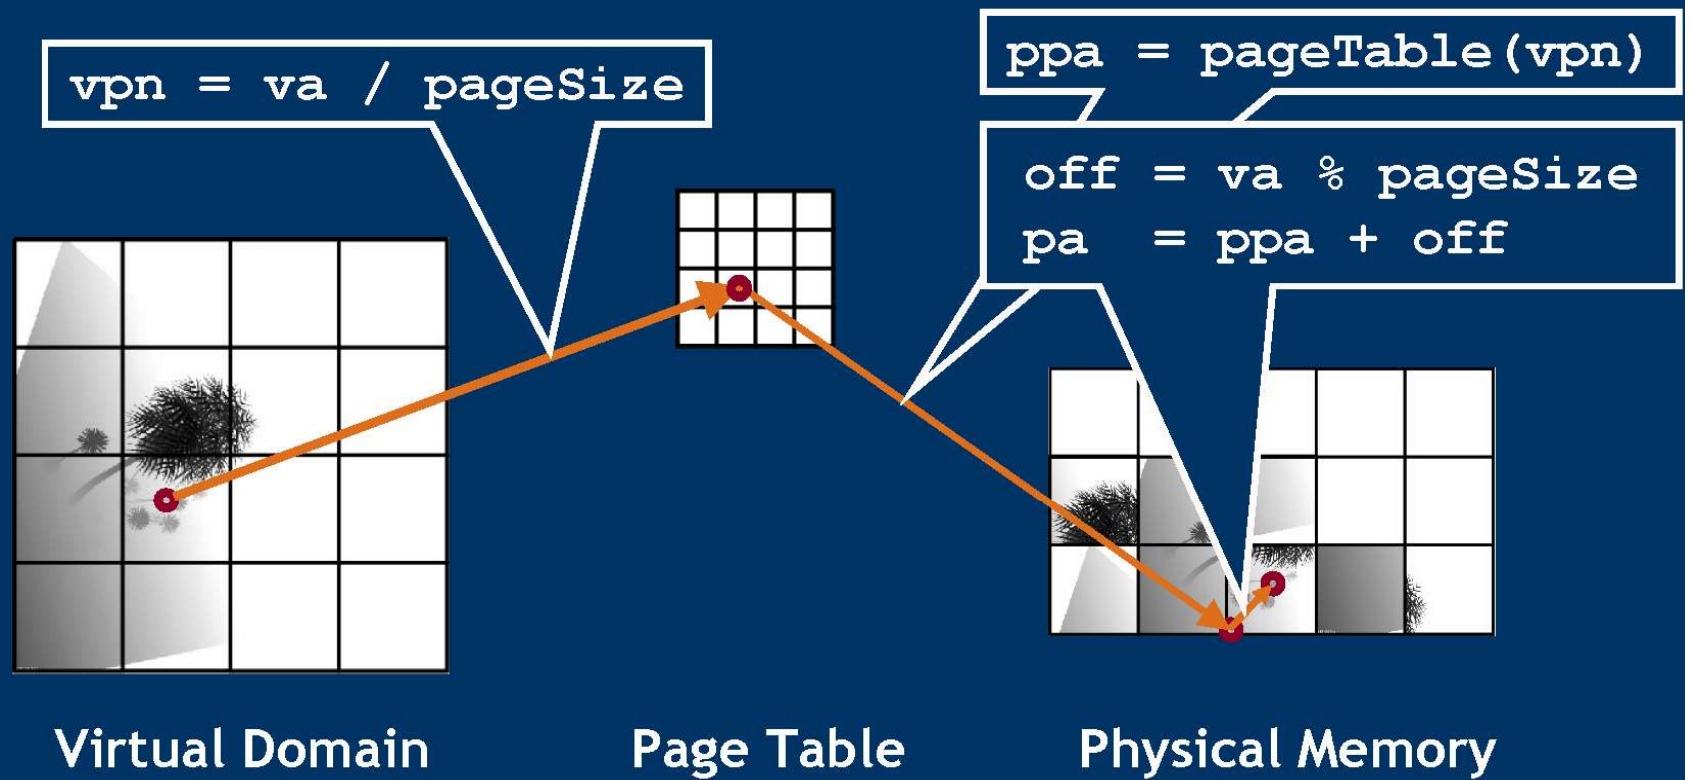
- On CPUs: Fernando et al., ACM SIGGRAPH 2001

### **Resolution-Matched Shadow Maps**

- On GPUs: Aaron Lefohn et al., ACM Transactions on Graphics 2007

# ASM Data Structure (Adaptive Shadow Maps)

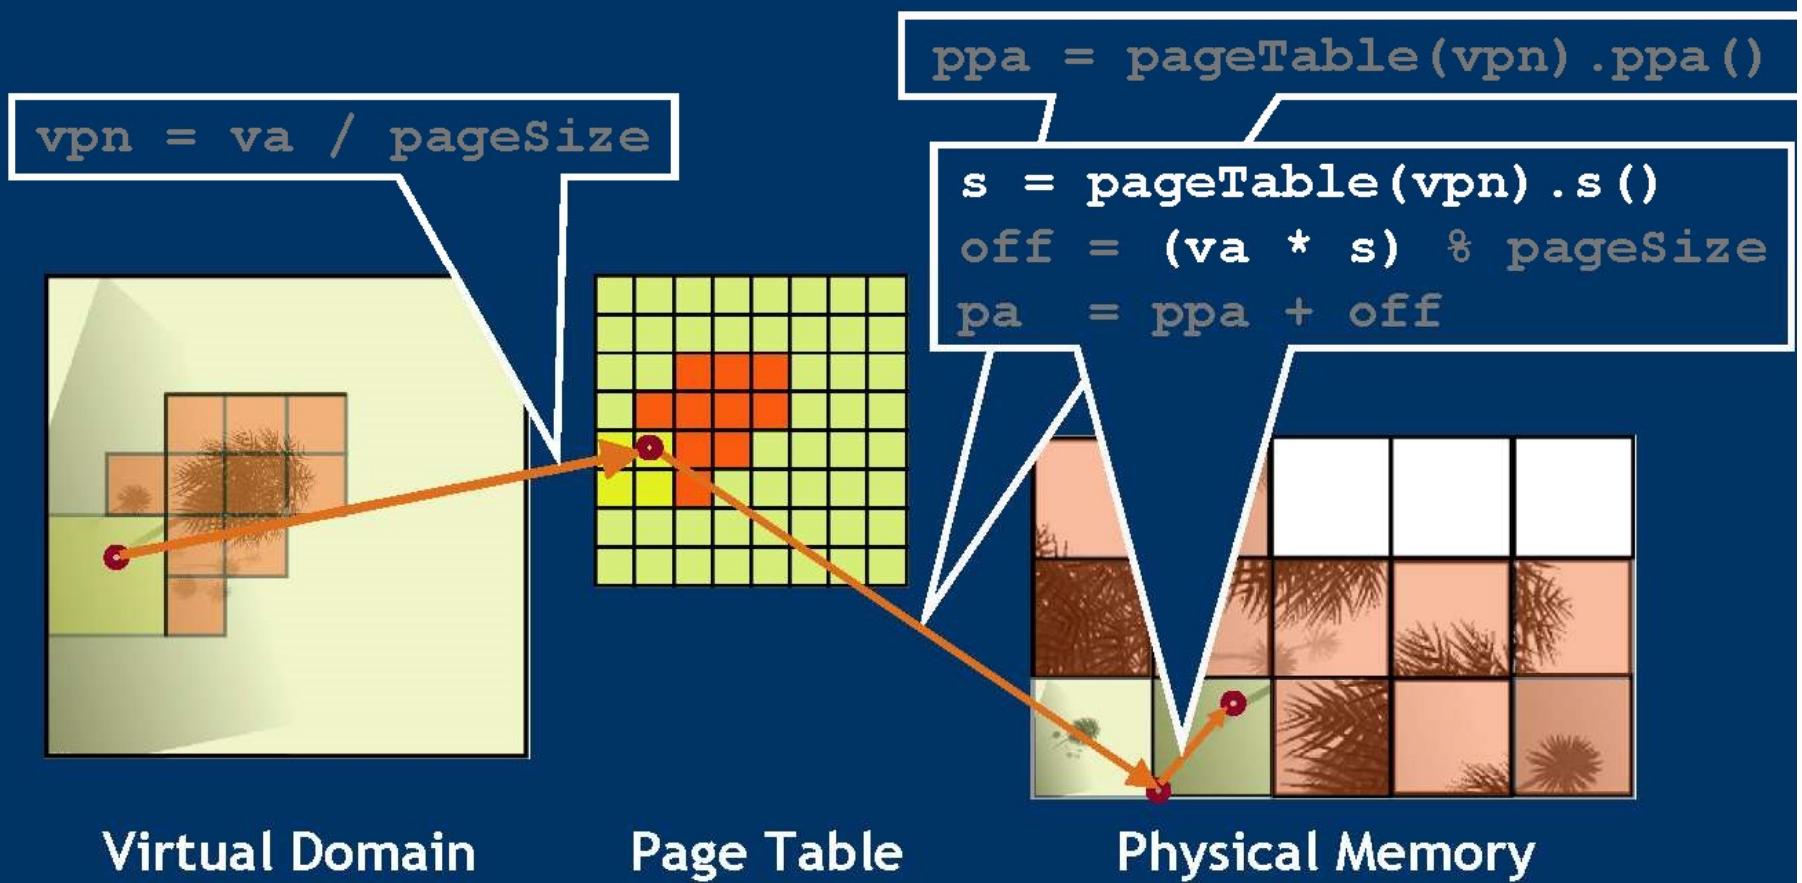
- Page table example



# ASM Data Structure (Adaptive Shadow Maps)

- Adaptive Page Table

- Map multiple virtual pages to single physical page



# Virtual Texturing



## Example #3:

**id Tech 5 Megatextures, id Software**

**Rage**

- Virtual Texturing in Software and Hardware, van Waveren et al., SIGGRAPH 2012 course notes + slides

[http://www.jurajobert.com/data/Virtual\\_Texturing\\_in\\_Software\\_and\\_Hardware\\_course\\_notes.pdf](http://www.jurajobert.com/data/Virtual_Texturing_in_Software_and_Hardware_course_notes.pdf)

<http://www.mrelusive.com/publications/papers/Software-Virtual-Textures.pdf>

[http://www.mrelusive.com/publications/presentations/2013\\_siggraph/hq\\_sw\\_hw\\_vts\\_12.pdf](http://www.mrelusive.com/publications/presentations/2013_siggraph/hq_sw_hw_vts_12.pdf)

# Virtual Texturing



Rage / id Tech 5 (id Software)

# Virtual Texturing

- Unique, very large virtual textures key to id tech 5 rendering
- Full description beyond the scope of this talk



# Virtual Texturing

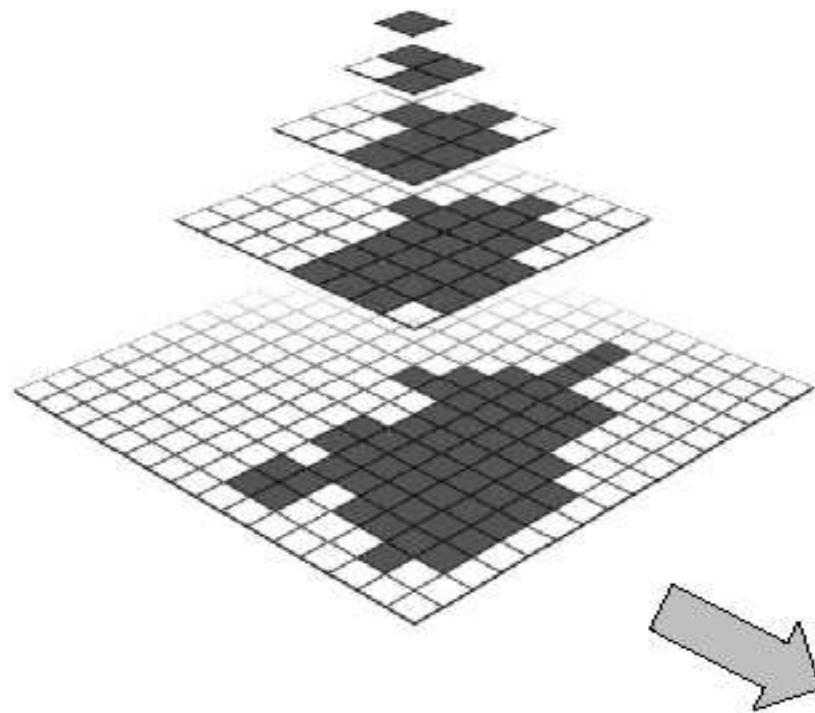


# Virtual Texturing

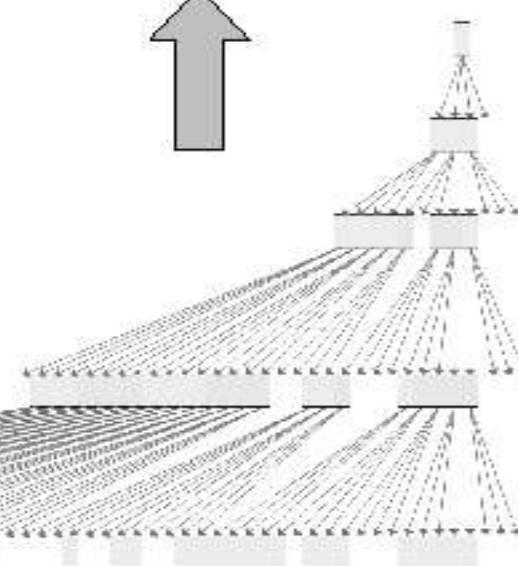
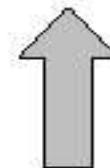
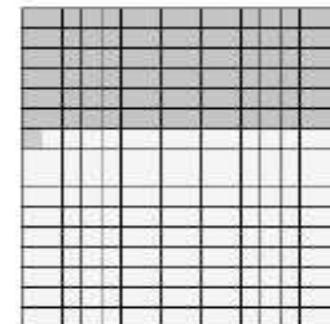


# Virtual Texturing

Texture Pyramid with Sparse Page Residency



Physical Page Texture

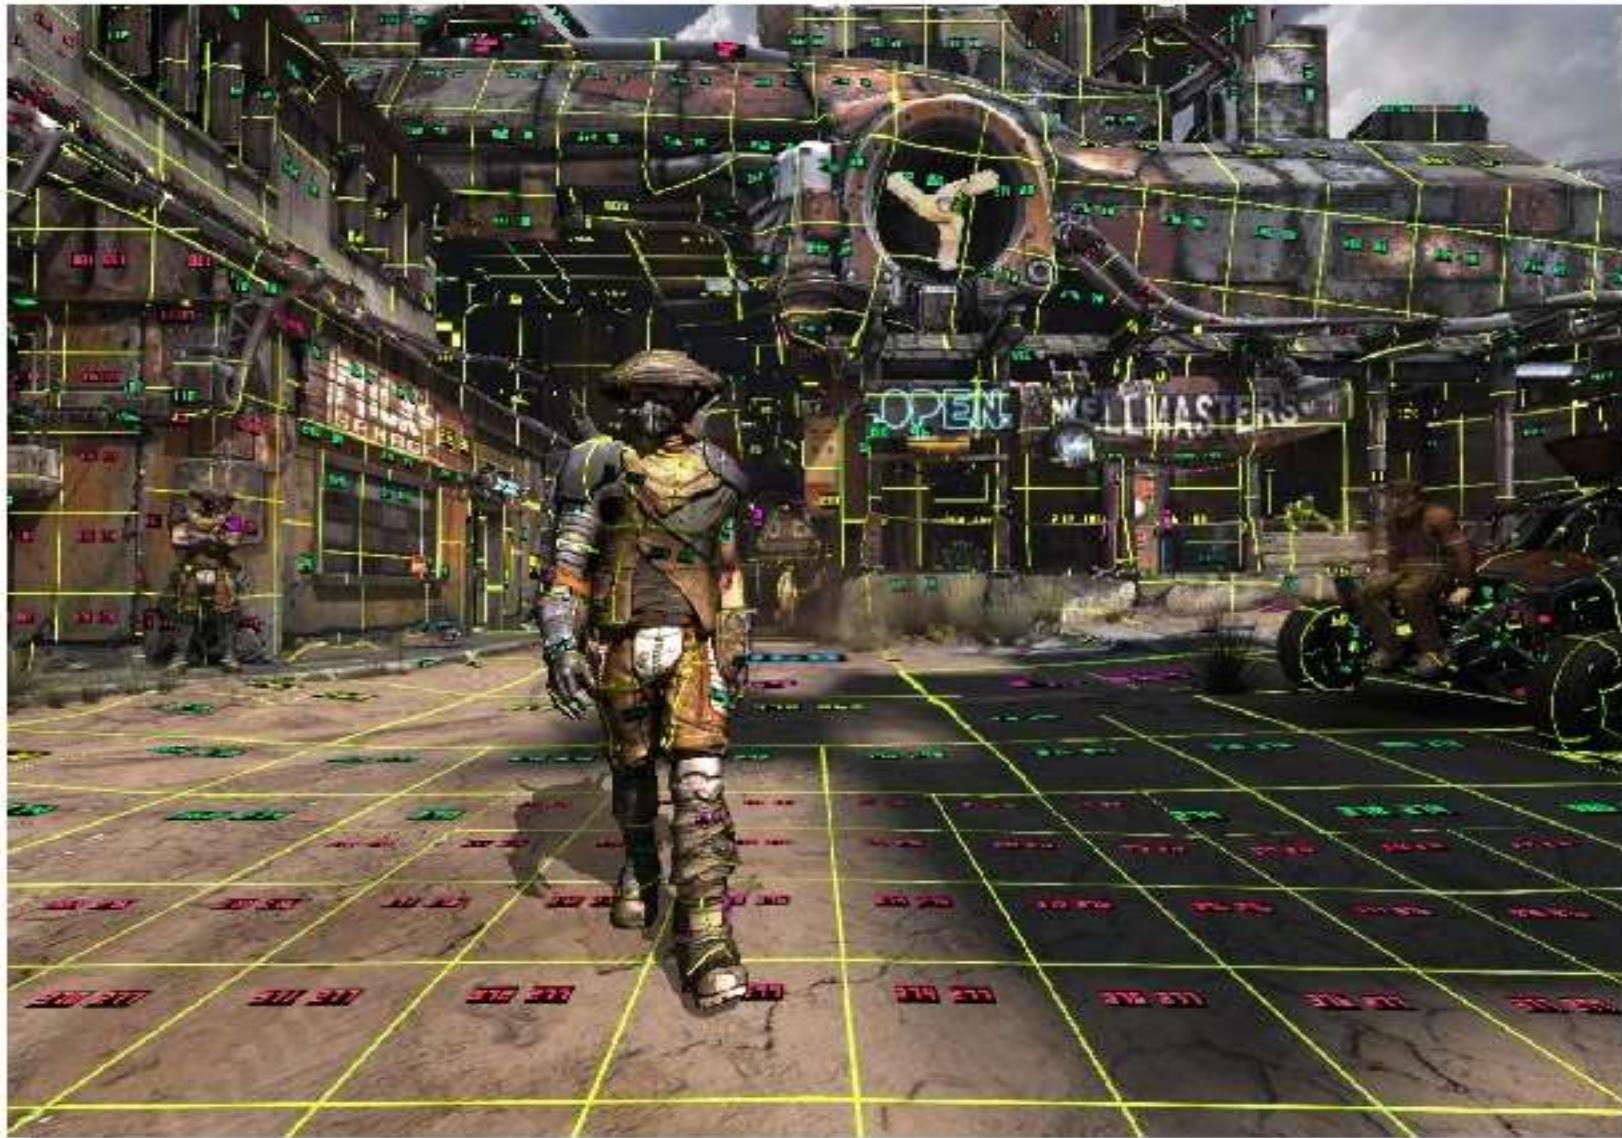


Quad-tree of Sparse Texture Pyramid

# Virtual Texturing



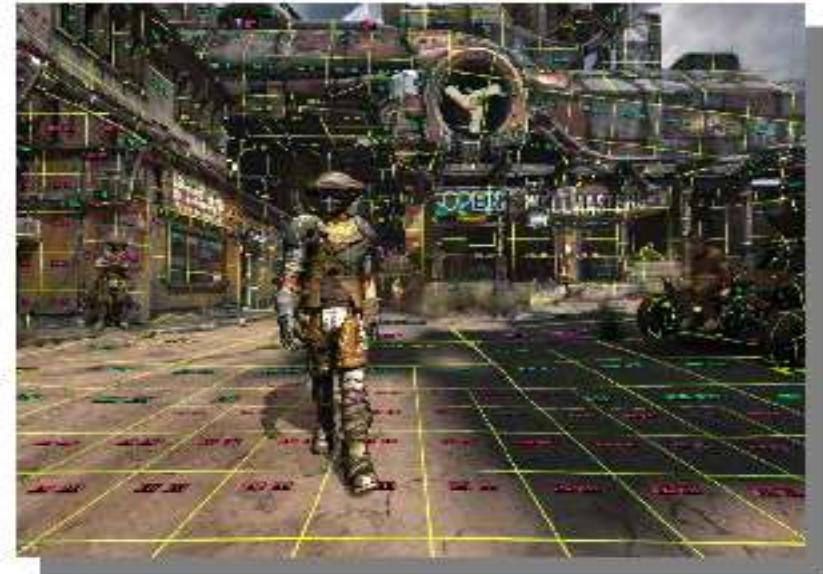
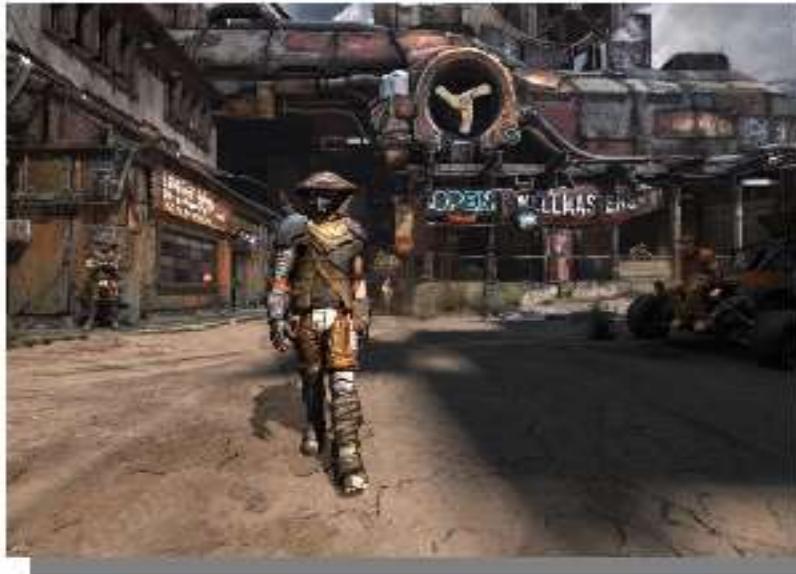
# Virtual Texturing

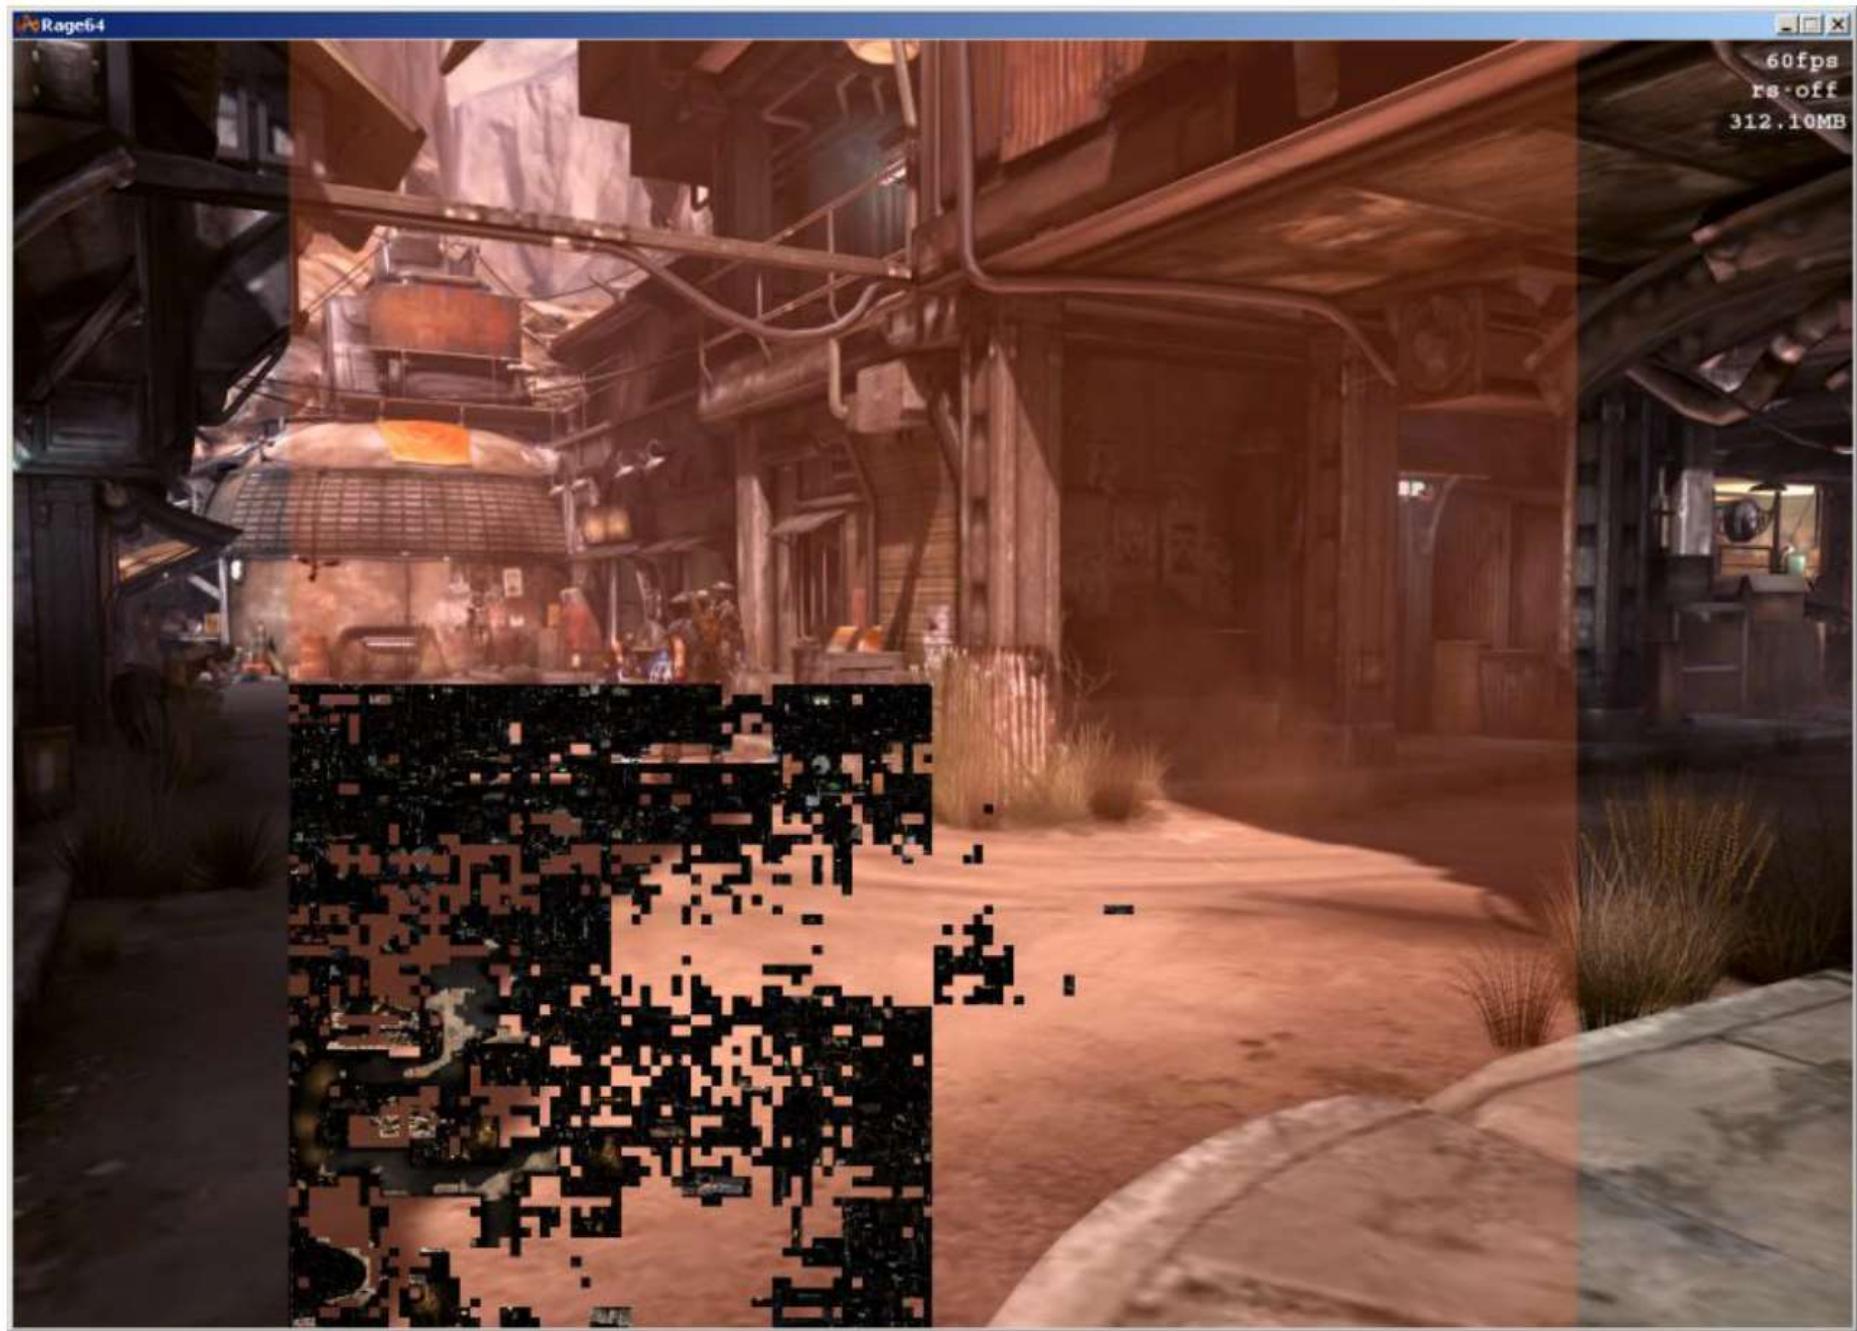


# Virtual Texturing

A few interesting issues...

- Texture filtering
- Thrashing due to physical memory oversubscription
- LOD transitions under high latency

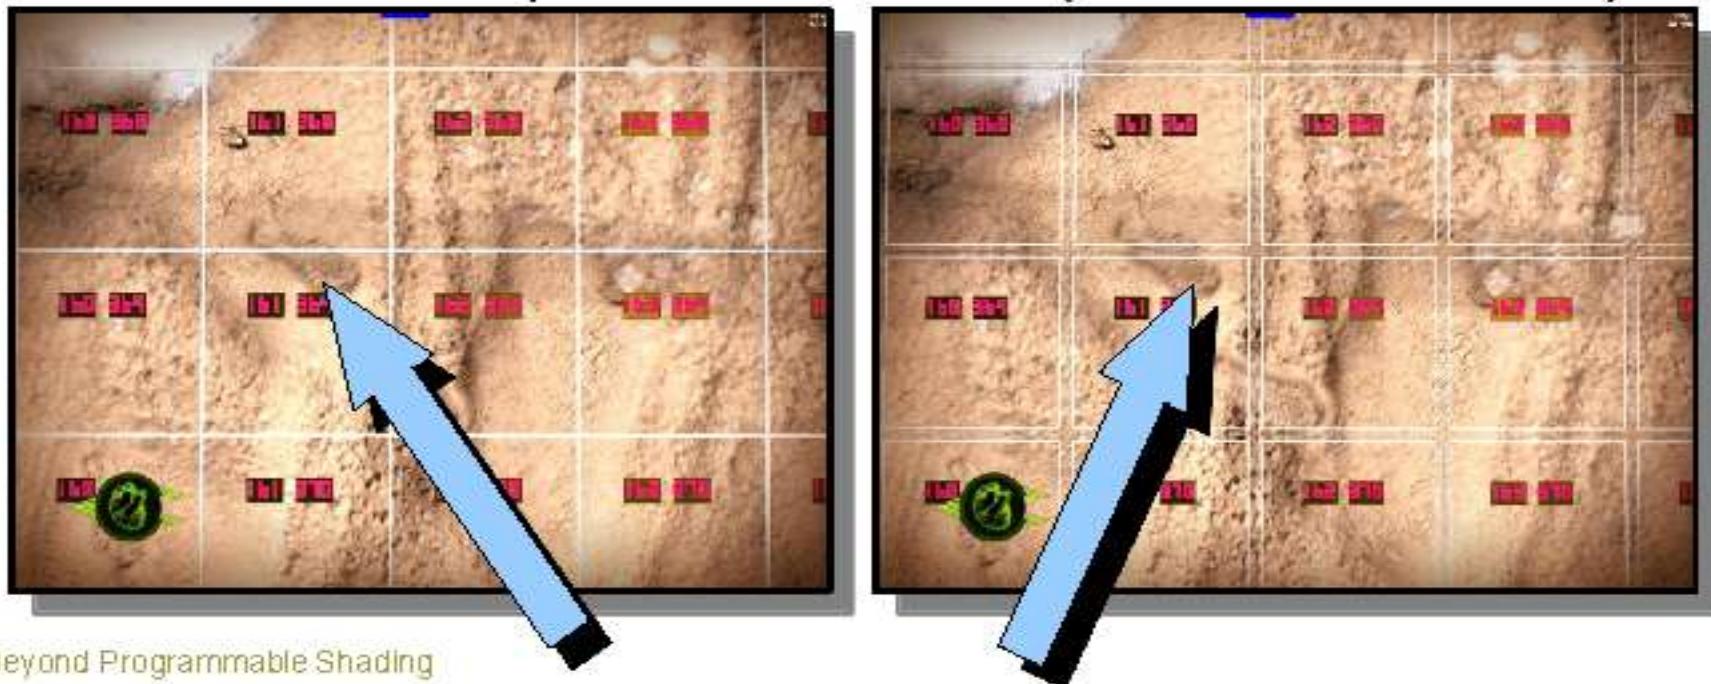




RAGE with PRTs (Image courtesy of id Software)

# Virtual Texturing - Filtering

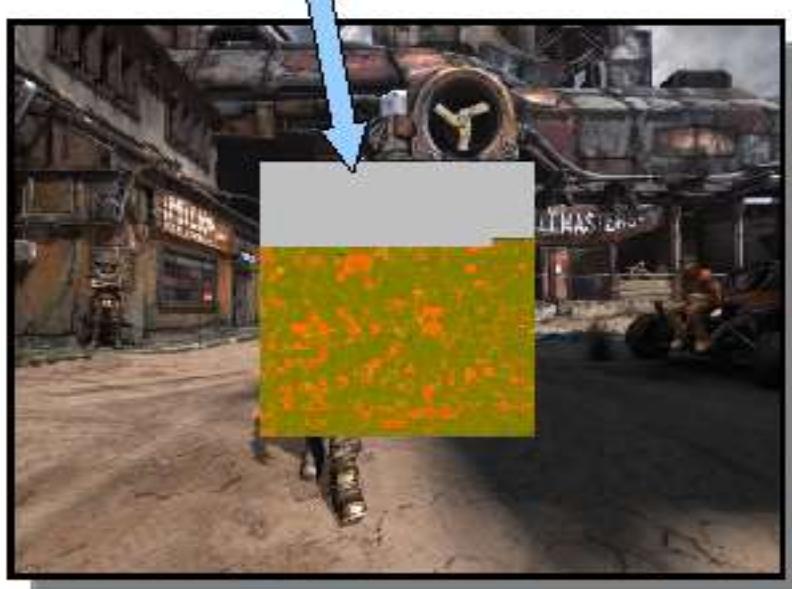
- We tried no filtering at all
- We tried bilinear filtering without borders
- Bilinear filtering with border works well
- Trilinear filtering reasonably but still expensive
- Anisotropic filtering possible via TXD (texgrad)
  - 4-texel border necessary (max aniso = 4)
  - TEX with implicit derivs ok too (on some hardware)



# Virtual Texturing - Thrashing

- Sometimes you need more physical pages than you have
- With conventional virtual memory, you must thrash
- With virtual texturing, you can globally adjust feedback LOD bias until working set fits

32 x 32 pages



1024 Physical Pages

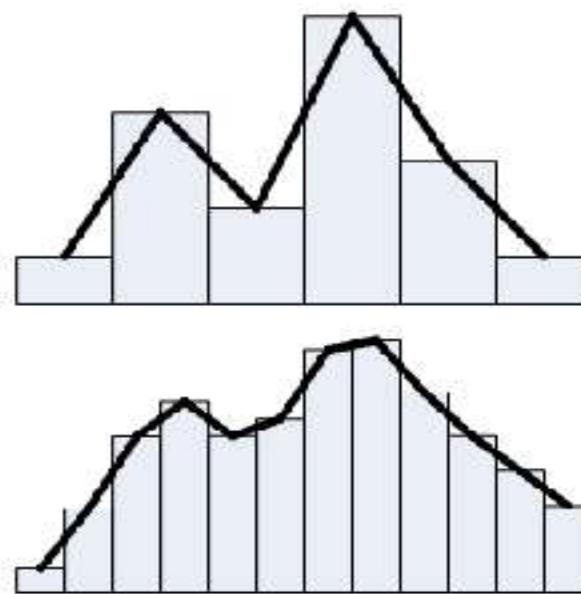
8x8 pages



64 Physical Pages

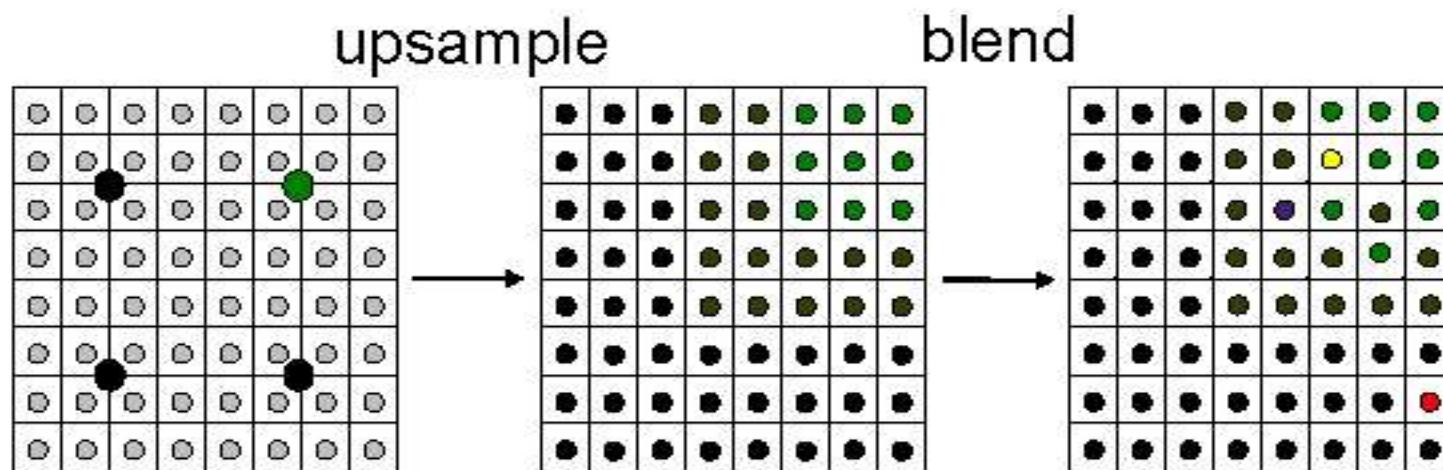
# Virtual Texturing – LOD Snap

- Latency between first need and availability can be high
  - Especially if optical disk read required ( $>100$  msec seek!)
- Visible snap happens when magnified texture changes LOD
- If we used trilinear filtering, blending in detail would be easy
- Instead continuously update physical pages with blended data



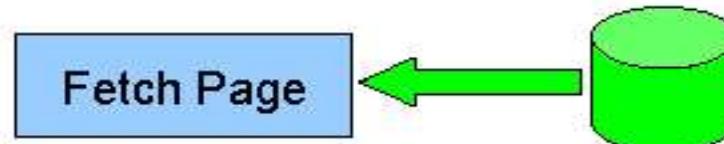
# Virtual Texturing - LOD Snap

- Upsample coarse page immediately
- Then blend in finer data when available



# Virtual Texturing - Management

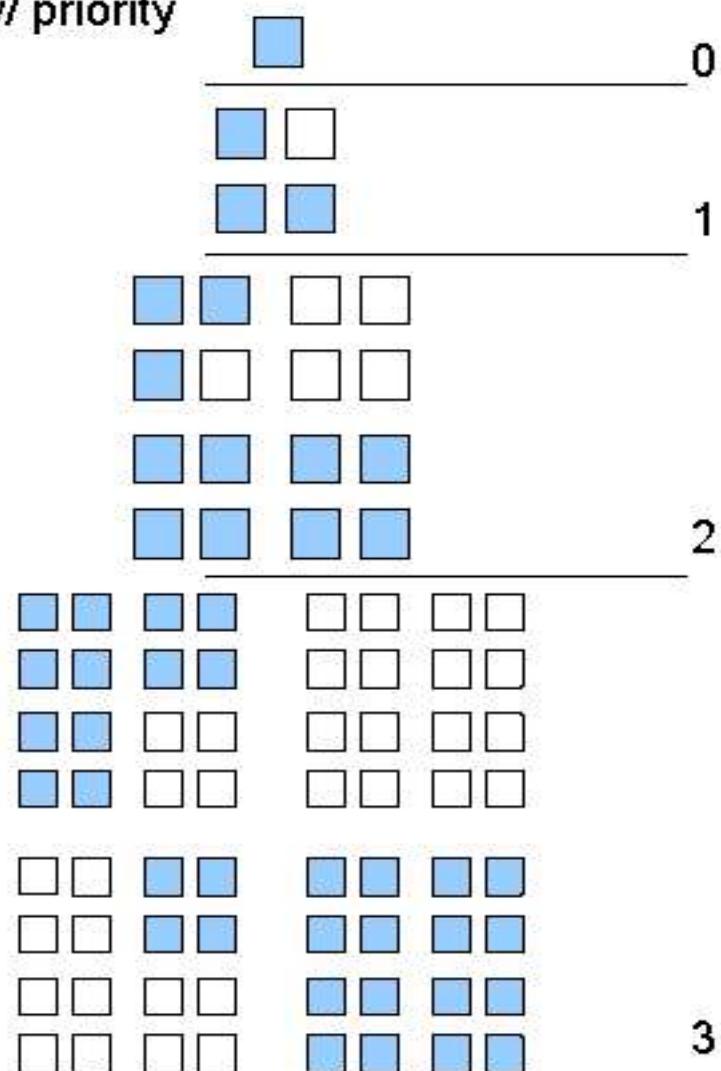
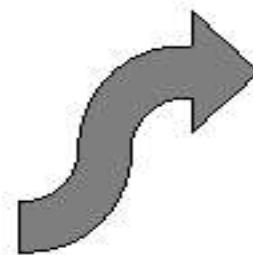
- Analysis tells us what pages we need
- We fetch what we can



- But this is a real-time app... so no blocking allowed
- Cache handles hits, schedules misses to load in background
- Resident pages managed independent of disk cache
- Physical pages organized as quad-tree per virtual texture
- Linked lists for free, LRU, and locked pages

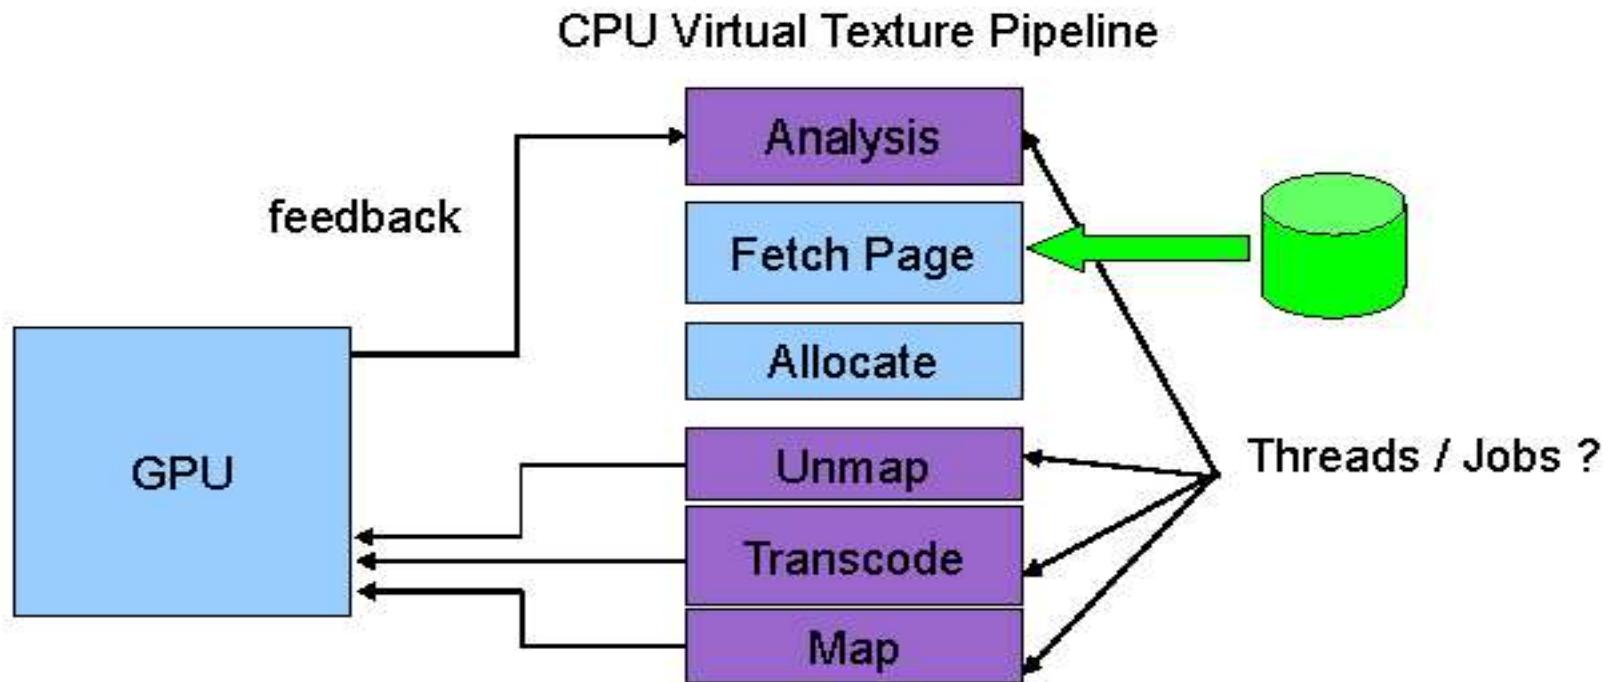
# Virtual Texturing - Feedback

- Feedback Analysis
  - Gen ~breadth-first quad-tree order w/ priority



# Virtual Texturing - Pipeline

- Compute intensive complex system with dependencies that we want to run in parallel on all the different platforms



# Virtual Texturing



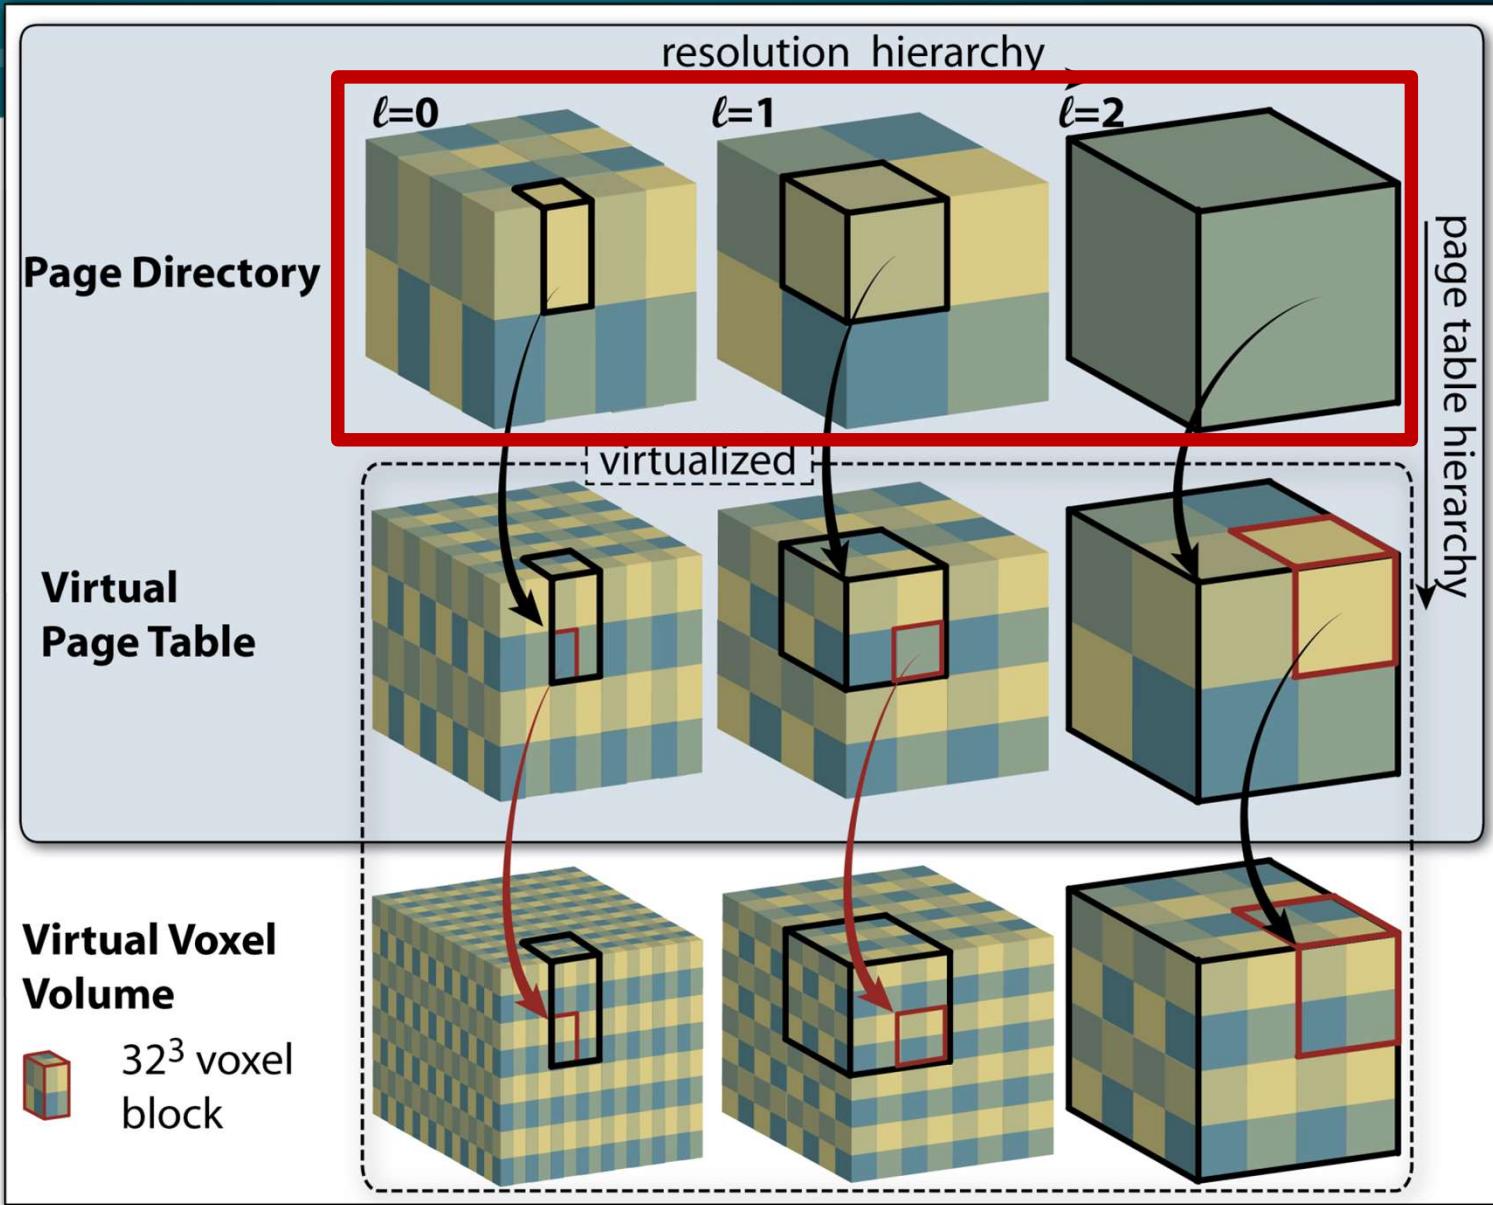
## Example #4:

### Petascale Volume Rendering

- Interactive Volume Exploration of Petascale Microscopy Data Streams Using a Visualization-Driven Virtual Memory Approach,  
Hadwiger et al., IEEE SciVis 2012

<http://dx.doi.org/10.1109/TVCG.2012.240>

# Petascale Volume Rendering



multi-resolution  
page directory

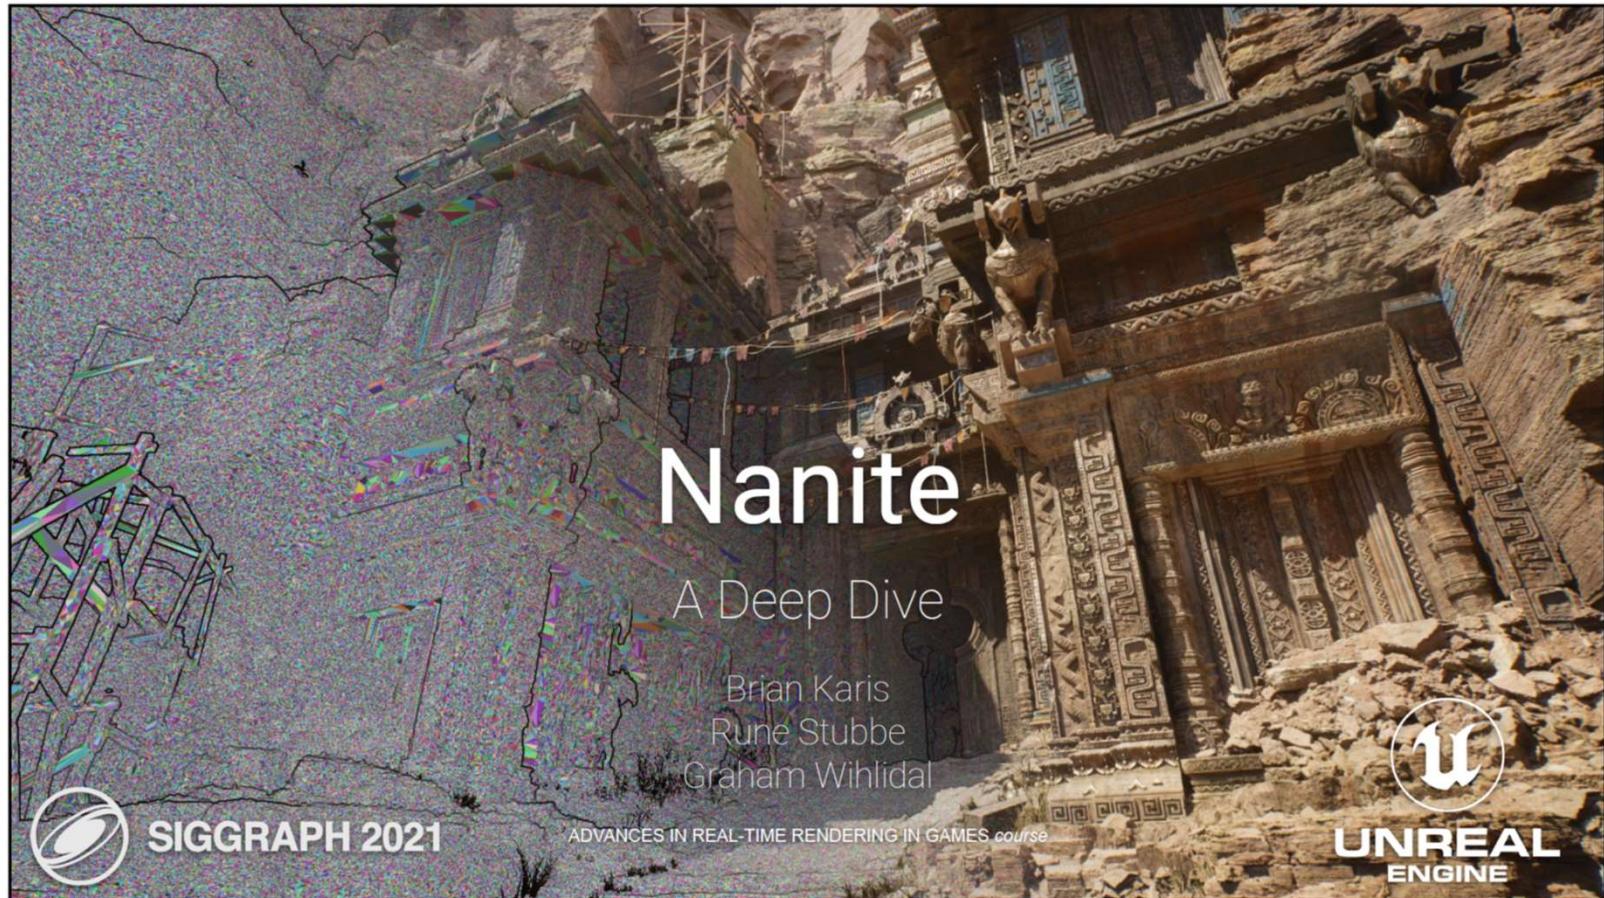


# **Virtual Geometry**

# Virtual Geometry



Unreal Engine 5 virtual geometry pipeline: Nanite  
Tech talk by Brian Karis, Epic Games





# Virtual Geometry

## The Dream

- Virtualize geometry like we did textures
- No more budgets
  - Polycount
  - Draw calls
  - Memory
- Directly use film quality source art
  - No manual optimization required
- No loss in quality



## Triangle cluster culling

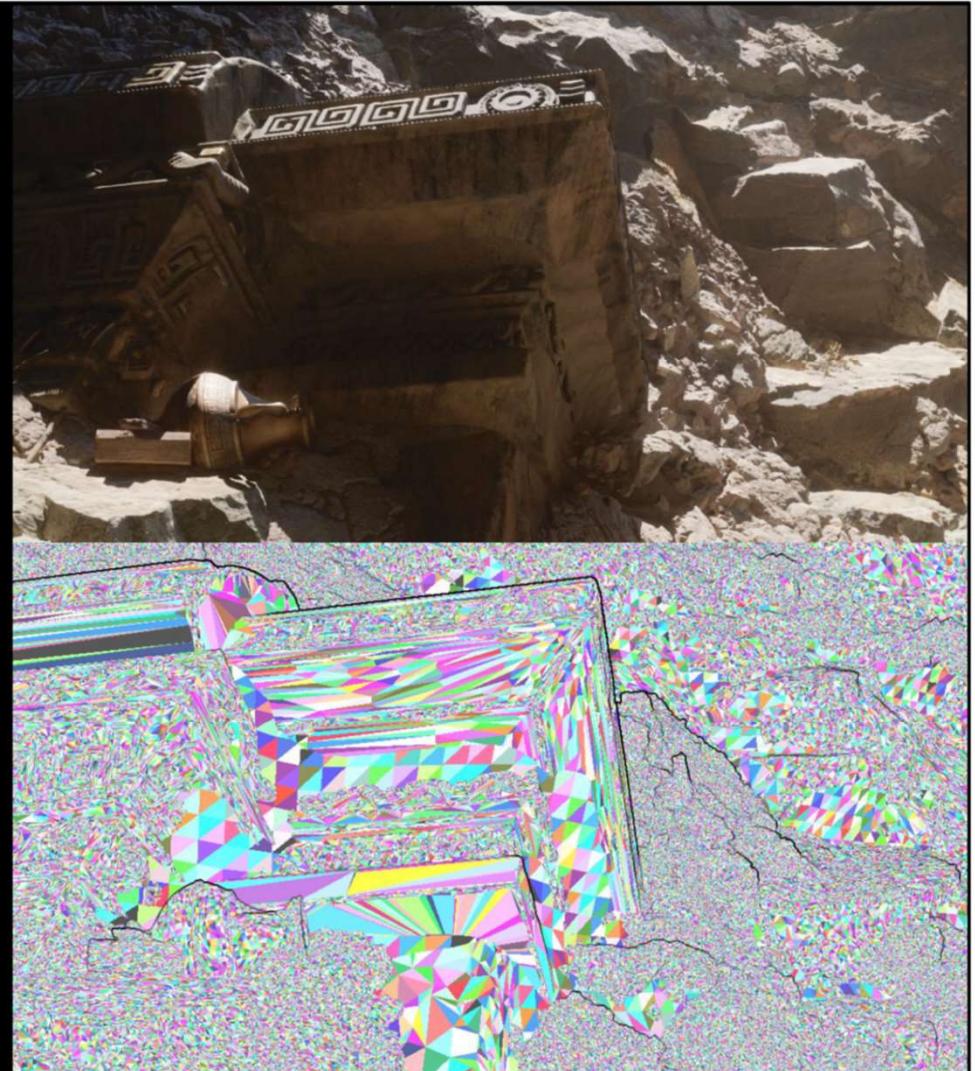
- Group triangles into clusters
  - Build bounding data for each cluster
- Cull clusters based on bounds
  - Frustum cull
  - Occlusion cull





## Pixel scale detail

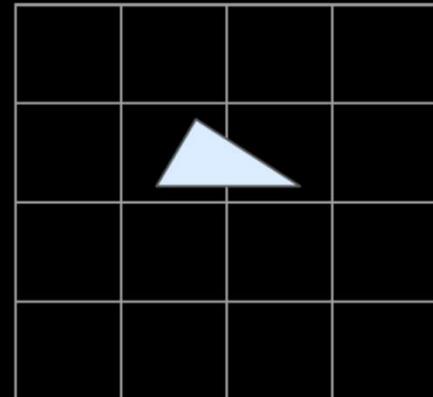
- Can we hit pixel scale detail with triangles > 1 pixel?
  - Depends how smooth
  - In general no
- We need to draw pixel sized triangles





## Tiny triangles

- Terrible for typical rasterizer
- Typical rasterizer:
  - Macro tile binning
  - Micro tile 4x4
  - Output 2x2 pixel quads
  - Highly parallel in pixels not triangles
- Modern GPUs setup 4 tris/clock max
  - Outputting SV\_PrimitiveID makes it even worse
- Can we beat the HW rasterizer in SW?





# Software Rasterization

3x faster!



# Micropoly software rasterizer

- 128 triangle clusters => threadgroup size 128
- 1 thread per vertex
  - Transform position
  - Store in groupshared
  - If more than 128 verts loop (max 2)
- 1 thread per triangle
  - Fetch indexes
  - Fetch transformed positions
  - Calculate edge equations and depth gradient
  - Calculate screen bounding rect
  - For all pixels in rect
    - If inside all edges then write pixel





# Hardware Rasterization

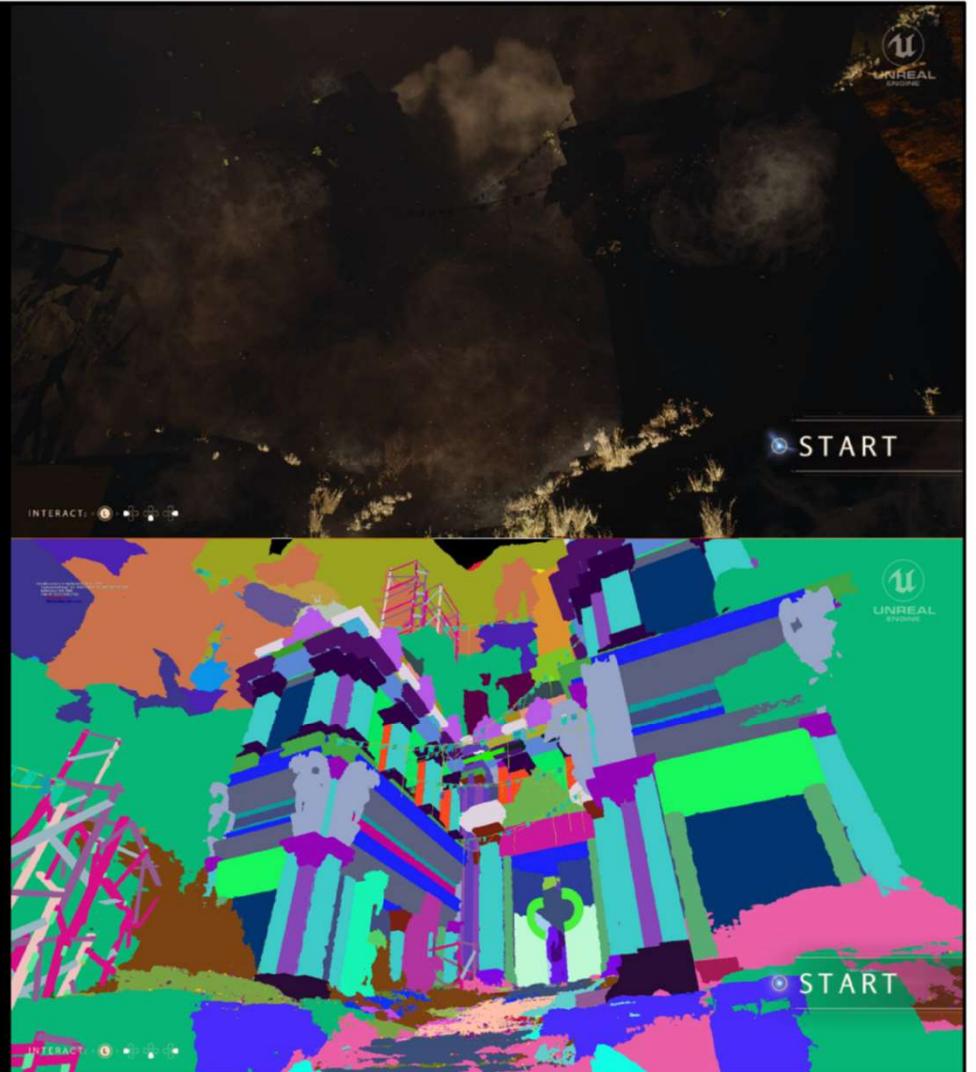
- What about big triangles?
  - Use HW rasterizer
- Choose SW or HW per cluster
- Also uses 64b atomic writes to UAV





# Material shading

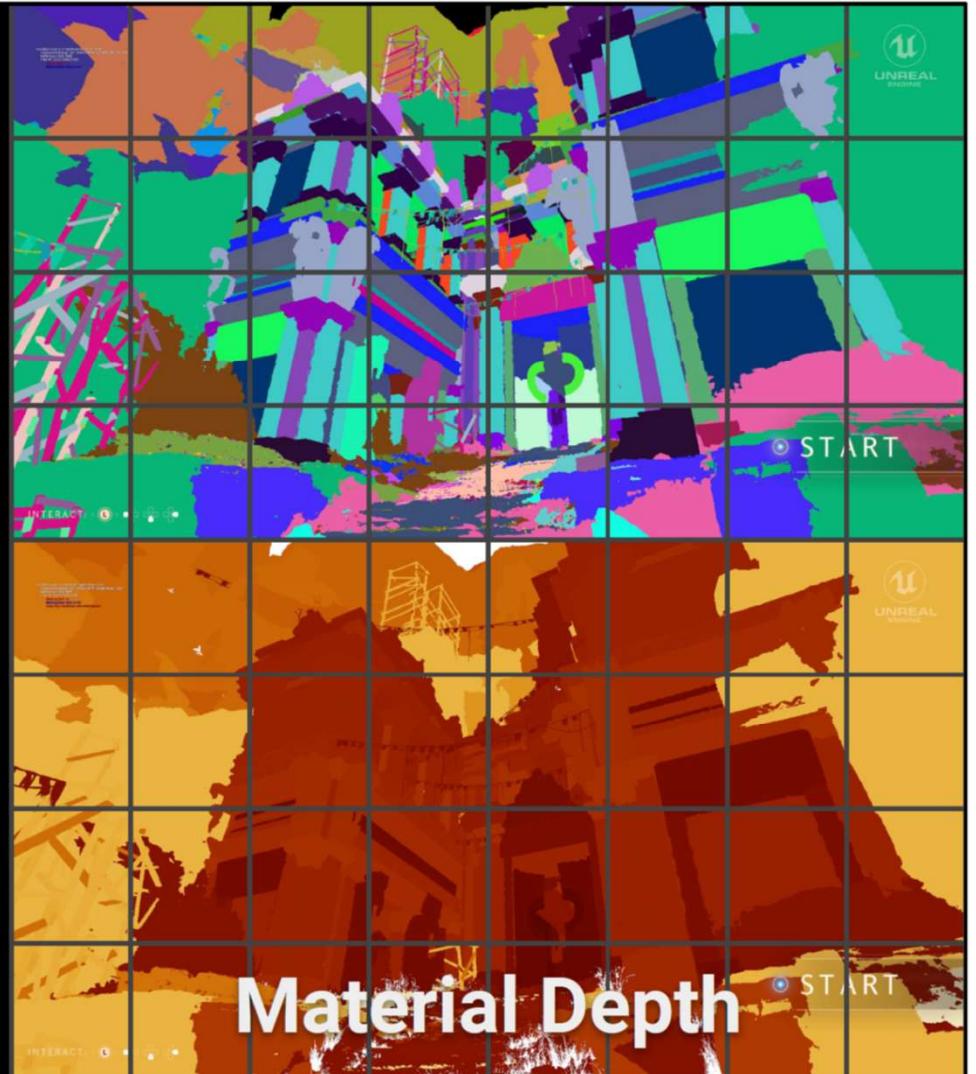
- Full screen quad per unique material
- Skip pixels not matching this material ID
- CPU unaware if some materials have no visible pixels
  - Material draw calls issued regardless
  - Unfortunate side effect of GPU driven
- How to do efficiently?
  - Don't test every pixel for matching material ID for every material pass





# Material culling

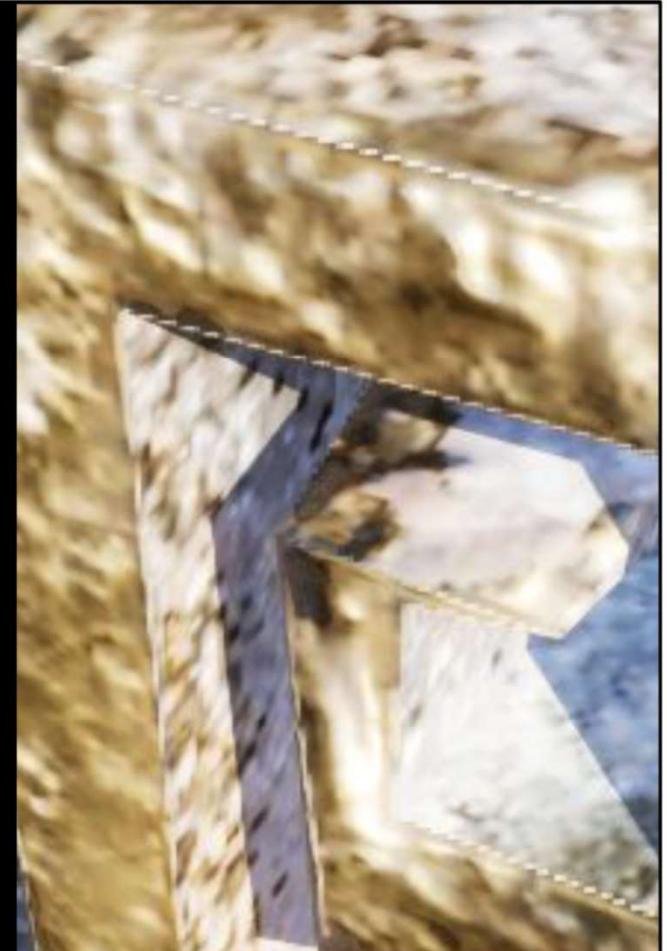
- Material covers small portion of the screen
  - HiZ handles this OK
  - We can do better
- Coarse tile classification / culling
  - Render 8x4 grid of tiles per material
  - Same shading approach as full screen quads
- Tile killed in vertex shader from 32b mask
  - $X=\text{NaN}$





# UV derivatives

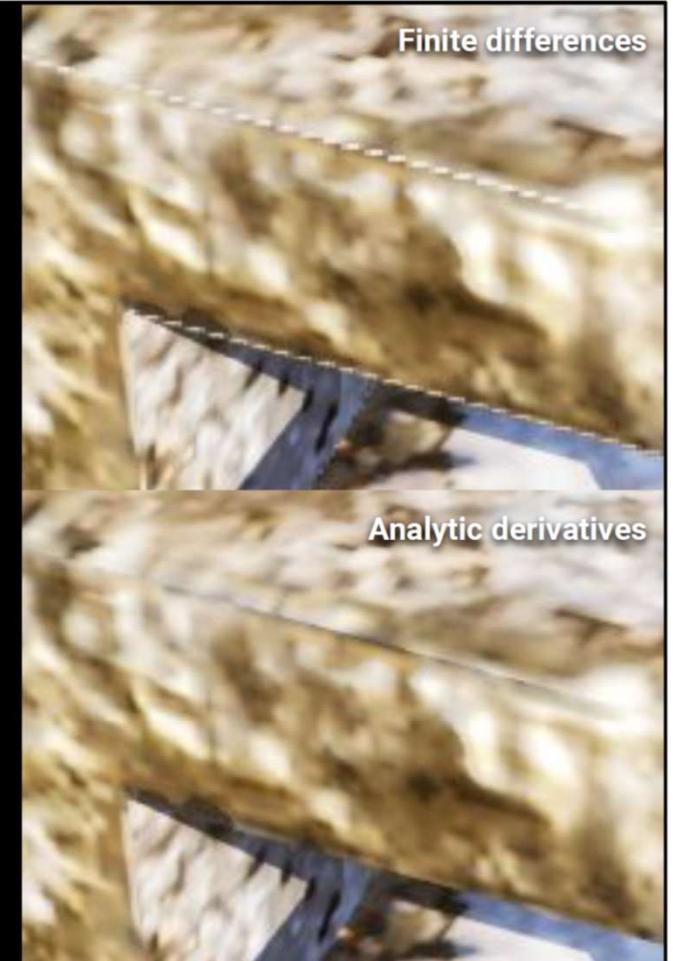
- Still a coherent pixel shader so we have finite difference derivatives
  - Pixel quads span
    - Triangles
  - Also span
    - Depth discontinuities
    - UV seams
    - Different objects
- ← Good!
- }
- Not good!





# Analytic derivatives

- Compute analytic derivatives
  - Attribute gradient across triangle
- Propagate through material node graph using chain rule
- If derivative can't be evaluated analytically
  - Fall back to finite differences
- Used to sample textures with SampleGrad
- Additional cost tiny
  - <2% overhead for material pass
  - Only affects calculations that affect texture sampling
  - Virtual texturing code already does SampleGrad





# Pipeline numbers



## Main pass

Instances pre-cull	896322
Instances post-cull	3668
Cluster node visits	39274
Cluster candidates	1536794
Visible clusters SW	184828
Visible clusters HW	6686

## Post pass

Instances pre-cull	102804
Instances post-cull	365
Cluster node visits	19139
Cluster candidates	458805
Visible clusters SW	7370
Visible clusters HW	536

## Total rasterized

Clusters	199,420
Triangles	25,041,711
Vertices	19,851,262



# Nanite shadows

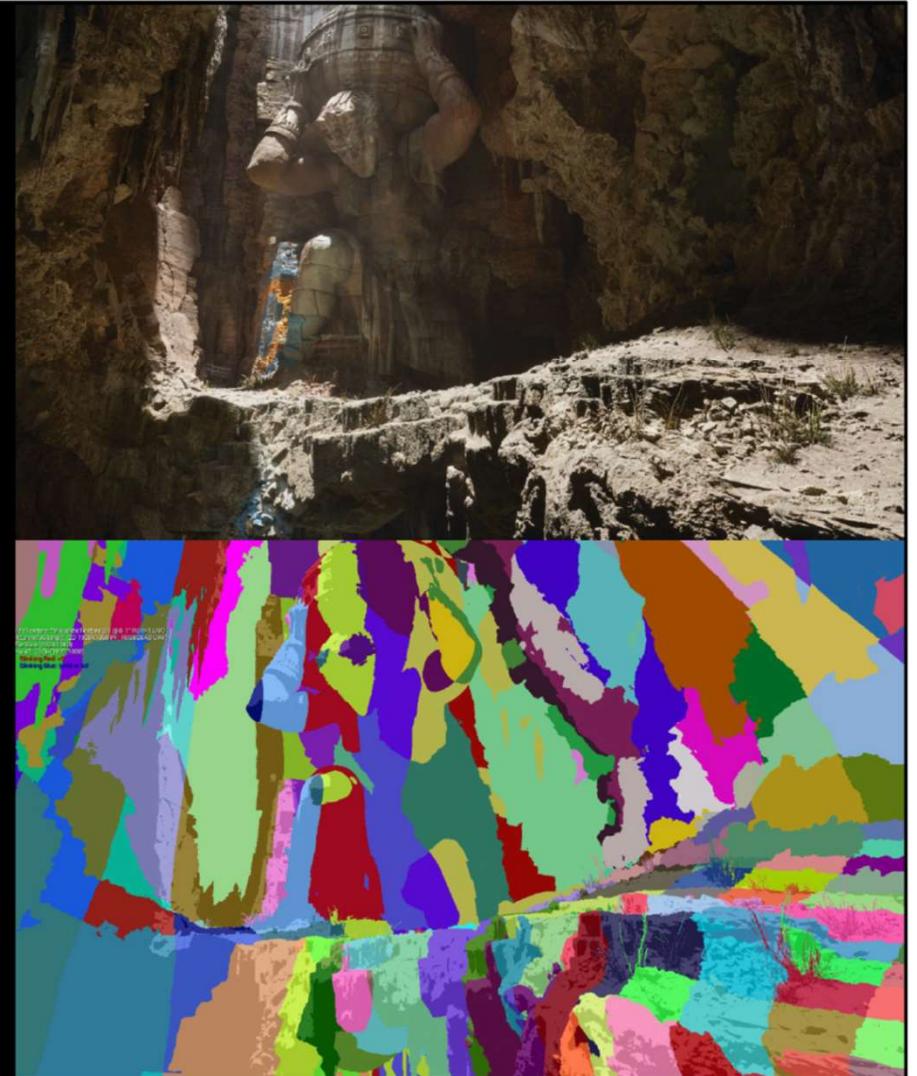
- Ray trace?
  - DXR isn't flexible enough
    - Complex LOD logic
    - Custom triangle encoding
    - No partial BVH updates
- Want a raster solution
  - Leverage all our other work
- Most lights don't move
  - Should cache as much as possible





# Virtual shadow maps

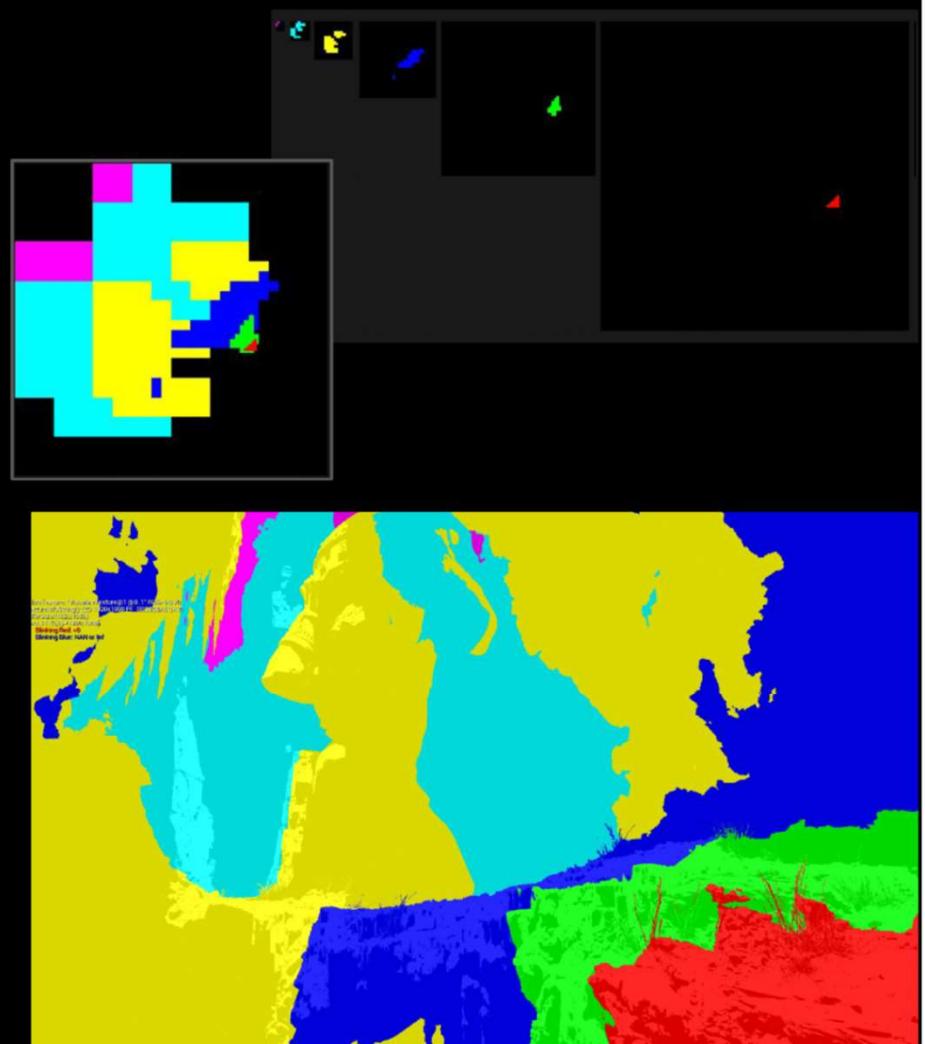
- Nanite enables new techniques
- 16k x 16k shadow maps everywhere
  - Spot: 1x projection
  - Point: 6x cube
  - Directional: Nx clipmaps
- Pick mip level where 1 texel = 1 pixel
- Only render the shadow map pixels that are visible
- Nanite culled and LODded to the detail required





# Virtual shadow maps

- Page size = 128 x 128
- Page table = 128 x 128, with mips
- Mark needed pages
  - Screen pixels project to shadow space
  - Pick mip level where 1 texel = 1 pixel
  - Mark that page
- Allocate physical pages for all needed
- If cached page already exists use that
  - And wasn't invalidated
  - Remove from needed page mask



Thank you.