

CS 380 - GPU and GPGPU Programming

Lecture 26: Programming Tensor Cores

Markus Hadwiger, KAUST

Reading Assignment #15++



Further suggested reading:

- Raihan et al., arXiv, Feb 2019, Modeling Deep Learning Accelerator Enabled GPUs
 - <https://arxiv.org/abs/1811.08309>
 - See also GPGPU-SIM: <http://www.gpgpu-sim.org/>
- CUTLASS 2.8 template library (last update Nov 2021)
 - <https://developer.nvidia.com/blog/cutlass-linear-algebra-cuda/>
 - <https://github.com/NVIDIA/cutlass>
- Register Cache: Caching for Warp-Centric CUDA Programs
 - <https://developer.nvidia.com/blog/register-cache-warp-cuda/>
- cuSPARSE library description in the CUDA SDK
- CUSP library: <http://cusplibrary.github.io/>
- Incomplete-LU and Cholesky Preconditioned Iterative Methods Using CUSPARSE and CUBLAS, Maxim Naumov
 - https://developer.download.nvidia.com/assets/cuda/files/psts_white_paper_final.pdf



Programming Tensor Cores

NVIDIA Volta SM

Multiprocessor: SM

- 64 FP32 + INT32 cores
- 32 FP64 cores
- 8 tensor cores
(FP16/FP32 mixed-precision)

4 partitions inside SM

- 16 FP32 + INT32 cores each
- 8 FP64 cores each
- 8 LD/ST units each
- 2 tensor cores each
- Each has: warp scheduler, dispatch unit, register file



NVIDIA Turing SM

Multiprocessor: SM

- 64 FP32 + INT32 cores
- 2 (!) FP64 cores
- 8 Turing tensor cores
(FP16/32, INT4/8 mixed-precision)
- 1 RT (ray tracing) core

4 partitions inside SM

- 16 FP32 + INT32 cores each
- 4 LD/ST units each
- 2 Turing tensor cores each
- Each has: warp scheduler,
dispatch unit, 16K register file



NVIDIA GA100 SM

Multiprocessor: SM

- 64 FP32 + 64 INT32 cores
- 32 FP64 cores
- 4 3rd gen tensor cores
- 1 2nd gen RT (ray tracing) core

4 partitions inside SM

- 16 FP32 + 16 INT32 cores
- 8 FP64 cores
- 8 LD/ST units each
- 1 3rd gen tensor core each
- Each has: warp scheduler, dispatch unit, 16K register file



NVIDIA GA102 SM

Multiprocessor: SM

- 128 (64+64) FP32 + 64 INT32 cores
- 2 (!) FP64 cores
- 4 3rd gen tensor cores
- 1 2nd gen RT (ray tracing) core

4 partitions inside SM

- 16+16 FP32 + 16 INT32 cores
- 4 LD/ST units each
- 1 3rd gen tensor core each
- Each has: warp scheduler, dispatch unit, 16K register file





Tensor Cores

Mixed-precision, fast matrix-matrix multiply and accumulate (mma)

$$\mathbf{D} = \left(\begin{array}{cccc} \mathbf{A}_{0,0} & \mathbf{A}_{0,1} & \mathbf{A}_{0,2} & \mathbf{A}_{0,3} \\ \mathbf{A}_{1,0} & \mathbf{A}_{1,1} & \mathbf{A}_{1,2} & \mathbf{A}_{1,3} \\ \mathbf{A}_{2,0} & \mathbf{A}_{2,1} & \mathbf{A}_{2,2} & \mathbf{A}_{2,3} \\ \mathbf{A}_{3,0} & \mathbf{A}_{3,1} & \mathbf{A}_{3,2} & \mathbf{A}_{3,3} \end{array} \right) \text{FP16 or FP32} \times \left(\begin{array}{cccc} \mathbf{B}_{0,0} & \mathbf{B}_{0,1} & \mathbf{B}_{0,2} & \mathbf{B}_{0,3} \\ \mathbf{B}_{1,0} & \mathbf{B}_{1,1} & \mathbf{B}_{1,2} & \mathbf{B}_{1,3} \\ \mathbf{B}_{2,0} & \mathbf{B}_{2,1} & \mathbf{B}_{2,2} & \mathbf{B}_{2,3} \\ \mathbf{B}_{3,0} & \mathbf{B}_{3,1} & \mathbf{B}_{3,2} & \mathbf{B}_{3,3} \end{array} \right) \text{FP16} + \left(\begin{array}{cccc} \mathbf{C}_{0,0} & \mathbf{C}_{0,1} & \mathbf{C}_{0,2} & \mathbf{C}_{0,3} \\ \mathbf{C}_{1,0} & \mathbf{C}_{1,1} & \mathbf{C}_{1,2} & \mathbf{C}_{1,3} \\ \mathbf{C}_{2,0} & \mathbf{C}_{2,1} & \mathbf{C}_{2,2} & \mathbf{C}_{2,3} \\ \mathbf{C}_{3,0} & \mathbf{C}_{3,1} & \mathbf{C}_{3,2} & \mathbf{C}_{3,3} \end{array} \right) \text{FP16 or FP32}$$

From this, build larger shapes (sizes), higher dimensionalities, ...


API currently only allows using larger shapes (16x16, ...) in warps (wmma)



Tensor Cores

Fused matrix multiply and accumulate



- Input matrices can be (at most) half-precision (FP16); (**Ampere has more!**)
- Accumulate can be FP16 or FP32; (**Ampere has more!**)



Ampere Tensor Cores: Mixed Precision



New in Ampere: TF32, BF16, FP64



plus FP64 (new in Ampere; GA100 only)


plus INT4/INT8/binary data types (experimental; introduced in Turing)

Ampere Tensor Cores: Sparsity Support



Sparse MMA instructions

2:4 structured sparsity





Low-level options

- CUDA C WMMA (warp-level matrix multiply and accumulate)
- PTX wmma and mma (needed for some features) instructions
- SASS hmma instructions (not documented)

High-level options

- NVIDIA CUTLASS (template abstractions for hi-perf matrix-multiplies)
- NVIDIA cuBLAS
- NVIDIA cuDNN
- Integration into TensorFlow, ...

CUDA C Warp Matrix Functions (WMMA)



Warp Level Matrix Multiply Accumulate (WMMA)

CUDA C Programming Guide (11.5), Appendix B.24

namespace **nvcuda::wmma** (and **nvcuda::wmma::experimental**)

```
template<typename Use, int m, int n, int k, typename T, typename Layout=void>
class fragment;

void load_matrix_sync(fragment<...> &a, const T* mptr, unsigned ldm);
void load_matrix_sync(fragment<...> &a, const T* mptr, unsigned ldm, layout_t layout);
void store_matrix_sync(T* mptr, const fragment<...> &a, unsigned ldm, layout_t layout);
void fill_fragment(fragment<...> &a, const T& v);
void mma_sync(fragment<...> &d, const fragment<...> &a, const fragment<...> &b, const fragment<...> &c, bool satf=false);
```

Concept of a matrix *fragment* (section of a matrix split across threads in a warp)

Dimensions **m,n,k**: **m X k matrix_a**; **k X n matrix_b**; **m X n accumulator**

CUDA C Warp Matrix Functions (WMMA)



Data types (\mathbf{T})

Volta, Turing, and Ampere:

wmma API splits
this into fragments

| Matrix A | Matrix B | Accumulator | Matrix Size (m-n-k) |
|----------------------------|----------------------------|---------------------|-----------------------|
| <code>__half</code> | <code>__half</code> | <code>float</code> | <code>16x16x16</code> |
| <code>__half</code> | <code>__half</code> | <code>float</code> | <code>32x8x16</code> |
| <code>__half</code> | <code>__half</code> | <code>float</code> | <code>8x32x16</code> |
| <code>__half</code> | <code>__half</code> | <code>__half</code> | <code>16x16x16</code> |
| <code>__half</code> | <code>__half</code> | <code>__half</code> | <code>32x8x16</code> |
| <code>__half</code> | <code>__half</code> | <code>__half</code> | <code>8x32x16</code> |
| <code>unsigned char</code> | <code>unsigned char</code> | <code>int</code> | <code>16x16x16</code> |
| <code>unsigned char</code> | <code>unsigned char</code> | <code>int</code> | <code>32x8x16</code> |
| <code>unsigned char</code> | <code>unsigned char</code> | <code>int</code> | <code>8x32x16</code> |
| <code>signed char</code> | <code>signed char</code> | <code>int</code> | <code>16x16x16</code> |
| <code>signed char</code> | <code>signed char</code> | <code>int</code> | <code>32x8x16</code> |
| <code>signed char</code> | <code>signed char</code> | <code>int</code> | <code>8x32x16</code> |



CUDA C Warp Matrix Functions (WMMA)

Data types (\mathbf{T})

wmma API splits
this into fragments

Alternate Floating Point support:

| Matrix A | Matrix B | Accumulator | Matrix Size (m-n-k) |
|------------------------------|------------------------------|-------------|---------------------|
| <code>__nv_bfloat16</code> | <code>__nv_bfloat16</code> | float | 16x16x16 |
| <code>__nv_bfloat16</code> | <code>__nv_bfloat16</code> | float | 32x8x16 |
| <code>__nv_bfloat16</code> | <code>__nv_bfloat16</code> | float | 8x32x16 |
| <code>precision::tf32</code> | <code>precision::tf32</code> | float | 16x16x8 |

Double Precision Support:

| Matrix A | Matrix B | Accumulator | Matrix Size (m-n-k) |
|---------------------|---------------------|---------------------|---------------------|
| <code>double</code> | <code>double</code> | <code>double</code> | 8x8x4 |

Experimental support for sub-byte operations:

Turing and Ampere:

| Matrix A | Matrix B | Accumulator | Matrix Size (m-n-k) |
|----------------------------|----------------------------|-------------|---------------------|
| <code>precision::u4</code> | <code>precision::u4</code> | int | 8x8x32 |
| <code>precision::s4</code> | <code>precision::s4</code> | int | 8x8x32 |
| <code>precision::b1</code> | <code>precision::b1</code> | int | 8x8x128 |

CUDA C Warp Matrix Functions (WMMA)



Warp Level Matrix Multiply Accumulate (WMMA)

CUDA C Programming Guide (11.5), Appendix B.24

```
#include <mma.h>

using namespace nvcuda;

__global__ void wmma_ker(half *a, half *b, float *c) {
    // Declare the fragments
    wmma::fragment<wmma::matrix_a, 16, 16, 16, half, wmma::col_major> a_frag;
    wmma::fragment<wmma::matrix_b, 16, 16, 16, half, wmma::row_major> b_frag;
    wmma::fragment<wmma::accumulator, 16, 16, 16, float> c_frag;

    // Initialize the output to zero
    wmma::fill_fragment(c_frag, 0.0f);

    // Load the inputs
    wmma::load_matrix_sync(a_frag, a, 16);
    wmma::load_matrix_sync(b_frag, b, 16);

    // Perform the matrix multiplication
    wmma::mma_sync(c_frag, a_frag, b_frag, c_frag);

    // Store the output
    wmma::store_matrix_sync(c, c_frag, 16, wmma::mem_row_major);
}
```



PTX WMMA and MMA Instructions

PTX ISA 7.5, Section 9.7.13

| Instruction | Sparsity | Multiplicand Data-type | Shape | PTX ISA version |
|-------------|----------|---|--|---------------------------------------|
| wmma | Dense | Floating-point - .f16 | .m16n16k16, .m8n32k16, and .m32n8k16 | PTX ISA version 6.0 |
| wmma | Dense | Alternate floating-point format - .bf16 | .m16n16k16, .m8n32k16, and .m32n8k16 | PTX ISA version 7.0 |
| wmma | Dense | Alternate floating-point format - .tf32 | .m16n16k8 | PTX ISA version 7.0 |
| wmma | Dense | Integer - .u8/.s8 | .m16n16k16, .m8n32k16, and .m32n8k16 | PTX ISA version 6.3 |
| wmma | Dense | Sub-byte integer - .u4/.s4 | .m8n8k32 | PTX ISA version 6.3 (preview feature) |
| wmma | Dense | Single-bit - .b1 | .m8n8k128 | PTX ISA version 6.3 (preview feature) |



PTX WMMA and MMA Instructions

PTX ISA 7.5

| Instruction | Sparsity | Multiplicand Data-type | Shape | PTX ISA version |
|-------------|----------|---|---------------------------------------|---------------------|
| mma | Dense | Floating-point - .f64 | .m8n8k4 | PTX ISA version 7.0 |
| mma | Dense | Floating-point - .f16 | .m8n8k4 | PTX ISA version 6.4 |
| | | | .m16n8k8 | PTX ISA version 6.5 |
| | | | .m16n8k16 | PTX ISA version 7.0 |
| mma | Dense | Alternate floating-point format - .bf16 | .m16n8k8 and .m16n8k16 | PTX ISA version 7.0 |
| mma | Dense | Alternate floating-point format - .tf32 | .m16n8k4 and .m16n8k8 | PTX ISA version 7.0 |
| mma | Dense | Integer - .u8/.s8 | .m8n8k16 | PTX ISA version 6.5 |
| | | | .m16n8k16 and .m16n8k32 | PTX ISA version 7.0 |
| | | | .m8n8k32 | PTX ISA version 6.5 |
| mma | Dense | Sub-byte integer - .u4/.s4 | .m16n8k32 and .m16n8k64 | PTX ISA version 7.0 |
| | | | .m8n8k128, .m16n8k128, and .m16n8k256 | PTX ISA version 7.0 |
| | | | .m16n8k16 and .m16n8k32 | PTX ISA version 7.1 |
| mma | Sparse | Floating-point - .f16 | .m16n8k16 and .m16n8k32 | PTX ISA version 7.1 |
| mma | Sparse | Alternate floating-point format - .bf16 | .m16n8k16 and .m16n8k32 | PTX ISA version 7.1 |
| mma | Sparse | Alternate floating-point format - .tf32 | .m16n8k8 and .m16n8k16 | PTX ISA version 7.1 |
| mma | Sparse | Integer - .u8/.s8 | .m16n8k32 and .m16n8k64 | PTX ISA version 7.1 |
| mma | Sparse | Sub-byte integer - .u4/.s4 | .m16n8k64 and .m16n8k128 | PTX ISA version 7.1 |



PTX WMMA and MMA Instructions

Load and store: wmma

wmma.load

Collectively load a matrix from memory for WMMA

Syntax

Floating point format .f16 loads:

```
wmma.load.a.sync.aligned.layout.shape{.ss}.atype r, [p] {, stride};  
wmma.load.b.sync.aligned.layout.shape{.ss}.btype r, [p] {, stride};  
wmma.load.c.sync.aligned.layout.shape{.ss}.ctype r, [p] {, stride};  
  
.layout = {.row, .col};  
.shape = {.m16n16k16, .m8n32k16, .m32n8k16};  
.ss = {.global, .shared};  
.atype = {.f16, .s8, .u8};  
.btype = {.f16, .s8, .u8};  
.ctype = {.f16, .f32, .s32};
```

Alternate floating point format .bf16 loads:

```
wmma.load.a.sync.aligned.layout.shape{.ss}.atype r, [p] {, stride};  
wmma.load.b.sync.aligned.layout.shape{.ss}.btype r, [p] {, stride};  
wmma.load.c.sync.aligned.layout.shape{.ss}.ctype r, [p] {, stride};  
  
.layout = {.row, .col};  
.shape = {.m16n16k16, .m8n32k16, .m32n8k16};  
.ss = {.global, .shared};  
.atype = {.bf16 };  
.btype = {.bf16 };  
.ctype = {.f32 };
```

Alternate floating point format .tf32 loads:

```
wmma.load.a.sync.aligned.layout.shape{.ss}.atype r, [p] {, stride};  
wmma.load.b.sync.aligned.layout.shape{.ss}.btype r, [p] {, stride};  
wmma.load.c.sync.aligned.layout.shape{.ss}.ctype r, [p] {, stride};  
  
.layout = {.row, .col};  
.shape = {.m16n16k8 };  
.ss = {.global, .shared};  
.atype = {.tf32 };  
.btype = {.tf32 };  
.ctype = {.f32 };
```



PTX WMMA and MMA Instructions

Load and store: wmma

wmma.load

Collectively load a matrix from memory for WMMA

Syntax

Double precision Floating point .f64 loads:

```
wmma.load.a.sync.aligned.layout.shape{.ss}.atype r, [p] {, stride}
wmma.load.b.sync.aligned.layout.shape{.ss}.btype r, [p] {, stride}
wmma.load.c.sync.aligned.layout.shape{.ss}.ctype r, [p] {, stride}
.layout = {.row, .col};
.shape = {.m8n8k4 };
.ss = {.global, .shared};
.atype = {.f64 };
.btype = {.f64 };
.ctype = {.f64 };
```

Sub-byte loads:

```
wmma.load.a.sync.aligned.row.shape{.ss}.atype r, [p] {, stride}
wmma.load.b.sync.aligned.col.shape{.ss}.btype r, [p] {, stride}
wmma.load.c.sync.aligned.layout.shape{.ss}.ctype r, [p] {, stride}
.layout = {.row, .col};
.shape = {.m8n8k32};
.ss = {.global, .shared};
.atype = {.s4, .u4};
.btype = {.s4, .u4};
.ctype = {.s32};
```

Single-bit loads:

```
wmma.load.a.sync.aligned.row.shape{.ss}.atype r, [p] {, stride}
wmma.load.b.sync.aligned.col.shape{.ss}.btype r, [p] {, stride}
wmma.load.c.sync.aligned.layout.shape{.ss}.ctype r, [p] {, stride}
.layout = {.row, .col};
.shape = {.m8n8k128};
.ss = {.global, .shared};
.atype = {.b1};
.btype = {.b1};
.ctype = {.s32};
```



PTX WMMA and MMA Instructions

wmma example

```
.global .align 32 .f16 A[256], B[256];
.global .align 32 .f32 C[256], D[256];
.reg .b32 a<8> b<8> c<8> d<8>;

wmma.load.a.sync.aligned.m16n16k16.global.row.f16
    {a0, a1, a2, a3, a4, a5, a6, a7}, [A];
wmma.load.b.sync.aligned.m16n16k16.global.col.f16
    {b0, b1, b2, b3, b4, b5, b6, b7}, [B];

wmma.load.c.sync.aligned.m16n16k16.global.row.f32
    {c0, c1, c2, c3, c4, c5, c6, c7}, [C];

wmma.mma.sync.aligned.m16n16k16.row.col.f32.f32
    {d0, d1, d2, d3, d4, d5, d6, d7},
    {a0, a1, a2, a3, a4, a5, a6, a7},
    {b0, b1, b2, b3, b4, b5, b6, b7},
    {c0, c1, c2, c3, c4, c5, c6, c7};

wmma.store.d.sync.aligned.m16n16k16.global.col.f32
    [D], {d0, d1, d2, d3, d4, d5, d6, d7};
```



PTX WMMA and MMA Instructions

mma: fixed assignments of matrix fragments to registers in each thread of warp

9.7.13.4.2. Matrix Fragments for mma.m8n8k4 with .f64 floating point type

A warp executing `mma.m8n8k4` with `.f64` floating point type will compute an MMA operation of shape `.m8n8k4`.

Elements of the matrix are distributed across the threads in a warp so each thread of the warp holds a fragment of the matrix.

- ▶ Multiplicand A:

| .atype | Fragment | Elements (low to high) |
|-------------------|--|------------------------|
| <code>.f64</code> | A vector expression containing a single <code>.f64</code> register, containing single <code>.f64</code> element from the matrix A. | <code>a0</code> |

| Row\Col | 0 | 1 | 2 | 3 |
|---------|--------|--------|--------|--------|
| 0 | T0:a0 | T1:a0 | T2:a0 | T3:a0 |
| 1 | T4:a0 | T5:a0 | T6:a0 | T7:a0 |
| 2 | | | | |
| .. | | | | |
| 7 | T28:a0 | T29:a0 | T30:a0 | T31:a0 |

%laneid:{fragments}



PTX WMMA and MMA Instructions

mma: fixed assignments of matrix fragments to registers in each thread of warp

9.7.13.4.1. Matrix Fragments for mma.m8n8k4 with .f16 floating point type

A warp executing `mma.m8n8k4` with `.f16` floating point type will compute 4 MMA operations of shape `.m8n8k4`.

Elements of 4 matrices need to be distributed across the threads in a warp. The following table shows distribution of matrices for MMA operations.

| MMA Computation | Threads participating in MMA computation |
|-------------------|--|
| MMA computation 1 | Threads with <code>%laneid</code> 0-3 (low group) and 16-19 (high group) |
| MMA computation 2 | Threads with <code>%laneid</code> 4-7 (low group) and 20-23 (high group) |
| MMA computation 3 | Threads with <code>%laneid</code> 8-11 (low group) and 24-27 (high group) |
| MMA computation 4 | Threads with <code>%laneid</code> 12-15 (low group) and 28-31 (high group) |

- Multiplicand A:

| .atype | Fragment | Elements (low to high) |
|-------------------|---|------------------------|
| <code>.f16</code> | A vector expression containing two <code>.f16x2</code> registers, with each register containing two <code>.f16</code> elements from the matrix A. | a0, a1, a2, a3 |

► Row Major:

MMA computation 1

| Row\Col | 0 | 1 | 2 | 3 |
|---------|-------------------------|---|---|---|
| 0 | T0 : { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 3 | T3: { a0, a1, a2, a3 } | | | |
| 4 | T16: { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 7 | T19: { a0, a1, a2, a3 } | | | |

MMA computation 3

| Row\Col | 0 | 1 | 2 | 3 |
|---------|-------------------------|---|---|---|
| 0 | T8 : { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 3 | T11: { a0, a1, a2, a3 } | | | |
| 4 | T24: { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 7 | T27: { a0, a1, a2, a3 } | | | |

MMA computation 2

| Row\Col | 0 | 1 | 2 | 3 |
|---------|-------------------------|---|---|---|
| 0 | T4 : { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 3 | T7: { a0, a1, a2, a3 } | | | |
| 4 | T20: { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 7 | T23: { a0, a1, a2, a3 } | | | |

MMA computation 4

| Row\Col | 0 | 1 | 2 | 3 |
|---------|--------------------------|---|---|---|
| 0 | T12 : { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 3 | T15: { a0, a1, a2, a3 } | | | |
| 4 | T28: { a0, a1, a2, a3 } | | | |
| .. | | ↓ | | |
| 7 | T31: { a0, a1, a2, a3 } | | | |

`%laneid:{fragments}`



PTX WMMA and MMA Instructions

mma: fixed assignments of matrix fragments to registers in each thread of warp

9.7.13.4.1. Matrix Fragments for mma.m8n8k4 with .f16 floating point type

A warp executing `mma.m8n8k4` with `.f16` floating point type will compute 4 MMA operations of shape `.m8n8k4`.

Elements of 4 matrices need to be distributed across the threads in a warp. The following table shows distribution of matrices for MMA operations.

| MMA Computation | Threads participating in MMA computation |
|-------------------|--|
| MMA computation 1 | Threads with <code>%laneid</code> 0-3 (low group) and 16-19 (high group) |
| MMA computation 2 | Threads with <code>%laneid</code> 4-7 (low group) and 20-23 (high group) |
| MMA computation 3 | Threads with <code>%laneid</code> 8-11 (low group) and 24-27 (high group) |
| MMA computation 4 | Threads with <code>%laneid</code> 12-15 (low group) and 28-31 (high group) |

► `.ctype` is `.f16`

| MMA computation 1 | | | | | | | |
|-------------------|--|---|---|---|---|---|---|
| Row\Col | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| 0 | T0 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 3 | T3 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| 4 | T16 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 7 | T19 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |

| MMA computation 3 | | | | | | | |
|-------------------|--|---|---|---|---|---|---|
| Row\Col | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| 0 | T8 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 3 | T11 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| 4 | T24 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 7 | T27 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |

| MMA computation 2 | | | | | | | |
|-------------------|--|---|---|---|---|---|---|
| Row\Col | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| 0 | T4 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 3 | T7 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| 4 | T20 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 7 | T23 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |

| MMA computation 4 | | | | | | | |
|-------------------|--|---|---|---|---|---|---|
| Row\Col | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| 0 | T12 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 3 | T15 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| 4 | T28 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |
| .. | | | | | | | |
| 7 | T31 : { c0, c1, c2, c3, c4, c5, c6, c7 } | | | | | | |

► `.ctype` is `.f32`

| MMA computation 1 | | | | | | | |
|-------------------|------------------|------------------|------------------|------------------|---|---|---|
| R\c | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| 0 | T0 : { c0, c1 } | T2 : { c0, c1 } | T0 : { c4, c5 } | T2 : { c4, c5 } | | | |
| 1 | T1 : { c0, c1 } | T3 : { c0, c1 } | T1 : { c4, c5 } | T3 : { c4, c5 } | | | |
| 2 | T0 : { c2, c3 } | T2 : { c2, c3 } | T0 : { c6, c7 } | T2 : { c6, c7 } | | | |
| 3 | T1 : { c2, c3 } | T3 : { c2, c3 } | T1 : { c6, c7 } | T3 : { c6, c7 } | | | |
| 4 | T16 : { c0, c1 } | T18 : { c0, c1 } | T16 : { c4, c5 } | T18 : { c4, c5 } | | | |
| 5 | T17 : { c0, c1 } | T19 : { c0, c1 } | T17 : { c4, c5 } | T19 : { c4, c5 } | | | |
| 6 | T16 : { c2, c3 } | T18 : { c2, c3 } | T16 : { c6, c7 } | T18 : { c6, c7 } | | | |
| 7 | T17 : { c2, c3 } | T19 : { c2, c3 } | T17 : { c6, c7 } | T19 : { c6, c7 } | | | |

► Accumulators C (or D):

| .ctype / .dtype | Fragment | Elements (low to high) |
|-----------------|---|--------------------------------|
| .f16 | A vector expression containing four <code>.f16x2</code> registers, with each register containing two <code>.f16</code> elements from the matrix C (or D). | c0, c1, c2, c3, c4, c5, c6, c7 |
| .f32 | A vector expression of eight <code>.f32</code> registers. | |




PTX WMMA and MMA Instructions

Sparse matrices: mma.sp

9.7.13.5. Matrix multiply-accumulate operation using mma.sp instruction with sparse matrix A

This section describes warp-level `mma.sp` instruction with sparse matrix A. This variant of the `mma` operation can be used when A is a structured sparse matrix with 50% zeros in each row distributed in a shape-specific granularity. For an $M \times N \times K$ sparse `mma.sp` operation, the $M \times K$ matrix A is packed into $M \times K / 2$ elements. For each K-wide row of matrix A, 50% elements are zeros and the remaining K/2 non-zero elements are packed in the operand representing matrix A. The mapping of these K/2 elements to the corresponding K-wide row is provided explicitly as metadata.





PTX WMMA and MMA Instructions

Load and store: mma ldmatrix

Warp-wide load matrix instruction

```
// Load a single 8x8 matrix using 64-bit addressing
.reg .b64 addr;
.reg .b32 d;
ldmatrix.sync.aligned.m8n8.x1.shared.b16 {d}, [addr];

// Load two 8x8 matrices in column-major format
.reg .b64 addr;
.reg .b32 d<2>;
ldmatrix.sync.aligned.m8n8.x2.trans.shared.b16 {d0, d1}, [addr];

// Load four 8x8 matrices
.reg .b64 addr;
.reg .b32 d<4>;
ldmatrix.sync.aligned.m8n8.x4.b16 {d0, d1, d2, d3}, [addr];
```




PTX WMMA to SASS

Raihan et al., 2019

Get SASS code from cuobjdump disassembly

Micro-benchmarking



PTX WMMA to SASS



Raihan et al., 2019

Get SASS code from cuobjdump disassembly

| | Cumulative Clock Cycles |
|------|----------------------------|
| SET1 | 10 |
| | 12 |
| | 14 |
| | 18 |
| SET2 | 20 |
| | 22 |
| | 24 |
| | 28 |
| SET3 | 30 |
| | 32 |
| | 34 |
| | 38 |
| SET4 | 40 |
| | 42 |
| | 44 |
| | 54 |

HMMA.884.F32.F32.STEP0 R8, R24.reuse.COL, R22.reuse.ROW, R8;
HMMA.884.F32.F32.STEP1 R10, R24.reuse.COL, R22.reuse.ROW, R10;
HMMA.884.F32.F32.STEP2 R4, R24.reuse.COL, R22.reuse.ROW, R4;
HMMA.884.F32.F32.STEP3 R6, R24.COL, R22.ROW, R6;
HMMA.884.F32.F32.STEP0 R8, R20.reuse.COL, R18.reuse.ROW, R8;
HMMA.884.F32.F32.STEP1 R10, R20.reuse.COL, R18.reuse.ROW, R10;
HMMA.884.F32.F32.STEP2 R4, R20.reuse.COL, R18.reuse.ROW, R4;
HMMA.884.F32.F32.STEP3 R6, R20.COL, R18.ROW, R6;
HMMA.884.F32.F32.STEP0 R8, R14.reuse.COL, R12.reuse.ROW, R8;
HMMA.884.F32.F32.STEP1 R10, R14.reuse.COL, R12.reuse.ROW, R10;
HMMA.884.F32.F32.STEP2 R4, R14.reuse.COL, R12.reuse.ROW, R4;
HMMA.884.F32.F32.STEP3 R6, R14.COL, R12.ROW, R6;
HMMA.884.F32.F32.STEP0 R8, R16.reuse.COL, R2.reuse.ROW, R8;
HMMA.884.F32.F32.STEP1 R10, R16.reuse.COL, R2.reuse.ROW, R10;
HMMA.884.F32.F32.STEP2 R4, R16.reuse.COL, R2.reuse.ROW, R4;
HMMA.884.F32.F32.STEP3 R6, R16.COL, R2.ROW, R6;

(a) Disassembled SASS instructions for Mixed precision mode

PTX WMMA to SASS



Raihan et al., 2019

Get SASS code from cuobjdump disassembly

| | Cumulative Clock Cycles |
|--|----------------------------|
| SET1 | 12 |
| | 21 |
| HMMA.884.F16.F16.STEP0 R4, R22.reuse.T, R12.reuse.T, R4; | 25 |
| HMMA.884.F16.F16.STEP1 R6, R22.T, R12.T, R6; | 34 |
| SET2 | 38 |
| | 47 |
| HMMA.884.F16.F16.STEP0 R4, R16.reuse.T, R14.reuse.T, R4; | 51 |
| HMMA.884.F16.F16.STEP1 R6, R16.T, R14.T, R6; | 64 |
| SET3 | |
| | |
| HMMA.884.F16.F16.STEP0 R4, R18.reuse.T, R8.reuse.T, R4; | |
| HMMA.884.F16.F16.STEP1 R6, R18.T, R8.T, R6; | |
| SET4 | |
| | |
| HMMA.884.F16.F16.STEP0 R4, R2.reuse.T, R10.reuse.T, R4; | |
| HMMA.884.F16.F16.STEP1 R6, R2.T, R10.T, R6; | |

(b) Disassembled SASS instructions for FP16 mode

PTX WMMA to SASS



Raihan et al., 2019, reverse-engineered matrix fragment assignment




Figure 8: Distribution of operand matrix elements to threads for tensor cores in the RTX 2080 (Turing).

PTX WMMA to SASS



Raihan et al., 2019, reverse-engineered Tensor core microarchitecture





Figure 13: Proposed Tensor Core Microarchitecture



DEVELOPING CUDA KERNELS TO PUSH TENSOR CORES TO THE ABSOLUTE LIMIT ON NVIDIA A100

Andrew Kerr, May 21, 2020




NVIDIA AMPERE ARCHITECTURE

NVIDIA A100

New and Faster Tensor Core Operations

- Floating-point Tensor Core operations **8x** and **16x** faster than F32 CUDA Cores
- Integer Tensor Core operations **32x** and **64x** faster than F32 CUDA Cores
- New IEEE double-precision Tensor Cores **2x** faster than F64 CUDA Cores




Additional Data Types and Mode

- Bfloat16, double, Tensor Float 32

Asynchronous copy

- Copy directly into shared memory - deep software pipelines

Many additional new features - see “Inside NVIDIA Ampere Architecture”




PROGRAMMING NVIDIA AMPERE ARCHITECTURE

Deep Learning and Math Libraries using Tensor Cores (with CUDA kernels under the hood)

- cuDNN, cuBLAS, cuTENSOR, cuSOLVER, cuFFT, cuSPARSE
- “CUDNN V8: New Advances in Deep Learning Acceleration” (GTC 2020 - S21685)
- “How CUDA Math Libraries Can Help you Unleash the Power of the New NVIDIA A100 GPU” (GTC 2020 - S21681)
- “Inside the Compilers, Libraries and Tools for Accelerated Computing” (GTC 2020 - S21766)

CUDA C++ Device Code


- CUTLASS, CUDA Math API, CUB, Thrust, libcu++



PROGRAMMING NVIDIA AMPERE ARCHITECTURE with CUDA C++




This is a talk for CUDA programmers



CUTLASS

CUDA C++ Templates for Deep Learning and Linear Algebra



CUTLASS

What's new?

CUTLASS 2.2: optimal performance on NVIDIA Ampere Architecture

- Higher throughput Tensor Cores: more than 2x speedup for all data types
- New floating-point types: bfloat16, Tensor Float 32, double
- Deep software pipelines with `cp.async`: efficient and latency tolerant


CUTLASS 2.1

- Planar complex: complex-valued GEMMs with batching options targeting Volta and Turing Tensor Cores
- BLAS-style host side API

CUTLASS 2.0: significant refactoring using modern C++11 programming

- Efficient: particularly for Turing Tensor Cores
- Tensor Core programming model: reusable components for linear algebra kernels in CUDA
- Documentation, profiling tools, reference implementations, SDK examples, more..


<https://github.com/NVIDIA/cutlass>




CUTLASS PERFORMANCE ON NVIDIA AMPERE ARCHITECTURE

CUTLASS 2.2 - CUDA 11 Toolkit - NVIDIA A100

Mixed Precision Floating Point





Double Precision Floating Point



m=3456, n=4096

Mixed Precision Integer



A complex network graph is displayed as a background, consisting of numerous small white and green circular nodes connected by thin gray lines. The graph is highly interconnected, forming several large clusters and radiating from central hubs.

TENSOR CORES ON NVIDIA AMPERE ARCHITECTURE

WHAT ARE TENSOR CORES?

Matrix operations: $D = \text{op}(A, B) + C$


- Matrix multiply-add
- XOR-POPC

Input Data types: A, B

- half, bfloat16, Tensor Float 32, double, int8, int4, bin1

Accumulation Data Types: C, D

- half, float, int32_t, double




WHAT ARE TENSOR CORES?

Matrix operations: $D = \text{op}(A, B) + C$

- Matrix multiply-add
- XOR-POPC

M -by- N -by- K matrix operation

- Warp-synchronous, collective operation
- 32 threads within warp collectively hold A, B, C, and D operands




NVIDIA AMPERE ARCHITECTURE - TENSOR CORE OPERATIONS

| PTX | Data Types (A * B + C) | Shape | Speedup on NVIDIA A100 (vs F32 CUDA cores) | Speedup on Turing* (vs F32 Cores) | Speedup on Volta* (vs F32 Cores) |
|---------------------------------------|---|-------------------------------|---|--------------------------------------|-------------------------------------|
| mma.sync.m16n8k16 mma.sync.m16n8k8 | F16 * F16 + F16 F16 * F16 + F32 BF16 * BF16 + F32 | 16-by-8-by-16 16-by-8-by-8 | 16x | 8x | 8x |
| mma.sync.m16n8k8 | TF32 * TF32 + F32 | 16-by-8-by-8 | 8x | N/A | N/A |
| mma.sync.m8n8k4 | F64 * F64 + F64 | 8-by-8-by-4 | 2x | N/A | N/A |
| mma.sync.m16n8k32 mma.sync.m8n8k16 | S8 * S8 + S32 | 16-by-8-by-32 8-by-8-by-16 | 32x | 16x | N/A |
| mma.sync.m16n8k64 | S4 * S4 + S32 | 16-by-8-by-64 | 64x | 32x | N/A |
| mma.sync.m16n8k256 | B1 ^ B1 + S32 | 16-by-8-by-256 | 256x | 128x | N/A |

<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#warp-level-matrix-instructions-mma-and-friends>

* Instructions with equivalent functionality for Turing and Volta differ in shape from the NVIDIA Ampere Architecture in several cases.


TENSOR CORE OPERATION: FUNDAMENTAL SHAPE



Warp-wide Tensor Core operation: 8-by-8-by-128b

$S8 * S8 + S32$

8-by-8-by-16





mma.sync.aligned
(via inline PTX)

```
int32_t      D[2];
uint32_t const A;
uint32_t const B;
int32_t const C[2];
```

```
// Example targets 8-by-8-by-16 Tensor Core operation
asm(
    "mma.sync.aligned.m8n8k16.row.col.s32.s8.s8.s32 "
    " { %0, %1 }, "
    " %2, "
    " %3, "
    " { %4, %5 }; "
    :
    "=r"(D[0]), "=r"(D[1])
    :
    "r"(A),     "r"(B),
    "r"(C[0]),  "r"(C[1])
);
```


EXPANDING THE M DIMENSION



Warp-wide Tensor Core operation: 16-by-8-by-128b

F16 * F16 + F32

16-by-8-by-8




mma.sync.aligned
(via inline PTX)

```
float      D[4];
uint32_t const A[2];
uint32_t const B;
float      const C[4];
```


```
// Example targets 16-by-8-by-8 Tensor Core operation
asm(
    "mma.sync.aligned.m16n8k8.row.col.f32.f16.f16.f32 "
    " { %0, %1, %2, %3 }, "
    " { %4, %5}, "
    " %6, "
    " { %7, %8, %9, %10 };"
    :
    "=f"(D[0]), "=f"(D[1]), "=f"(D[2]), "=f"(D[3])
    :
    "r"(A[0]), "r"(A[1]),
    "r"(B),
    "f"(C[0]), "f"(C[1])
);
```

<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#warp-level-matrix-instructions-mma-and-friends>

EXPANDING THE K DIMENSION




Warp-wide Tensor Core operation: 16-by-8-by-256b



F16 * F16 + F32

16-by-8-by-16



mma.sync.aligned (via inline PTX)


```
float      D[4];
uint32_t const A[4];
uint32_t const B[2];
float      const C[4];
```

```
// Example targets 16-by-8-by-32 Tensor Core operation
asm(
    "mma.sync.aligned.m16n8k16.row.col.f32.f16.f16.f32 "
    " { %0, %1, %2, %3 },      "
    " { %4, %5, %6, %7 },      "
    " { %8, %9 },              "
    " { %10, %11, %12, %13 };"
    :
    "=f"(D[0]), "=f"(D[1]), "=f"(D[2]), "=f"(D[3])
    :
    "r"(A[0]), "r"(A[1]), "r"(A[2]), "r"(A[3]),
    "r"(B[0]), "r"(B[1]),
    "f"(C[0]), "f"(C[1]), "f"(C[2]), "f"(C[3])
);
```

<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#warp-level-matrix-instructions-mma-and-friends>

$S8 * S8 + S32$

16-by-8-by-32



mma.sync.aligned
(via inline PTX)


```
int32_t      D[4];
uint32_t const A[4];
uint32_t const B[2];
int32_t const C[4];
```

```
// Example targets 16-by-8-by-32 Tensor Core operation
asm(
    "mma.sync.aligned.m16n8k32.row.col.s32.s8.s8.s32 "
    " { %0, %1, %2, %3 },      "
    " { %4, %5, %6, %7 },      "
    " { %8, %9 },              "
    " { %10, %11, %12, %13 };"
    :
    "=r"(D[0]), "=r"(D[1]), "=r"(D[2]), "=r"(D[3])
    :
    "r"(A[0]), "r"(A[1]), "r"(A[2]), "r"(A[3]),
    "r"(B[0]), "r"(B[1]),
    "r"(C[0]), "r"(C[1]), "r"(C[2]), "r"(C[3])
);
```

<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#warp-level-matrix-instructions-mma-and-friends>

HALF-PRECISION : F16 * F16 + F16

16-by-8-by-16



`mma.sync.aligned`
(via inline PTX)


```
uint32_t      D[2]; // two registers needed (vs. four)
uint32_t const A[4];
uint32_t const B[2];
uint32_t const C[2]; // two registers needed (vs. four)
```

```
// Example targets 16-by-8-by-16 Tensor Core operation
asm(
    "mma.sync.aligned.m16n8k16.row.col.f16.f16.f16.f16 "
    " { %0, %1},           "
    " { %2, %3, %4, %5 },   "
    " { %6, %7 },           "
    " { %8, %9 };          "
    : "=r"(D[0]), "=r"(D[1])
    :
    "r"(A[0]), "r"(A[1]), "r"(A[2]), "r"(A[3]),
    "r"(B[0]), "r"(B[1]),
    "r"(C[0]), "r"(C[1])
);
```

<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html#warp-level-matrix-instructions-mma-and-friends>

DOUBLE-PRECISION: F64 * F64 + F64

8-by-8-by-4




mma.sync.aligned
(via inline PTX)

```
uint64_t      D[2];    // two 64-bit accumulators
uint64_t const A;      // one 64-bit element for A operand
uint64_t const B;      // one 64-bit element for B operand
uint64_t const C[2];   // two 64-bit accumulators

// Example targets 8-by-8-by-4 Tensor Core operation
asm(
    "mma.sync.aligned.m8n8k4.row.col.f64.f64.f64.f64 "
    " { %0, %1},   "
    " %2,          "
    " %3,          "
    " { %4, %5 }; "
    :
    "=l"(D[0]),  "=l"(D[1])
    :
    "l"(A),
    "l"(B),
    "l"(C[0]),  "l"(C[1])
);
```

CUTLASS: wraps PTX in template

m-by-n-by-k



`cutlass::arch::Mma`


```

/// Matrix multiply-add operation
template <
    /// Size of the matrix product (concept: GemmShape)
    typename Shape,
    /// Number of threads participating
    int kThreads,
    /// Data type of A elements
    typename ElementA,
    /// Layout of A matrix (concept: MatrixLayout)
    typename LayoutA,
    /// Data type of B elements
    typename ElementB,
    /// Layout of B matrix (concept: MatrixLayout)
    typename LayoutB,
    /// Element type of C matrix
    typename ElementC,
    /// Layout of C matrix (concept: MatrixLayout)
    typename LayoutC,
    /// Inner product operator
    typename Operator
>
struct Mma;
```

https://github.com/NVIDIA/cutlass/blob/master/include/cutlass/arch/mma_sm80.h

CUTLASS: wraps PTX in template


16-by-8-by-16



`cutlass::arch::Mma`

```
__global__ void kernel() {  
  
    // arrays containing logical elements  
    Array<half_t, 8> A;  
    Array<half_t, 4> B;  
    Array< float, 4> C;  
  
    // define the appropriate matrix operation  
    arch::Mma< GemmShape<16, 8, 16>, 32, ... > mma;  
  
    // in-place matrix multiply-accumulate  
    mma(C, A, B, C);  
  
    ...  
}
```

https://github.com/NVIDIA/cutlass/blob/master/include/cutlass/arch/mma_sm80.h



EFFICIENT DATA MOVEMENT FOR TENSOR CORES


HELLO WORLD: TENSOR CORES

Map each thread to coordinates of the matrix operation

Load inputs from memory

Perform the matrix operation

Store the result to memory



CUDA example

```
__global__ void tensor_core_example_8x8x16(
    int32_t           *D,
    uint32_t const   *A,
    uint32_t const   *B,
    int32_t const   *C) {

    // Compute the coordinates of accesses to A and B matrices
    int outer = threadIdx.x / 4;      // m or n dimension
    int inner = threadIdx.x % 4;      // k dimension

    // Compute the coordinates for the accumulator matrices
    int c_row = threadIdx.x / 4;
    int c_col = 2 * (threadIdx.x % 4);

    // Compute linear offsets into each matrix
    int ab_idx = outer * 4 + inner;
    int cd_idx = c_row * 8 + c_col;

    // Issue Tensor Core operation
    asm(
        "mma.sync.aligned.m8n8k16.row.col.s32.s8.s32 "
        " { %0, %1 }, "
        " %2, "
        " %3, "
        " { %4, %5 }; "
        :
        "=r"(D[cd_idx]), "=r"(D[cd_idx + 1])
        :
        "r"(A[ab_idx]),
        "r"(B[ab_idx]),
        "r"(C[cd_idx]), "r"(C[cd_idx + 1])
    );
}
```

PERFORMANCE IMPLICATIONS

Load A and B inputs from memory: $2 \times 4B$ per thread

Perform one Tensor Core operation: 2048 flops per warp

2048 flops require 256 B of loaded data

→ 8 flops/byte

NVIDIA A100 Specifications:

- 624 TFLOP/s (INT8)
 - 1.6 TB/s (HBM2)
- 400 flops/byte

8 flops/byte * 1.6 TB/s → 12 TFLOP/s

This kernel is global memory bandwidth limited.

CUDA example

```
__global__ void tensor_core_example_8x8x16(
    int32_t          *D,
    uint32_t const   *A,
    uint32_t const   *B,
    int32_t const   *C) {

    // Compute the coordinates of accesses to A and B matrices
    int outer = threadIdx.x / 4;      // m or n dimension
    int inner = threadIdx.x % 4;      // k dimension


    // Compute the coordinates for the accumulator matrices
    int c_row = threadIdx.x / 4;
    int c_col = 2 * (threadIdx.x % 4);

    // Compute linear offsets into each matrix
    int ab_idx = outer * 4 + inner;
    int cd_idx = c_row * 8 + c_col;

    // Issue Tensor Core operation
    asm(
        "mma.sync.aligned.m8n8k16.row.col.s32.s8.s32 "
        " { %0, %1 }, "
        " %2,
        " %3,
        " { %4, %5 }; "
        :
        "=r"(D[cd_idx]), "+r"(D[cd_idx + 1])
    );
    : "r"(A[ab_idx]),
      "r"(B[ab_idx]),
      "r"(C[cd_idx]),
      "r"(C[cd_idx + 1])
    );
}
```

FEEDING THE DATA PATH

Efficient storing and loading through Shared Memory




Tiled, hierarchical model: reuse data in Shared Memory and in Registers

See [CUTLASS GTC 2018](#) talk for more details about this model.

FEEDING THE DATA PATH

Move data from Global Memory to Tensor Cores as efficiently as possible


- Latency-tolerant pipeline from Global Memory
- Conflict-free Shared Memory stores
- Conflict-free Shared Memory loads



ASYNCHRONOUS COPY: EFFICIENT PIPELINES


New NVIDIA Ampere Architecture feature: cp.async

- Asynchronous copy directly from Global to Shared Memory
- See “*Inside the NVIDIA Ampere Architecture*” for more details (GTC 2020 - S21730)



Enables efficient software pipelines


- Minimizes data movement: L2 → L1 → RF → SMEM becomes L2 → SMEM
- Saves registers: RF no longer needed to hold the results of long-latency load instructions
- Indirection: fetch several stages in advance for greater latency tolerance




FEEDING THE DATA PATH

Move data from Global Memory to Tensor Cores as efficiently as possible

- Latency-tolerant pipeline from Global Memory
- Conflict-free Shared Memory stores
- Conflict-free Shared Memory loads



GLOBAL MEMORY TO TENSOR CORES



LDMATRIX: FETCH TENSOR CORE OPERANDS


PTX instruction to load a matrix from Shared Memory

Each thread supplies a pointer to 128b row of data in Shared Memory

Each 128b row is broadcast to groups of four threads

(potentially different threads than the one supplying the pointer)

Data matches arrangement of inputs to Tensor Core operations



Shared Memory

Tensor Cores

Shared Memory
Pointers

| | | | | | |
|-----------|---|-----|-----|-----|-----|
| T0 | → | T0 | T1 | T2 | T3 |
| T1 | → | T4 | T5 | T6 | T7 |
| T2 | → | T8 | T9 | T10 | T11 |
| T3 | → | T12 | T13 | T14 | T15 |
| T4 | → | T16 | T17 | T18 | T19 |
| T5 | → | T20 | T21 | T22 | T23 |
| T6 | → | T24 | T25 | T26 | T27 |
| T7 | → | T28 | T29 | T30 | T31 |

| | | | | | |
|-----------|---|-----|-----|-----|-----|
| T8 | → | T0 | T1 | T2 | T3 |
| T9 | → | T4 | T5 | T6 | T7 |
| T10 | → | T8 | T9 | T10 | T11 |
| T11 | → | T12 | T13 | T14 | T15 |
| T12 | → | T16 | T17 | T18 | T19 |
| T13 | → | T20 | T21 | T22 | T23 |
| T14 | → | T24 | T25 | T26 | T27 |
| T15 | → | T28 | T29 | T30 | T31 |

Shared Memory
Pointers

| | | | | | |
|------------|---|-----|-----|-----|-----|
| T16 | → | T0 | T1 | T2 | T3 |
| T17 | → | T4 | T5 | T6 | T7 |
| T18 | → | T8 | T9 | T10 | T11 |
| T19 | → | T12 | T13 | T14 | T15 |
| T20 | → | T16 | T17 | T18 | T19 |
| T21 | → | T20 | T21 | T22 | T23 |
| T22 | → | T24 | T25 | T26 | T27 |
| T23 | → | T28 | T29 | T30 | T31 |

| | | | | | |
|------------|---|-----|-----|-----|-----|
| T24 | → | T0 | T1 | T2 | T3 |
| T25 | → | T4 | T5 | T6 | T7 |
| T26 | → | T8 | T9 | T10 | T11 |
| T27 | → | T12 | T13 | T14 | T15 |
| T28 | → | T16 | T17 | T18 | T19 |
| T29 | → | T20 | T21 | T22 | T23 |
| T30 | → | T24 | T25 | T26 | T27 |
| T31 | → | T28 | T29 | T30 | T31 |

| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |
| T0 | T1 | T2 | T3 | | | | |
| T4 | T5 | T6 | T7 | | | | |
| T8 | T9 | T10 | T11 | | | | |
| T12 | T13 | T14 | T15 | | | | |
| T16 | T17 | T18 | T19 | | | | |
| T20 | T21 | T22 | T23 | | | | |
| T24 | T25 | T26 | T27 | | | | |
| T28 | T29 | T30 | T31 | | | | |
| T0 | T1 | T2 | T3 | | | | |
| T4 | T5 | T6 | T7 | | | | |
| T8 | T9 | T10 | T11 | | | | |
| T12 | T13 | T14 | T15 | | | | |
| T16 | T17 | T18 | T19 | | | | |
| T20 | T21 | T22 | T23 | | | | |
| T24 | T25 | T26 | T27 | | | | |
| T28 | T29 | T30 | T31 | | | | |

LDMATRIX: PTX INSTRUCTION

PTX instruction to load a matrix from SMEM

Each thread supplies a pointer to 128b row of data in Shared Memory

Each 128b row is broadcast to groups of four threads


(potentially different threads than the one supplying the pointer)

Data matches arrangement of inputs to Tensor Core operations



```
// Inline PTX assembly for ldmatrix

uint32_t R[4];
uint32_t smem_ptr;

asm volatile (
    "ldmatrix.sync.aligned.x4.m8n8.shared.b16 "
    "{%, %1, %2, %3}, [%4];"
    :
    "=r"(R[0]), "=r"(R[1]), "=r"(R[2]), "=r"(R[3])
    :
    "r"(smem_ptr)
);
```



GLOBAL MEMORY TO TENSOR CORES



NVIDIA AMPERE ARCHITECTURE - SHARED MEMORY BANK TIMING

Bank conflicts between threads in the same phase


4B words are accessed in 1 phase

8B words are accessed in 2 phases:

- Process addresses of the first 16 threads in a warp
- Process addresses of the second 16 threads in a warp

16B words are accessed in 4 phases:

- Each phase processes **8 consecutive threads** of a warp





128 bit access size

Slide borrowed from: Guillaume Thomas-Collignon and Paulius Micikevicius. "Volta Architecture and performance optimization." GTC 2018.

<http://on-demand.gputechconf.com/gtc/2018/presentation/s81006-volta-architecture-and-performance-optimization.pdf>

GLOBAL MEMORY TO TENSOR CORES



Bank conflict on either store or load
from Shared Memory

GLOBAL TO SHARED MEMORY

Load from Global Memory


Load
(128 bits per thread)

| | | | | | | | | | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|--|--|--|--|--|--|--|--|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 | | | | | | | | |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 | | | | | | | | |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 | | | | | | | | |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 | | | | | | | | |

Store
(128 bits per thread)

Store to Shared Memory

| | | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|--|--|--|--|--|--|--|--|--|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 | | | | | | | | | |
| T9 | T8 | T11 | T10 | T13 | T12 | T15 | T14 | | | | | | | | | |
| T18 | T19 | T16 | T17 | T22 | T23 | T20 | T21 | | | | | | | | | |
| T27 | T26 | T25 | T24 | T31 | T30 | T29 | T28 | | | | | | | | | |




Permuted Shared Memory layout


XOR function maps thread index to Shared Memory location

GLOBAL TO SHARED MEMORY

Load from Global Memory




Load
(128 bits per thread)



Store to Shared Memory

Store
(128 bits per thread)



Phase 0: T0 .. T7

Phase 1: T8 .. T15

Phase 2: T16 .. T23


Phase 3: T24 .. T31

GLOBAL TO SHARED MEMORY

Load from Global Memory

| | | | | | | | | | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|--|--|--|--|--|--|--|--|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 | | | | | | | | |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 | | | | | | | | |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 | | | | | | | | |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 | | | | | | | | |

Load
(128 bits per thread)



Store to Shared Memory

| | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|--|--|--|--|--|--|--|--|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 | | | | | | | | |
| T9 | T8 | T11 | T10 | T13 | T12 | T15 | T14 | | | | | | | | |
| T18 | T19 | T16 | T17 | T22 | T23 | T20 | T21 | | | | | | | | |
| T27 | T26 | T25 | T24 | T31 | T30 | T29 | T28 | | | | | | | | |

Store
(128 bits per thread)

Phase 0: T0 .. T7

Phase 1: T8 .. T15

Phase 2: T16 .. T23


Phase 3: T24 .. T31

GLOBAL TO SHARED MEMORY

Load from Global Memory

| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |

Load
(128 bits per thread)



Store to Shared Memory

Store
(128 bits per thread)

| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 |
| T9 | T8 | T11 | T10 | T13 | T12 | T15 | T14 |
| T18 | T19 | T16 | T17 | T22 | T23 | T20 | T21 |
| T27 | T26 | T25 | T24 | T31 | T30 | T29 | T28 |

Phase 0: T0 .. T7

Phase 1: T8 .. T15

Phase 2: T16 .. T23

Phase 3: T24 .. T31

GLOBAL TO SHARED MEMORY


Load from Global Memory

Load (128 bits per thread)

Store to Shared Memory

| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 |
| T9 | T8 | T11 | T10 | T13 | T12 | T15 | T14 |
| T18 | T19 | T16 | T17 | T22 | T23 | T20 | T21 |
| T27 | T26 | T25 | T24 | T31 | T30 | T29 | T28 |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

Store (128 bits per thread)



Phase 0: T0 .. T7

Phase 1: T8 .. T15


Phase 2: T16 .. T23

Phase 3: T24 .. T31

FEEDING THE DATA PATH

Move data from Global Memory to Tensor Cores as efficiently as possible

- Latency-tolerant pipeline from Global Memory
- Conflict-free Shared Memory stores
- **Conflict-free Shared Memory loads**



LOADING FROM SHARED MEMORY TO REGISTERS

Logical view of threadblock tile

| | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 | T20 | T21 | T22 | T23 | T24 | T25 | T26 | T27 | T28 | T29 | T30 | T31 |
| | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | |

Shared Memory Pointers

| | |
|-----|---|
| T0 | → |
| T1 | → |
| T2 | → |
| T3 | → |
| T4 | → |
| T5 | → |
| T6 | → |
| T7 | → |
| T8 | → |
| T9 | → |
| T10 | → |
| T11 | → |
| T12 | → |
| T13 | → |
| T14 | → |
| T15 | → |

Shared Memory Pointers

| | | | | | |
|-----|---|-----|-----|-----|-----|
| T16 | → | T0 | T1 | T2 | T3 |
| T17 | → | T4 | T5 | T6 | T7 |
| T18 | → | T8 | T9 | T10 | T11 |
| T19 | → | T12 | T13 | T14 | T15 |
| T20 | → | T16 | T17 | T18 | T19 |
| T21 | → | T20 | T21 | T22 | T23 |
| T22 | → | T24 | T25 | T26 | T27 |
| T23 | → | T28 | T29 | T30 | T31 |
| T24 | → | T0 | T1 | T2 | T3 |
| T25 | → | T4 | T5 | T6 | T7 |
| T26 | → | T8 | T9 | T10 | T11 |
| T27 | → | T12 | T13 | T14 | T15 |
| T28 | → | T16 | T17 | T18 | T19 |
| T29 | → | T20 | T21 | T22 | T23 |
| T30 | → | T24 | T25 | T26 | T27 |
| T31 | → | T28 | T29 | T30 | T31 |

Load Matrix from Shared Memory

| | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|--|-----|-----|--|--|-----|-----|--|--|--|--|--|
| T0 | T16 | | | | T1 | T17 | | | | | | | | | |
| T18 | T2 | | | | T19 | T3 | | | | | | | | | |
| | | T4 | T20 | | | | | | T5 | T21 | | | | | |
| | | T22 | T6 | | | | | | T23 | T7 | | | | | |
| T8 | T24 | | | | T9 | T25 | | | | | | | | | |
| T26 | T10 | | | | T27 | T11 | | | | | | | | | |
| | | T12 | T28 | | | | | | T13 | T29 | | | | | |
| | | T30 | T14 | | | | | | T31 | T15 | | | | | |

LOADING FROM SHARED MEMORY TO REGISTERS

Logical view of threadblock tile

| | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 | T20 | T21 | T22 | T23 | T24 | T25 | T26 | T27 | T28 | T29 | T30 | T31 |
| | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | |

Shared Memory Pointers

| | |
|-----|---|
| T0 | → |
| T1 | → |
| T2 | → |
| T3 | → |
| T4 | → |
| T5 | → |
| T6 | → |
| T7 | → |
| T8 | → |
| T9 | → |
| T10 | → |
| T11 | → |
| T12 | → |
| T13 | → |
| T14 | → |
| T15 | → |

Shared Memory Pointers

| | | | |
|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 |
| T4 | T5 | T6 | T7 |
| T8 | T9 | T10 | T11 |
| T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 |
| T20 | T21 | T22 | T23 |
| T24 | T25 | T26 | T27 |
| T28 | T29 | T30 | T31 |

Load Matrix from Shared Memory

| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T16 | | | T1 | T17 | | |
| T18 | T2 | | | T19 | T3 | | |
| | | T4 | T20 | | | T5 | T21 |
| | | T22 | T6 | | | T23 | T7 |
| T8 | T24 | | | T9 | T25 | | |
| T26 | T10 | | | T27 | T11 | | |
| | | T12 | T28 | | | T13 | T29 |
| | | T30 | T14 | | | T31 | T15 |

| | |
|-----|---|
| T8 | → |
| T9 | → |
| T10 | → |
| T11 | → |
| T12 | → |
| T13 | → |
| T14 | → |
| T15 | → |

| | |
|-----|---|
| T16 | → |
| T17 | → |
| T18 | → |
| T19 | → |
| T20 | → |
| T21 | → |
| T22 | → |
| T23 | → |
| T24 | → |
| T25 | → |
| T26 | → |
| T27 | → |
| T28 | → |
| T29 | → |
| T30 | → |
| T31 | → |

| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |

| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |

| | | | |
|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 |
| T4 | T5 | T6 | T7 |
| T8 | T9 | T10 | T11 |
| T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 |
| T20 | T21 | T22 | T23 |
| T24 | T25 | T26 | T27 |
| T28 | T29 | T30 | T31 |

| | | | |
|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 |
| T4 | T5 | T6 | T7 |
| T8 | T9 | T10 | T11 |
| T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 |
| T20 | T21 | T22 | T23 |
| T24 | T25 | T26 | T27 |
| T28 | T29 | T30 | T31 |

LOADING FROM SHARED MEMORY TO REGISTERS

Logical view of threadblock tile

| | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 | T20 | T21 | T22 | T23 | T24 | T25 | T26 | T27 | T28 | T29 | T30 | T31 |
| | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | |

Shared Memory Pointers

| | |
|-----|---|
| T0 | → |
| T1 | → |
| T2 | → |
| T3 | → |
| T4 | → |
| T5 | → |
| T6 | → |
| T7 | → |
| T8 | → |
| T9 | → |
| T10 | → |
| T11 | → |
| T12 | → |
| T13 | → |
| T14 | → |
| T15 | → |

Shared Memory Pointers

| | |
|-----|---|
| T16 | → |
| T17 | → |
| T18 | → |
| T19 | → |
| T20 | → |
| T21 | → |
| T22 | → |
| T23 | → |
| T24 | → |
| T25 | → |
| T26 | → |
| T27 | → |
| T28 | → |
| T29 | → |
| T30 | → |
| T31 | → |

| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |

| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |

Load Matrix from Shared Memory

| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T16 | | | T1 | T17 | | |
| T18 | T2 | | | T19 | T3 | | |
| | | T4 | T20 | | | T5 | T21 |
| | | | | | | | |
| | | T22 | T6 | | | T23 | T7 |
| T8 | T24 | | | T9 | T25 | | |
| T26 | T10 | | | T27 | T11 | | |
| | | T12 | T28 | | | T13 | T29 |
| | | | | | | | |
| | | T30 | T14 | | | T31 | T15 |

| | |
|-----|---|
| T16 | → |
| T17 | → |
| T18 | → |
| T19 | → |
| T20 | → |
| T21 | → |
| T22 | → |
| T23 | → |
| T24 | → |
| T25 | → |
| T26 | → |
| T27 | → |
| T28 | → |
| T29 | → |
| T30 | → |
| T31 | → |

| | | | |
|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 |
| T4 | T5 | T6 | T7 |
| T8 | T9 | T10 | T11 |
| T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 |
| T20 | T21 | T22 | T23 |
| T24 | T25 | T26 | T27 |
| T28 | T29 | T30 | T31 |

LOADING FROM SHARED MEMORY TO REGISTERS

Logical view of threadblock tile

| | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 | T20 | T21 | T22 | T23 | T24 | T25 | T26 | T27 | T28 | T29 | T30 | T31 |
| | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | |

Shared Memory Pointers

| | |
|-----|---|
| T0 | → |
| T1 | → |
| T2 | → |
| T3 | → |
| T4 | → |
| T5 | → |
| T6 | → |
| T7 | → |
| T8 | → |
| T9 | → |
| T10 | → |
| T11 | → |
| T12 | → |
| T13 | → |
| T14 | → |
| T15 | → |

Shared Memory Pointers

| | | | |
|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 |
| T4 | T5 | T6 | T7 |
| T8 | T9 | T10 | T11 |
| T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 |
| T20 | T21 | T22 | T23 |
| T24 | T25 | T26 | T27 |
| T28 | T29 | T30 | T31 |

Load Matrix from Shared Memory

| | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T16 | | | | T1 | T17 | | |
| T18 | T2 | | | | T19 | T3 | | |
| | | T4 | T20 | | | | T5 | T21 |
| | | | | T6 | | | T7 | |
| T8 | T24 | | | | T9 | T25 | | |
| T26 | T10 | | | | T27 | T11 | | |
| | | T12 | T28 | | | | T13 | T29 |
| | | | | T14 | | | T31 | T15 |
| | | | | | | | | |

| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |


| | | | | | | | |
|----|----|-----|-----|-----|-----|-----|-----|
| T0 | T4 | T8 | T12 | T16 | T20 | T24 | T28 |
| T1 | T5 | T9 | T13 | T17 | T21 | T25 | T29 |
| T2 | T6 | T10 | T14 | T18 | T22 | T26 | T30 |
| T3 | T7 | T11 | T15 | T19 | T23 | T27 | T31 |

| | | | | | |
|-----|---|-----|-----|-----|-----|
| T16 | → | T0 | T1 | T2 | T3 |
| T17 | → | T4 | T5 | T6 | T7 |
| T18 | → | T8 | T9 | T10 | T11 |
| T19 | → | T12 | T13 | T14 | T15 |
| T20 | → | T16 | T17 | T18 | T19 |
| T21 | → | T20 | T21 | T22 | T23 |
| T22 | → | T24 | T25 | T26 | T27 |
| T23 | → | T28 | T29 | T30 | T31 |



| | | | | | |
|-----|---|-----|-----|-----|-----|
| T24 | → | T0 | T1 | T2 | T3 |
| T25 | → | T4 | T5 | T6 | T7 |
| T26 | → | T8 | T9 | T10 | T11 |
| T27 | → | T12 | T13 | T14 | T15 |
| T28 | → | T16 | T17 | T18 | T19 |
| T29 | → | T20 | T21 | T22 | T23 |
| T30 | → | T24 | T25 | T26 | T27 |
| T31 | → | T28 | T29 | T30 | T31 |

| | | | |
|-----|-----|-----|-----|
| T0 | T1 | T2 | T3 |
| T4 | T5 | T6 | T7 |
| T8 | T9 | T10 | T11 |
| T12 | T13 | T14 | T15 |
| T16 | T17 | T18 | T19 |
| T20 | T21 | T22 | T23 |
| T24 | T25 | T26 | T27 |
| T28 | T29 | T30 | T31 |

ADVANCING TO NEXT K GROUP



ADVANCING TO NEXT K GROUP



| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| T0 | T16 | | | T1 | T17 | | |
| T18 | T2 | | | T19 | T3 | | |
| | | T4 | T20 | | | T5 | T21 |
| | | T22 | T6 | | | T23 | T7 |
| T8 | T24 | | | T9 | T25 | | |
| T26 | T10 | | | T27 | T11 | | |
| | | T12 | T28 | | | T13 | T29 |
| | | T30 | T14 | | | T31 | T15 |


| | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|
| | | T0 | T16 | | | T1 | T17 |
| | | T18 | T2 | | | T19 | T3 |
| T4 | T20 | | | T5 | T21 | | |
| T22 | T6 | | | T23 | T7 | | |
| | | T8 | T24 | | | T9 | T25 |
| | | T26 | T10 | | | T27 | T11 |
| T12 | T28 | | | T13 | T29 | | |
| T30 | T14 | | | T31 | T15 | | |



smem_ptr = row_idx * 8 + column_idx;


smem_ptr = smem_ptr ^ 2;

LOADING FROM SHARED MEMORY TO REGISTERS




Phase 0

Logical view of threadblock tile





K=16..31

Load Matrix from Shared Memory




LOADING FROM SHARED MEMORY TO REGISTERS




K=16..31

Phase 1





Load Matrix from Shared Memory



LOADING FROM SHARED MEMORY TO REGISTERS


Phase 2

Logical view of threadblock tile



K=16..31



Load Matrix from Shared Memory



LOADING FROM SHARED MEMORY TO REGISTERS


Phase 3

Logical view of threadblock tile



K=16..31


Load Matrix from Shared Memory




CUTLASS

CUDA C++ Templates as an Optimal Abstraction Layer for Tensor Cores

- Latency-tolerant pipeline from Global Memory
- Conflict-free Shared Memory stores
- Conflict-free Shared Memory loads



CUTLASS: OPTIMAL ABSTRACTION FOR TENSOR CORES



```
using Mma = cutlass::gemm::warp::DefaultMmaTensorOp<
    GemmShape<64, 64, 16>,
    half_t, LayoutA,
    half_t, LayoutB,
    float, RowMajor
>;  
  
__shared__ ElementA smem_buffer_A[Mma::Shape::kM * GemmK];
__shared__ ElementB smem_buffer_B[Mma::Shape::kN * GemmK];  
  
// Construct iterators into SMEM tiles
Mma::IteratorA iter_A({smem_buffer_A, lda}, thread_id);
Mma::IteratorB iter_B({smem_buffer_B, ldb}, thread_id);  
  
Mma::FragmentA frag_A;
Mma::FragmentB frag_B;
Mma::FragmentC accum;  
  
Mma mma;  
  
accum.clear();  
  
#pragma unroll 1
for (int k = 0; k < GemmK; k += Mma::Shape::kK) {  
  
    iter_A.load(frag_A); // Load fragments from A and B matrices
    iter_B.load(frag_B);  
  
    ++iter_A; ++iter_B; // Advance along GEMM K to next tile in A
                         // and B matrices  
  
    mma(accum, frag_A, frag_B, accum);
}
```

CUTLASS: OPTIMAL ABSTRACTION FOR TENSOR CORES

Tile Iterator Constructors:

Initialize pointers into permuted Shared Memory buffers

```
using Mma = cutlass::gemm::warp::DefaultMmaTensorOp<
    GemmShape<64, 64, 16>,
    half_t, LayoutA,                                // GEMM A operand
    half_t, LayoutB,                                // GEMM B operand
    float, RowMajor                                // GEMM C operand
>;
```

```
__shared__ ElementA smem_buffer_A[Mma::Shape::kM * GemmK];
__shared__ ElementB smem_buffer_B[Mma::Shape::kN * GemmK];
```

```
// Construct iterators into SMEM tiles
Mma::IteratorA iter_A({smem_buffer_A, lda}, thread_id);
Mma::IteratorB iter_B({smem_buffer_B, ldb}, thread_id);
```

```
Mma::FragmentA frag_A;
Mma::FragmentB frag_B;
Mma::FragmentC accum;
```

```
Mma mma;
```

```
accum.clear();
```

```
#pragma unroll 1
for (int k = 0; k < GemmK; k += Mma::Shape::kK) {
```

```
    iter_A.load(frag_A); // Load fragments from A and B matrices
    iter_B.load(frag_B);
```

```
    ++iter_A; ++iter_B; // Advance along GEMM K to next tile in A
                         // and B matrices
```

```
    mma(accum, frag_A, frag_B, accum);
```

Fragments:

Register-backed arrays holding each thread's data

Tile Iterator:

load() - Fetches data from permuted Shared Memory buffers

operator++() - advances to the next logical matrix in SMEM

Warp-level matrix multiply:

Decomposes a large matrix multiply into Tensor Core operations

Thank you.

- Md Aamir Raihan, Negar Goli, Tor Aamodt
- Andrew Kerr et al., NVIDIA
- NVIDIA