

CS 380 - GPU and GPGPU Programming

Lecture 6: GPU Architecture, Pt. 4

Markus Hadwiger, KAUST

Reading Assignment #3 (until Sep 16)



Read (required):

- Programming Massively Parallel Processors book, 4th edition,
Chapter 4 (*Compute architecture and scheduling*)
- NVIDIA CUDA C++ Programming Guide (current: v12.6, Aug 29, 2024):
*Read **Chapter 5.6*** (Compute Capability);
*Read **Chapter 19.1*** (Compute Capabilities);
*Browse all of **Chapter 19*** (Compute Capabilities)
*Browse all of **Chapter 8.2*** (Maximize Utilization) and
Chapter 8.4 (Maximize Instruction Throughput)

https://docs.nvidia.com/cuda/pdf/CUDA_C_Programming_Guide.pdf

Where this is going...




Summary: three key ideas for high-throughput execution

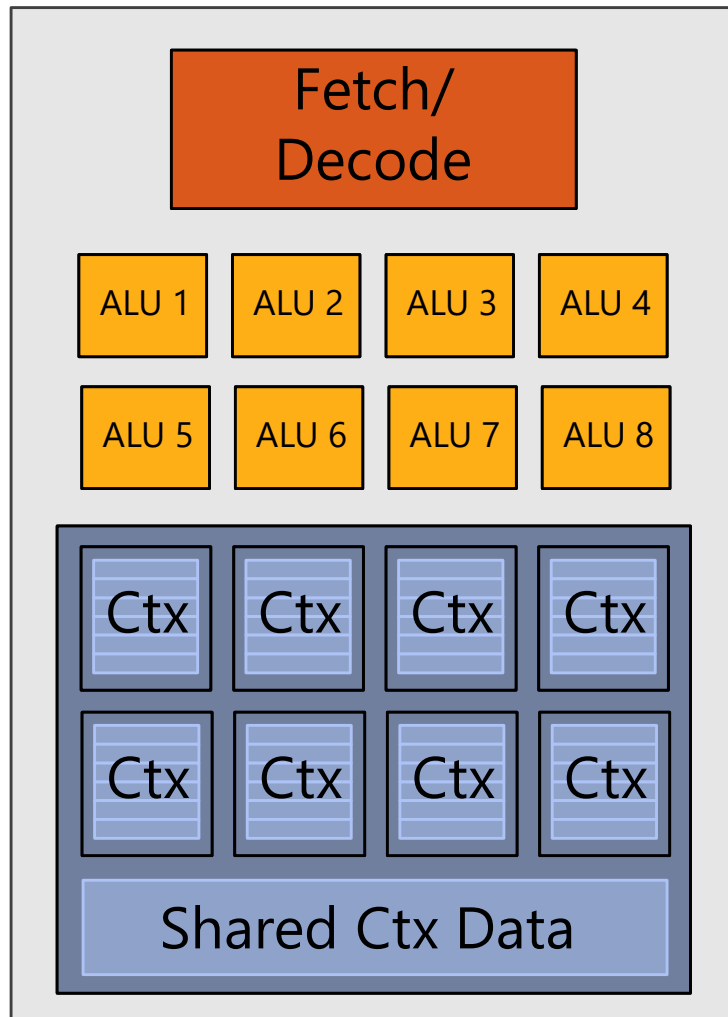
1. Use many “slimmed down cores,” run them in parallel
2. Pack cores full of ALUs (by sharing instruction stream overhead across groups of fragments)
 - Option 1: Explicit SIMD vector instructions
 - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
 - When one group stalls, work on another group

**GPUs are here!
(usually)**

Idea #1: Slim down



Idea #2: Add ALUs (sharing inst. stream)



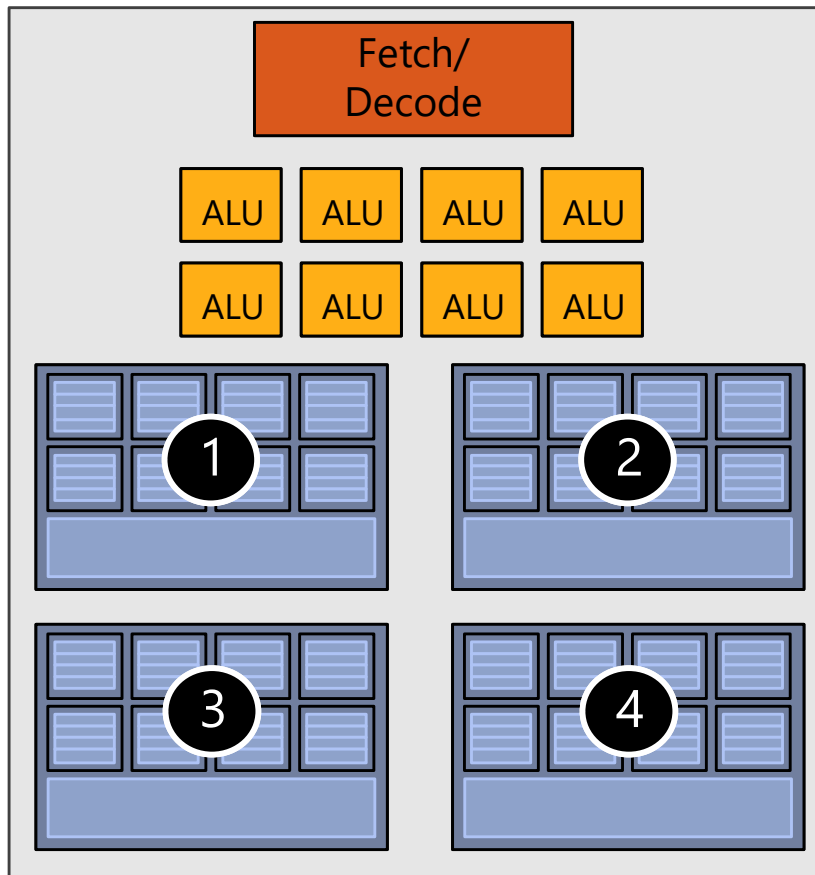
Idea #2:

Amortize cost/complexity of managing an instruction stream across many ALUs

SIMD processing

(or **SIMT**, SPMD)

Idea #3: Store multiple group contexts



Idea #3:

Interleave execution of groups of threads

(the number of groups is *not fixed*, but depends on the context storage requirements of a given kernel!)

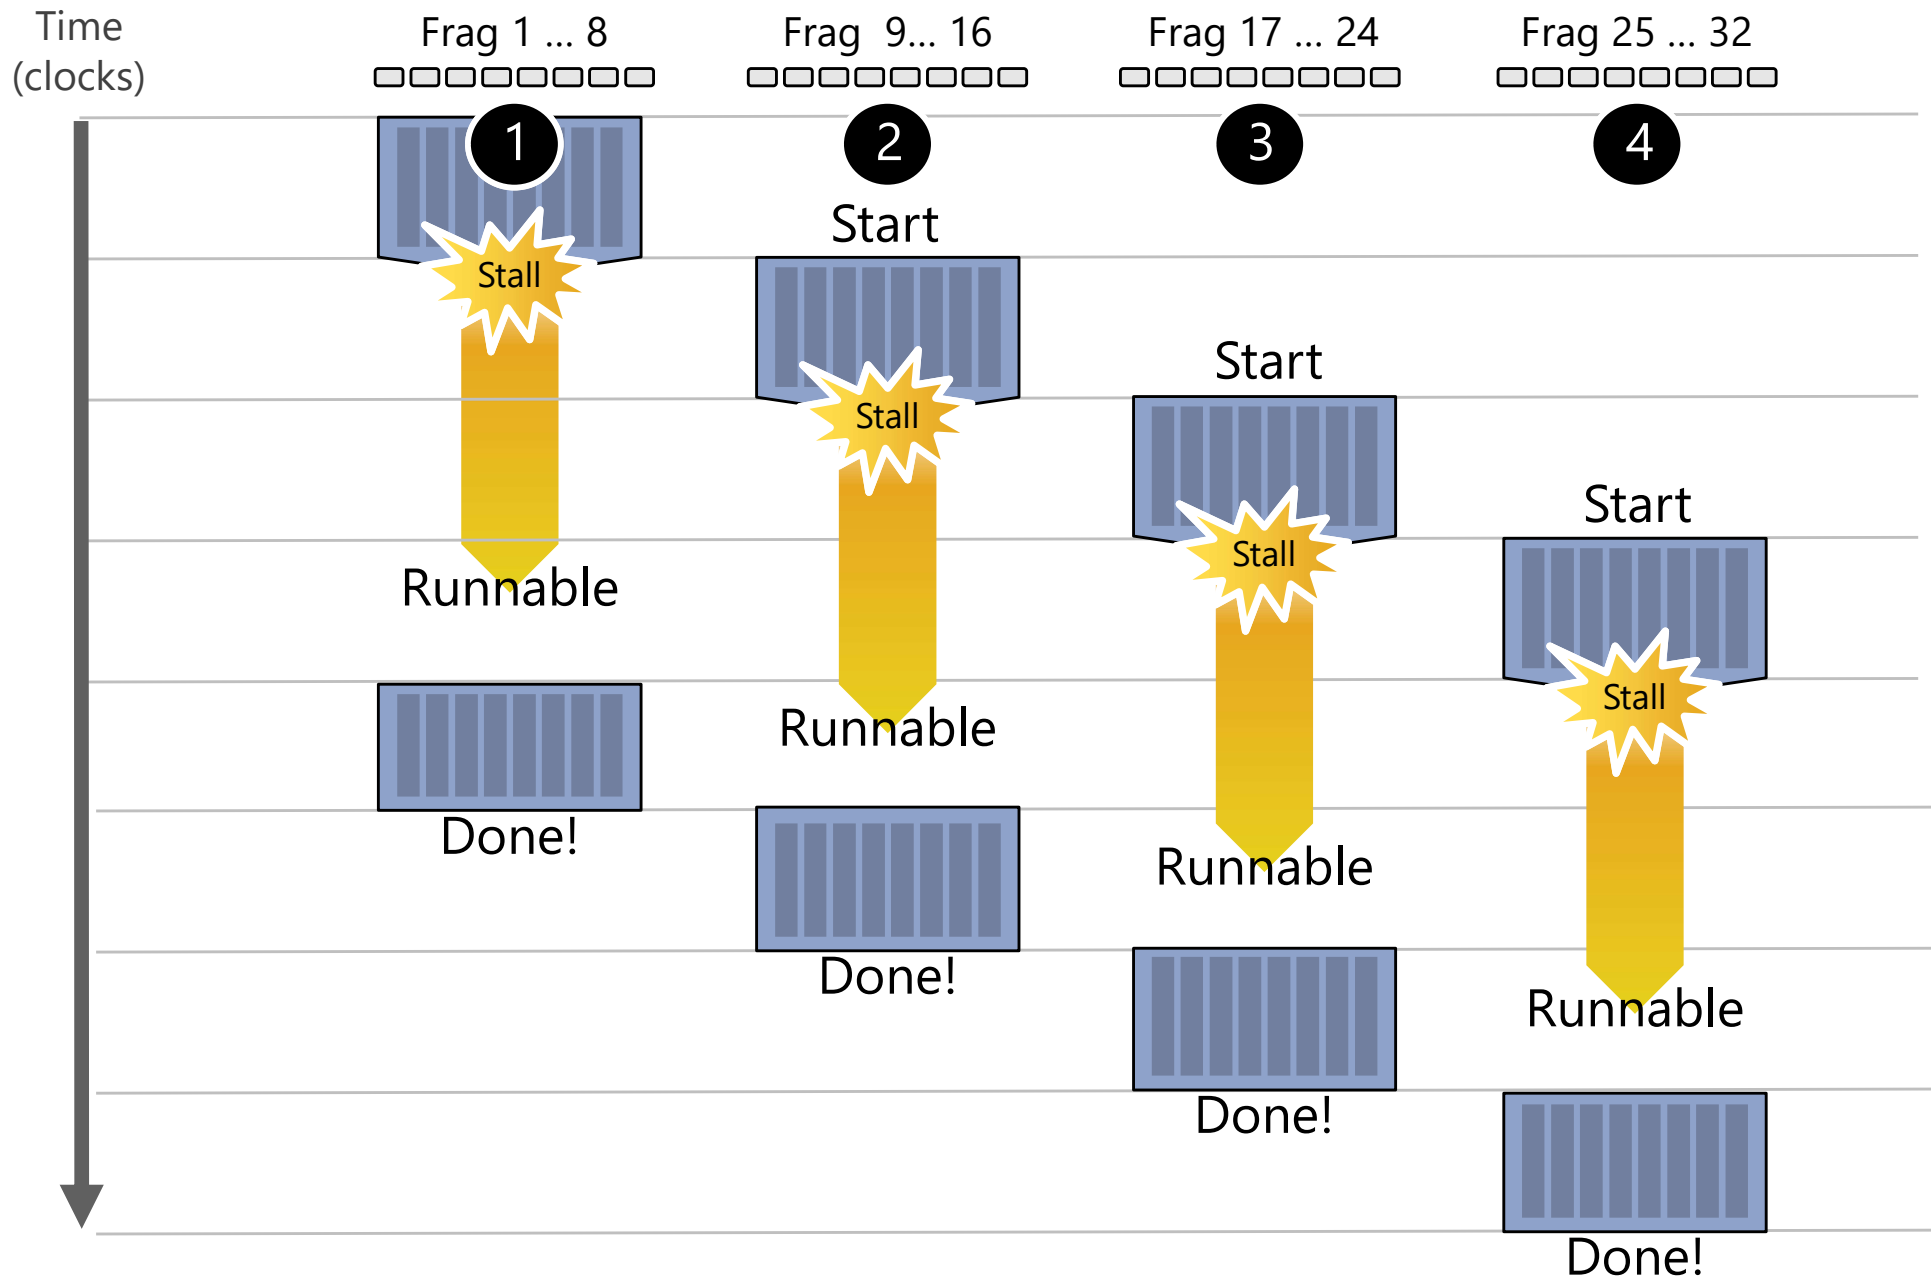
Idea #3: Interleave execution of groups

But we have **LOTS** of independent fragments.

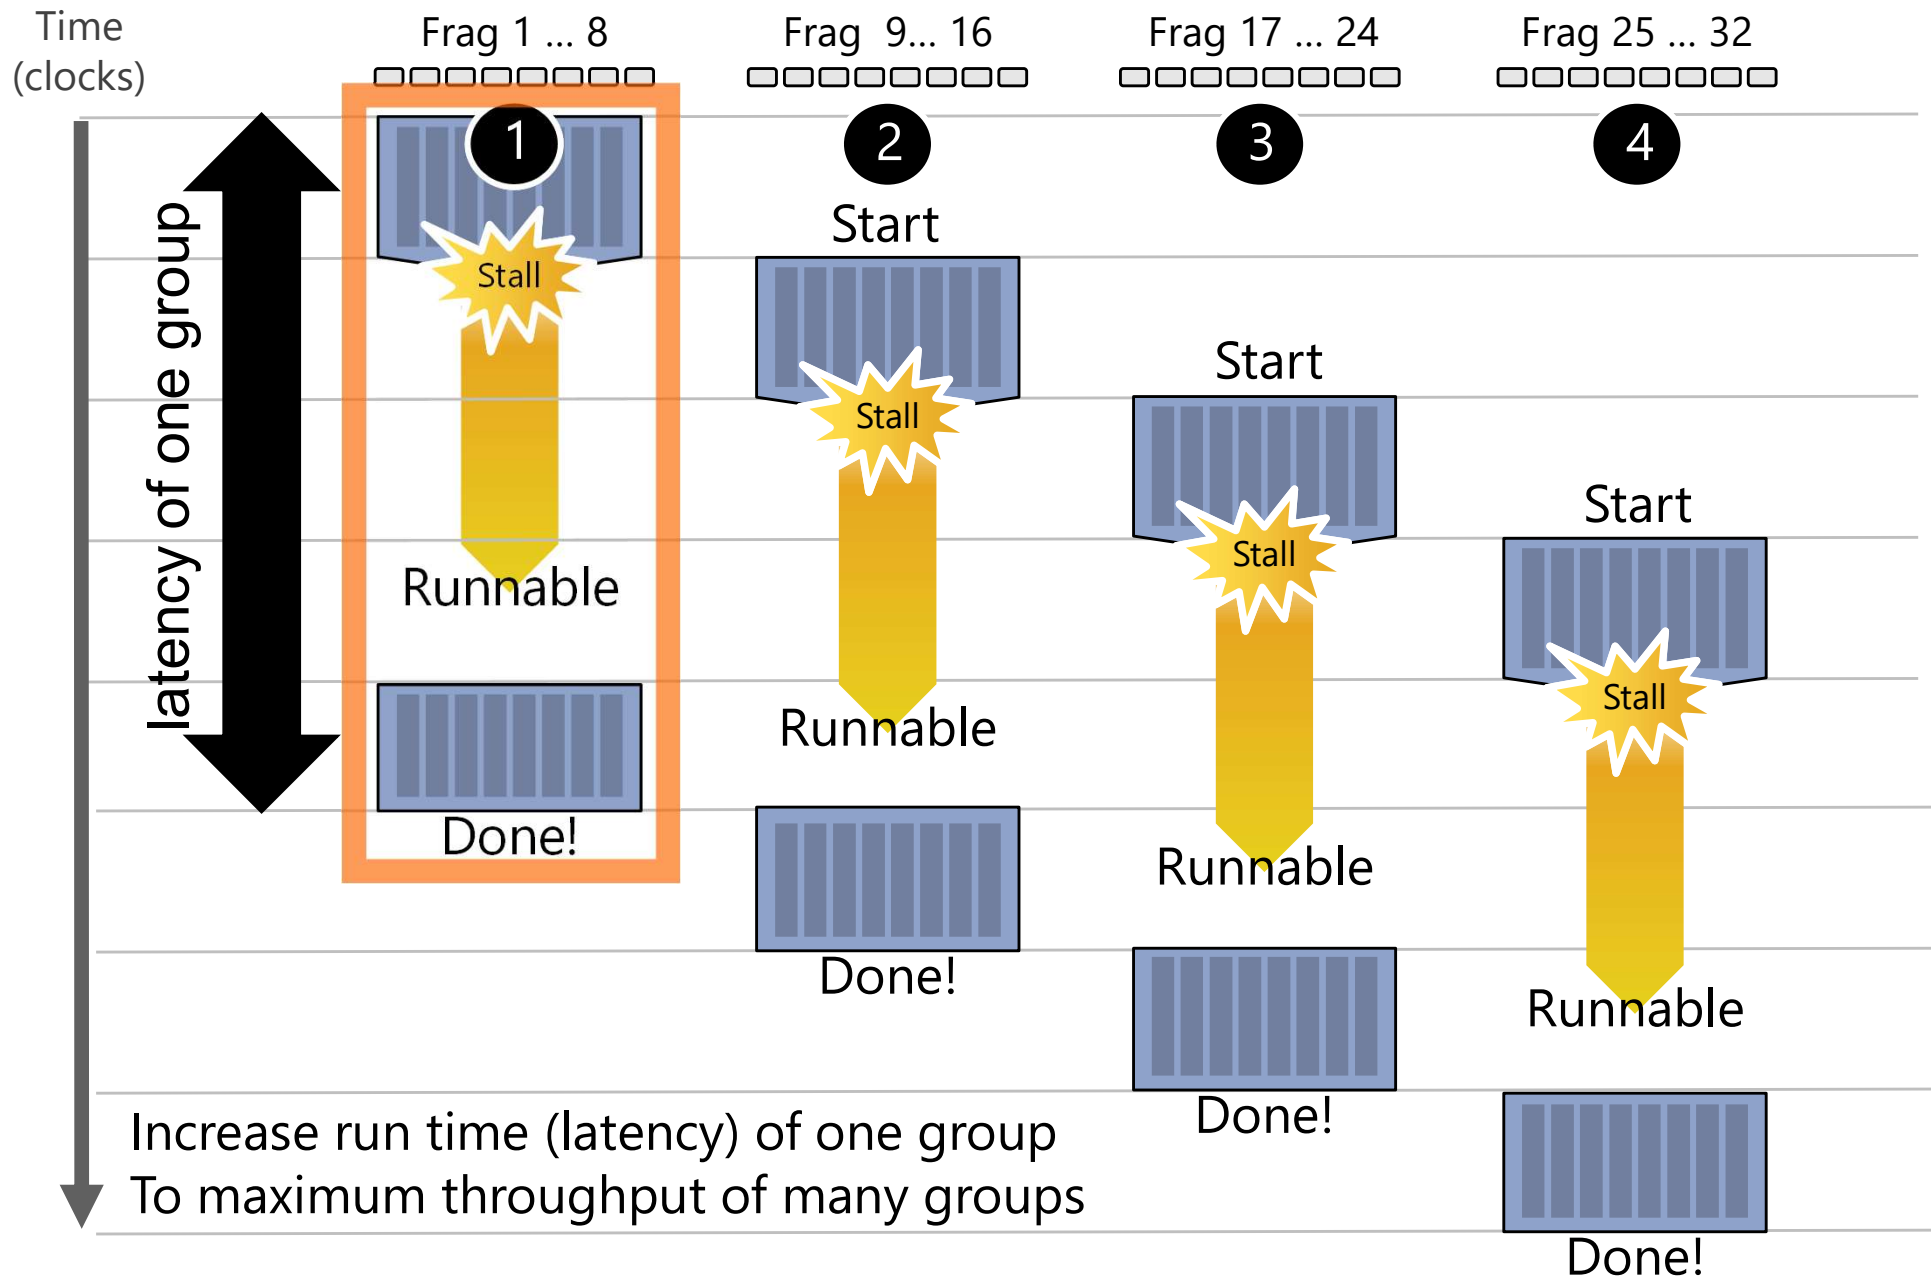
Idea #3:

Interleave processing of many fragments on a single core to avoid stalls caused by high latency operations.


Hiding shader stalls



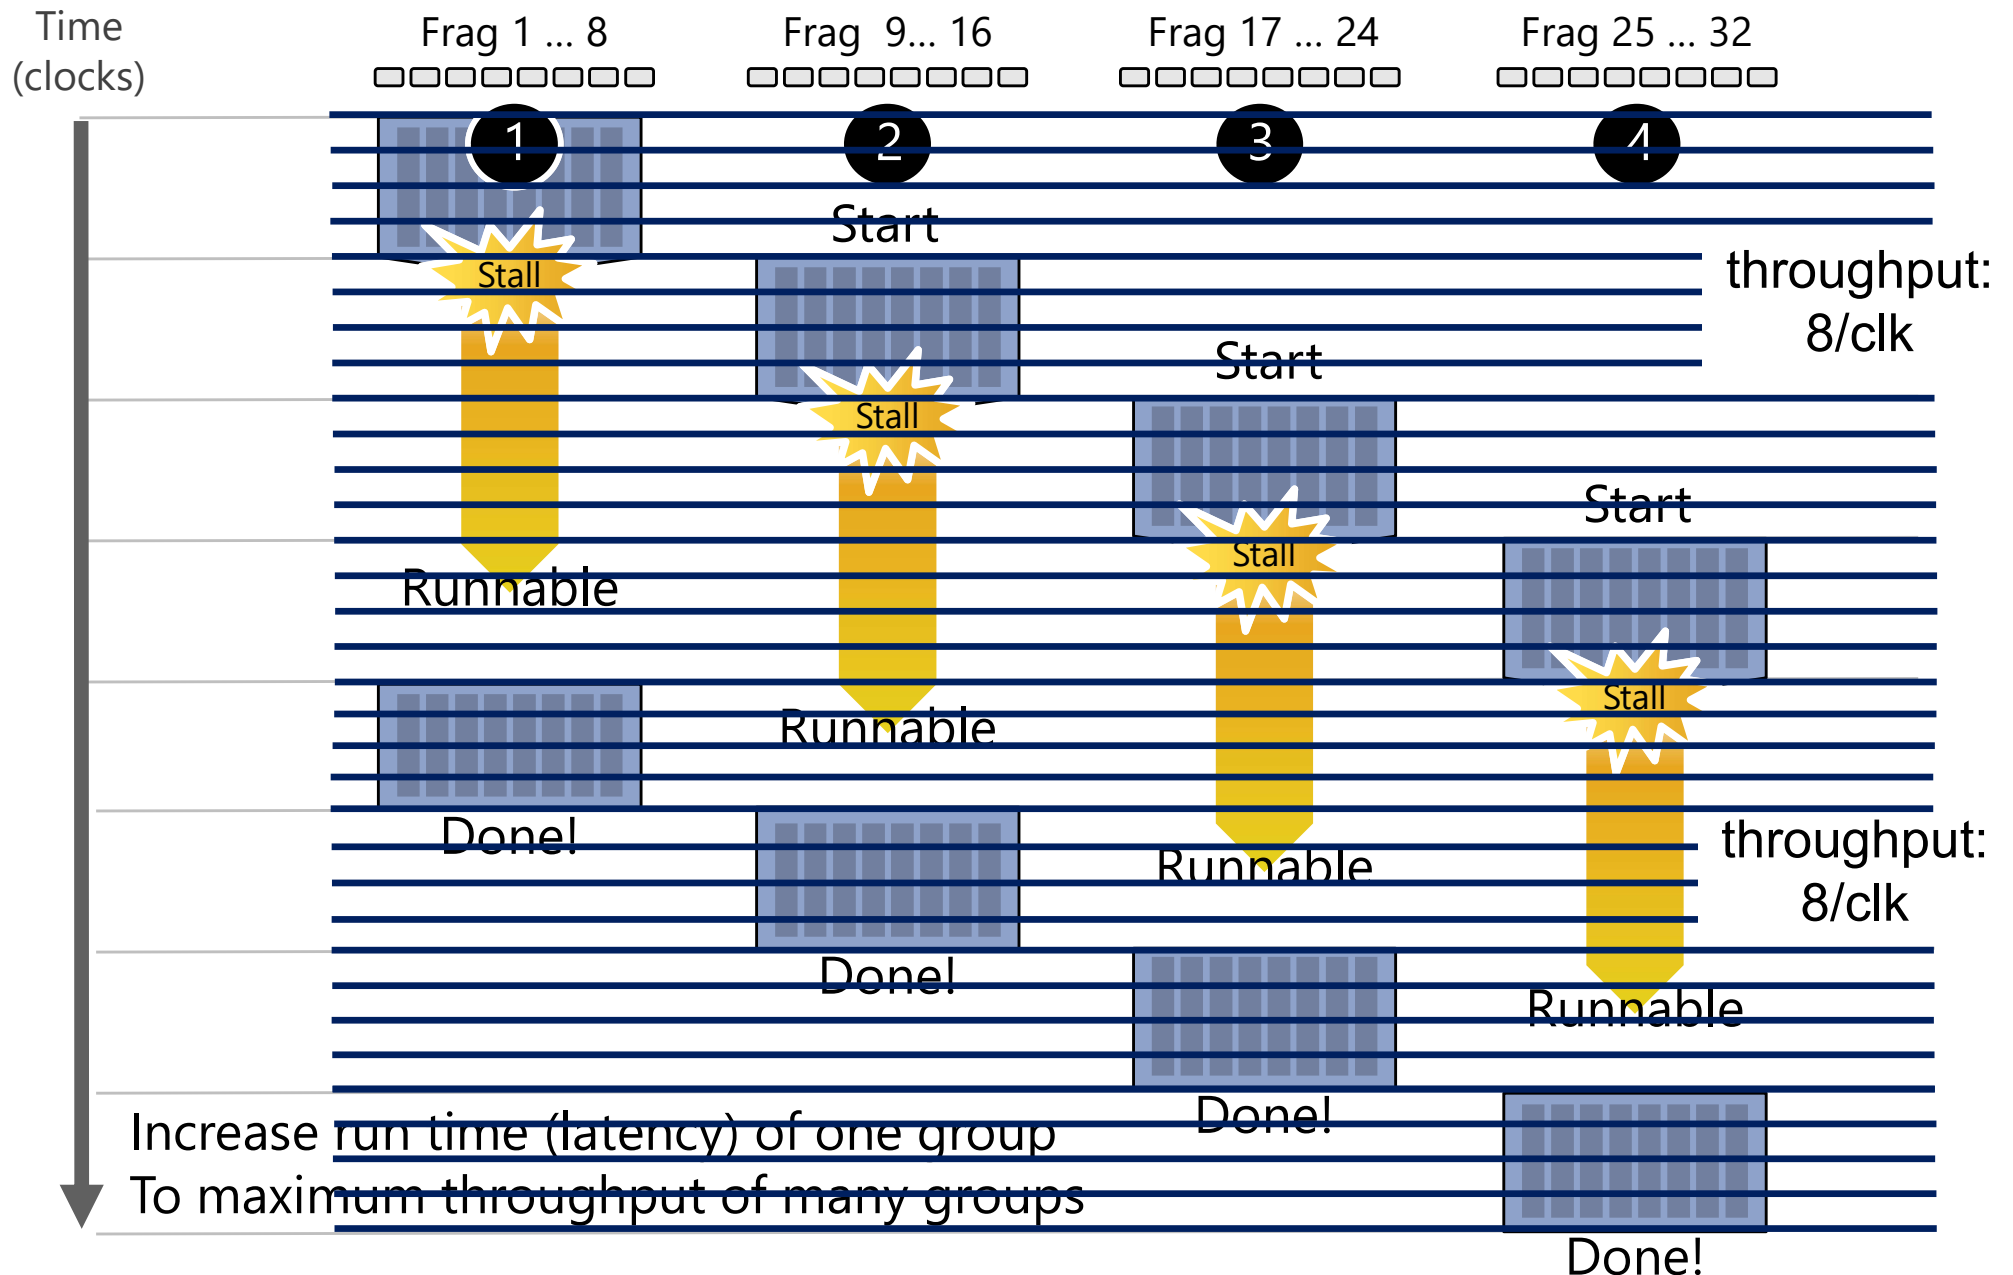
Hiding shader stalls




Throughput! (4 groups of threads)



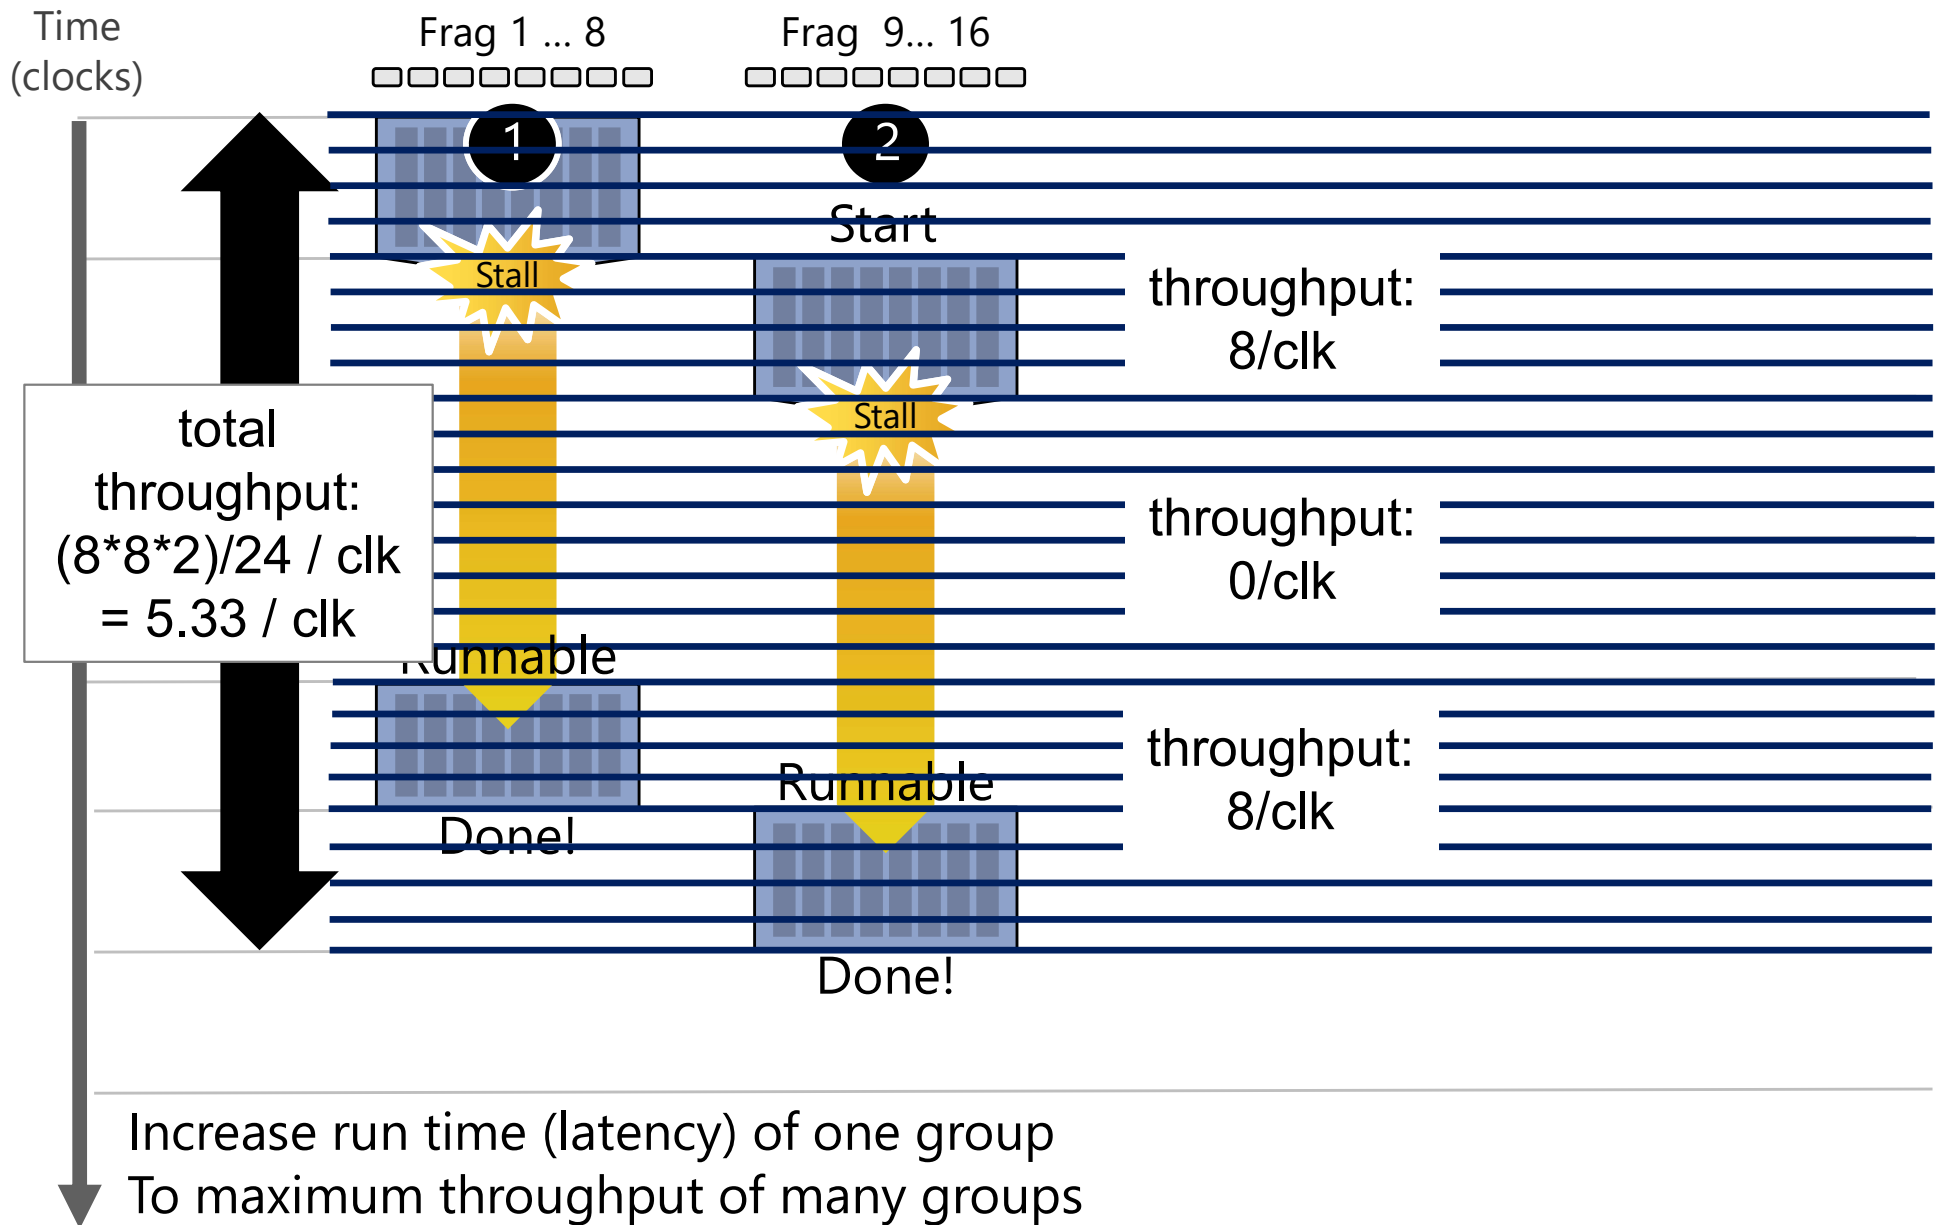
Throughput! (4 groups of threads)




Throughput! (4 groups of threads)



Throughput! (2 groups of threads)




Idea #3: Store multiple group contexts




Twenty small contexts (few regs/thread)

(maximal latency hiding ability)




Twelve medium contexts (more regs/th.)



Four large contexts (many regs/thread)

(low latency hiding ability)



Complete GPU

16 cores


8 mul-add _[mad] ALUs per core
(8*16 = **128** total)

16 simultaneous
instruction streams

64 (4*16) concurrent (but
interleaved) instruction streams

512 (8*4*16) concurrent
fragments (resident threads)

= **256 GFLOPs** (@ 1GHz)
(**128** * 2 _[mad] * 1G)



Complete GPU

16 cores


8 mul-add _[mad] ALUs per core
(8*16 = **128** total)

16 simultaneous
instruction streams

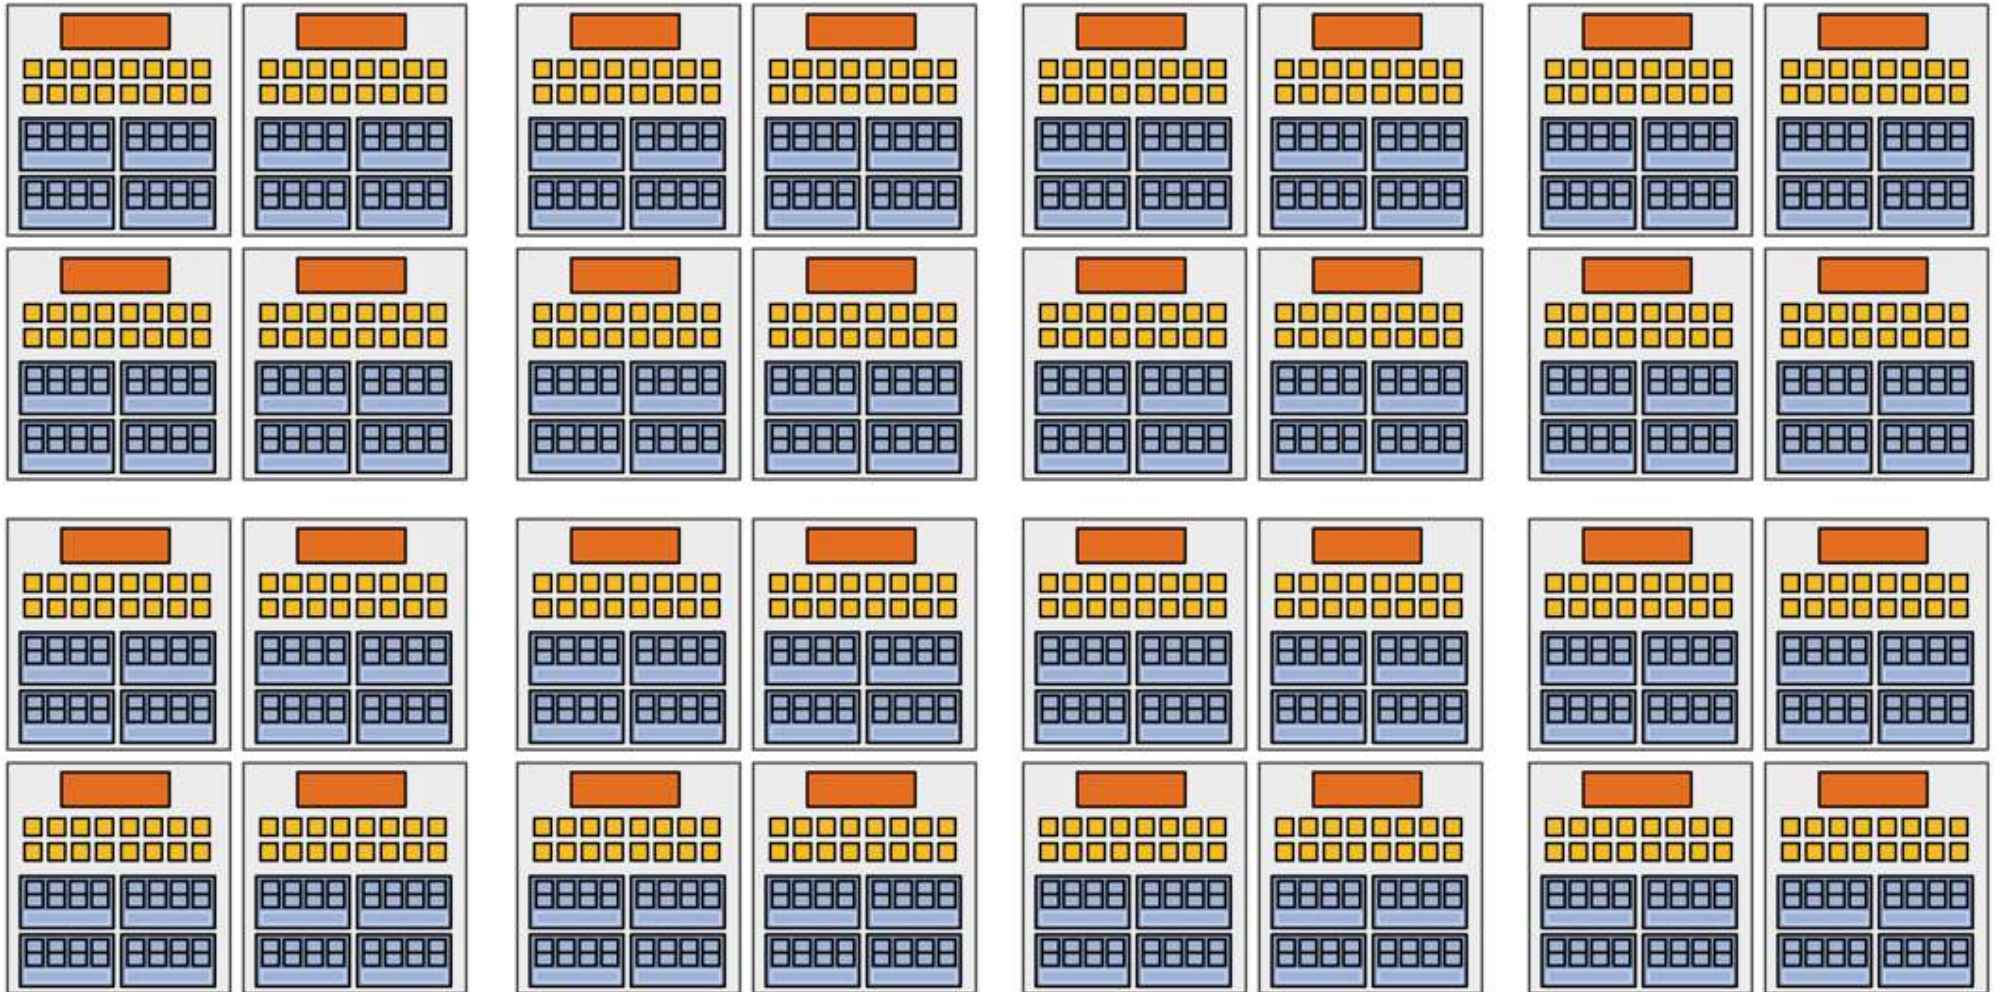
64 (4*16) concurrent (but
interleaved) instruction streams

512 (8*4*16) concurrent
fragments (resident threads)

= **256 GFLOPs** (@ 1GHz)
(**128** * 2 _[mad] * 1G)



"Enthusiast" GPU (Some time ago :)



32 cores, 16 ALUs per core (512 total) = 1 TFLOP (@ 1 GHz)

Where We've Arrived...



Summary: three key ideas for high-throughput execution

1. Use many “slimmed down cores,” run them in parallel
2. Pack cores full of ALUs (by sharing instruction stream overhead across groups of fragments)
 - Option 1: Explicit SIMD vector instructions
 - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
 - When one group stalls, work on another group

GPUs are here!
(usually)

GPU Architecture: Real Architectures

NVIDIA Architectures (since first CUDA GPU)



Tesla [CC 1.x]: 2007-2009

- G80, G9x: 2007 (Geforce 8800, ...)
GT200: 2008/2009 (GTX 280, ...)

Fermi [CC 2.x]: 2010 (2011, 2012, 2013, ...)

- GF100, ... (GTX 480, ...)
GF104, ... (GTX 460, ...)
GF110, ... (GTX 580, ...)

Kepler [CC 3.x]: 2012 (2013, 2014, 2016, ...)

- GK104, ... (GTX 680, ...)
GK110, ... (GTX 780, GTX Titan, ...)

Maxwell [CC 5.x]: 2015

- GM107, ... (GTX 750Ti, ...)
GM204, ... (GTX 980, Titan X, ...)

Pascal [CC 6.x]: 2016 (2017, 2018, 2021, 2022, ...)

- GP100 (Tesla P100, ...)
- GP10x: x=2,4,6,7,8, ...
(GTX 1060, 1070, 1080, Titan X *Pascal*, Titan Xp, ...)

Volta [CC 7.0, 7.2]: 2017/2018

- GV100, ...
(Tesla V100, Titan V, Quadro GV100, ...)

Turing [CC 7.5]: 2018/2019

- TU102, TU104, TU106, TU116, TU117, ...
(Titan RTX, RTX 2070, 2080 (Ti), GTX 1650, 1660, ...)

Ampere [CC 8.0, 8.6, 8.7]: 2020

- GA100, GA102, GA104, GA106, ...
(A100, RTX 3070, 3080, 3090 (Ti), RTX A6000, ...)

Hopper [CC 9.0], Ada Lovelace [CC 8.9]: 2022/23

- GH100, AD102, AD103, AD104, ...
(H100, L40, RTX 4080 (12/16 GB), 4090, RTX 6000, ...)

Blackwell [CC 10.0]: *coming in 2024/25*

- GB200/GB202, GB20x, ...?
(RTX 5080/5090, GB200 NVL72, HGX B100/200, ...?)

Interlude: PTX vs. SASS Code (1)



PTX is virtual machine ISA

SASS is actual machine ISA

For disassembly:

cuobjdump / nvdiasm

See [CUDA_Binary_Uutilities.pdf](#)

For debugging (and code inspection) see:

<https://developer.nvidia.com/nsight-visual-studio-edition>

```
Address: _Z9matrixMulPfS_S_ii

0x00002ed0      MOV R1, R1;
72:
73:      // Index of the last sub-matrix of A processed by the block
74:      int aEnd  = aBegin + wA - 1;
0x00002ed8 [0083] ld.param.s32    %r14, [__cudaparm__Z9matrixMulPfS_S_ii_wA];
0x00002ed8      MVI R0, 0x1c;
0x00002ee0      R2A A1, R0;
0x00002ee8      MOV R0, g [A1+0x0];
0x00002ef0 [0084] mov.s32    %r15, %r13;
0x00002ef0      MOV32 R1, R1;
0x00002ef4 [0085] add.s32    %r16, %r14, %r15;
0x00002ef4      IADD32 R0, R0, R1;
0x00002ef8 [0086] sub.s32    %r17, %r16, 1;
0x00002ef8      IADD32I R8, R0, 0xffffffff;
0x00002f00 [0087] mov.s32    %r18, %r17;
0x00002f00      MOV R8, R8;
75:
76:      // Step size used to iterate through the sub-matrices of A
77:      int aStep = BLOCK_SIZE;
0x00002f08 [0089] mov.s32    %r19, 16;
0x00002f08      MVI R9, 0x10;
0x00002f10 [0090] mov.s32    %r20, %r19;
0x00002f10      MOV32 R9, R9;
78:
79:      // Index of the first sub-matrix of B processed by the block
80:      int bBegin = BLOCK_SIZE * bx;
0x00002f14 [0092] mov.s32    %r21, %r2;
0x00002f14      MOV32 R4, R4;
0x00002f18 [0093] mul.lo.s32    %r22, %r21, 16;
0x00002f18      IMUL.U16.U16 R0, R4L, R31H;
0x00002f20      IMAD32I.U16 R0, R4H, 0x10, R0;
0x00002f28      SHL R2, R0, 0x10;
0x00002f30      IMAD32I.U16 R2, R4L, 0x10, R2;
0x00002f38 [0094] mov.s32    %r23, %r22;
0x00002f38      MOV R2, R2;
```


Interlude: PTX vs. SASS Code (2)



Note

- Size of instructions (here: 16 bytes)
- **MUFU.RCP** computing FP32 reciprocal on SFU (there is no SASS division: division is an algorithm comprising simpler instructions)
- This is debug code: redundant register moves not (yet) removed by optimizer in assembler (*result of virtual PTX registers being mapped to same physical register*)
- ...

(SASS on Ampere)

```
Disassembly  x bicubicTexture_kernel.cuh  bicubicTexture_cuda.cu  binomialOptions_kernel.cu  simpleCUFFT.cu  oceanFFT_kernel.cu  matrixMulC

Address: h0

Viewing Options
--- D:/development/CUDA_Samples/git_work/cuda-samples/Samples/5_Domain_Specific/bicubicTexture/bicubicTexture_kernel.cuh
__device__ float h0(float a) {
0x00000004300bbfe00          IADD3 R1, R1, -0x18, RZ
0x00000004300bbfe10          S2R R0, SR_LMEMHIOFF
0x00000004300bbfe20          ISETP.GE.U32.AND P0, PT, R1, R0, PT
0x00000004300bbfe30          @P0 BRA 0x4300bbfe50
0x00000004300bbfe40          BPT.TRAP 0x1
0x00000004300bbfe50          STL [R1+0x14], R21
0x00000004300bbfe60          STL [R1+0x10], R20
0x00000004300bbfe70          STL [R1+0xc], R18
0x00000004300bbfe80          STL [R1+0x8], R17
0x00000004300bbfe90          STL [R1+0x4], R16
0x00000004300bbfea0          STL [R1], R2
0x00000004300bbfeb0          BMOV.32.CLEAR R18, B6
0x00000004300bbfec0          MOV R4, R4
0x00000004300bbfed0          MOV R4, R4
0x00000004300bbfee0          MOV R17, R4
0x00000004300bbfef0          return -1.0f + w1(a) / (w0(a) + w1(a)) + 0.5f;
0x00000004300bbff00          MOV R4, R17
0x00000004300bbff10          MOV R20, 0x0
0x00000004300bbff20          MOV R21, 0x0
0x00000004300bbff30          CALL.ABS.NOINC 0x0
0x00000004300bbff40          MOV R0, R4
0x00000004300bbff50          MOV R4, R17
0x00000004300bbff60          MOV R16, R0
0x00000004300bbff70          MOV R20, 0x0
0x00000004300bbff80          MOV R21, 0x0
0x00000004300bbff90          CALL.ABS.NOINC 0x0
0x00000004300bbffa0          MOV R0, R4
0x00000004300bbffb0          MOV R4, R17
0x00000004300bbffc0          MOV R2, R0
0x00000004300bbffd0          MOV R20, 0x0
0x00000004300bbffe0          MOV R21, 0x0
0x00000004300bbfff0          CALL.ABS.NOINC 0x0
0x00000004300bc0000          MOV R4, R4
0x00000004300bc0010          FADD R4, R2, R4
0x00000004300bc0020          MOV R0, R16
0x00000004300bc0030          MOV R4, R4
0x00000004300bc0040          MOV R0, R0
0x00000004300bc0050          MOV R4, R4
0x00000004300bc0060          MOV R3, R4
0x00000004300bc0070          MUFU.RCP R5, R3
0x00000004300bc0080          FADD R3, -RZ, -R3
0x00000004300bc0090          MOV R3, R3
0x00000004300bc00a0          MOV R6, 0x3f800000
0x00000004300bc00b0          FFMA R6, R3, R5, R6
0x00000004300bc00c0          FCHK P0, R0, R4
0x00000004300bc00d0          FFMA R6, R5, R6, R5
0x00000004300bc00e0          MOV R5, RZ
0x00000004300bc00f0          MOV R0, R0

71 // h0 and h1 are the two offset functions
72 __device__ float h0(float a) {
73 // note +0.5 offset to compensate for CUDA linear filtering convention
74 return -1.0f + w1(a) / (w0(a) + w1(a)) + 0.5f;
75 }
76
77 __device__ float h1(float a) { return 1.0f + w3(a) / (w2(a) + w3(a)) + 0.5f; }
78
79
```

Interlude: PTX vs. SASS Code (2)



Note

- Size of instructions (here: 16 bytes)
- **MUFU.RCP** computing FP32 reciprocal on SFU (there is no SASS division: division is an algorithm comprising simpler instructions)
- This is debug code: redundant register moves not (yet) removed by optimizer in assembler (*result of virtual PTX registers being mapped to same physical register*)
- ...

(SASS on Ampere)

```
Disassembly  x bicubicTexture_kernel.cuh  bicubicTexture_cuda.cu  binomialOptions_kernel.cu  simpleCUFFT.cu  oceanFFT_kernel.cu  matrixMulC

Address: h0

Viewing Options
--- D:/development/CUDA_Samples/git_work/cuda-samples/Samples/5_Domain_Specific/bicubicTexture/bicubicTexture_kernel.cuh
__device__ float h0(float a) {
0x00000004300bbfe00          IADD3 R1, R1, -0x18, RZ
0x00000004300bbfe10          S2R R0, SR_LMEMHIOFF
0x00000004300bbfe20          ISETP.GE.U32.AND P0, PT, R1, R0, PT
0x00000004300bbfe30          @P0 BRA 0x4300bbfe50
0x00000004300bbfe40          BPT.TRAP 0x1
0x00000004300bbfe50          STL [R1+0x14], R21
0x00000004300bbfe60          STL [R1+0x10], R20
0x00000004300bbfe70          STL [R1+0xc], R18
0x00000004300bbfe80          STL [R1+0x8], R17
0x00000004300bbfe90          STL [R1+0x4], R16
0x00000004300bbfea0          STL [R1], R2
0x00000004300bbfeb0          BMOV.32.CLEAR R18, B6
0x00000004300bbfec0          MOV R4, R4
0x00000004300bbfed0          MOV R4, R4
0x00000004300bbfee0          MOV R17, R4
0x00000004300bbfef0          MOV R4, R17
0x00000004300bbff00          MOV R20, 0x0
0x00000004300bbff10          MOV R21, 0x0
0x00000004300bbff20          CALL.ABS.NOINC 0x0
0x00000004300bbff30          MOV R0, R4
0x00000004300bbff40          MOV R4, R17
0x00000004300bbff50          MOV R16, R0
0x00000004300bbff60          MOV R20, 0x0
0x00000004300bbff70          MOV R21, 0x0
0x00000004300bbff80          CALL.ABS.NOINC 0x0
0x00000004300bbff90          MOV R0, R4
0x00000004300bbffa0          MOV R4, R17
0x00000004300bbffb0          MOV R2, R0
0x00000004300bbffc0          MOV R20, 0x0
0x00000004300bbffd0          MOV R21, 0x0
0x00000004300bbffe0          CALL.ABS.NOINC 0x0
0x00000004300bbfff0          MOV R4, R4
0x00000004300bc0000          FADD R4, R2, R4
0x00000004300bc0010          MOV R0, R16
0x00000004300bc0020          MOV R4, R4
0x00000004300bc0030          MOV R0, R0
0x00000004300bc0040          MOV R4, R4
0x00000004300bc0050          MOV R3, R4
0x00000004300bc0060          MUFU.RCP R5, R3
0x00000004300bc0070          FADD R3, -RZ, -R3
0x00000004300bc0080          MOV R3, R3
0x00000004300bc0090          MOV R6, 0x3f800000
0x00000004300bc00a0          FFMA R6, R3, R5, R6
0x00000004300bc00b0          FCHK P0, R0, R4
0x00000004300bc00c0          FFMA R6, R5, R6, R5
0x00000004300bc00d0          MOV R5, RZ
0x00000004300bc00e0          MOV R0, R0
}

// h0 and h1 are the two offset functions
__device__ float h0(float a) {
// note +0.5 offset to compensate for CUDA linear filtering convention
return -1.0f + w1(a) / (w0(a) + w1(a)) + 0.5f;
}

__device__ float h1(float a) { return 1.0f + w3(a) / (w2(a) + w3(a)) + 0.5f; }
```

- Size of instructions
(here: 16 bytes)
- **MUFU.RCP** computing
FP32 reciprocal on SFU
(there is no SASS division:
division is an algorithm
comprising simpler instructions)
- This is debug code:
redundant register moves
not (yet) removed by
optimizer in assembler
*(result of virtual PTX
registers being mapped
to same physical register)*
-

[illegible]

Example: “Scalar” GF100

Main concept here:


There is one instruction dispatcher
(dispatch unit / fetch/decode unit)
per warp scheduler
(warp selector)

Details later...

Ignore less important subtleties...


GF100 has two warp schedulers, not one,
and each 32-thread instruction is executed
over two clock cycles, not one, etc.


Caveat on NVIDIA diagrams: if two dispatchers per warp scheduler are shown, it still doesn't mean that the ALU pipeline is “superscalar” (often, the second dispatcher dispatches to a *non-ALU* pipeline)
... need to look at CUDA programming guide info, also given
in our tables in row “# ALU dispatch / warp sched.”



Example: “Superscalar” ALUs in SM Architecture

NVIDIA Kepler GK104 architecture SMX unit (one “core”)



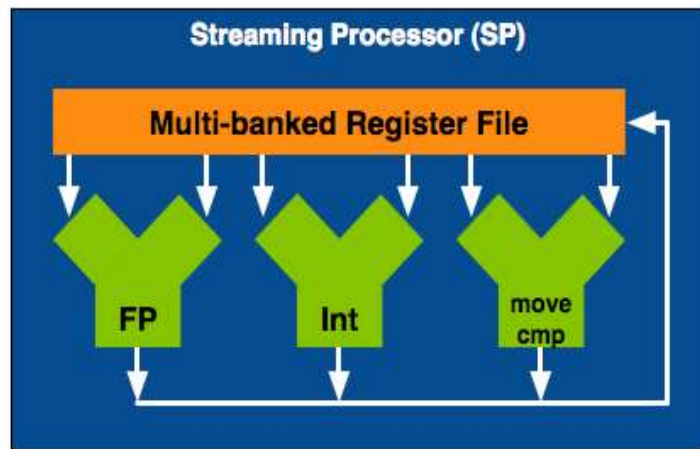


 = SIMD function unit,
control shared across 32 units
(1 MUL-ADD per clock)

 = "special" SIMD function unit,
control shared across 32 units
(operations like sin/cos)

 = SIMD load/store unit
(handles warp loads/stores, gathers/scatters)

NVIDIA Tesla Architecture (not the Tesla product line!), G80: 2007, GT200: 2008/2009



Courtesy AnandTech


G80: first CUDA GPU!
Multiprocessor: SM (CC 1.x)

- Streaming Processor (SP) [or: CUDA core; or: FP32 / FP64 / INT32 core, ...]
- Streaming Multiprocessor (SM)
- Texture/Processing Cluster (TPC)


NVIDIA Tesla Architecture (not the Tesla product line!), G80: 2007, GT200: 2008/2009



- G80/G92: $8 \text{ TPCs} * (2 * 8 \text{ SPs}) = 128 \text{ SPs}$ [= CUDA cores]
- GT200: $10 \text{ TPCs} * (3 * 8 \text{ SPs}) = 240 \text{ SPs}$ [= CUDA cores]
- **Arithmetic intensity** has increased (num. of ALUs vs. texture units)



G80 / G92



GT200

Courtesy AnandTech

NVIDIA Ampere GA100 Architecture (2020)



GA 100 (A100 Tensor Core GPU)

Full GPU: 128 SMs (in 8 GPCs/64 TPCs)



NVIDIA Hopper GH100 Architecture (2022)



GH 100 (H100 Tensor Core GPU)

Full GPU: 144 SMs (in 8 GPCs/72 TPCs)

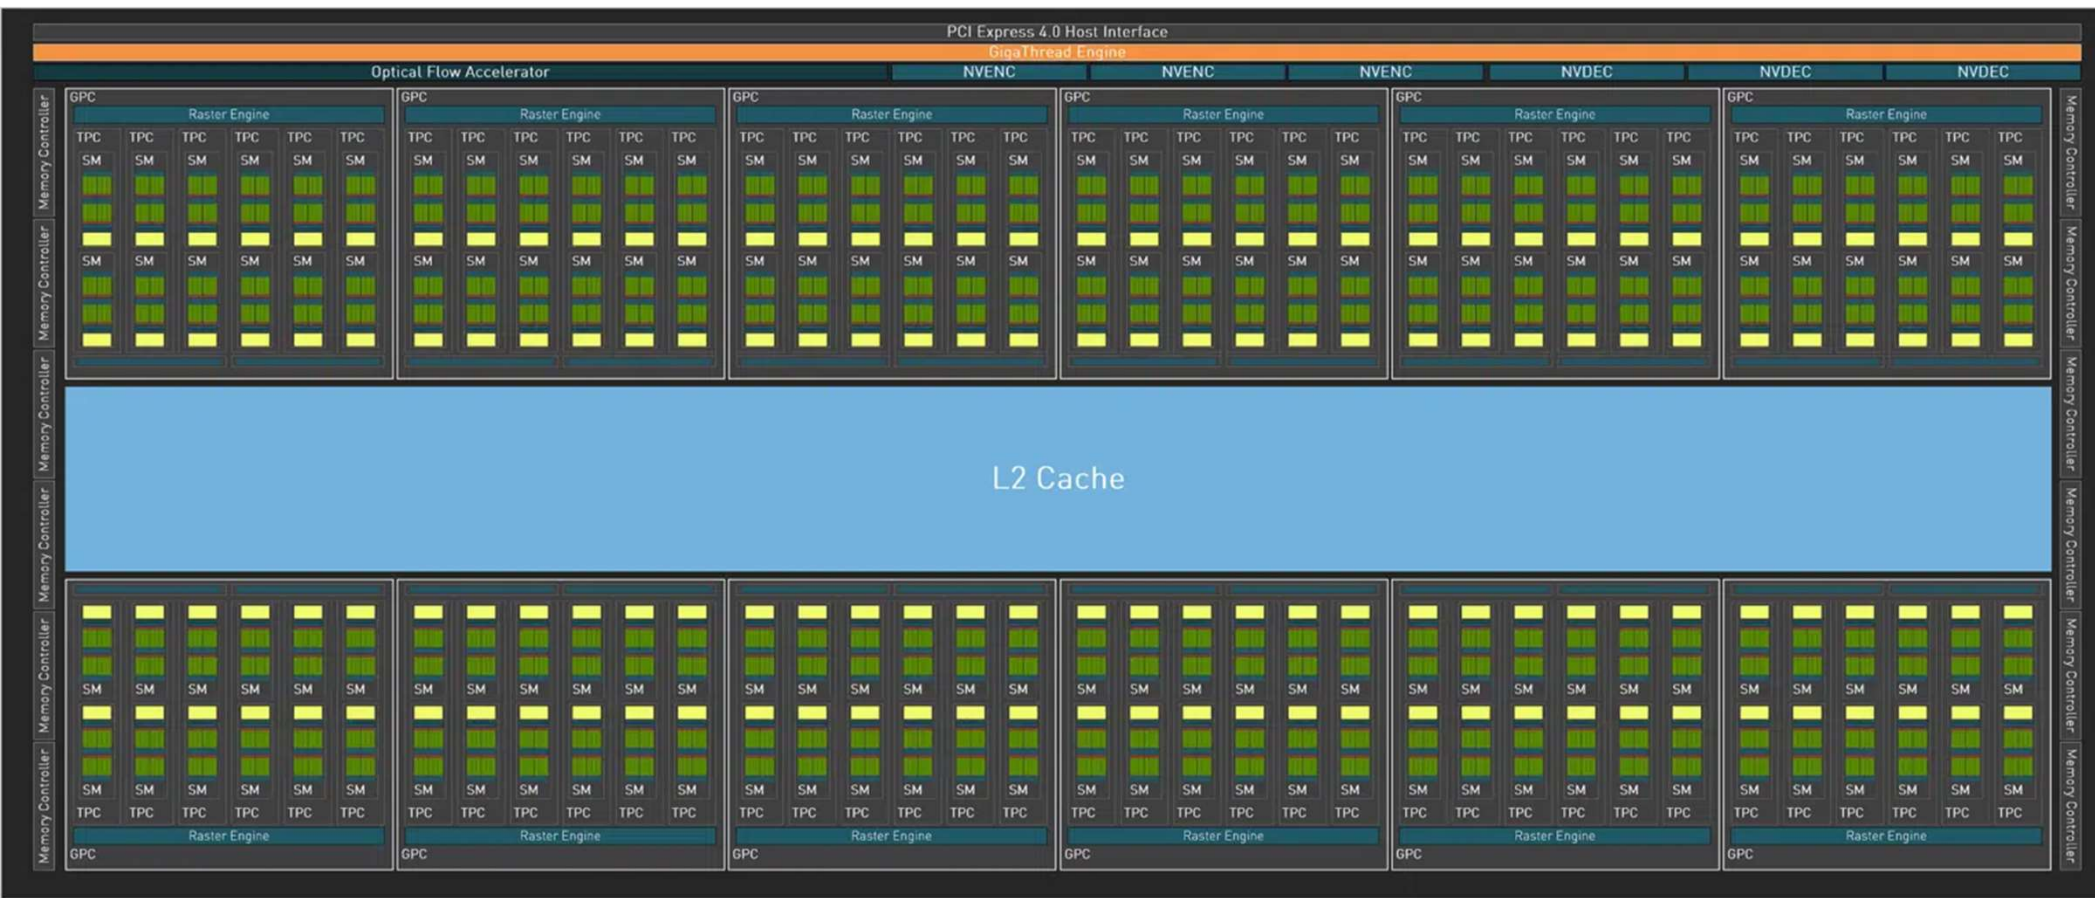


NVIDIA Ada Lovelace AD10x Architecture (2022)



Full AD 10x

Full GPU: 144 SMs (in 12 GPCs/72 TPCs)



NVIDIA Ada Lovelace AD102 Architecture (2022)



AD 102 (RTX 4090, ...)

Full RTX 4090: 128 SMs (in 11 GPCs/64 TPCs)



Thank you.