

CS 380 - GPU and GPGPU Programming

Lecture 26: GPU Texturing, Pt. 3

Markus Hadwiger, KAUST

Reading Assignment #10 (until Nov 11)



Read (required):

- Interpolation for Polygon Texture Mapping and Shading,
Paul Heckbert and Henry Moreton

<https://www.rh.cmu.edu/publications/interpolation-for-polygon-texture-mapping-and-shading/>

- Homogeneous Coordinates

https://en.wikipedia.org/wiki/Homogeneous_coordinates

Read (optional; highly recommended!):

- MIP-Map Level Selection for Texture Mapping

<https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=765326>

Next Lectures



Lecture 27: Thu, Nov 7: Vulkan tutorial #2


Lecture 28: Mon, Nov 11: 10:00-11:30 (on Zoom)

Lecture 29: Thu, Nov 14: 10:00-11:30 (on Zoom)

Lecture 30: Mon, Nov 18: Quiz #3

GPU Texturing

Texturing: General Approach



Texture space (u, v)




Object space (x_O, y_O, z_O)



Image Space (x_I, y_I)

Parametrization

Rendering
(Projection etc.)



Texture Mapping

2D (3D) Texture Space

| Texture Transformation

2D Object Parameters

| Parameterization

3D Object Space

| Model Transformation


3D World Space

| Viewing Transformation

3D Camera Space

| Projection

2D Image Space



Texture Mapping Polygons


Forward transformation: linear projective map

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} s \\ t \\ r \end{bmatrix}$$



Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Perspective Texture Mapping



linear interpolation
in object space $\frac{ax_1 + bx_2}{aw_1 + bw_2} \neq a \frac{x_1}{w_1} + b \frac{x_2}{w_2}$ linear interpolation
in screen space



$$a = b = 0.5$$



Perspective-correct linear interpolation

Only projected values interpolate correctly, so project A

- Linearly interpolate A_1/w_1 and A_2/w_2

Also interpolate $1/w_1$ and $1/w_2$

- These also interpolate linearly in screen space

Divide interpolants at each sample point to recover A

- $(A/w) / (1/w) = A$
- Division is expensive (more than add or multiply), so
 - Recover w for the sample point (reciprocate), and
 - Multiply each projected attribute by w

Barycentric triangle parameterization:

$$A = \frac{aA_1/w_1 + bA_2/w_2 + cA_3/w_3}{a/w_1 + b/w_2 + c/w_3} \quad a + b + c = 1$$


Perspective-Correct Interpolation Recipe





$$r_i(x, y) = \frac{r_i(x, y) / w(x, y)}{1 / w(x, y)}$$

- (1) Associate a record containing the n parameters of interest (r_1, r_2, \dots, r_n) with each vertex of the polygon.
- (2) For each vertex, transform object space coordinates to homogeneous screen space using 4×4 object to screen matrix, yielding the values (xw, yw, zw, w) .
- (3) Clip the polygon against plane equations for each of the six sides of the viewing frustum, linearly interpolating all the parameters when new vertices are created.
- (4) At each vertex, divide the homogeneous screen coordinates, the parameters r_i , and the number 1 by w to construct the variable list $(x, y, z, s_1, s_2, \dots, s_{n+1})$, where $s_i = r_i / w$ for $i \leq n$, $s_{n+1} = 1 / w$.
- (5) Scan convert in screen space by linear interpolation of all parameters, at each pixel computing $r_i = s_i / s_{n+1}$ for each of the n parameters; use these values for shading.

- Want to simulate a beamer
 - ... or a flashlight, or a slide projector
- Precursor to shadows
- Interesting mathematics:
2 perspective
projections involved!
- Easy to program!






Projective Shadows in Doom 3



- What about **homogeneous** texture coords?
- Need to do perspective divide also for projector!
 - $(s, t, q) \rightarrow (s/q, t/q)$ for every fragment
- How does OpenGL do that?
 - Needs to be perspective correct as well!
 - Trick: interpolate $(s/w, t/w, r/w, q/w)$
 - $(s/w) / (q/w) = s/q$ etc. at every fragment
- Remember: s, t, r, q are equivalent to x, y, z, w in projector space! $\rightarrow r/q = \text{projector depth!}$




- Apply multiple textures in one pass
- *Integral* part of programmable shading
 - e.g. diffuse texture map + gloss map
 - e.g. diffuse texture map + light map
- Performance issues
 - How many textures are free?
 - How many are available



- Used in virtually every commercial game
- Precalculate diffuse lighting on static objects
 - Only low resolution necessary
 - Diffuse lighting is view independent!
- Advantages:
 - No runtime lighting necessary
 - VERY fast!
 - Can take global effects (shadows, color bleeds) into account






Original LM texels




Bilinear Filtering


■ Why premultiplication is bad...



Full Size Texture
(with Lightmap)



Tiled Surface Texture
plus Lightmap



→ use tileable surface textures and low resolution lightmaps






Original scene




Light-mapped

- Precomputation based on non-realtime methods
 - Radiosity
 - Ray tracing
 - Monte Carlo Integration
 - Path tracing
 - Photon mapping






Lightmap



mapped



Original scene



Light-mapped

Interpolation #2



Interpolation Type + Purpose #2:

Interpolation of Samples in Texture Space

(Multi-Linear Interpolation)

- Spatial layout
 - Cartesian grids: 1D, 2D, 3D, 2D_ARRAY, ...
 - Cube maps, ...
- Formats (too many), e.g. OpenGL
 - GL_LUMINANCE16_ALPHA16
 - GL_RGB8, GL_RGBA8, ...: integer texture formats
 - GL_RGB16F, GL_RGBA32F, ...: float texture formats
 - compressed formats, high dynamic range formats, ...
- External (CPU) format vs. internal (GPU) format
 - OpenGL driver converts from external to internal

for Vulkan, see `vkImageView`

for Vulkan, see `vkImage`
and `vkImageView`

use `VK_IMAGE_TILING_OPTIMAL`
for `VkImageCreateInfo::tiling`



Magnification (Bi-linear Filtering Example)



Original image




Nearest neighbor




Bi-linear filtering

Nearest-Neighbor vs. Bi-Linear Interpolation




nearest-neighbor



bi-linear

wikipedia





Bilinear patch (courtesy J. Han)

Bi-Linear Interpolation



Consider area between 2x2 adjacent samples (e.g., pixel centers)

Example #2: 1 at top-left and bottom-right, 0 at bottom-left, 0.5 at top-right



Bi-Linear Interpolation



Consider area between 2x2 adjacent samples (e.g., pixel centers):

Given any (fractional) position


$$\alpha_1 := x_1 - \lfloor x_1 \rfloor \quad \alpha_1 \in [0.0, 1.0)$$

$$\alpha_2 := x_2 - \lfloor x_2 \rfloor \quad \alpha_2 \in [0.0, 1.0)$$

and 2x2 sample values

$$\begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix}$$

Compute: $f(\alpha_1, \alpha_2)$



Bi-Linear Interpolation



Consider area between 2x2 adjacent samples (e.g., pixel centers):

Given any (fractional) position


$$\alpha_1 := x_1 - \lfloor x_1 \rfloor \quad \alpha_1 \in [0.0, 1.0)$$

$$\alpha_2 := x_2 - \lfloor x_2 \rfloor \quad \alpha_2 \in [0.0, 1.0)$$

and 2x2 sample values

$$\begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix}$$

Compute: $f(\alpha_1, \alpha_2)$



Bi-Linear Interpolation



Weights in 2x2 format:

$$\begin{bmatrix} \alpha_2 \\ (1 - \alpha_2) \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) & \alpha_1 \end{bmatrix} = \begin{bmatrix} (1 - \alpha_1)\alpha_2 & \alpha_1\alpha_2 \\ (1 - \alpha_1)(1 - \alpha_2) & \alpha_1(1 - \alpha_2) \end{bmatrix}$$

Interpolate function at (fractional) position (α_1, α_2) :

$$f(\alpha_1, \alpha_2) = \begin{bmatrix} \alpha_2 & (1 - \alpha_2) \end{bmatrix} \begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix}$$

Bi-Linear Interpolation



Interpolate function at (fractional) position (α_1, α_2) :

$$\begin{aligned} f(\alpha_1, \alpha_2) &= \begin{bmatrix} \alpha_2 & (1 - \alpha_2) \end{bmatrix} \begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix} \\ &= \begin{bmatrix} \alpha_2 & (1 - \alpha_2) \end{bmatrix} \begin{bmatrix} (1 - \alpha_1)v_{01} + \alpha_1 v_{11} \\ (1 - \alpha_1)v_{00} + \alpha_1 v_{10} \end{bmatrix} \\ &= \begin{bmatrix} \alpha_2 v_{01} + (1 - \alpha_2)v_{00} & \alpha_2 v_{11} + (1 - \alpha_2)v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix} \end{aligned}$$

Bi-Linear Interpolation



Interpolate function at (fractional) position (α_1, α_2) :

$$f(\alpha_1, \alpha_2) = \begin{bmatrix} \alpha_2 & (1 - \alpha_2) \end{bmatrix} \begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix}$$

$$= (1 - \alpha_1)(1 - \alpha_2)v_{00} + \alpha_1(1 - \alpha_2)v_{10} + (1 - \alpha_1)\alpha_2v_{01} + \alpha_1\alpha_2v_{11}$$

$$= v_{00} + \alpha_1(v_{10} - v_{00}) + \alpha_2(v_{01} - v_{00}) + \alpha_1\alpha_2(v_{00} + v_{11} - v_{10} - v_{01})$$






REALLY IMPORTANT:


this is a different thing (for a different purpose)
than the linear (or, in perspective, rational-linear)
interpolation of texture coordinates!!

Texture Minification

- Problem: One pixel in image space covers many texels



- Caused by *undersampling*: texture information is lost



Texture space


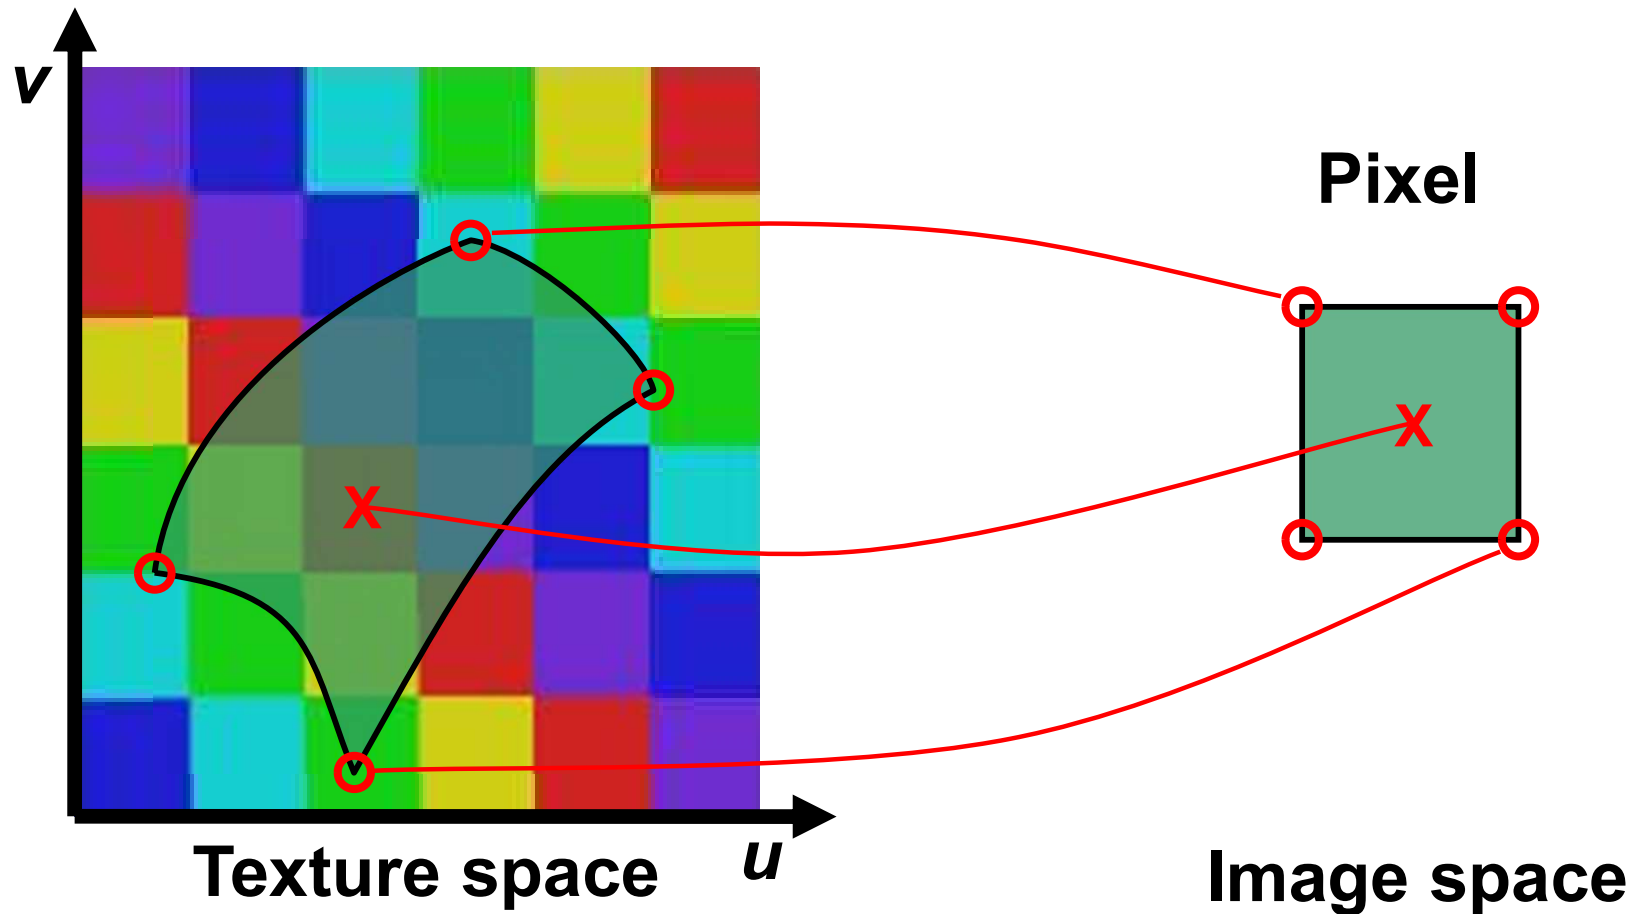




Image space

- A good pixel value is the weighted mean of the pixel area projected into texture space




- MIP Mapping (“Multum In Parvo”)
 - Texture size is reduced by factors of 2 (*downsampling* = “many things in a small place”)
 - Simple (4 pixel average) and memory efficient
 - Last image is only ONE texel



- MIP Mapping (“Multum In Parvo”)
 - Texture size is reduced by factors of 2 (*downsampling* = “many things in a small place”)
 - Simple (4 pixel average) and memory efficient
 - Last image is only ONE texel




geometric series:

$$\begin{aligned} a + ar + ar^2 + ar^3 + \dots + ar^{n-1} &= \\ &= \sum_{k=0}^{n-1} ar^k = a \left(\frac{1 - r^n}{1 - r} \right) \end{aligned}$$





Texture Anti-Aliasing: MIP Mapping

- MIP Mapping Algorithm
- $D := \text{ld}(\max(d_1, d_2))$ "Mip Map level"
- $T_0 := \text{value from texture } D_0 = \text{trunc}(D)$
 - Use *bilinear interpolation*



MIP-Map Level Computation



- Use the partial derivatives of texture coordinates with respect to screen space coordinates
- This is the Jacobian matrix
- Area of parallelogram is the absolute value of the Jacobian determinant (the *Jacobian*)

$$\begin{pmatrix} \partial u / \partial x & \partial u / \partial y \\ \partial v / \partial x & \partial v / \partial y \end{pmatrix} = \begin{pmatrix} s_x & s_y \\ t_x & t_y \end{pmatrix}$$

MIP-Map Level Computation (OpenGL)



- OpenGL 4.6 core specification, pp. 251-264

(3D tex coords!)

$$\lambda_{base}(x, y) = \log_2[\rho(x, y)]$$

$$\rho = \max \left\{ \sqrt{\left(\frac{\partial u}{\partial x}\right)^2 + \left(\frac{\partial v}{\partial x}\right)^2 + \left(\frac{\partial w}{\partial x}\right)^2}, \sqrt{\left(\frac{\partial u}{\partial y}\right)^2 + \left(\frac{\partial v}{\partial y}\right)^2 + \left(\frac{\partial w}{\partial y}\right)^2} \right\}$$


Does not use area of parallelogram but greater hypotenuse [Heckbert, 1983]

- Approximation without square-roots

$$m_u = \max \left\{ \left| \frac{\partial u}{\partial x} \right|, \left| \frac{\partial u}{\partial y} \right| \right\} \quad m_v = \max \left\{ \left| \frac{\partial v}{\partial x} \right|, \left| \frac{\partial v}{\partial y} \right| \right\} \quad m_w = \max \left\{ \left| \frac{\partial w}{\partial x} \right|, \left| \frac{\partial w}{\partial y} \right| \right\}$$


$$\max\{m_u, m_v, m_w\} \leq f(x, y) \leq m_u + m_v + m_w$$

MIP-Map Level Interpolation




- Level of detail value is fractional!
- Use fractional part to blend (lin.) between two adjacent mipmap levels



- Trilinear interpolation:
 - $T_1 :=$ value from texture $D_1 = D_0 + 1$ (bilinear interpolation)
 - Pixel value $:= (D_1 - D) \cdot T_0 + (D - D_0) \cdot T_1$
 - Linear interpolation between successive MIP Maps
 - Avoids "Mip banding" (but doubles texture lookups)



- Other example for bilinear vs. trilinear filtering

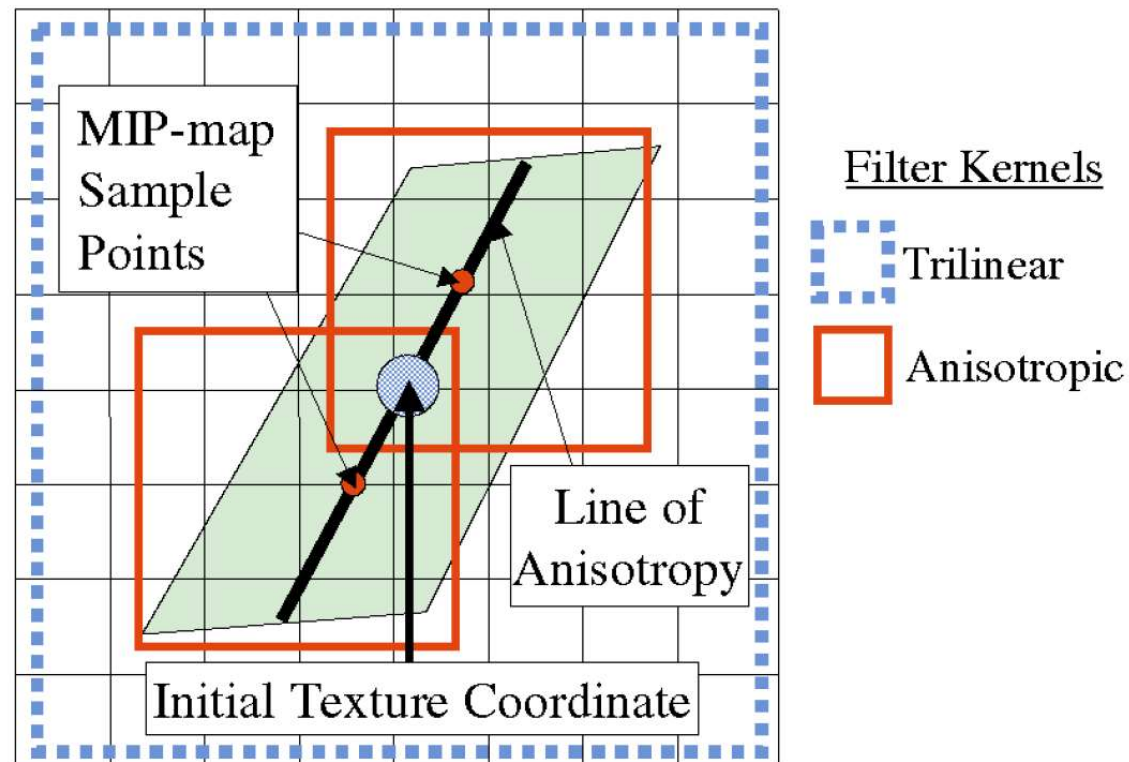


- Anisotropic filtering
 - View-dependent filter kernel
 - Implementation: *summed area table*, "*RIP Mapping*", *footprint assembly*, *elliptical weighted average* (EWA)




Texture space

Anisotropic Filtering: Footprint Assembly



■ Example



- Basically, everything done in hardware
- `gluBuild2DMipmaps()` generates MIPmaps
- Set parameters in `glTexParameter()`
 - `GL_TEXTURE_MAG_FILTER: GL_NEAREST, GL_LINEAR, ...`
 - `GL_TEXTURE_MIN_FILTER: GL_LINEAR_MIPMAP_NEAREST`
- Anisotropic filtering is an extension:
 - `GL_EXT_texture_filter_anisotropic`
 - Number of samples can be varied (4x,8x,16x)
 - Vendor specific support and extensions

for Vulkan, see `vkSampler`,
`VkSamplerCreateInfo::magFilter`, `VkSamplerCreateInfo::minFilter`,
`VK_FILTER_NEAREST`, `VK_FILTER_LINEAR`,
`VkSamplerCreateInfo::mipmapMode`,
`VK_SAMPLER_MIPMAP_MODE_NEAREST`, `VK_SAMPLER_MIPMAP_MODE_LINEAR`, ...



Thank you.