

# CS 380 - GPU and GPGPU Programming

## Lecture 22: GPU Texturing, Pt. 4

Markus Hadwiger, KAUST

# Reading Assignment #12 (until Nov 24)



## Read (required):

- Look at Vulkan *sparse resources*, especially *sparse partially-resident images*
  - <https://docs.vulkan.org/spec/latest/chapters/sparsemem.html>
- Read about shadow mapping
  - [https://en.wikipedia.org/wiki/Shadow\\_mapping](https://en.wikipedia.org/wiki/Shadow_mapping)
- Look at Unreal Engine 5 virtual texturing
  - <https://dev.epicgames.com/documentation/en-us/unreal-engine/virtual-texturing-in-unreal-engine/>
- Look at Unreal Engine 5 MegaLights
  - <https://dev.epicgames.com/documentation/en-us/unreal-engine/megalights-in-unreal-engine/>


## Read (optional):

- CUDA Warp-Level Primitives
  - <https://developer.nvidia.com/blog/using-cuda-warp-level-primitives/>
- Warp-aggregated atomics
  - <https://developer.nvidia.com/blog/cuda-pro-tip-optimized-filtering-warp-aggregated-atomics/>

# GPU Texturing




# 2D Texture Mapping



For each fragment:  
interpolate the  
texture coordinates  
**(barycentric)**  
**Or:**

**Use arbitrary, computed coordinates**



**Texture-Lookup:**  
interpolate the  
texture data  
**(bi-linear)**  
**Or:**

**Nearest-neighbor for “array lookup”**

# Interpolation #1



# Interpolation Type + Purpose #1: **Interpolation of Texture Coordinates**


*(Linear / Rational-Linear Interpolation)*

# Linear Interpolation / Convex Combinations



Linear interpolation in 1D:

$$f(\alpha) = (1 - \alpha)v_1 + \alpha v_2$$



Line embedded in 2D (linear interpolation of vertex coordinates/attributes):

$$f(\alpha_1, \alpha_2) = \alpha_1 v_1 + \alpha_2 v_2$$

$$\alpha_1 + \alpha_2 = 1$$

$$f(\alpha) = v_1 + \alpha(v_2 - v_1)$$

$$\alpha = \alpha_2$$

Line segment:  $\alpha_1, \alpha_2 \geq 0$  ( $\rightarrow$  convex combination)

Compare to line parameterization  
with parameter t:

$$v(t) = v_1 + t(v_2 - v_1)$$

# Linear Interpolation / Convex Combinations




**Linear** combination ( $n$ -dim. space):

$$\alpha_1 v_1 + \alpha_2 v_2 + \dots + \alpha_n v_n = \sum_{i=1}^n \alpha_i v_i$$

**Affine** combination: Restrict to  $(n - 1)$ -dim. subspace:

$$\alpha_1 + \alpha_2 + \dots + \alpha_n = \sum_{i=1}^n \alpha_i = 1$$



**Convex** combination:  $\alpha_i \geq 0$

(restrict to simplex in subspace)

# Linear Interpolation / Convex Combinations




$$\alpha_1 v_1 + \alpha_2 v_2 + \dots + \alpha_n v_n = \sum_{i=1}^n \alpha_i v_i$$

$$\alpha_1 + \alpha_2 + \dots + \alpha_n = \sum_{i=1}^n \alpha_i = 1$$

Re-parameterize to get affine coordinates:

$$\begin{aligned}\alpha_1 v_1 + \alpha_2 v_2 + \alpha_3 v_3 &= \\ \tilde{\alpha}_1(v_2 - v_1) + \tilde{\alpha}_2(v_3 - v_1) + v_1 &\\ \tilde{\alpha}_1 &= \alpha_2 \\ \tilde{\alpha}_2 &= \alpha_3\end{aligned}$$



# Linear Interpolation / Convex Combinations



The weights  $\alpha_i$  are the (normalized) barycentric coordinates



→ linear attribute interpolation in simplex

$$\alpha_1 v_1 + \alpha_2 v_2 + \dots + \alpha_n v_n = \sum_{i=1}^n \alpha_i v_i$$

$$\alpha_1 + \alpha_2 + \dots + \alpha_n = \sum_{i=1}^n \alpha_i = 1$$

$$\alpha_i \geq 0$$

attribute interpolation



spatial position  
interpolation

wikipedia



# Homogeneous Coordinates (1)



## Projective geometry

- (Real) projective spaces  $\mathbf{RP}^n$ :  
Real projective line  $\mathbf{RP}^1$ , real projective plane  $\mathbf{RP}^2$ , ...
- A point in  $\mathbf{RP}^n$  is a line through the origin (i.e., all the scalar multiples of the same vector) in an  $(n+1)$ -dimensional (real) vector space



## Homogeneous coordinates of 2D projective point in $\mathbf{RP}^2$

- Coordinates differing only by a non-zero factor  $\lambda$  map to the same point  
 $(\lambda x, \lambda y, \lambda)$       dividing out the  $\lambda$  gives  $(x, y, 1)$ , corresponding to  $(x, y)$  in  $\mathbf{R}^2$
- Coordinates with last component = 0 map to “points at infinity”  
 $(\lambda x, \lambda y, 0)$       division by last component not allowed; but again this is the same point if it only differs by a scalar factor, e.g., this is the same point as  $(x, y, 0)$



# Homogeneous Coordinates (2)

## Examples of usage

- Translation (with translation vector  $\vec{b}$ )
- Affine transformations (linear transformation + translation)

$$\vec{y} = A\vec{x} + \vec{b}.$$

- With homogeneous coordinates:

$$\left[ \begin{array}{c} \vec{y} \\ 1 \end{array} \right] = \left[ \begin{array}{cc|c} & A & \vec{b} \\ 0 & \dots & 0 \end{array} \right] \left[ \begin{array}{c} \vec{x} \\ 1 \end{array} \right]$$

- Setting the last coordinate = 1 and the last row of the matrix to [ 0, ..., 0, 1 ] results in translation of the point  $\vec{x}$  (via addition of translation vector  $\vec{b}$ )
- The matrix above is a linear map, but because it is one dimension higher, it does not have to move the origin in the  $(n+1)$ -dimensional space for translation




# Homogeneous Coordinates (3)

## Examples of usage

- Projection (e.g., OpenGL projection matrices)


$$\begin{bmatrix} \frac{2}{right-left} & 0 & 0 & -\frac{right+left}{right-left} \\ 0 & \frac{2}{top-bottom} & 0 & -\frac{top+bottom}{top-bottom} \\ 0 & 0 & \frac{-2}{far-near} & -\frac{far+near}{far-near} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

orthographic



$$\begin{bmatrix} \frac{z_{near}}{width/2} & 0.0 & \frac{left+right}{width/2} & 0.0 \\ 0.0 & \frac{z_{near}}{height/2} & \frac{top+bottom}{height/2} & 0.0 \\ 0.0 & 0.0 & -\frac{z_{far}+z_{near}}{z_{far}-z_{near}} & \frac{2z_{far}z_{near}}{z_{far}-z_{near}} \\ 0.0 & 0.0 & -1.0 & 0.0 \end{bmatrix}$$

perspective



# Texture Mapping

---

2D (3D) Texture Space

| Texture Transformation

2D Object Parameters

| Parameterization

3D Object Space

| Model Transformation


3D World Space

| Viewing Transformation

3D Camera Space

| Projection

2D Image Space



# Texture Mapping Polygons

---

Forward transformation: linear projective map


$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} s \\ t \\ r \end{bmatrix}$$

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

# Incorrect attribute interpolation

---



$A' \neq A !$

# Linear interpolation

---

Compute intermediate attribute value

- Along a line:  $A = aA_1 + bA_2, \quad a+b=1$
- On a plane:  $A = aA_1 + bA_2 + cA_3, \quad a+b+c=1$

Only projected values interpolate linearly in screen space (straight lines project to straight lines)


- $x$  and  $y$  are projected (divided by  $w$ )
- Attribute values are not naturally projected

Choice for attribute interpolation in screen space


- Interpolate unprojected values
  - Cheap and easy to do, but gives wrong values
  - Sometimes OK for color, but
  - Never acceptable for texture coordinates
- Do it right

# Linear Perspective


---




**Correct Linear Perspective**



**Incorrect Perspective**





# Perspective Texture Mapping



linear interpolation  
in object space

$$\frac{ax_1 + bx_2}{aw_1 + bw_2} \neq a\frac{x_1}{w_1} + b\frac{x_2}{w_2}$$

linear interpolation  
in screen space



$$a = b = 0.5$$



# Perspective-correct linear interpolation

---

Only projected values interpolate correctly, so project A

- Linearly interpolate  $A_1/w_1$  and  $A_2/w_2$

Also interpolate  $1/w_1$  and  $1/w_2$

- These also interpolate linearly in screen space

Divide interpolants at each sample point to recover A


- $(A/w) / (1/w) = A$
- Division is expensive (more than add or multiply), so
  - Recover w for the sample point (reciprocate), and
  - Multiply each projected attribute by w

Barycentric triangle parameterization:

$$A = \frac{aA_1/w_1 + bA_2/w_2 + cA_3/w_3}{a/w_1 + b/w_2 + c/w_3} \quad a + b + c = 1$$

# Perspective Texture Mapping

- Solution: interpolate  $(s/w, t/w, 1/w)$
- $(s/w) / (1/w) = s$  etc. at every fragment



Heckbert and Moreton



# Perspective-Correct Interpolation Recipe



$$r_i(x, y) = \frac{r_i(x, y)/w(x, y)}{1/w(x, y)}$$

- (1) Associate a record containing the  $n$  parameters of interest  $(r_1, r_2, \dots, r_n)$  with each vertex of the polygon.
- (2) For each vertex, transform object space coordinates to homogeneous screen space using  $4 \times 4$  object to screen matrix, yielding the values  $(xw, yw, zw, w)$ .
- (3) Clip the polygon against plane equations for each of the six sides of the viewing frustum, linearly interpolating all the parameters when new vertices are created.
- (4) At each vertex, divide the homogeneous screen coordinates, the parameters  $r_i$ , and the number 1 by  $w$  to construct the variable list  $(x, y, z, s_1, s_2, \dots, s_{n+1})$ , where  $s_i = r_i/w$  for  $i \leq n$ ,  $s_{n+1} = 1/w$ .
- (5) Scan convert in screen space by linear interpolation of all parameters, at each pixel computing  $r_i = s_i/s_{n+1}$  for each of the  $n$  parameters; use these values for shading.



# Projective Map vs. Interpolation Recipe (1)

In general (see previous slides),  
we had the projective map:

Let's rename and rewrite this as:

$$\begin{bmatrix} s \\ t \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_{world} \\ y_{world} \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix}$$

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

For homogeneous points  
we can also divide by w:

Coordinates on the right become  
screen space coordinates!

$$\begin{bmatrix} s \\ t \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix},$$

$$\begin{bmatrix} s/w \\ t/w \\ q/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}.$$



# Projective Map vs. Interpolation Recipe (2)

In general (see previous slides),  
we had the projective map:

Let's rename and rewrite this as:

$$\begin{bmatrix} s \\ t \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_{world} \\ y_{world} \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix}$$

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

For homogeneous points  
we can also divide by w:

Coordinates on the right become  
screen space coordinates!

$$\begin{bmatrix} s \\ t \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix},$$

(special case  $q = 1$ )

$$\boxed{\begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}}.$$

# Projective Map vs. Interpolation Recipe (3)



In general (see previous slides),  
we had the projective map:

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Now consider scanline interpolation:  
(barycentric interpolation is linear along any line: here, horizontal line)

$$\begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x + \Delta x \\ y \\ 1 \end{bmatrix},$$


$$\begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} + \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} \Delta x \\ 0 \\ 0 \end{bmatrix}$$

$$\Delta_x \begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} \Delta x \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} a \cdot \Delta x \\ d \cdot \Delta x \\ g \cdot \Delta x \end{bmatrix} = \begin{bmatrix} a \\ d \\ g \end{bmatrix}$$



$$(\Delta x = 1)$$

# Projective Texture Mapping

- Want to simulate a beamer
  - ... or a flashlight, or a slide projector
- Precursor to shadows
- Interesting mathematics:  
2 perspective projections involved!
- Easy to program!



# Projective Texture Mapping



# Projective Shadows in Doom 3



- What about **homogeneous texture coords**?
- Need to do perspective divide also for projector!
  - $(s, t, q) \rightarrow (s/q, t/q)$  for every fragment
- How does OpenGL do that?
  - Needs to be perspective correct as well!
  - Trick: interpolate  $(s/w, t/w, r/w, q/w)$
  - $(s/w) / (q/w) = s/q$  etc. at every fragment
- Remember:  $s, t, r, q$  are equivalent to  $x, y, z, w$  in projector space!  $\rightarrow r/q = \text{projector depth!}$



- Apply multiple textures in one pass
- *Integral* part of programmable shading
  - e.g. diffuse texture map + gloss map
  - e.g. diffuse texture map + light map
- Performance issues
  - How many textures are free?
  - How many are available




# Example: Light Mapping

- Used in virtually every commercial game
- Precalculate diffuse lighting on static objects
  - Only low resolution necessary
  - Diffuse lighting is view independent!
- Advantages:
  - No runtime lighting necessary
    - VERY fast!
  - Can take global effects (shadows, color bleeds) into account



# Light Mapping




Original LM texels




Bilinear Filtering



## ■ Why premultiplication is bad...



Full Size Texture  
(with Lightmap)



Tiled Surface Texture  
plus Lightmap

→ use tileable surface textures and low  
resolution lightmaps



# Light Mapping



Original scene



Light-mapped



# Example: Light Mapping


- Precomputation based on non-realtimemethods
  - Radiosity
  - Ray tracing
    - Monte Carlo Integration
    - Path tracing
    - Photon mapping



# Light Mapping



Lightmap



mapped



# Light Mapping



Original scene

Light-mapped



Thank you.