

CS 380 - GPU and GPGPU Programming

Lecture 23: GPU Texturing, Pt. 5

Markus Hadwiger, KAUST

Reading Assignment #12 (until Nov 24)



Read (required):

- Look at Vulkan *sparse resources*, especially *sparse partially-resident images*
 - <https://docs.vulkan.org/spec/latest/chapters/sparsemem.html>
- Read about shadow mapping
 - https://en.wikipedia.org/wiki/Shadow_mapping
- Look at Unreal Engine 5 virtual texturing
 - <https://dev.epicgames.com/documentation/en-us/unreal-engine/virtual-texturing-in-unreal-engine/>
- Look at Unreal Engine 5 MegaLights
 - <https://dev.epicgames.com/documentation/en-us/unreal-engine/megalights-in-unreal-engine/>

Read (optional):

- CUDA Warp-Level Primitives
 - <https://developer.nvidia.com/blog/using-cuda-warp-level-primitives/>
- Warp-aggregated atomics
 - <https://developer.nvidia.com/blog/cuda-pro-tip-optimized-filtering-warp-aggregated-atomics/>

GPU Texturing

Interpolation #1



Interpolation Type + Purpose #1: **Interpolation of Texture Coordinates**

(Linear / Rational-Linear Interpolation)

Homogeneous Coordinates (1)



Projective geometry

- (Real) projective spaces \mathbf{RP}^n :
Real projective line \mathbf{RP}^1 , real projective plane \mathbf{RP}^2 , ...
- A point in \mathbf{RP}^n is a line through the origin (i.e., all the scalar multiples of the same vector) in an $(n+1)$ -dimensional (real) vector space



Homogeneous coordinates of 2D projective point in \mathbf{RP}^2

- Coordinates differing only by a non-zero factor λ map to the same point
 $(\lambda x, \lambda y, \lambda)$ dividing out the λ gives $(x, y, 1)$, corresponding to (x, y) in \mathbf{R}^2
- Coordinates with last component = 0 map to “points at infinity”
 $(\lambda x, \lambda y, 0)$ division by last component not allowed; but again this is the same point if it only differs by a scalar factor, e.g., this is the same point as $(x, y, 0)$

Texture Mapping

2D (3D) Texture Space

| Texture Transformation

2D Object Parameters

| Parameterization

3D Object Space

| Model Transformation


3D World Space

| Viewing Transformation

3D Camera Space

| Projection

2D Image Space



Texture Mapping Polygons

Forward transformation: linear projective map

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} s \\ t \\ r \end{bmatrix}$$

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Linear interpolation

Compute intermediate attribute value

- Along a line: $A = aA_1 + bA_2, \quad a+b=1$
- On a plane: $A = aA_1 + bA_2 + cA_3, \quad a+b+c=1$

Only projected values interpolate linearly in screen space (straight lines project to straight lines)

- x and y are projected (divided by w)
- Attribute values are not naturally projected

Choice for attribute interpolation in screen space

- Interpolate unprojected values
 - Cheap and easy to do, but gives wrong values
 - Sometimes OK for color, but
 - Never acceptable for texture coordinates
- Do it right

Perspective-correct linear interpolation

Only projected values interpolate correctly, so project A

- Linearly interpolate A_1/w_1 and A_2/w_2

Also interpolate $1/w_1$ and $1/w_2$

- These also interpolate linearly in screen space

Divide interpolants at each sample point to recover A


- $(A/w) / (1/w) = A$
- Division is expensive (more than add or multiply), so
 - Recover w for the sample point (reciprocate), and
 - Multiply each projected attribute by w

Barycentric triangle parameterization:

$$A = \frac{aA_1/w_1 + bA_2/w_2 + cA_3/w_3}{a/w_1 + b/w_2 + c/w_3} \quad a + b + c = 1$$

Perspective Texture Mapping

- Solution: interpolate $(s/w, t/w, 1/w)$
- $(s/w) / (1/w) = s$ etc. at every fragment



Heckbert and Moreton



Perspective-Correct Interpolation Recipe



$$r_i(x, y) = \frac{r_i(x, y)/w(x, y)}{1/w(x, y)}$$

- (1) Associate a record containing the n parameters of interest (r_1, r_2, \dots, r_n) with each vertex of the polygon.
- (2) For each vertex, transform object space coordinates to homogeneous screen space using 4×4 object to screen matrix, yielding the values (xw, yw, zw, w) .
- (3) Clip the polygon against plane equations for each of the six sides of the viewing frustum, linearly interpolating all the parameters when new vertices are created.
- (4) At each vertex, divide the homogeneous screen coordinates, the parameters r_i , and the number 1 by w to construct the variable list $(x, y, z, s_1, s_2, \dots, s_{n+1})$, where $s_i = r_i/w$ for $i \leq n$, $s_{n+1} = 1/w$.
- (5) Scan convert in screen space by linear interpolation of all parameters, at each pixel computing $r_i = s_i/s_{n+1}$ for each of the n parameters; use these values for shading.



Projective Map vs. Interpolation Recipe (1)

In general (see previous slides),
we had the projective map:

Let's rename and rewrite this as:

$$\begin{bmatrix} s \\ t \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_{world} \\ y_{world} \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix}$$

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

For homogeneous points
we can also divide by w:

Coordinates on the right become
screen space coordinates!

$$\begin{bmatrix} s \\ t \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix},$$

$$\begin{bmatrix} s/w \\ t/w \\ q/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}.$$



Projective Map vs. Interpolation Recipe (2)

In general (see previous slides),
we had the projective map:

Let's rename and rewrite this as:

$$\begin{bmatrix} s \\ t \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_{world} \\ y_{world} \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix}$$

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

For homogeneous points
we can also divide by w:

Coordinates on the right become
screen space coordinates!

$$\begin{bmatrix} s \\ t \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \cdot w \\ y \cdot w \\ w \end{bmatrix},$$

(special case $q = 1$)

$$\boxed{\begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}}.$$

Projective Map vs. Interpolation Recipe (3)



In general (see previous slides),
we had the projective map:

Backward transformation: linear projective map

$$\begin{bmatrix} s \\ t \\ r \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}^{-1} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Now consider scanline interpolation:
(barycentric interpolation is linear along any line: here, horizontal line)

$$\begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x + \Delta x \\ y \\ 1 \end{bmatrix},$$

$$\begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} + \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} \Delta x \\ 0 \\ 0 \end{bmatrix}$$

$$\Delta_x \begin{bmatrix} s/w \\ t/w \\ 1/w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} \Delta x \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} a \cdot \Delta x \\ d \cdot \Delta x \\ g \cdot \Delta x \end{bmatrix} = \begin{bmatrix} a \\ d \\ g \end{bmatrix}$$

$$(\Delta x = 1)$$

Interpolation #2



Interpolation Type + Purpose #2: **Interpolation of Samples in Texture Space**

(Multi-Linear Interpolation)

Magnification (Bi-linear Filtering Example)



Original image




Nearest neighbor




Bi-linear filtering



Nearest-Neighbor vs. Bi-Linear Interpolation




nearest-neighbor



wikipedia

bi-linear





Bilinear patch (courtesy J. Han)



Bi-Linear Interpolation

Consider area between 2x2 adjacent samples (e.g., pixel centers)

Example #2: 1 at top-left and bottom-right, 0 at bottom-left, 0.5 at top-right





Bi-Linear Interpolation

Consider area between 2x2 adjacent samples (e.g., pixel centers):

Given any (fractional) position


$$\alpha_1 := x_1 - \lfloor x_1 \rfloor \quad \alpha_1 \in [0.0, 1.0)$$

$$\alpha_2 := x_2 - \lfloor x_2 \rfloor \quad \alpha_2 \in [0.0, 1.0)$$

and 2x2 sample values

$$\begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix}$$

Compute: $f(\alpha_1, \alpha_2)$





Bi-Linear Interpolation

Consider area between 2x2 adjacent samples (e.g., pixel centers):

Given any (fractional) position


$$\alpha_1 := x_1 - \lfloor x_1 \rfloor \quad \alpha_1 \in [0.0, 1.0)$$

$$\alpha_2 := x_2 - \lfloor x_2 \rfloor \quad \alpha_2 \in [0.0, 1.0)$$

and 2x2 sample values

$$\begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix}$$

Compute: $f(\alpha_1, \alpha_2)$





Bi-Linear Interpolation

Weights in 2x2 format:

$$\begin{bmatrix} \alpha_2 \\ (1 - \alpha_2) \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) & \alpha_1 \end{bmatrix} = \begin{bmatrix} (1 - \alpha_1)\alpha_2 & \alpha_1\alpha_2 \\ (1 - \alpha_1)(1 - \alpha_2) & \alpha_1(1 - \alpha_2) \end{bmatrix}$$

Interpolate function at (fractional) position (α_1, α_2) :

$$f(\alpha_1, \alpha_2) = [\alpha_2 \quad (1 - \alpha_2)] \begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix}$$



Bi-Linear Interpolation

Interpolate function at (fractional) position (α_1, α_2) :

$$f(\alpha_1, \alpha_2) = [\alpha_2 \quad (1 - \alpha_2)] \begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix}$$

$$= [\alpha_2 \quad (1 - \alpha_2)] \begin{bmatrix} (1 - \alpha_1)v_{01} + \alpha_1 v_{11} \\ (1 - \alpha_1)v_{00} + \alpha_1 v_{10} \end{bmatrix}$$

$$= [\alpha_2 v_{01} + (1 - \alpha_2)v_{00} \quad \alpha_2 v_{11} + (1 - \alpha_2)v_{10}] \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix}$$



Bi-Linear Interpolation

Interpolate function at (fractional) position (α_1, α_2) :

$$f(\alpha_1, \alpha_2) = [\alpha_2 \quad (1 - \alpha_2)] \begin{bmatrix} v_{01} & v_{11} \\ v_{00} & v_{10} \end{bmatrix} \begin{bmatrix} (1 - \alpha_1) \\ \alpha_1 \end{bmatrix}$$

$$= (1 - \alpha_1)(1 - \alpha_2)v_{00} + \alpha_1(1 - \alpha_2)v_{10} + (1 - \alpha_1)\alpha_2v_{01} + \alpha_1\alpha_2v_{11}$$

$$= v_{00} + \alpha_1(v_{10} - v_{00}) + \alpha_2(v_{01} - v_{00}) + \alpha_1\alpha_2(v_{00} + v_{11} - v_{10} - v_{01})$$






REALLY IMPORTANT:

this is a different thing (for a different purpose)
than the linear (or, in perspective, rational-linear)
interpolation of texture coordinates!!

Texture Minification


Texture Aliasing: Minification

- Problem: One pixel in image space covers many texels




Texture Aliasing: Minification

- Caused by *undersampling*: texture information is lost





Texture Anti-Aliasing: Minification

- A good pixel value is the weighted mean of the pixel area projected into texture space



Texture Anti-Aliasing: MIP Mapping

- MIP Mapping (“Multum In Parvo”)
 - Texture size is reduced by factors of 2 (*downsampling* = "many things in a small place")
 - Simple (4 pixel average) and memory efficient
 - Last image is only ONE texel




Texture Anti-Aliasing: MIP Mapping

- MIP Mapping (“Multum In Parvo”)
 - Texture size is reduced by factors of 2 (*downsampling* = "many things in a small place")
 - Simple (4 pixel average) and memory efficient
 - Last image is only ONE texel

geometric series:



$$a + ar + ar^2 + ar^3 + \cdots + ar^{n-1} =$$

$$= \sum_{k=0}^{n-1} ar^k = a \left(\frac{1 - r^n}{1 - r} \right)$$






Texture Anti-Aliasing: MIP Mapping

- MIP Mapping Algorithm
- $D := ld(\max(d_1, d_2))$ "Mip Map level"
- $T_0 := \text{value from texture } D_0 = \text{trunc}(D)$
 - Use bilinear interpolation





MIP-Map Level Computation



- Use the partial derivatives of texture coordinates with respect to screen space coordinates
- This is the Jacobian matrix
- Area of parallelogram is the absolute value of the Jacobian determinant (the *Jacobian*)

$$\begin{pmatrix} \partial u / \partial x & \partial u / \partial y \\ \partial v / \partial x & \partial v / \partial y \end{pmatrix} = \begin{pmatrix} s_x & s_y \\ t_x & t_y \end{pmatrix}$$



MIP-Map Level Computation (OpenGL)

- OpenGL 4.6 core specification, pp. 251-264
(3D tex coords!)

$$\lambda_{base}(x, y) = \log_2[\rho(x, y)]$$

$$\rho = \max \left\{ \sqrt{\left(\frac{\partial u}{\partial x}\right)^2 + \left(\frac{\partial v}{\partial x}\right)^2 + \left(\frac{\partial w}{\partial x}\right)^2}, \sqrt{\left(\frac{\partial u}{\partial y}\right)^2 + \left(\frac{\partial v}{\partial y}\right)^2 + \left(\frac{\partial w}{\partial y}\right)^2} \right\}$$

Does not use area of parallelogram but greater hypotenuse [Heckbert, 1983]


- Approximation without square-roots

$$m_u = \max \left\{ \left| \frac{\partial u}{\partial x} \right|, \left| \frac{\partial u}{\partial y} \right| \right\} \quad m_v = \max \left\{ \left| \frac{\partial v}{\partial x} \right|, \left| \frac{\partial v}{\partial y} \right| \right\} \quad m_w = \max \left\{ \left| \frac{\partial w}{\partial x} \right|, \left| \frac{\partial w}{\partial y} \right| \right\}$$

$$\max\{m_u, m_v, m_w\} \leq f(x, y) \leq m_u + m_v + m_w$$



MIP-Map Level Interpolation




- Level of detail value is fractional!
- Use fractional part to blend (lin.) between two adjacent mipmap levels

Texture Anti-Aliasing: MIP Mapping


■ Trilinear interpolation:

- $T_1 :=$ value from texture $D_1 = D_0 + 1$ (bilin.interpolation)
- Pixel value := $(D_1 - D) \cdot T_0 + (D - D_0) \cdot T_1$
 - Linear interpolation between successive MIP Maps
- Avoids "Mip banding" (but doubles texture lookups)





Texture Anti-Aliasing: MIP Mapping

- Other example for bilinear vs. trilinear filtering



Anti-Aliasing: Anisotropic Filtering

- Anisotropic filtering
 - View-dependent filter kernel
 - Implementation: *summed area table*, "*RIP Mapping*", *footprint assembly*, *elliptical weighted average* (EWA)




Texture space






Anisotropic Filtering: Footprint Assembly



Anti-Aliasing: Anisotropic Filtering

■ Example



- Basically, everything done in hardware
- gluBuild2DMipmaps () generates MIPmaps
- Set parameters in glTexParameter()
 - GL_TEXTURE_MAG_FILTER: GL_NEAREST, GL_LINEAR, ...
 - GL_TEXTURE_MIN_FILTER: GL_LINEAR_MIPMAP_NEAREST
- Anisotropic filtering is an extension:
 - GL_EXT_texture_filter_anisotropic
 - Number of samples can be varied (4x,8x,16x)
 - Vendor specific support and extensions

for Vulkan, see `vkSampler`,
`VkSamplerCreateInfo::magFilter`, `VkSamplerCreateInfo::minFilter`,
`VK_FILTER_NEAREST`, `VK_FILTER_LINEAR`,
`VkSamplerCreateInfo::mipmapMode`,
`VK_SAMPLER_MIPMAP_MODE_NEAREST`, `VK_SAMPLER_MIPMAP_MODE_LINEAR`, ...



Thank you.