

# CS 380 - GPU and GPGPU Programming

## Lecture 24: GPU Virtual Texturing

Markus Hadwiger, KAUST

# Reading Assignment #13 (until Dec 1)



## Read (required):

- Look at Unreal Engine 5 virtual geometry system (Nanite)
  - <https://dev.epicgames.com/documentation/en-us/unreal-engine/nanite-virtualized-geometry-in-unreal-engine/>
- Look at Unreal Engine 5 Nanite tech talk by Brian Karis (*Journey to Nanite*) slides
  - [https://www.highperformancegraphics.org/slides22/Journey\\_to\\_Nanite.pdf](https://www.highperformancegraphics.org/slides22/Journey_to_Nanite.pdf)
- Look at Unreal Engine 5 Lumen tech talk by Daniel Wright et al. slides
  - <https://advances.realtimerendering.com/s2022/SIGGRAPH2022-Advances-Lumen-Wright%20et%20al.pdf>

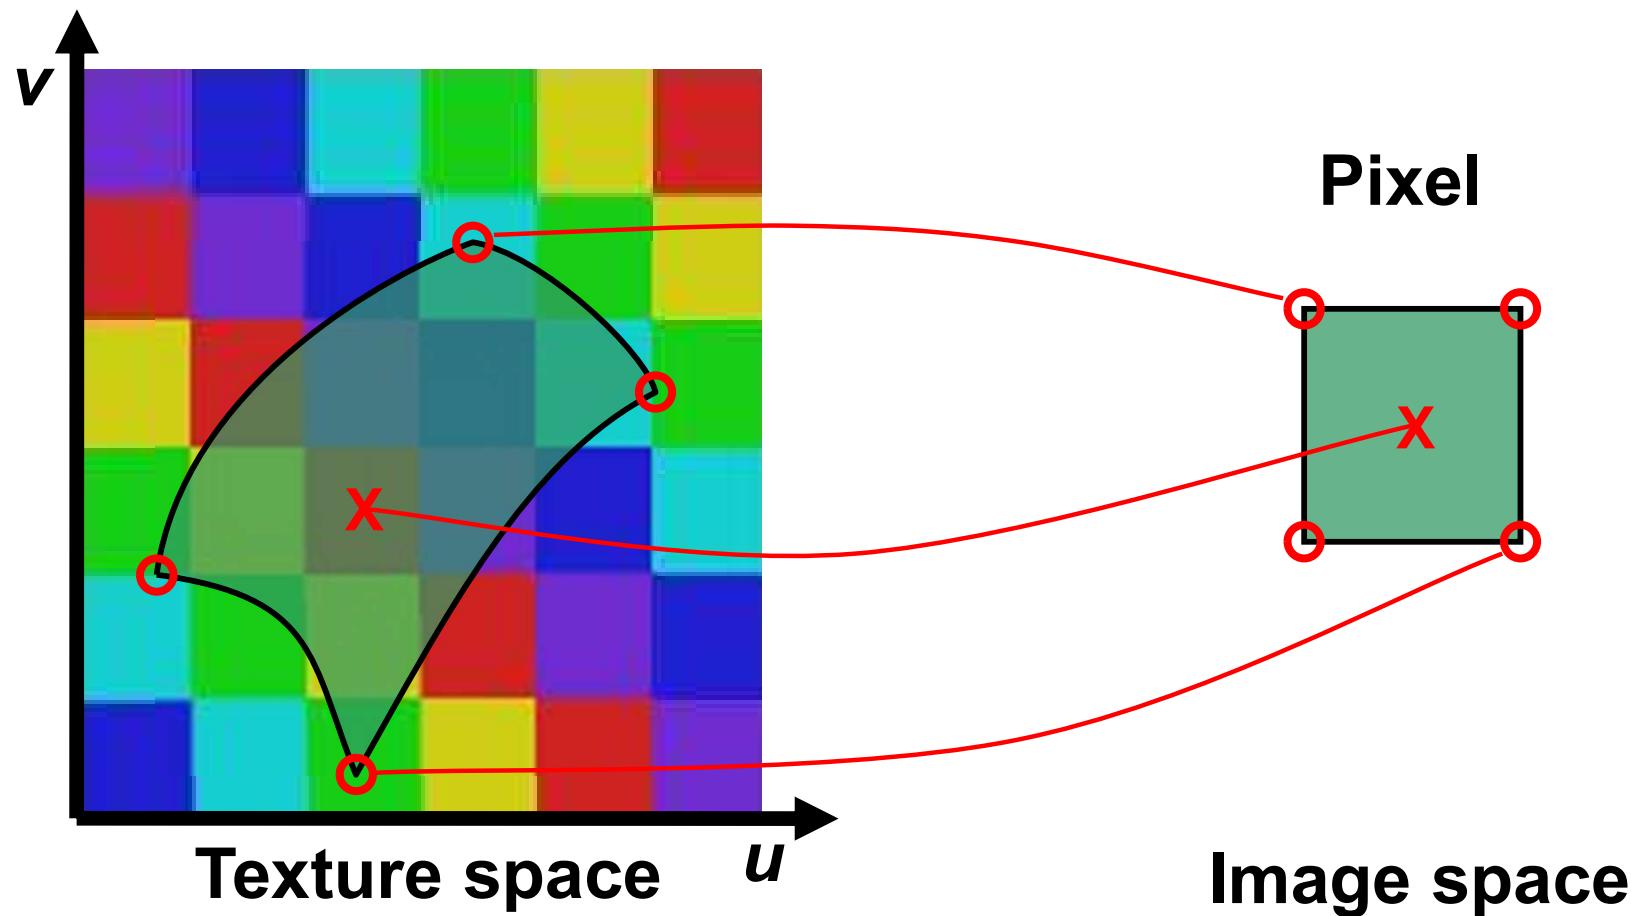
## Read (optional):

- Look at Unreal Engine 5 Nanite tech talk by Brian Karis (*Journey to Nanite*) talk
  - [https://www.youtube.com/watch?v=NRnj\\_1npORU](https://www.youtube.com/watch?v=NRnj_1npORU)

# Texture Minification

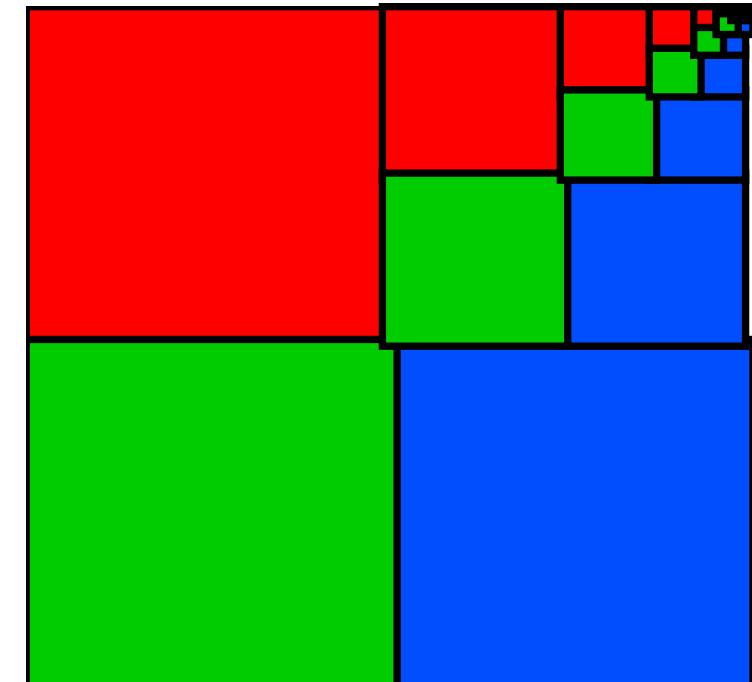
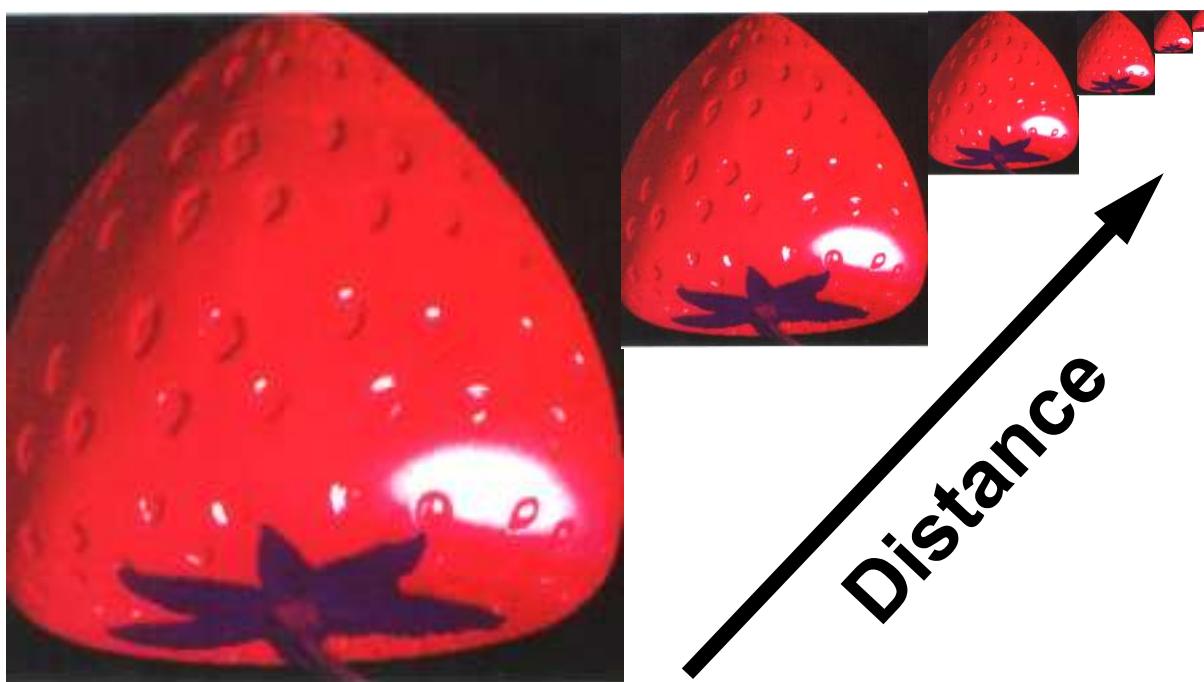
# Texture Anti-Aliasing: Minification

- A good pixel value is the weighted mean of the pixel area projected into texture space



# Texture Anti-Aliasing: MIP Mapping

- MIP Mapping (“Multum In Parvo”)
  - Texture size is reduced by factors of 2 (*downsampling* = "many things in a small place")
  - Simple (4 pixel average) and memory efficient
  - Last image is only ONE texel



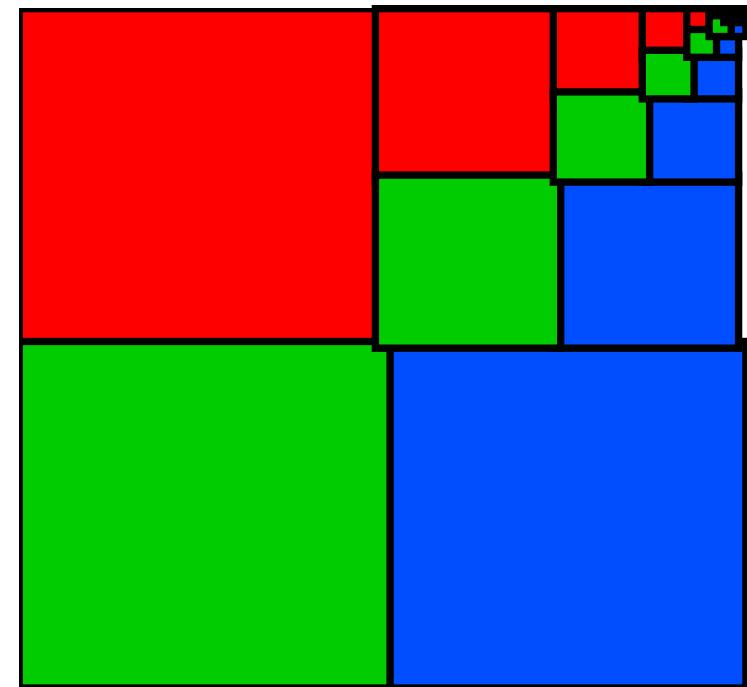
# Texture Anti-Aliasing: MIP Mapping

- MIP Mapping (“Multum In Parvo”)
  - Texture size is reduced by factors of 2 (*downsampling* = "many things in a small place")
  - Simple (4 pixel average) and memory efficient
  - Last image is only ONE texel

geometric series:

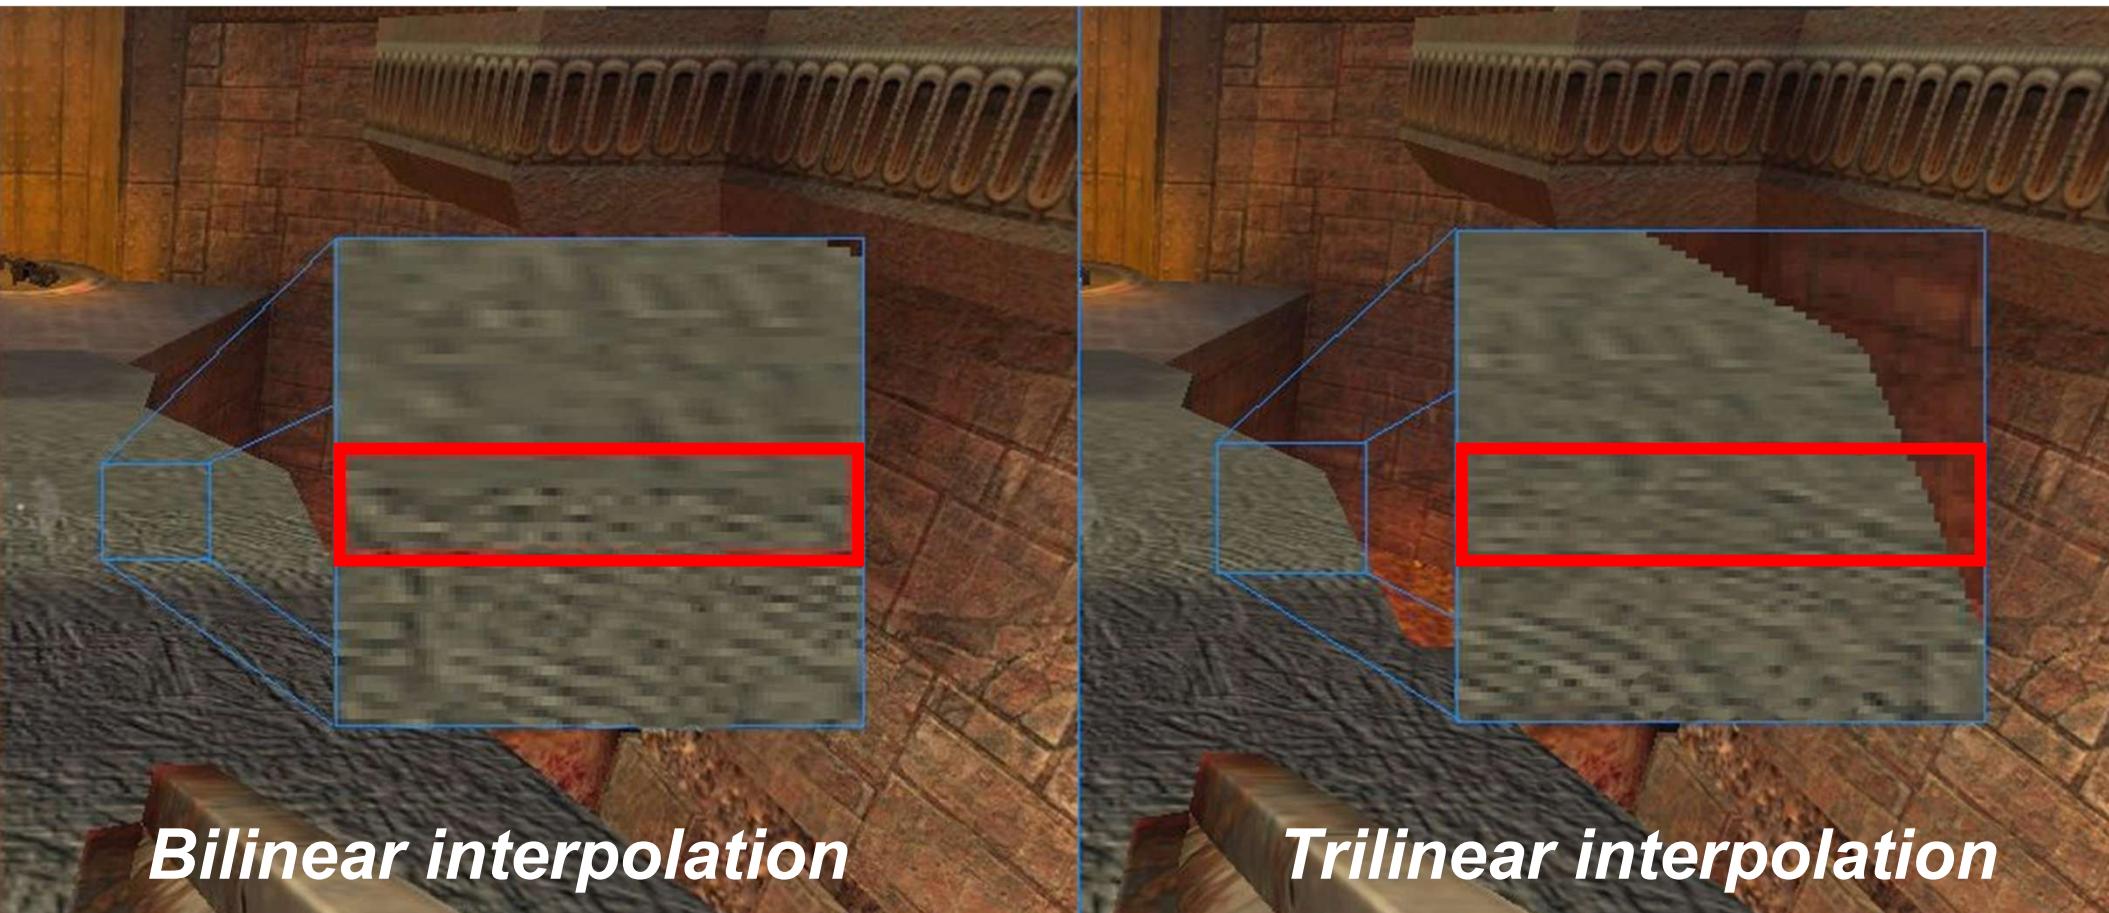
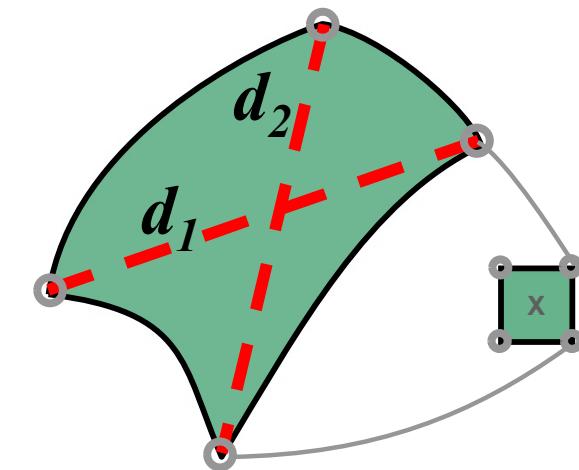
$$a + ar + ar^2 + ar^3 + \cdots + ar^{n-1} =$$

$$= \sum_{k=0}^{n-1} ar^k = a \left( \frac{1 - r^n}{1 - r} \right)$$



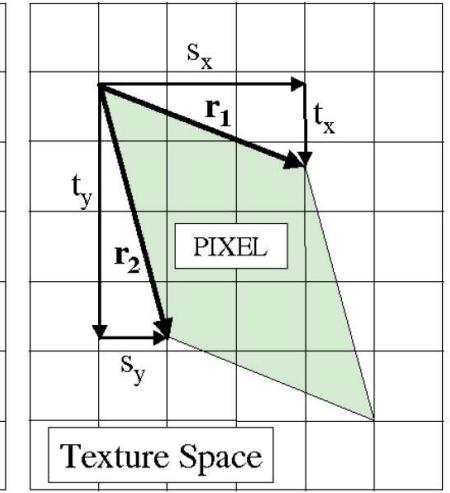
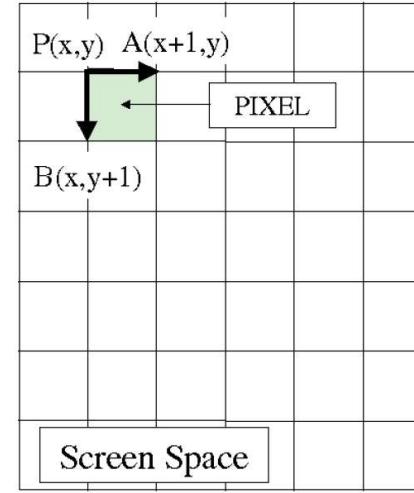
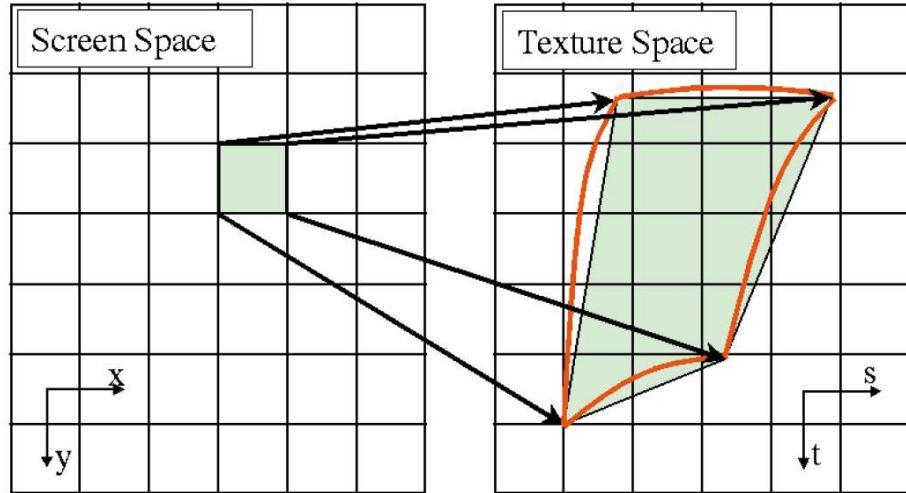
# Texture Anti-Aliasing: MIP Mapping

- MIP Mapping Algorithm
- $D := ld(\max(d_1, d_2))$  "Mip Map level"
- $T_0 := \text{value from texture } D_0 = \text{trunc}(D)$ 
  - Use bilinear interpolation





# MIP-Map Level Computation



- Use the partial derivatives of texture coordinates with respect to screen space coordinates
- This is the Jacobian matrix
- Area of parallelogram is the absolute value of the Jacobian determinant (the *Jacobian*)

$$\begin{pmatrix} \partial u / \partial x & \partial u / \partial y \\ \partial v / \partial x & \partial v / \partial y \end{pmatrix} = \begin{pmatrix} s_x & s_y \\ t_x & t_y \end{pmatrix}$$

# MIP-Map Level Computation (OpenGL)



- OpenGL 4.6 core specification, pp. 251-264  
(3D tex coords!)

$$\lambda_{base}(x, y) = \log_2[\rho(x, y)]$$

$$\rho = \max \left\{ \sqrt{\left(\frac{\partial u}{\partial x}\right)^2 + \left(\frac{\partial v}{\partial x}\right)^2 + \left(\frac{\partial w}{\partial x}\right)^2}, \sqrt{\left(\frac{\partial u}{\partial y}\right)^2 + \left(\frac{\partial v}{\partial y}\right)^2 + \left(\frac{\partial w}{\partial y}\right)^2} \right\}$$

Does not use area of parallelogram but greater hypotenuse [Heckbert, 1983]

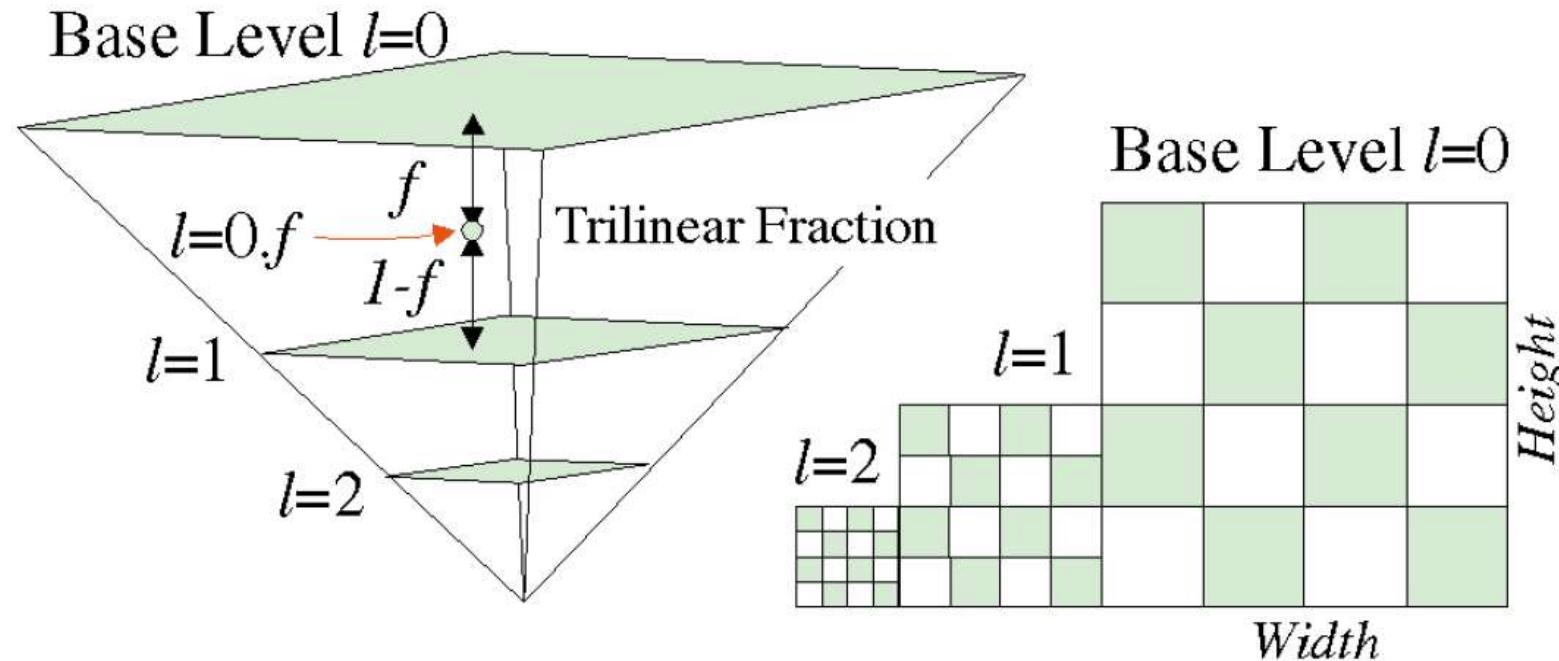
- Approximation without square-roots

$$m_u = \max \left\{ \left| \frac{\partial u}{\partial x} \right|, \left| \frac{\partial u}{\partial y} \right| \right\} \quad m_v = \max \left\{ \left| \frac{\partial v}{\partial x} \right|, \left| \frac{\partial v}{\partial y} \right| \right\} \quad m_w = \max \left\{ \left| \frac{\partial w}{\partial x} \right|, \left| \frac{\partial w}{\partial y} \right| \right\}$$

$$\max\{m_u, m_v, m_w\} \leq f(x, y) \leq m_u + m_v + m_w$$



# MIP-Map Level Interpolation

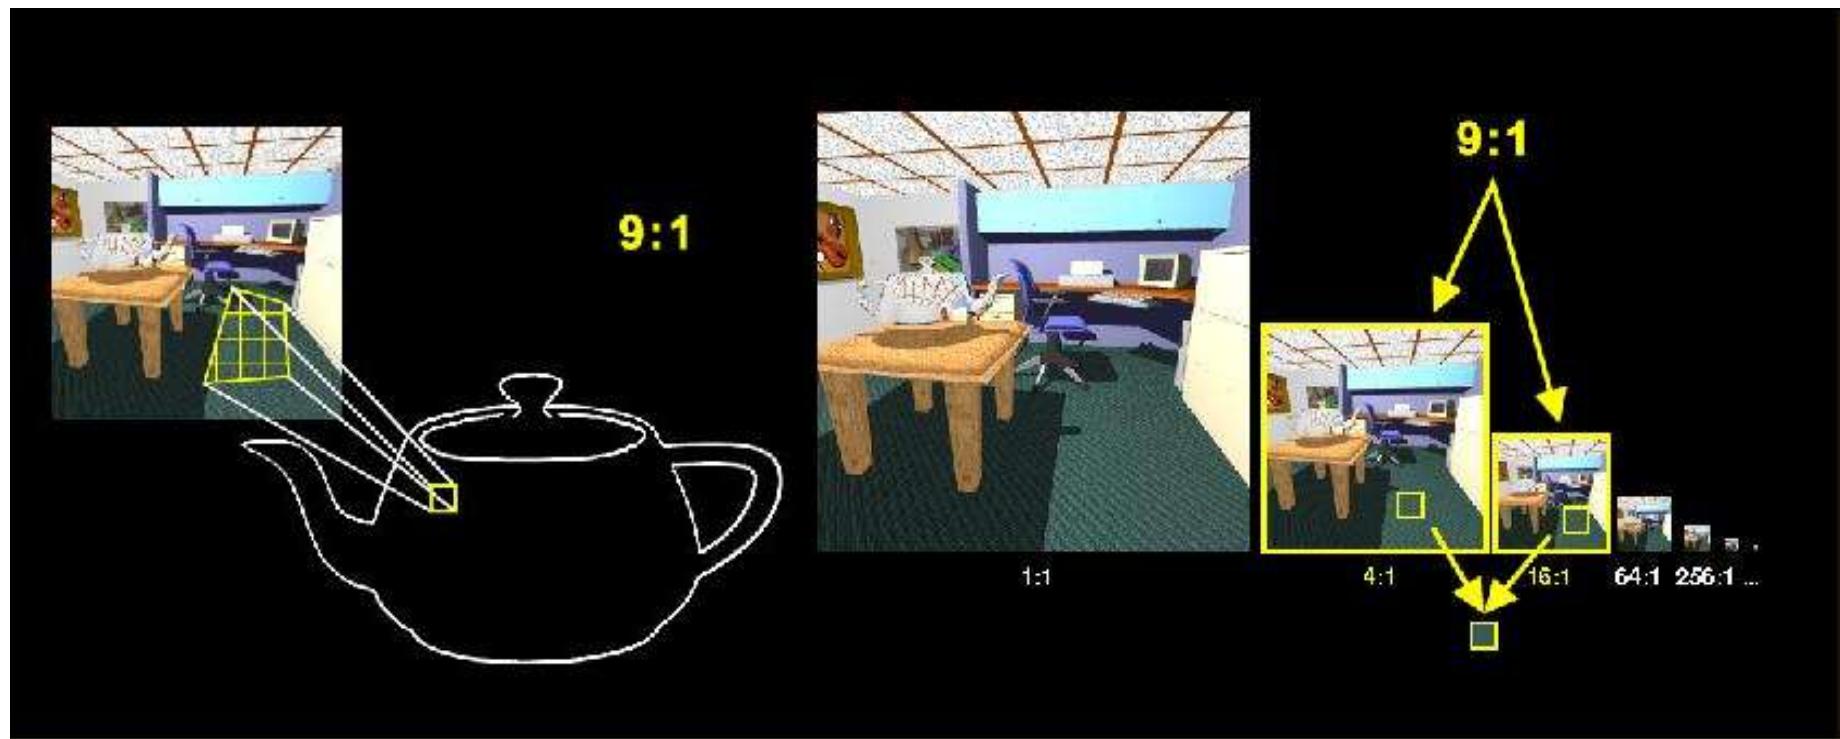


- Level of detail value is fractional!
- Use fractional part to blend (lin.) between two adjacent mipmap levels

# Texture Anti-Aliasing: MIP Mapping

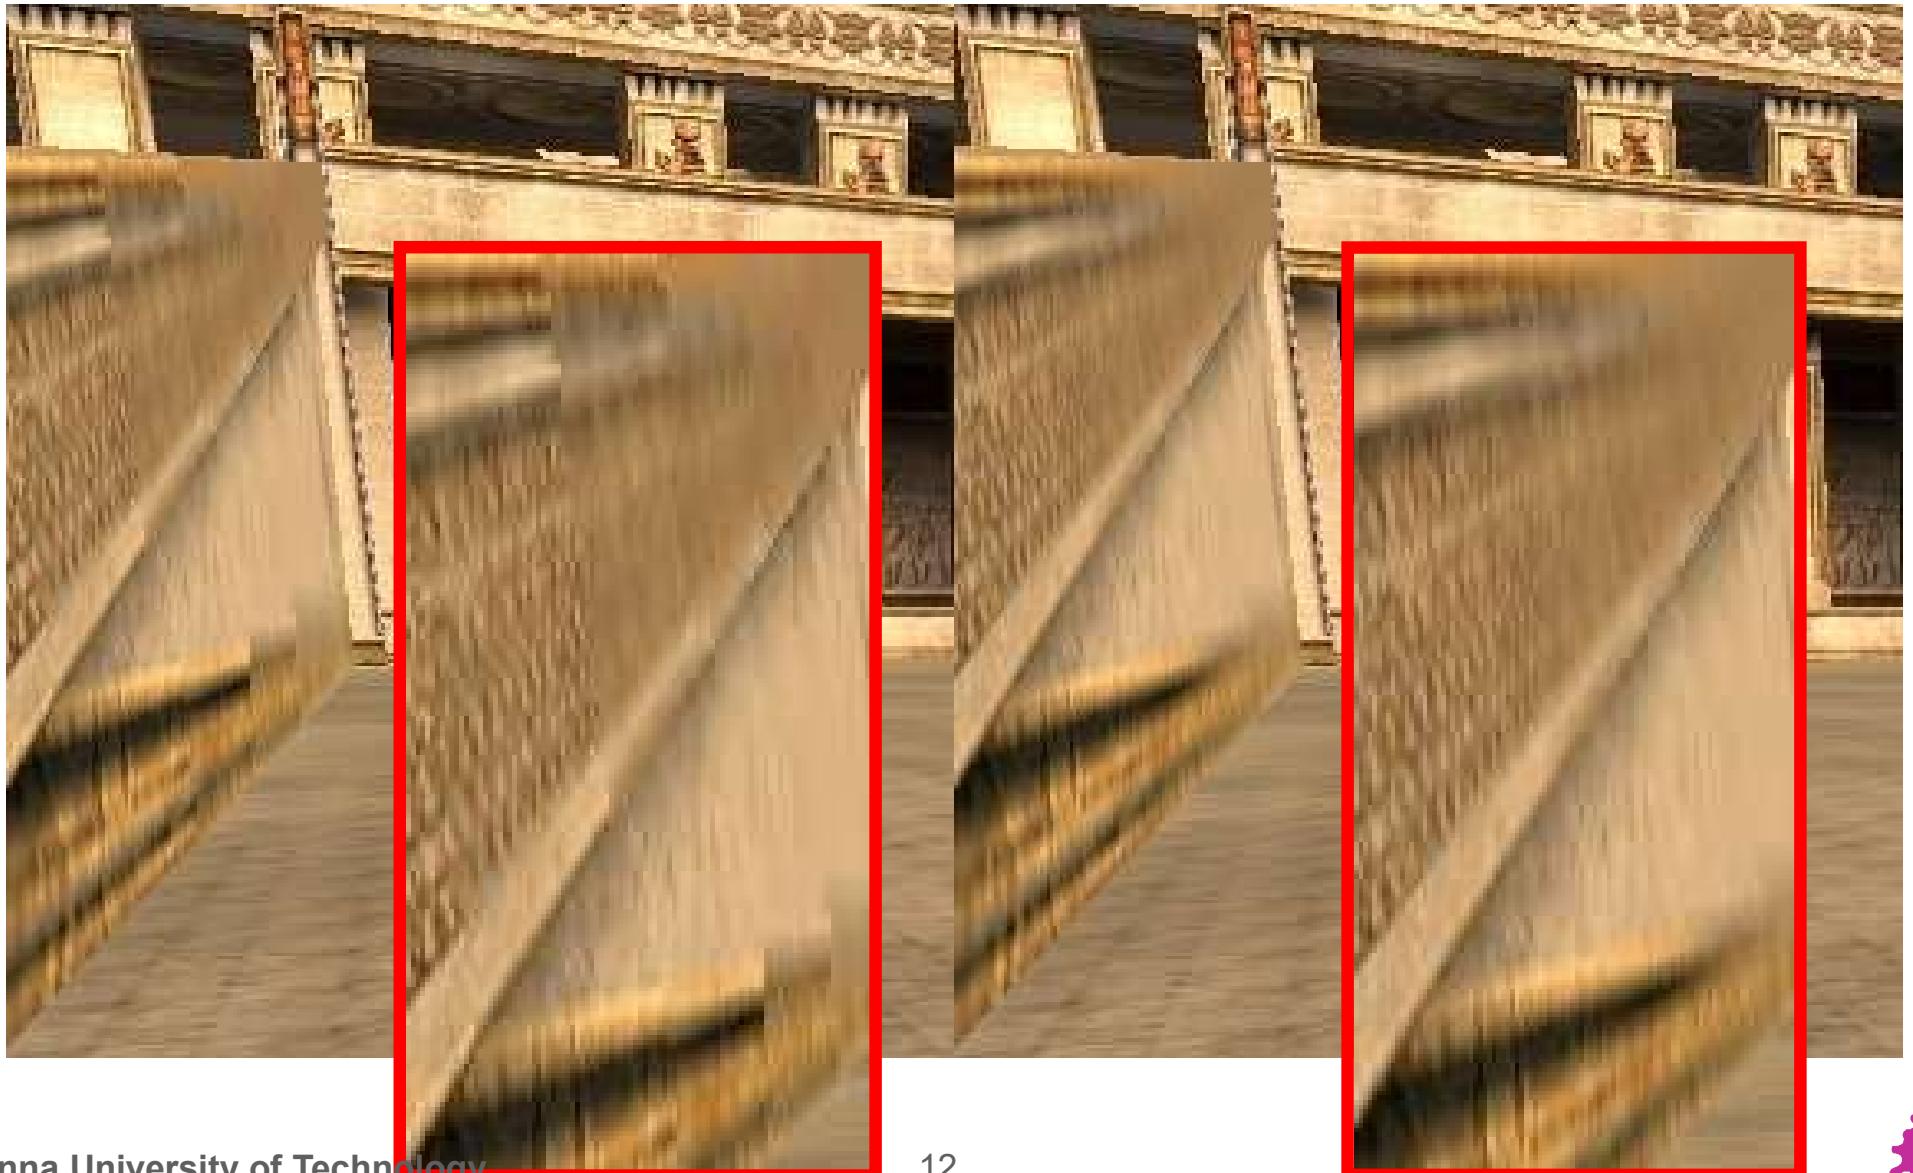
## ■ Trilinear interpolation:

- $T_1 :=$  value from texture  $D_1 = D_0 + 1$  (bilin.interpolation)
- Pixel value :=  $(D_1 - D) \cdot T_0 + (D - D_0) \cdot T_1$ 
  - Linear interpolation between successive MIP Maps
- Avoids "Mip banding" (but doubles texture lookups)



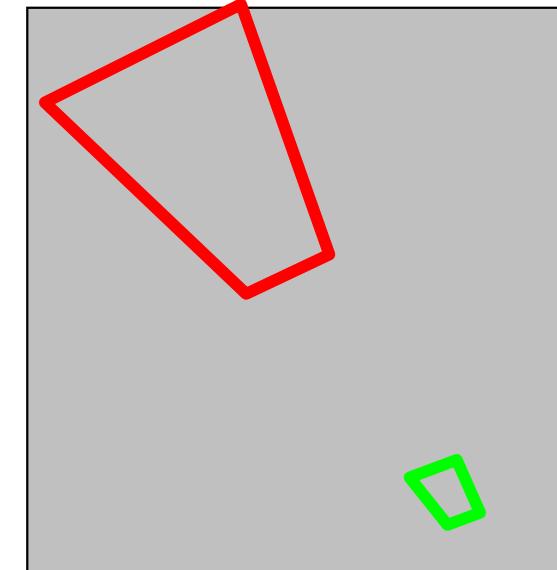
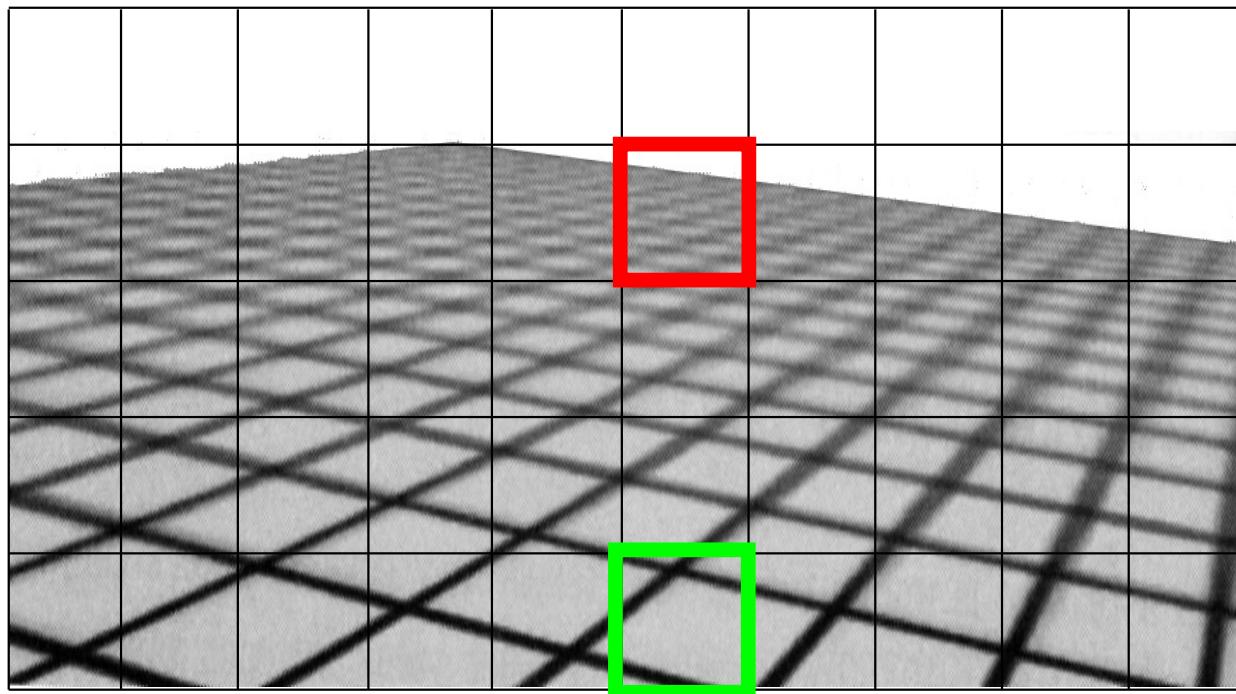
# Texture Anti-Aliasing: MIP Mapping

- Other example for bilinear vs. trilinear filtering



# Anti-Aliasing: Anisotropic Filtering

- Anisotropic filtering
  - View-dependent filter kernel
  - Implementation: *summed area table*, "*RIP Mapping*", *footprint assembly*, *elliptical weighted average* (EWA)

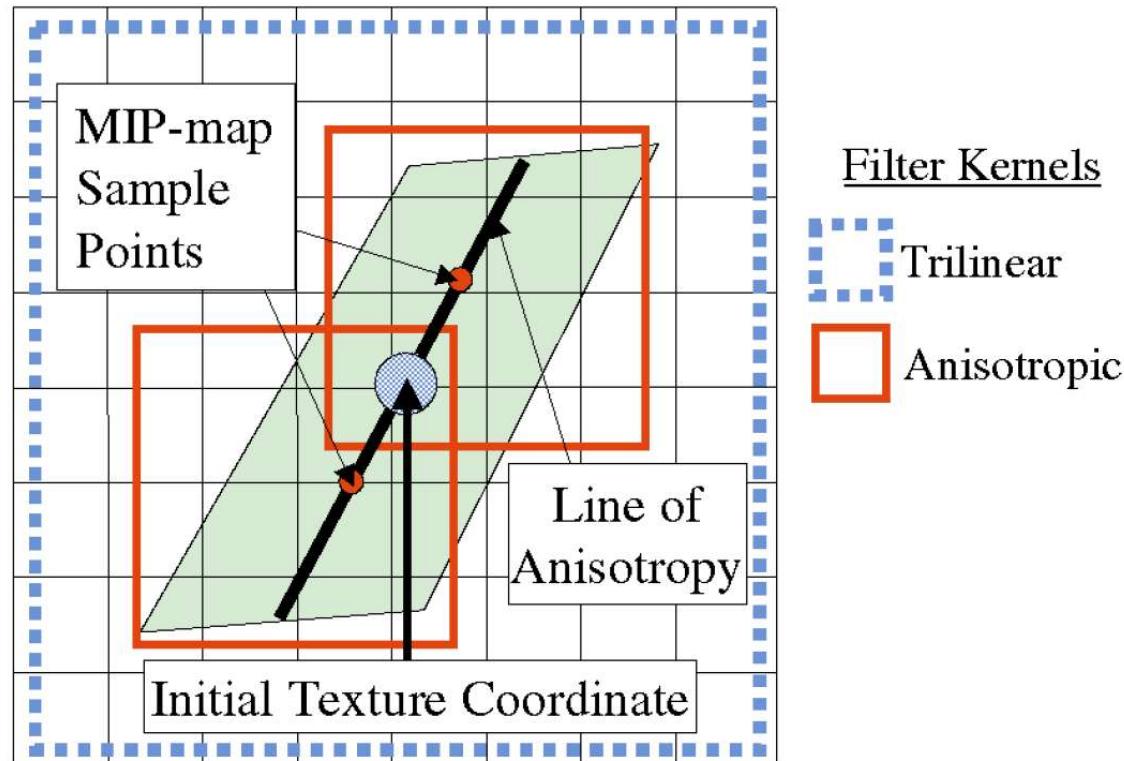


Texture space



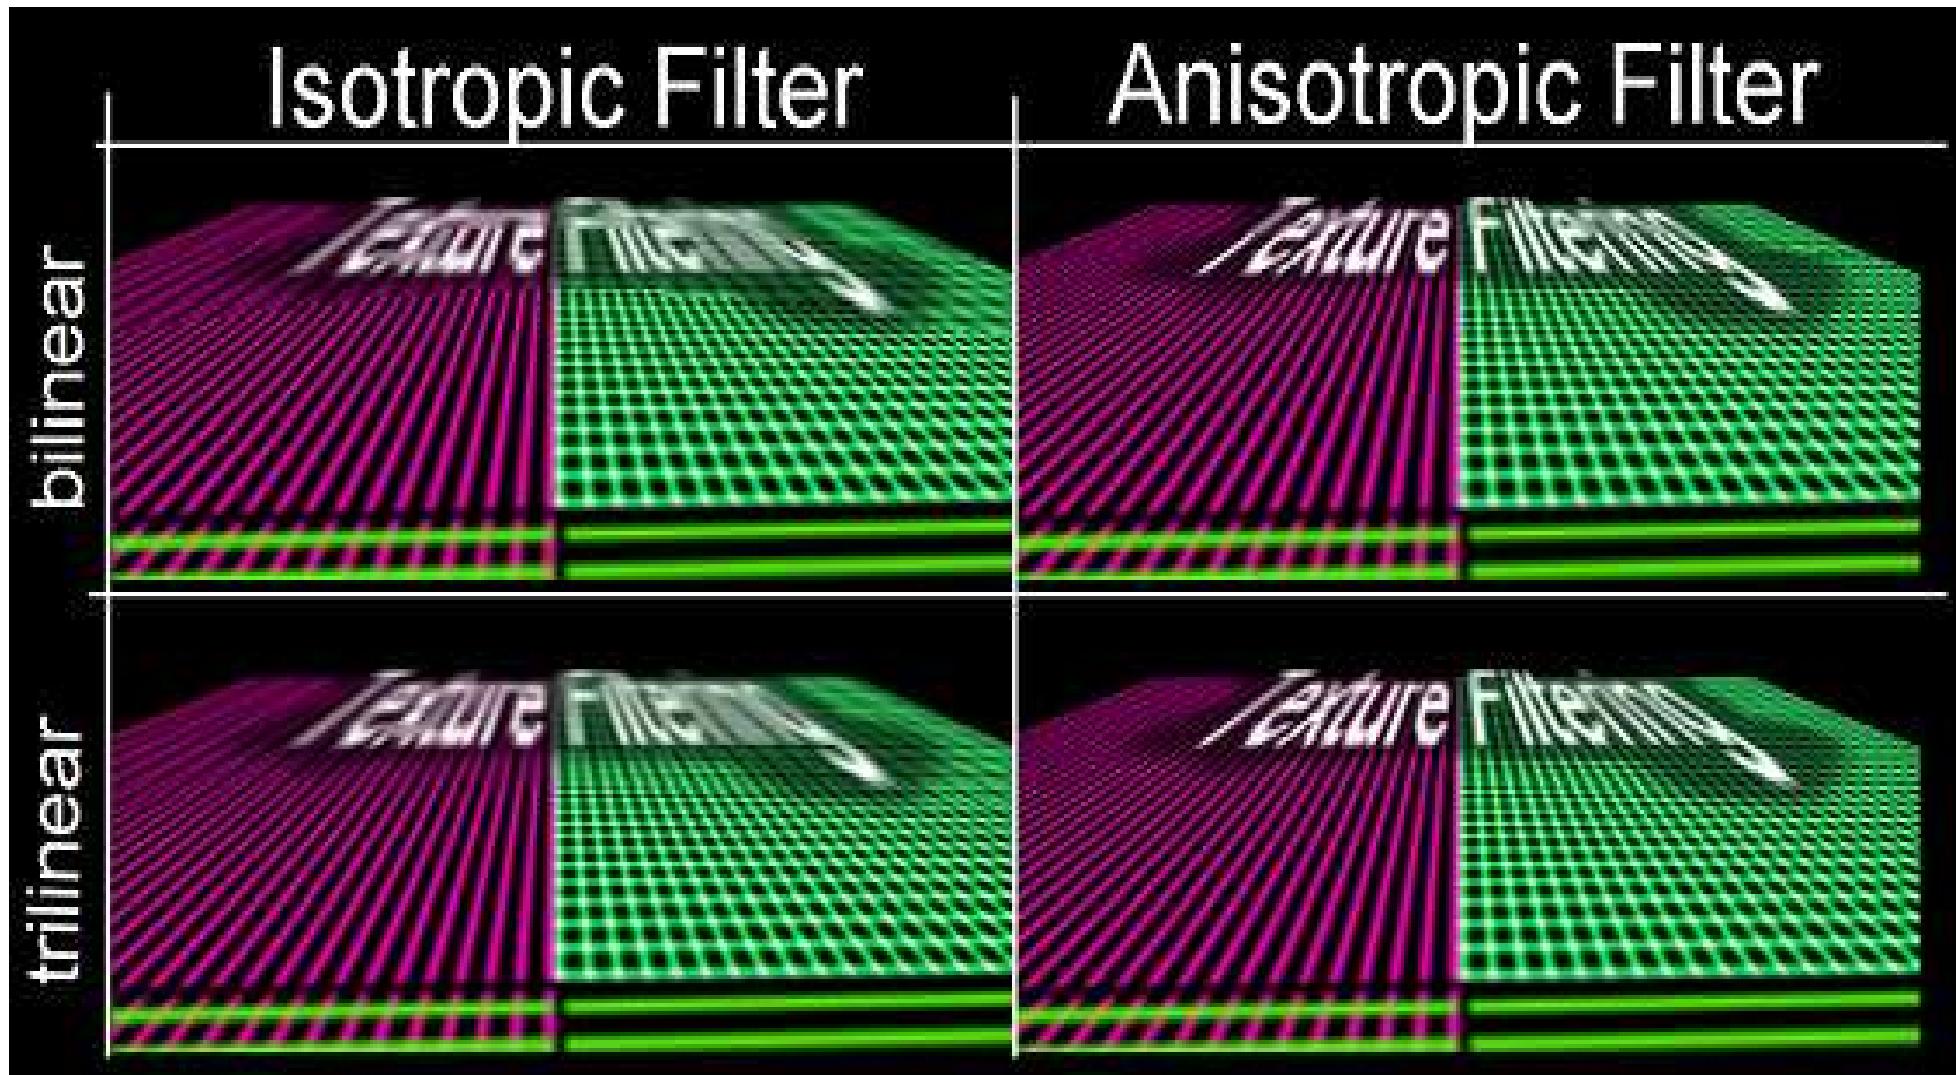


# Anisotropic Filtering: Footprint Assembly



# Anti-Aliasing: Anisotropic Filtering

## ■ Example



- Basically, everything done in hardware
- gluBuild2DMipmaps () generates MIPmaps
- Set parameters in glTexParameter()
  - GL\_TEXTURE\_MAG\_FILTER: GL\_NEAREST, GL\_LINEAR, ...
  - GL\_TEXTURE\_MIN\_FILTER: GL\_LINEAR\_MIPMAP\_NEAREST
- Anisotropic filtering is an extension:
  - GL\_EXT\_texture\_filter\_anisotropic
  - Number of samples can be varied (4x,8x,16x)
    - Vendor specific support and extensions

for Vulkan, see `vkSampler`,  
`VkSamplerCreateInfo::magFilter`, `VkSamplerCreateInfo::minFilter`,  
`VK_FILTER_NEAREST`, `VK_FILTER_LINEAR`,  
`VkSamplerCreateInfo::mipmapMode`,  
`VK_SAMPLER_MIPMAP_MODE_NEAREST`, `VK_SAMPLER_MIPMAP_MODE_LINEAR`, ...



# GPU Virtual Texturing

# Virtual Texturing



## Example #1:

### **ARB Sparse Textures (originally: AMD Partially Resident Textures)**

ARB\_sparse\_texture / ARB\_sparse\_texture2

[https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB\\_sparse\\_texture.txt](https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB_sparse_texture.txt)

[https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB\\_sparse\\_texture2.txt](https://www.khronos.org/registry/OpenGL/extensions/ARB/ARB_sparse_texture2.txt)

Hardware Virtual Texturing, Graham Sellers,  
from SIGGRAPH 2013 course “Rendering Massive Virtual Worlds”

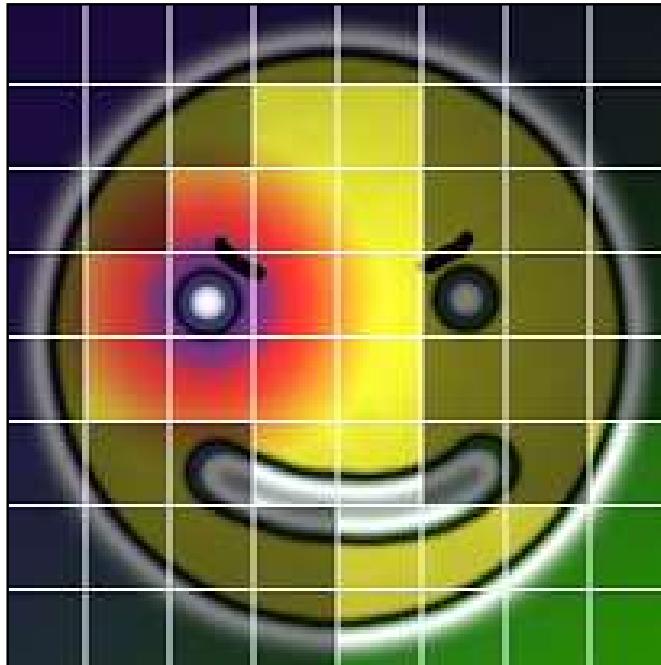
[https://cesiumjs.org/hosted-apps/massiveworlds/downloads/Graham/Hardware\\_Virtual\\_Textures.pptx](https://cesiumjs.org/hosted-apps/massiveworlds/downloads/Graham/Hardware_Virtual_Textures.pptx)



# Virtual Texturing

Divide texture up into tiles

- Commit only *used* tiles to memory
- Store data in separate physical texture



Virtual Texture

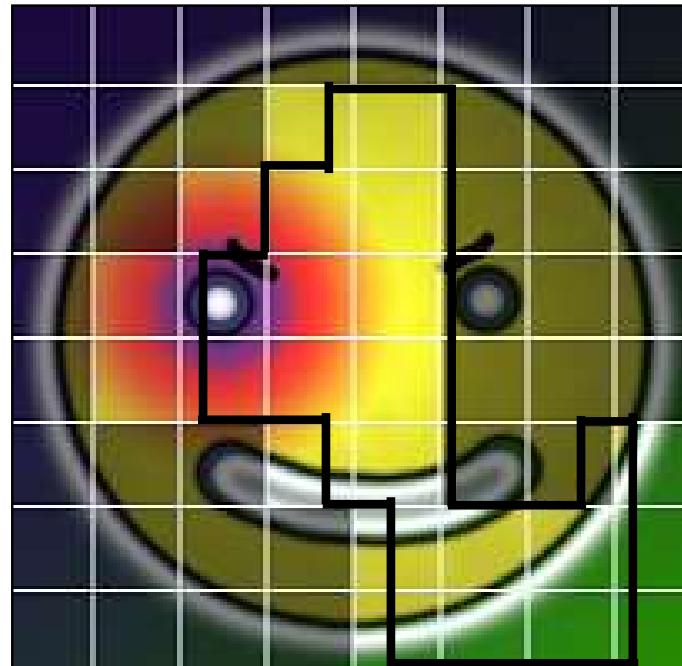


Physical Texture



# Virtual Texturing

Memory requirements set by number of resident tiles, not texture dimensions



RGBA8, 1024x1024, 64 tiles

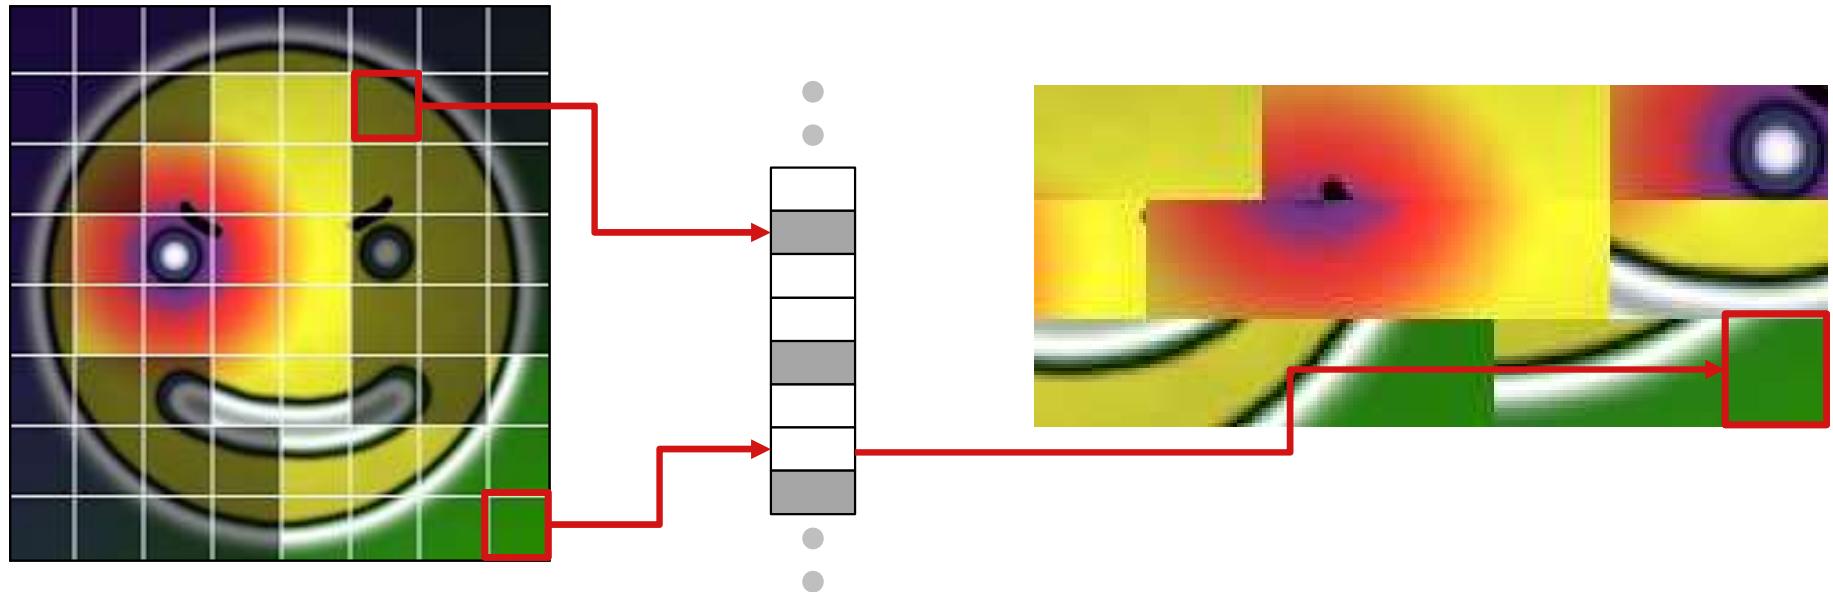
	<b>Virtual</b>	<b>Physical</b>
Memory	4096 kB	1536 kB



# Virtual Texturing

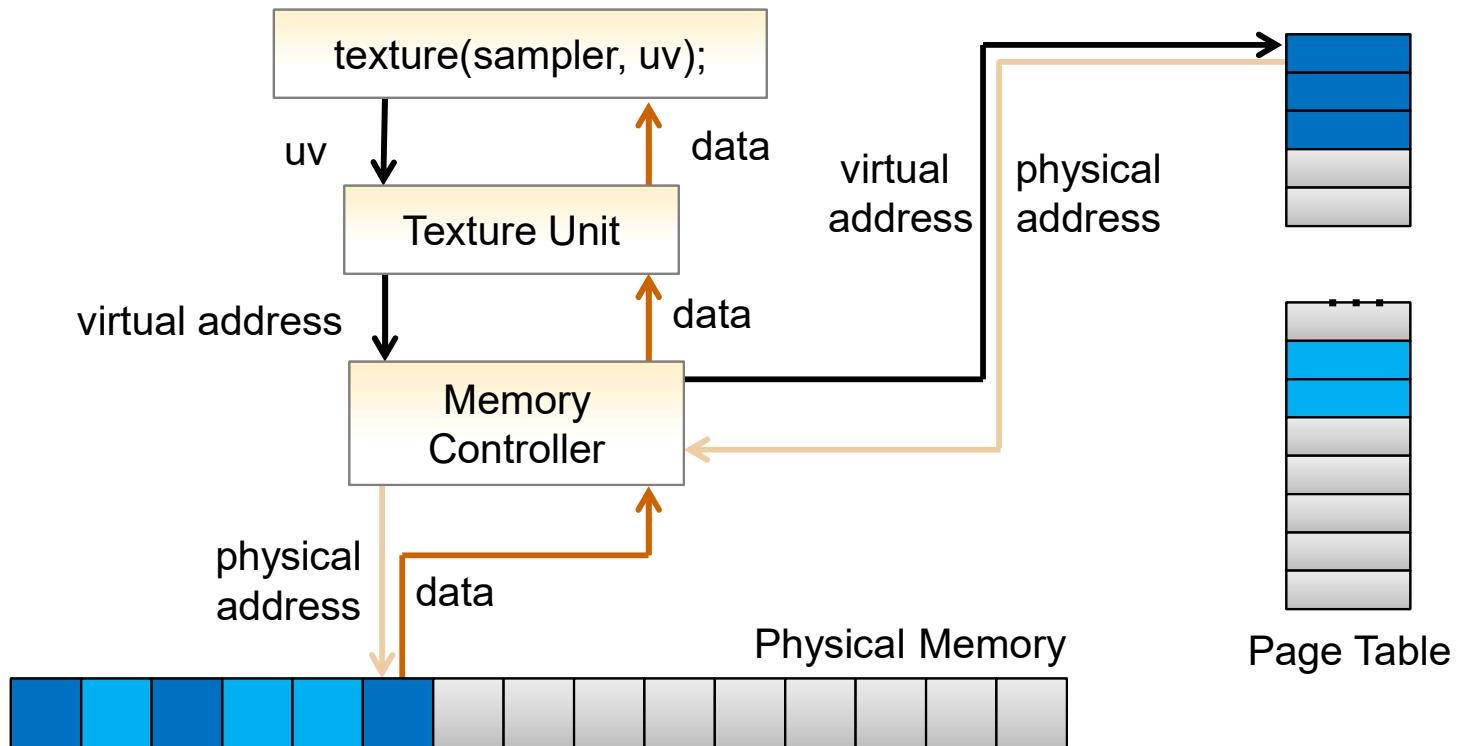
Use indirection table to map virtual to physical

- This is also known as a *page table*





# GPU Virtual Memory



# Summary (Shader vs. Full Hardware Support)



	SVTs	HVTs
Address translation	Shader code	HW page table
Filtering	HW + shader code	HW only
# of texture fetches	2, dependent	1
Supported formats	The ones implemented	All supported by HW
Supported texture types	The ones implemented	All supported by HW



## Example #2:

### Adaptive Shadow Maps (ASM)

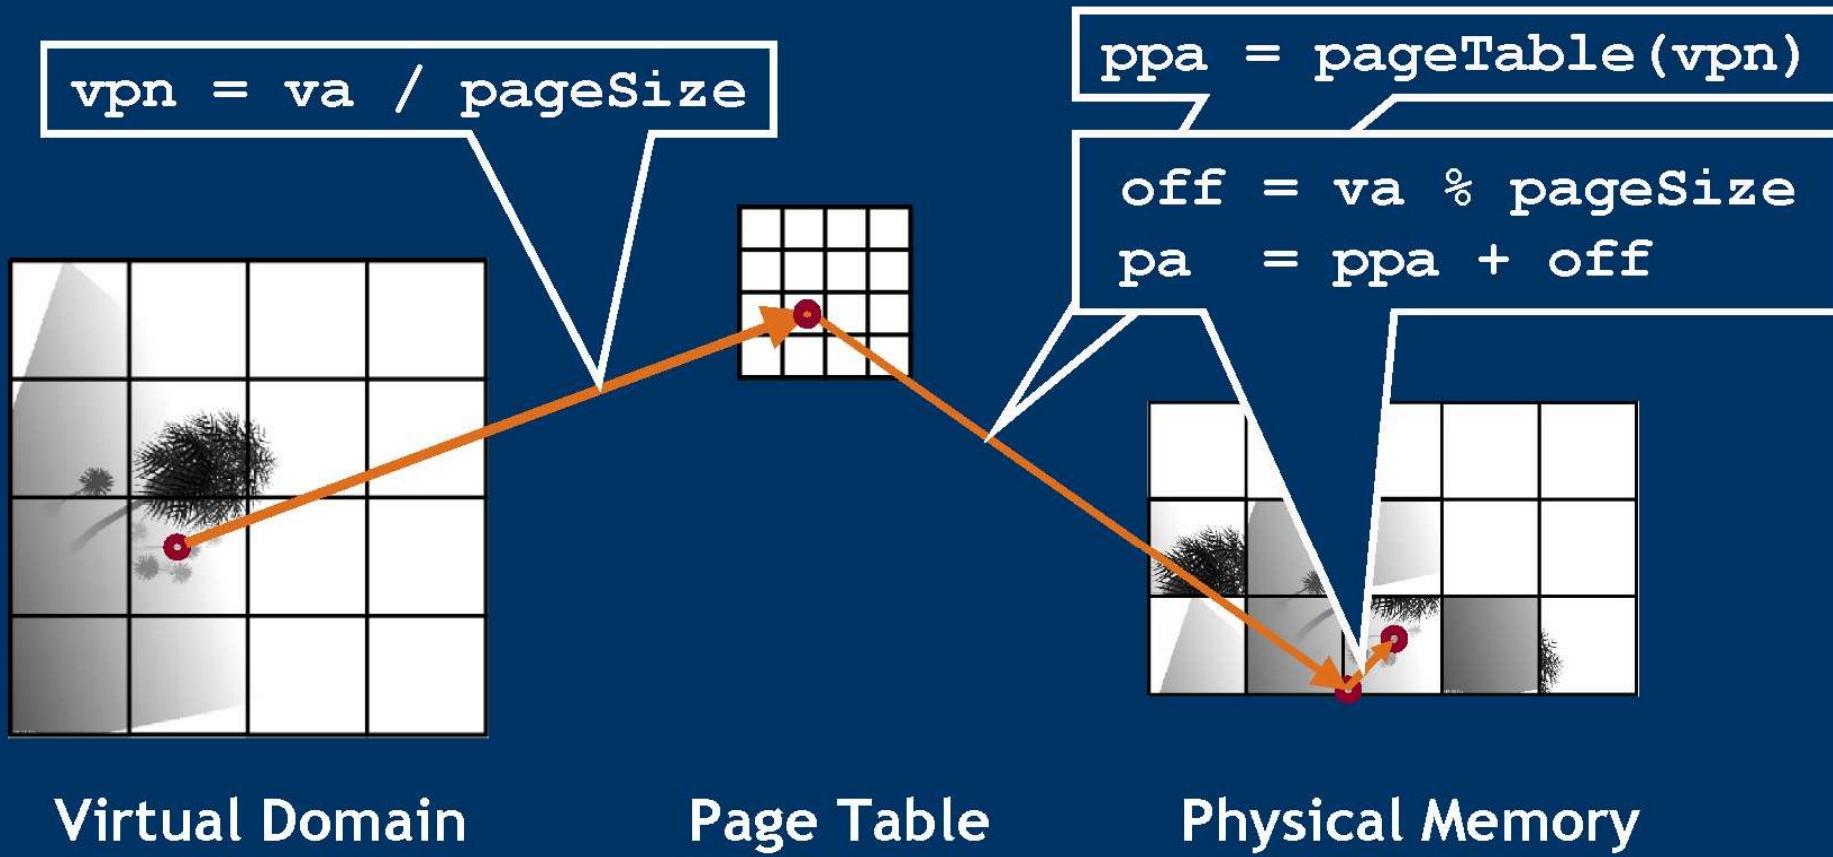
- On CPUs: Fernando et al., ACM SIGGRAPH 2001

### Resolution-Matched Shadow Maps

- On GPUs: Aaron Lefohn et al., ACM Transactions on Graphics 2007

# ASM Data Structure (Adaptive Shadow Maps)

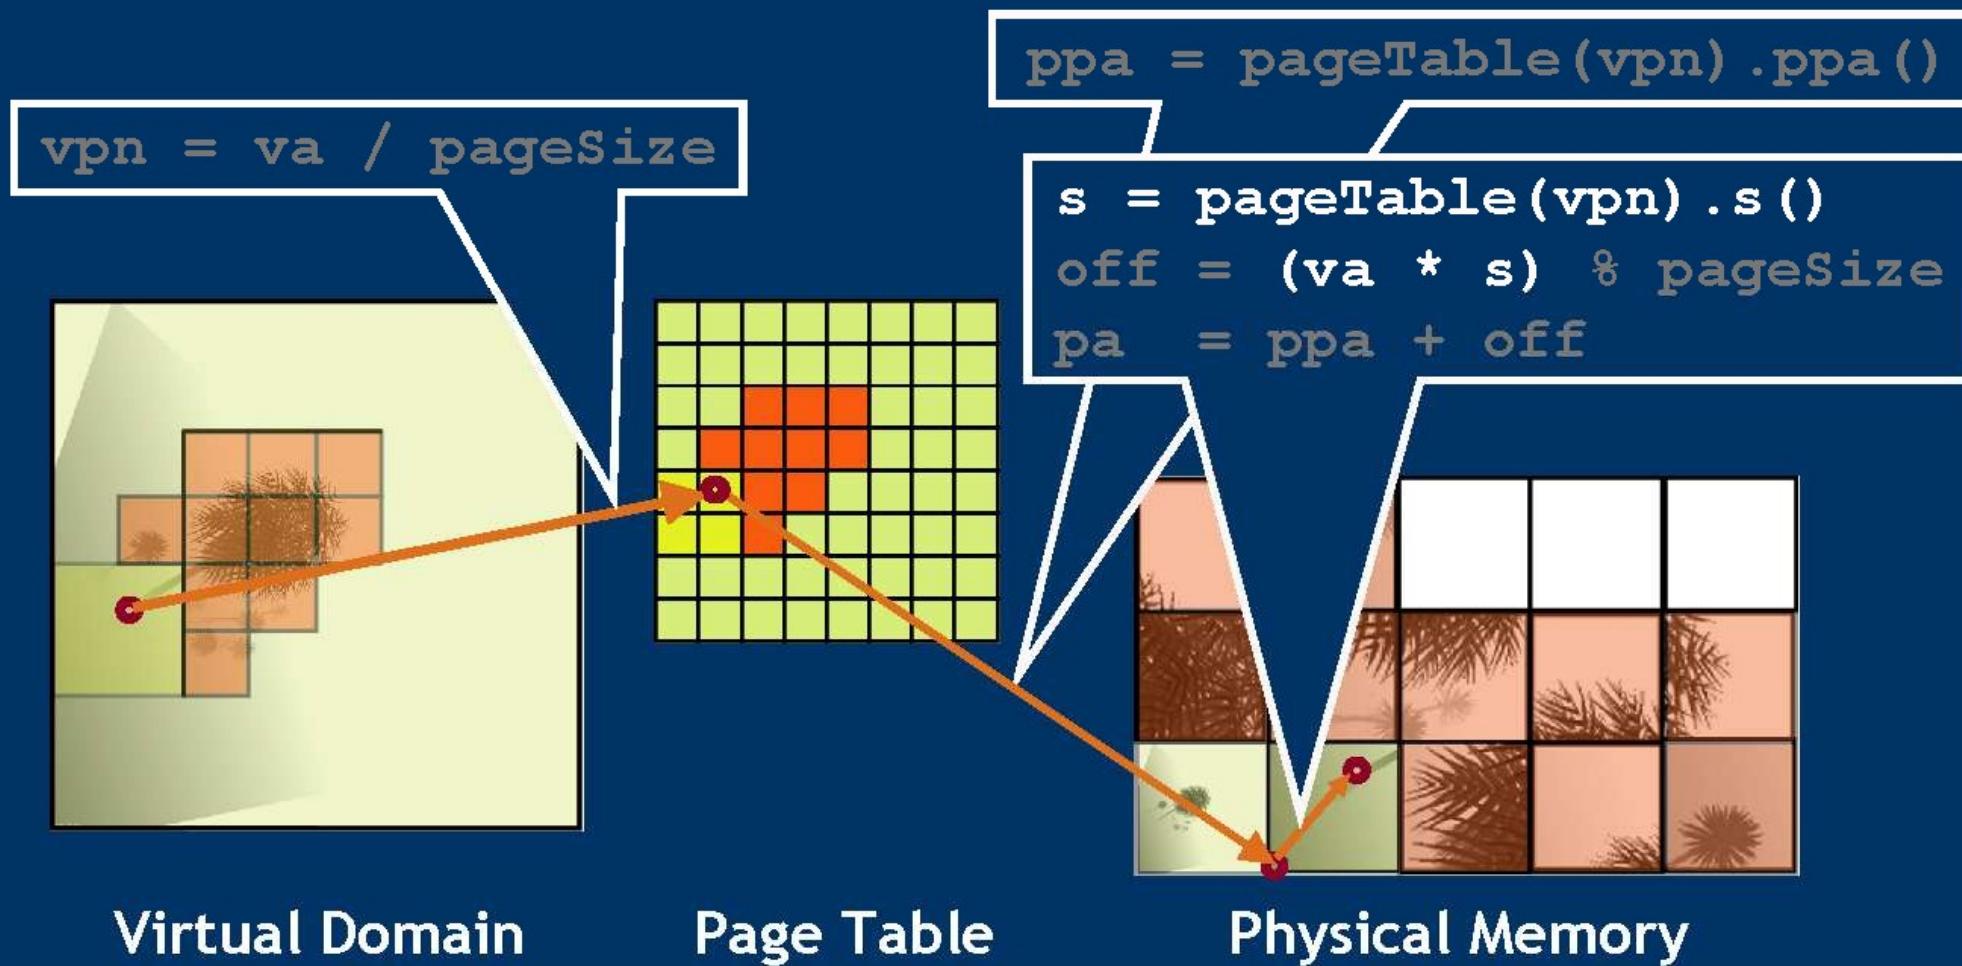
- Page table example



# ASM Data Structure (Adaptive Shadow Maps)

- Adaptive Page Table

- Map multiple virtual pages to single physical page



Virtual Domain

Page Table

Physical Memory





## Example #3:

**id Tech 5 Megatextures, id Software**

**Rage**

- Virtual Texturing in Software and Hardware, van Waveren et al., SIGGRAPH 2012 course notes + slides

[http://www.jurajjobert.com/data/Virtual\\_Texturing\\_in\\_Software\\_and\\_Hardware\\_course\\_notes.pdf](http://www.jurajjobert.com/data/Virtual_Texturing_in_Software_and_Hardware_course_notes.pdf)

<http://www.mrelusive.com/publications/papers/Software-Virtual-Textures.pdf>

[http://www.mrelusive.com/publications/presentations/2013\\_siggraph/hq\\_sw\\_hw\\_vts\\_12.pdf](http://www.mrelusive.com/publications/presentations/2013_siggraph/hq_sw_hw_vts_12.pdf)



# Virtual Texturing



Rage / id Tech 5 (id Software)

# Virtual Texturing

- Unique, very large virtual textures key to id tech 5 rendering
- Full description beyond the scope of this talk



# Virtual Texturing

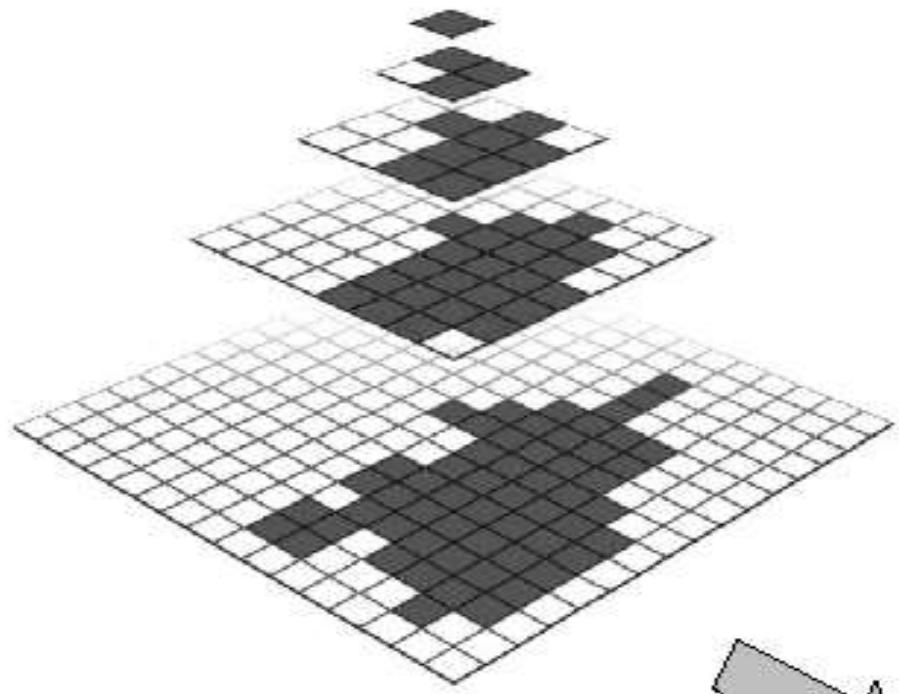


# Virtual Texturing

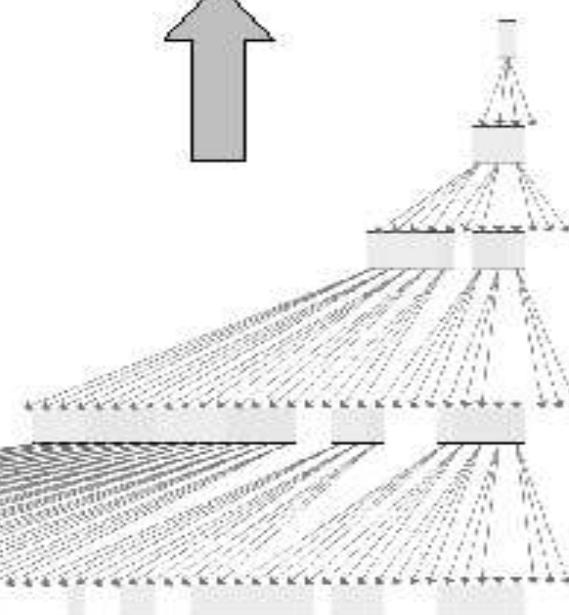
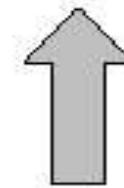
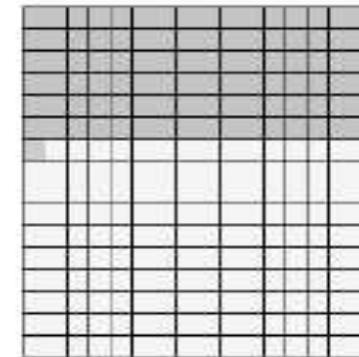


# Virtual Texturing

Texture Pyramid with Sparse Page Residency



Physical Page Texture



Quad-tree of Sparse Texture Pyramid

# Virtual Texturing



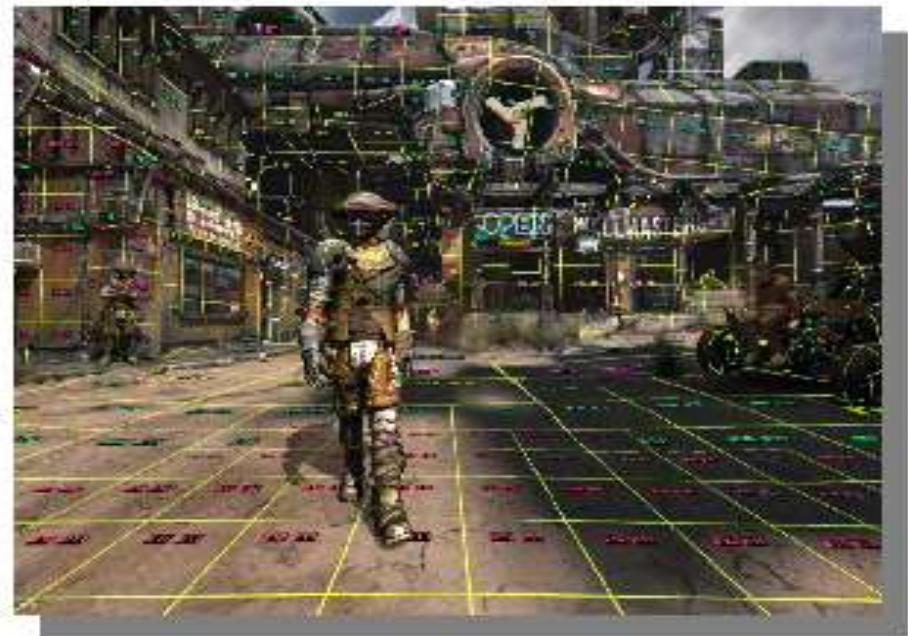
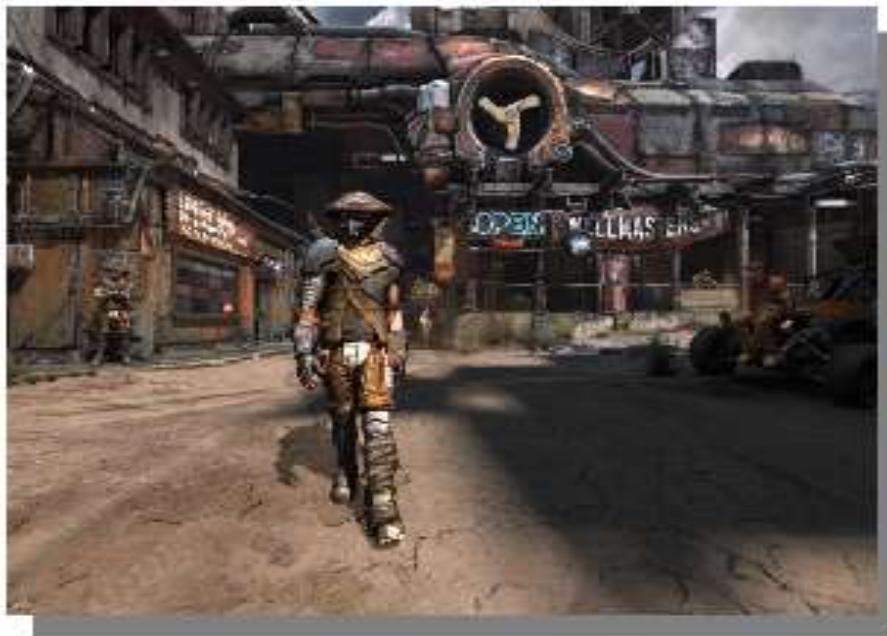
# Virtual Texturing



# Virtual Texturing

A few interesting issues...

- Texture filtering
- Thrashing due to physical memory oversubscription
- LOD transitions under high latency



Rage64

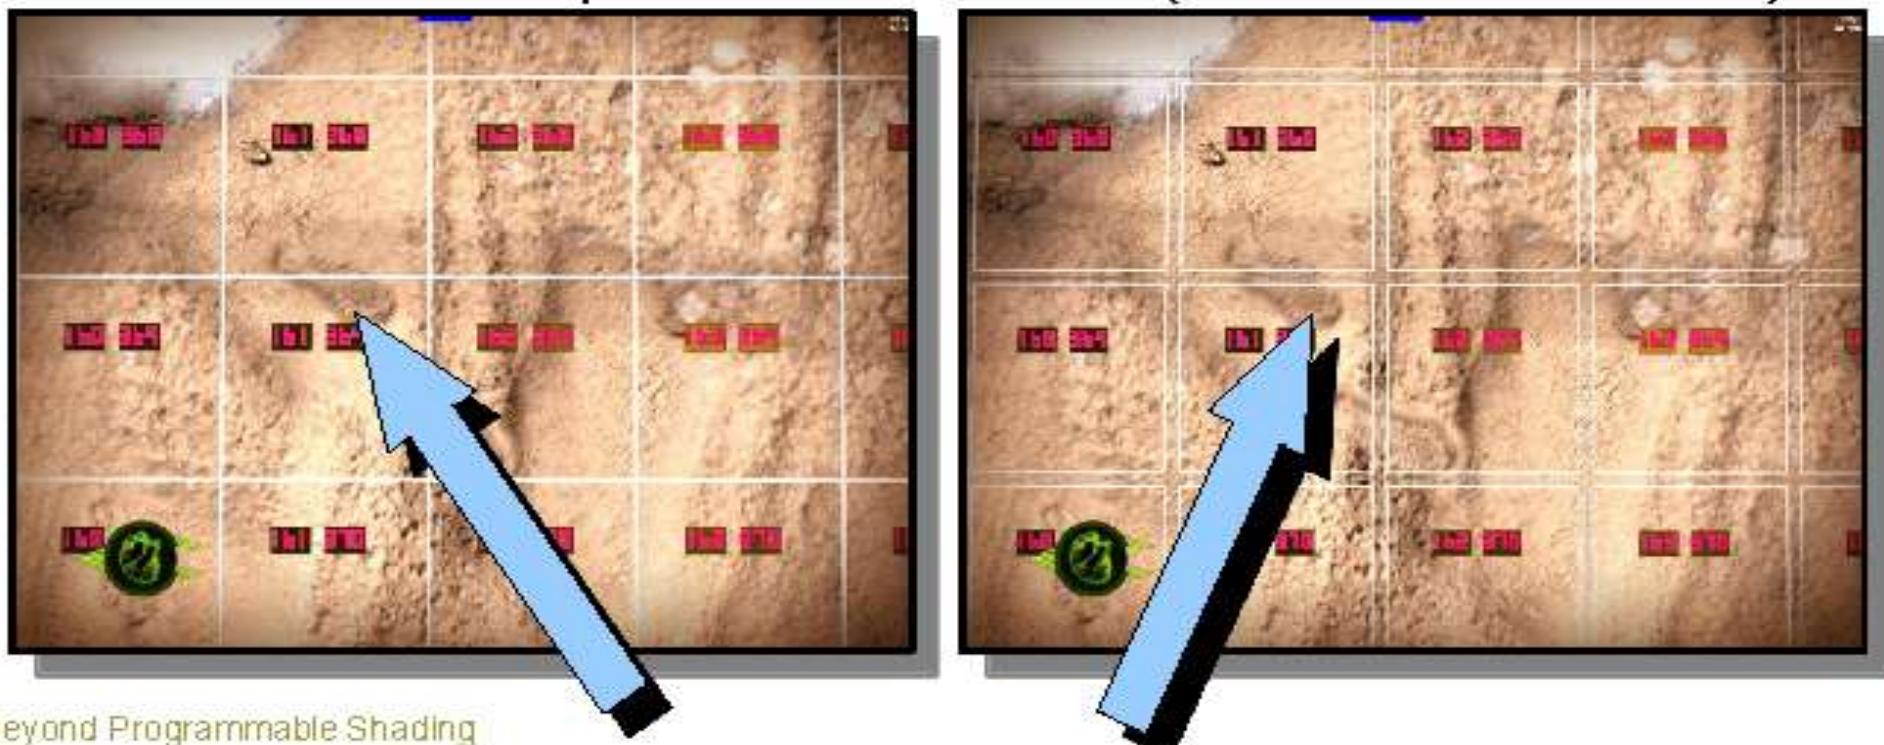
60fps  
rs-off  
312.10MB



RAGE with PRTs (Image courtesy of id Software)

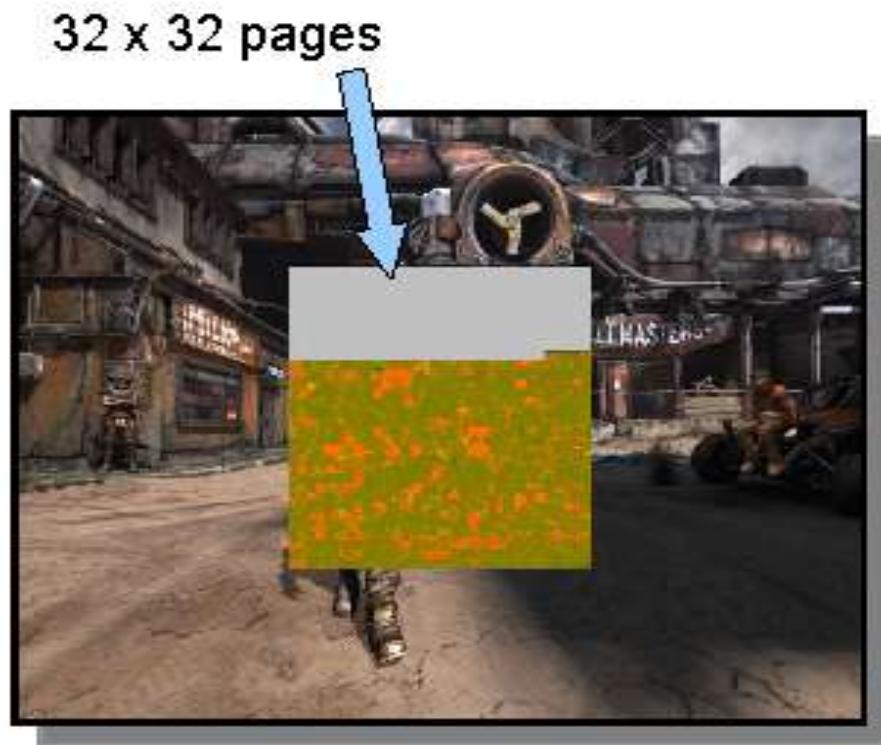
# Virtual Texturing - Filtering

- We tried no filtering at all
- We tried bilinear filtering without borders
- Bilinear filtering with border works well
- Trilinear filtering reasonably but still expensive
- Anisotropic filtering possible via TXD (texgrad)
  - 4-texel border necessary (max aniso = 4)
  - TEX with implicit derivs ok too (on some hardware)

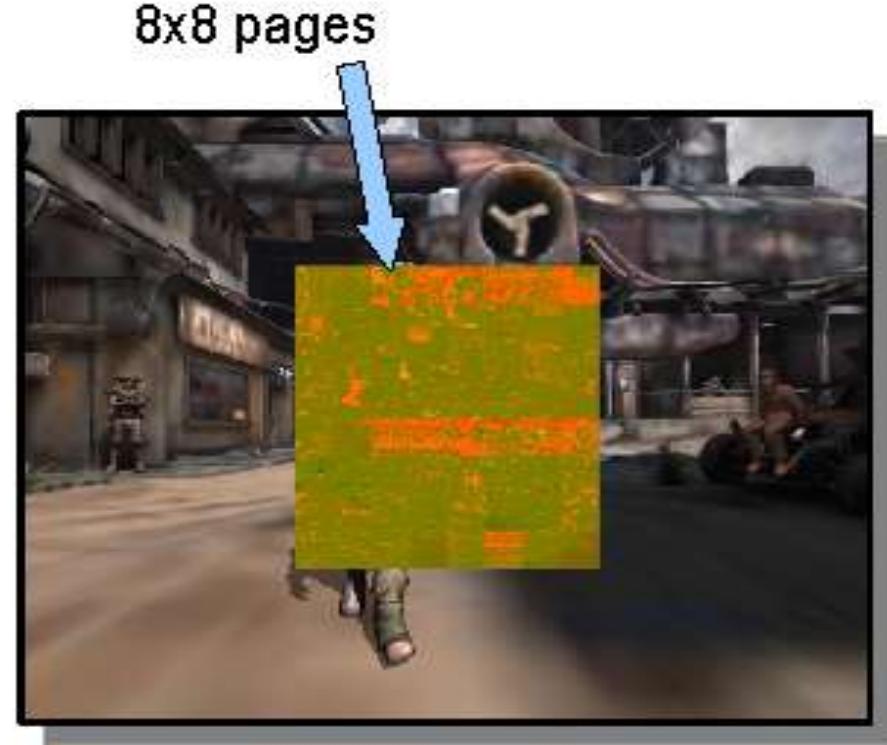


# Virtual Texturing - Thrashing

- Sometimes you need more physical pages than you have
- With conventional virtual memory, you must thrash
- With virtual texturing, you can globally adjust feedback LOD bias until working set fits



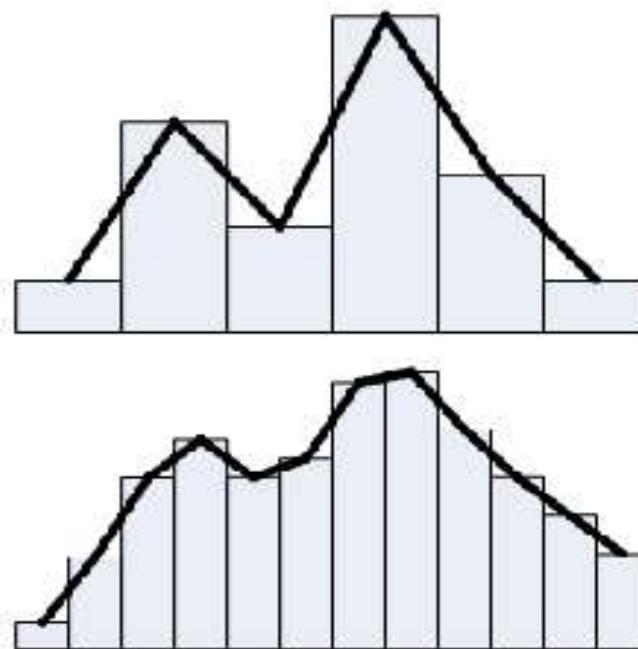
1024 Physical Pages



64 Physical Pages

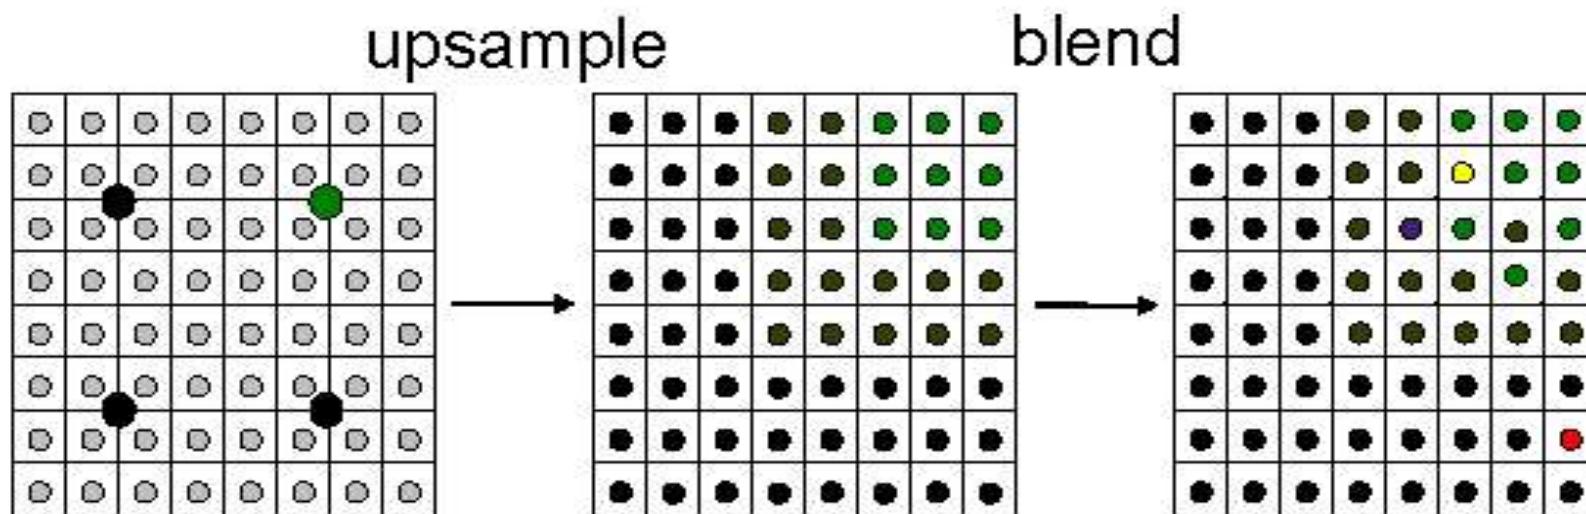
# Virtual Texturing – LOD Snap

- Latency between first need and availability can be high
  - Especially if optical disk read required (>100 msec seek!)
- Visible snap happens when magnified texture changes LOD
- If we used trilinear filtering, blending in detail would be easy
- Instead continuously update physical pages with blended data



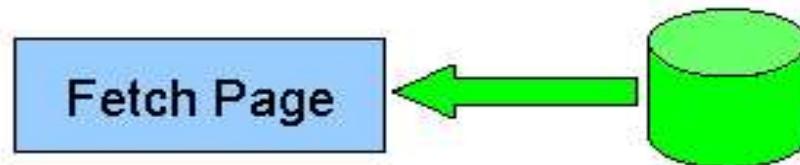
# Virtual Texturing – LOD Snap

- Upsample coarse page immediately
- Then blend in finer data when available



# Virtual Texturing - Management

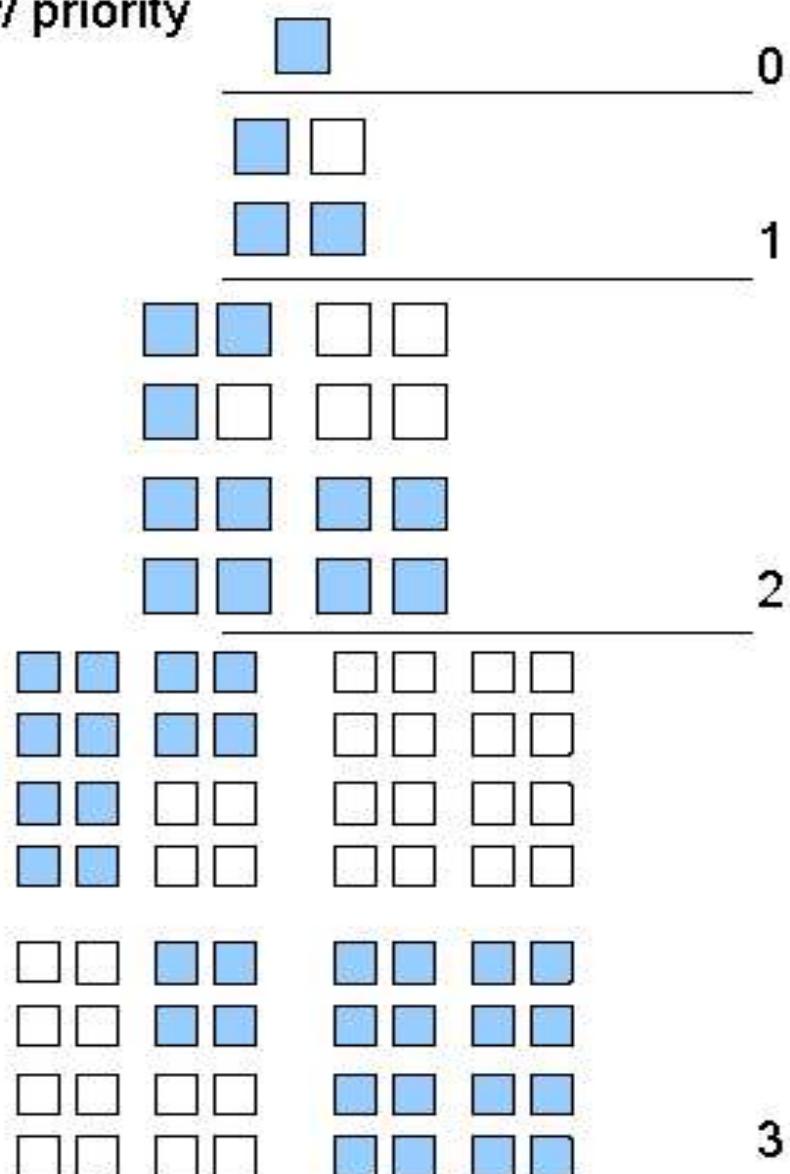
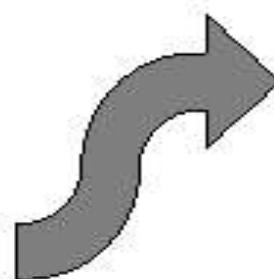
- Analysis tells us what pages we need
- We fetch what we can



- But this is a real-time app... so no blocking allowed
- Cache handles hits, schedules misses to load in background
- Resident pages managed independent of disk cache
- Physical pages organized as quad-tree per virtual texture
- Linked lists for free, LRU, and locked pages

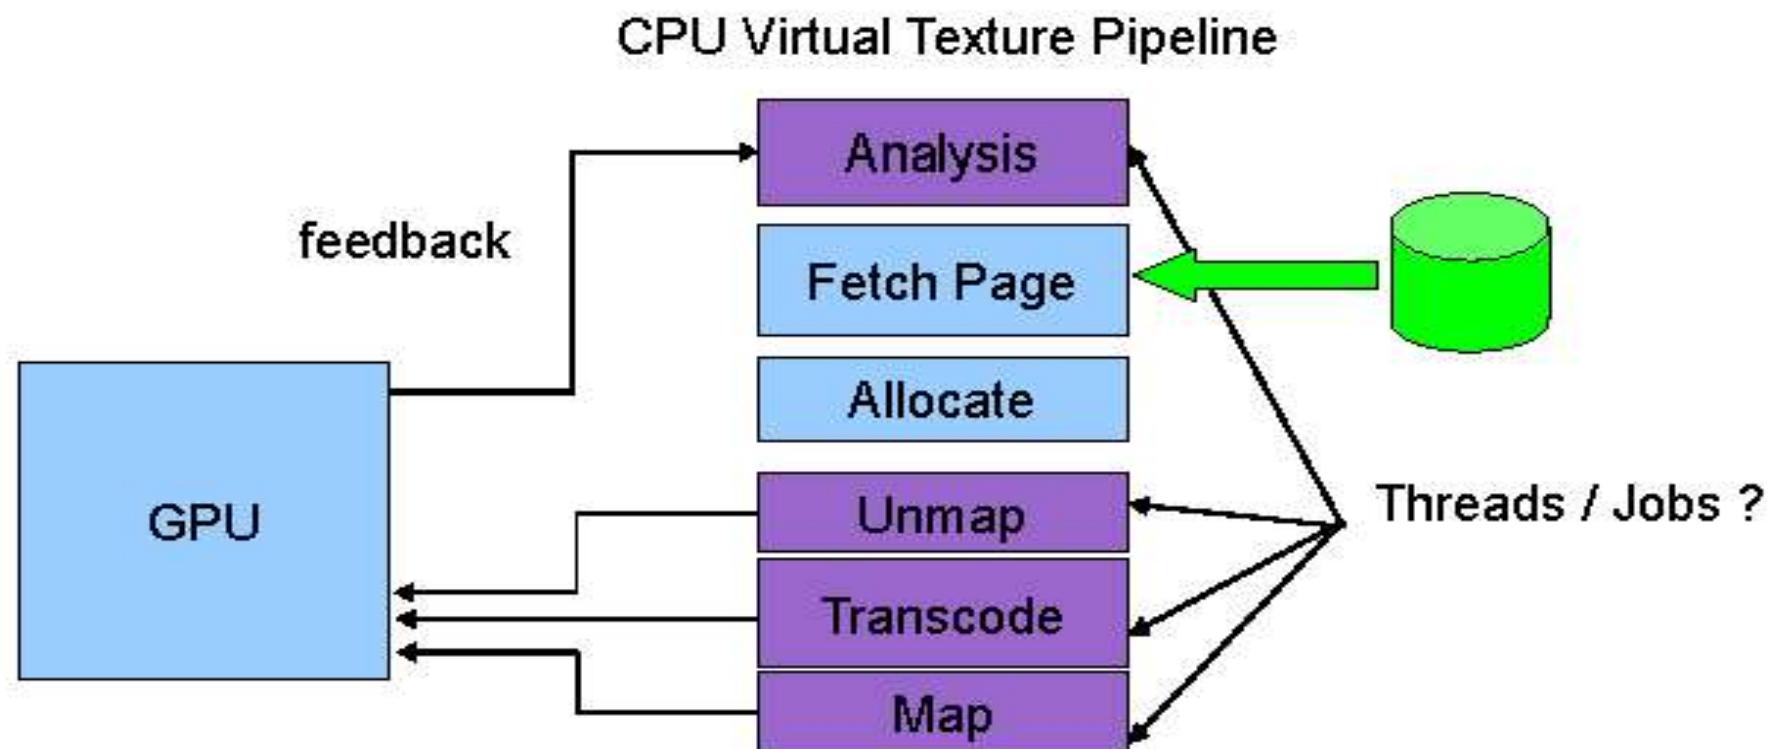
# Virtual Texturing - Feedback

- Feedback Analysis
  - Gen ~breadth-first quad-tree order w/ priority



# Virtual Texturing - Pipeline

- Compute intensive complex system with dependencies that we want to run in parallel on all the different platforms



# Virtual Texturing



## Example #4:

### Unreal Engine 5 Virtual Texturing

<https://dev.epicgames.com/documentation/en-us/unreal-engine/virtual-texturing-in-unreal-engine>

<https://dev.epicgames.com/documentation/en-us/unreal-engine/streaming-virtual-texturing-in-unreal-engine>

<https://dev.epicgames.com/documentation/en-us/unreal-engine/virtual-texture-memory-pools-in-unreal-engine>

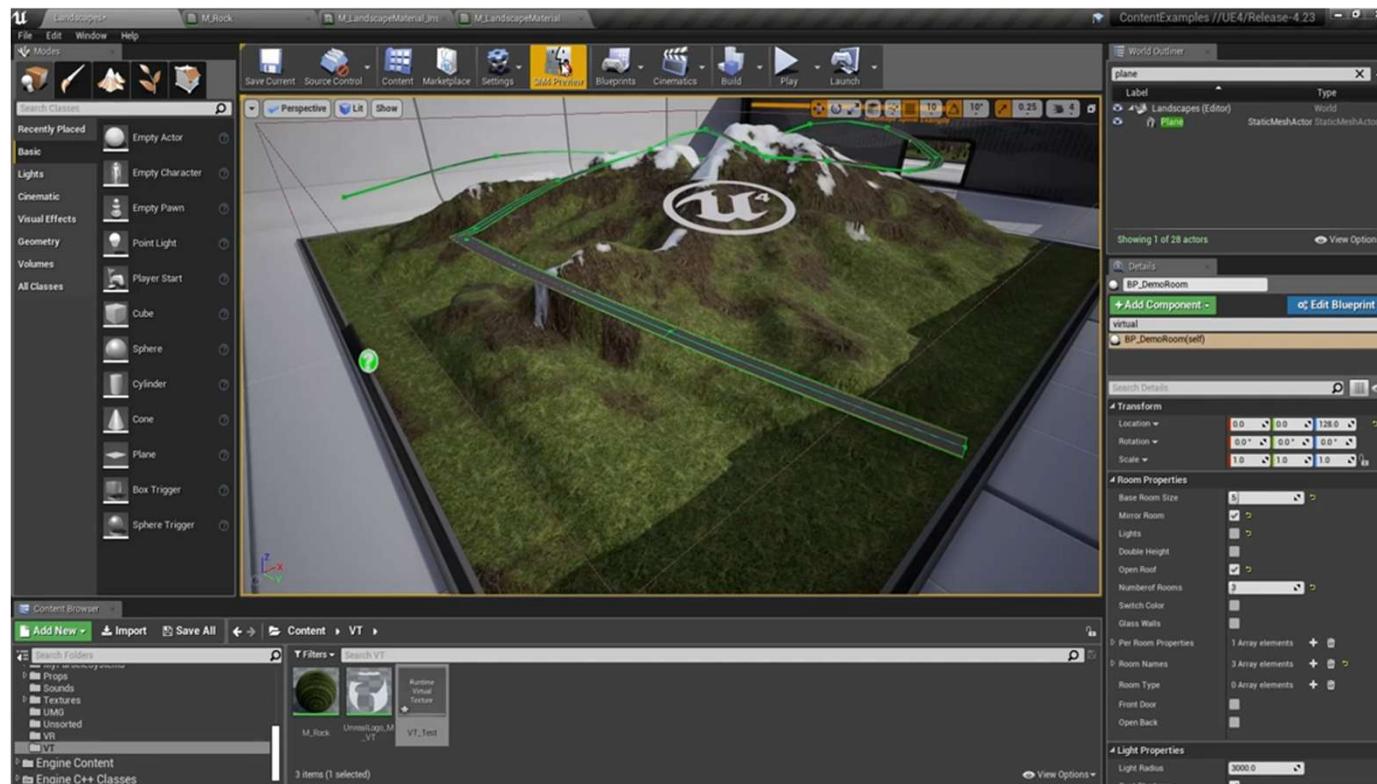
<https://dev.epicgames.com/documentation/en-us/unreal-engine/virtual-texturing-settings-and-properties-in-unreal-engine>

# Unreal Engine 5 Virtual Texturing



## Run-Time Virtual Texturing

- Procedurally generate hires texture content; materials, etc.

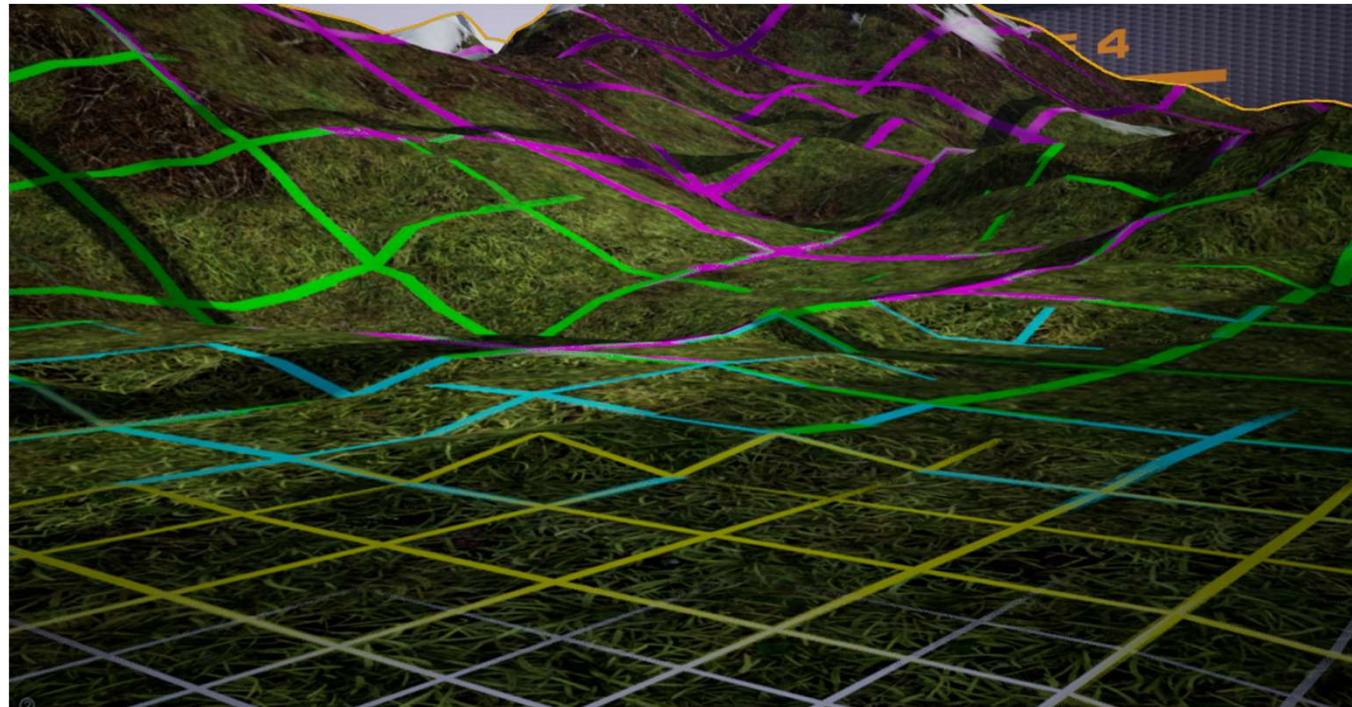




# Unreal Engine 5 Virtual Texturing

## Streaming Virtual Texturing

- Stream texture data from disk; default tile size 128, up to 4K tiles
- Also: pre-baked hires light maps, etc.



# Unreal Engine 5 Virtual Texture Memory Pools



One pool for virtual textures of each type/format

- Can use *residency mipmap bias* (mipmap LOD bias)



# Virtual Texturing



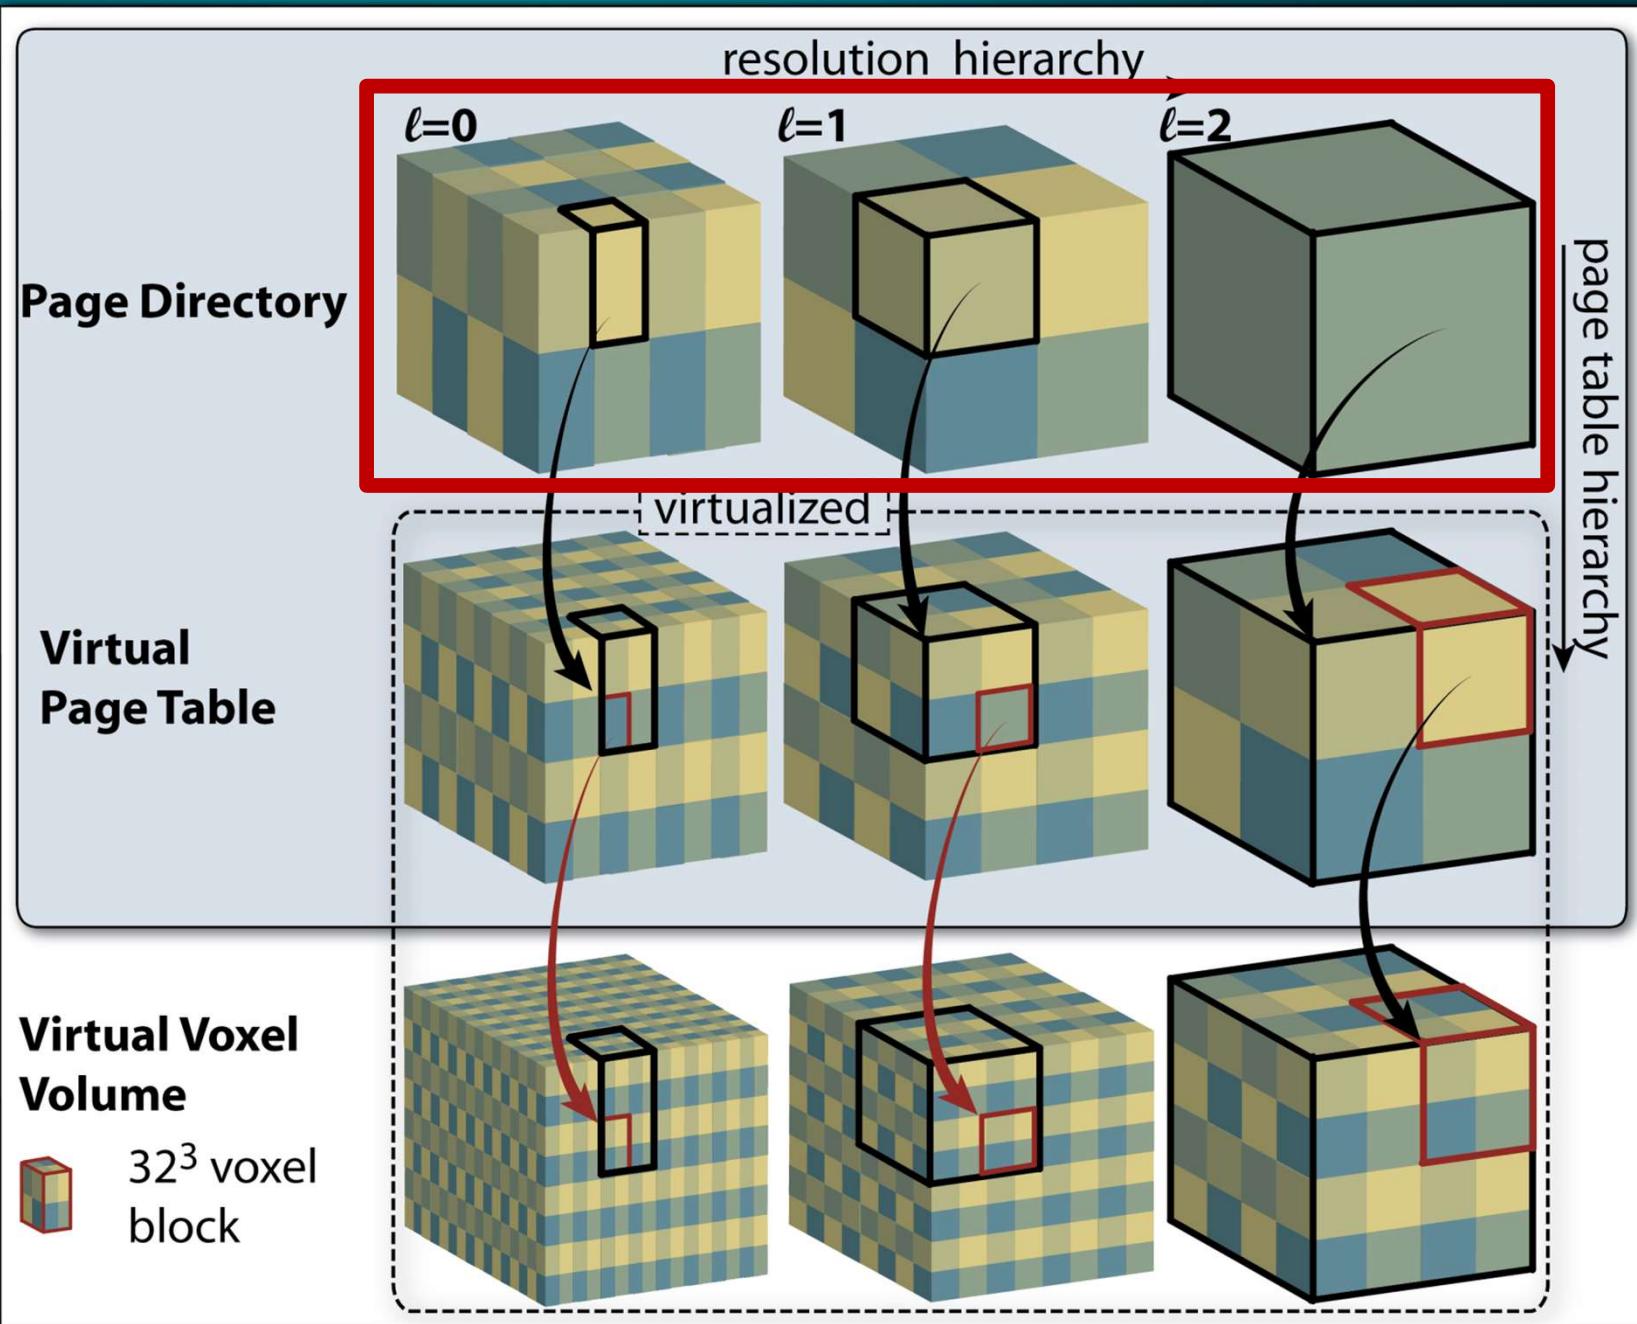
## Example #5:

### Petascale Volume Rendering

- Interactive Volume Exploration of Petascale Microscopy Data Streams Using a Visualization-Driven Virtual Memory Approach,  
Hadwiger et al., IEEE SciVis 2012

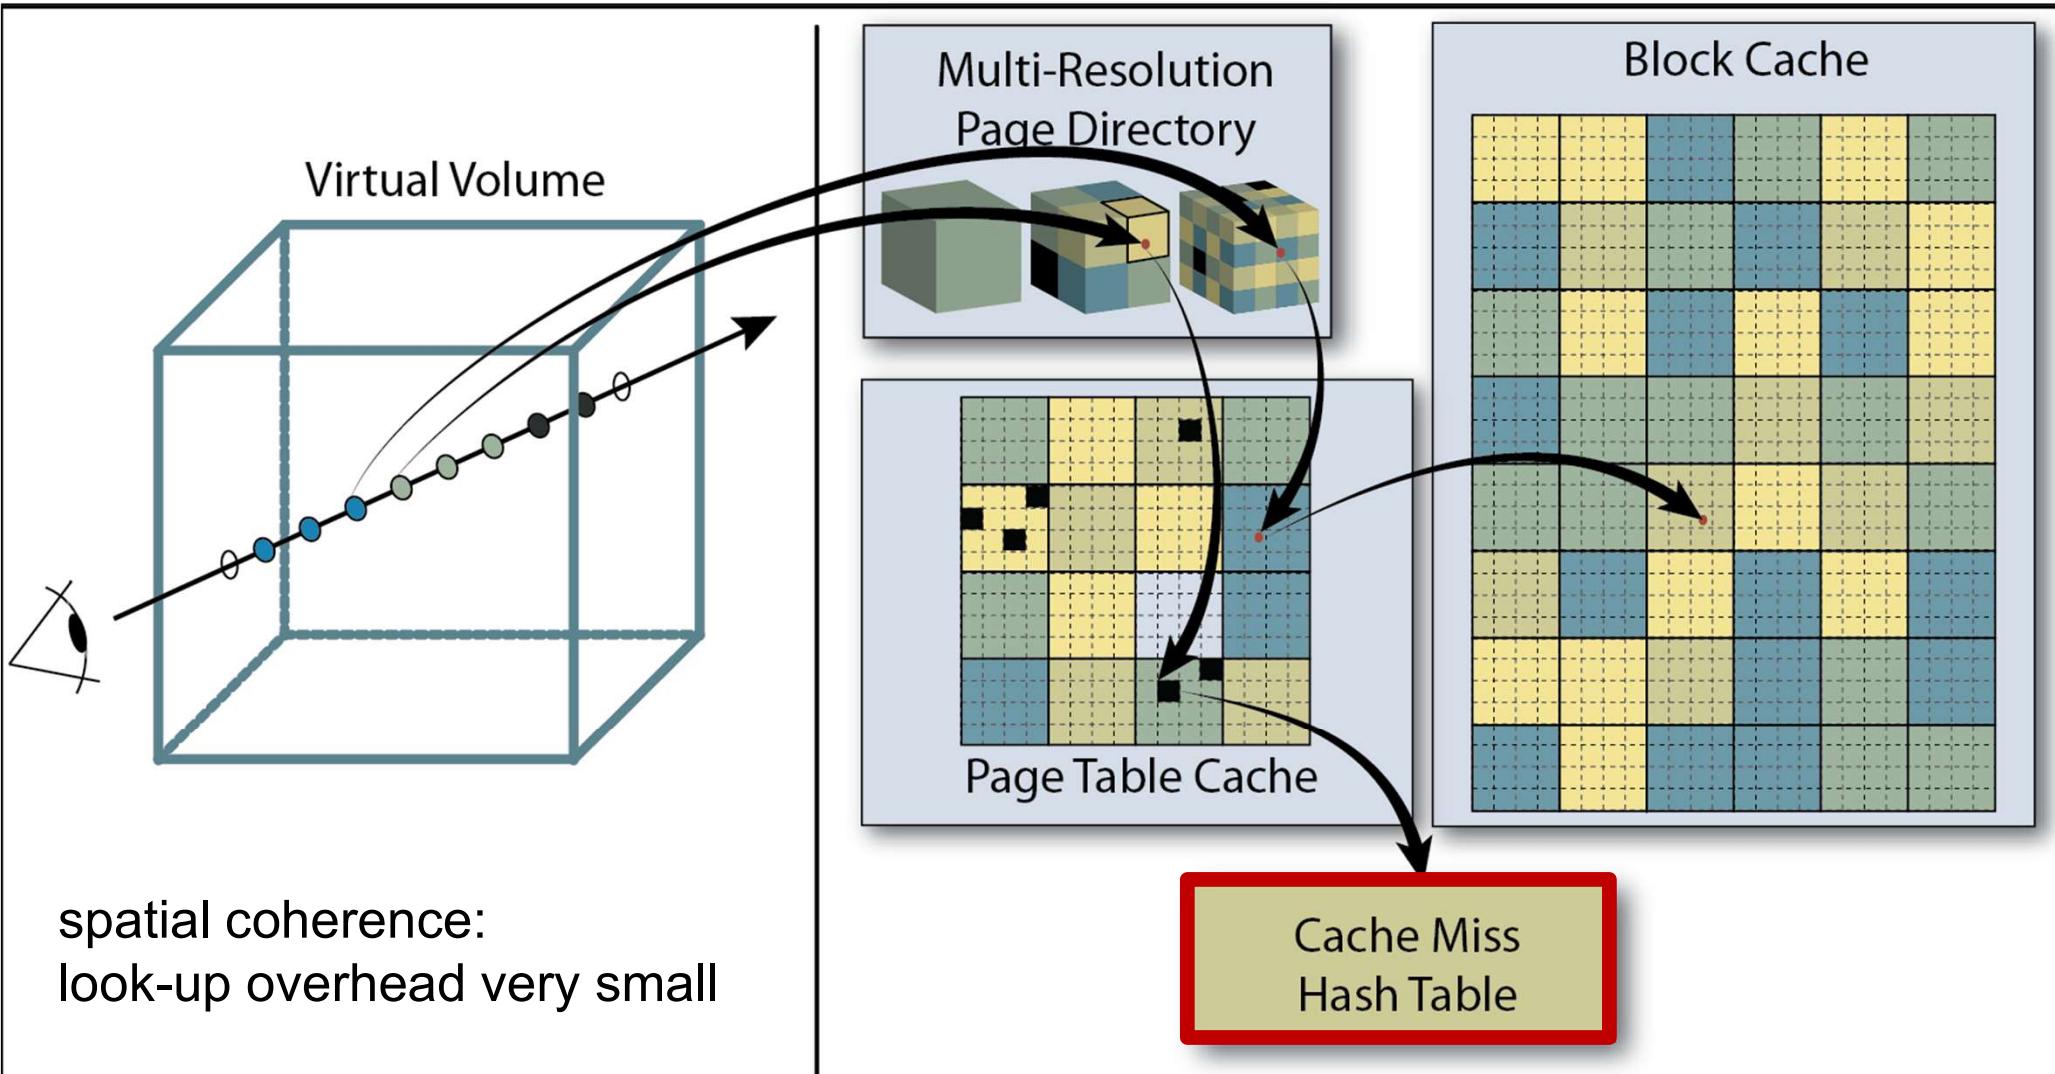
<http://dx.doi.org/10.1109/TVCG.2012.240>

# Petascale Volume Rendering





# Petascale Volume Rendering



Thank you.