

CS 380 - GPU and GPGPU Programming

Lecture 27: Neural Shading on GPUs

(based on NVIDIA Siggraph course)

Markus Hadwiger, KAUST



Reading Assignment #15++

Read (optional):

- SIGGRAPH 2025 course on Neural Shading, Bitterli et al. (NVIDIA)

<https://github.com/shader-slang/neural-shading-s25>

- Random-Access Neural Compression of Material Textures, Vaidyanathan et al. (NVIDIA), 2023

<https://arxiv.org/abs/2305.17105>

- Real-Time Neural Materials using Block-Compressed Features, Weinreich et al. (Ubisoft), 2024

<https://arxiv.org/abs/2311.16121>

- Neural Texture Block Compression, Fujieda and Harada (AMD), 2024

https://gpuopen.com/download/2024_NeuralTextureBCCompression.pdf

- Filtering After Shading With Stochastic Texture Filtering, Pharr et al. (NVIDIA), 2024

<https://arxiv.org/abs/2407.06107>

<https://research.nvidia.com/labs/rtr/publication/pharr2024stochtex/stochtex-slides.pdf>

- Improved Stochastic Texture Filtering Through Sample Reuse, Wronski et al. (NVIDIA), 2025

<https://arxiv.org/abs/2504.05562>

- Real-Time Neural Appearance Models, Zeltner et al. (NVIDIA), 2024

https://research.nvidia.com/labs/rtr/neural_appearance_models/

- The Neural Light Grid: A Scalable Production-Ready Learned Irradiance Volume, Iwanicki et al. (Activision), 2024

https://www.activision.com/cdn/research/Neural_Light_Grid.pdf

- Decoding Light: Neural Compression of Global Illumination, Cao (Lightspeed Studios), 2025

<https://gdcvault.com/play/1035521/Decoding-Light-Neural-Compression-for>

Quiz #3: Dec 11



Organization

- First 30 min of lecture
- No material (book, notes, ...) allowed

Content of questions

- Lectures (both actual lectures and slides)
- Reading assignments
- Programming assignments (algorithms, methods)
- Solve short practical examples

Neural Shading

WHAT IS NEURAL SHADING?



- Usually: Involves Neural Networks
- Anything that's trainable

Neural Shader

WHAT IS NEURAL SHADING?



- Usually: Involves Neural Networks
- Anything that's trainable

Neural **Shader**

- Runs in the graphics pipeline
- Part of your normal shader code

HOW DO NEURAL SHADERS HELP?



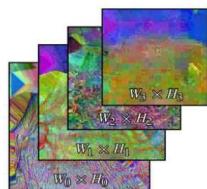
- Consumer GPUs ship with neural network accelerators
- While you render, these sit at 0% utilization
- These are untapped FLOPS!
- ...and very efficient FLOPS
- Accessible in the graphics pipeline with Cooperative Vectors



HOW DO NEURAL SHADERS HELP?



[Fujieda and Harada, 2024]



[Belcour and Benyoub, 2025]



[Vaidyanathan et al., 2023]



[Kuznetsov et al., 2021]



[Mullia et al., 2024]



[Zeltner et al., 2024]



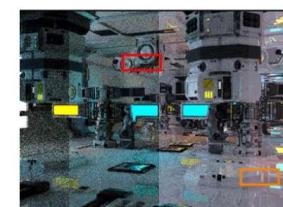
[Mildenhall et al., 2020]



[Müller et al., 2022]



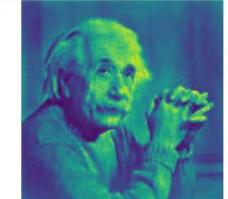
[Kerbl et al., 2023]



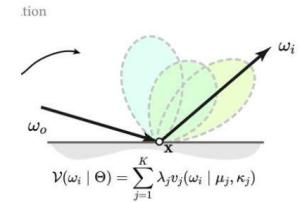
[Müller et al., 2021]



[Dereviannykh et al., 2024]



[Müller et al., 2019]



[Dong et al., 2023]

Compression

Materials

Geometry

Caching

Guiding

CAN WE LEARN BETTER MIP MAPS?

- Mipmaps reduce aliasing **and** improve coherency for distant surfaces
- Color maps downsample well using simple box filter
- Normal maps do not



CAN WE LEARN BETTER MIP MAPS?



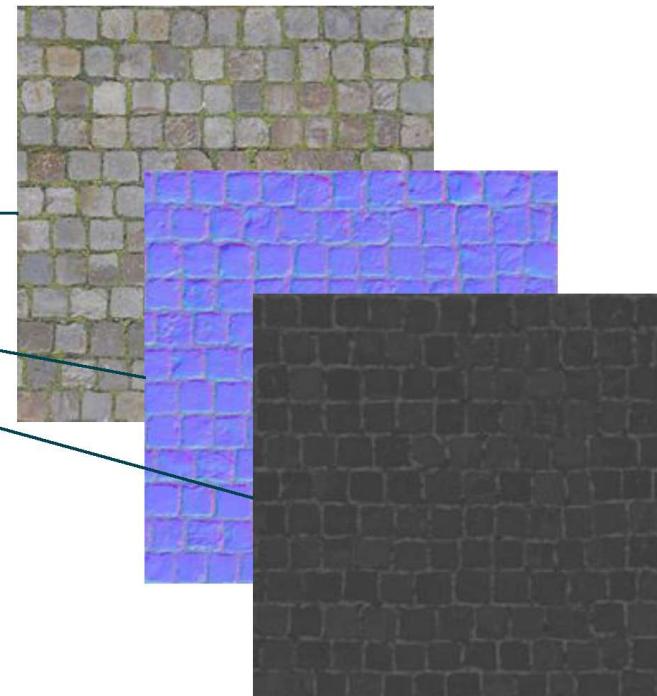
1. Render a beautiful PBR material
2. Generate + render low res with **Mip Maps**
3. Generate + render reference with **Supersampling**
4. Calculate the loss



BASIC SLANG PROGRAM

```
struct MaterialParameters
{
    Tensor<float3, 2> albedo;
    Tensor<float3, 2> normal;
    Tensor<float, 2> roughness;

    float3 get_albedo(int2 pixel)
    {
        return albedo.getv(pixel);
    }
    float3 get_normal(int2 pixel)
    {
        return normalize(normal.getv(pixel));
    }
    float get_roughness(int2 pixel)
    {
        return roughness.getv(pixel);
    }
};
```



BASIC SLANG PROGRAM



```
float3 render(int2 pixel, MaterialParameters material, float3 light_dir, float3 view_dir)
{
    // Bright white light
    float light_intensity = 5.0;

    // Sample very shiny BRDF (it rained today!)
    float3 brdf_sample = sample_BRDF(
        material.get_albedo(pixel),           // albedo color
        normalize(light_dir),                // light direction
        normalize(view_dir),                 // view direction
        material.get_normal(pixel),          // normal map sample
        0.05,                                // roughness
        0.0,                                 // metallic (no metal)
        1.0,                                // specular
    );

    // Combine light with BRDF sample to get pixel colour
    return brdf_sample * light_intensity;
}
```

BASIC PYTHON PROGRAM



```
# Create the app and load the slang module.
app = App(width=2048, height=2048, title="Mipmap Example")
module = spy.Module.load_from_file(app.device, "nsc_basicprogram.slang")

# Load some materials.
albedo_map = spy.Tensor.load_from_image(app.device, "PavingStones070_2K.diffuse.jpg", linearize=True)
normal_map = spy.Tensor.load_from_image(app.device, "PavingStones070_2K.normal.jpg", scale=2, offset=-1)
roughness_map = spy.Tensor.load_from_image(app.device, "PavingStones070_2K.roughness.jpg", grayscale=True)

while app.process_events():
    # Allocate a tensor for output + call the render function
    output = spy.Tensor.empty_like(albedo_map)
    module.render(pixel = spy.call_id(),
                  material = {
                      "albedo": albedo_map,
                      "normal": normal_map,
                      "roughness": roughness_map,
                  },
                  light_dir = spy.math.normalize(spy.float3(0.2, 0.2, 1.0)),
                  view_dir = spy.float3(0, 0, 1),
                  _result = output)

    # Blit tensor to screen.
    app.blit(output)
    app.present()
```

RESULT!

Conveniently Shiny Cobblestones!



DOWNSAMPLING

```
def downsample(source: spy.Tensor, steps: int) -> spy.Tensor:
    for i in range(steps):
        dest = spy.Tensor.empty(
            device=app.device,
            shape=(source.shape[0] // 2, source.shape[1] // 2),
            dtype=source.dtype)
        module.downsample(spy.call_id(), source, _result=dest)
        source = dest
    return source
```

```
float3 downsample(
    int2 pixel,
    Tensor<float3, 2> source)
{
    float3 res = 0;
    res += source.getv(pixel * 2 + int2(0, 0));
    res += source.getv(pixel * 2 + int2(1, 0));
    res += source.getv(pixel * 2 + int2(0, 1));
    res += source.getv(pixel * 2 + int2(1, 1));
    return res * 0.25;
}
```

DOWNSAMPLED INPUTS



```
# ... downsampled output generation/render here

# Quarter res rendered output BRDF from half res inputs.
lr_output = spy.Tensor.empty_like(output)
module.render(pixel = spy.call_id(),
              material = {
                  "albedo":    downsample(albedo_map, 2),
                  "normal":    downsample(normal_map, 2),
                  "roughness": downsample(roughness_map, 2),
              },
              light_dir = spy.math.normalize(spy.float3(0.2, 0.2, 1.0)),
              view_dir = spy.float3(0, 0, 1),
              _result = lr_output)

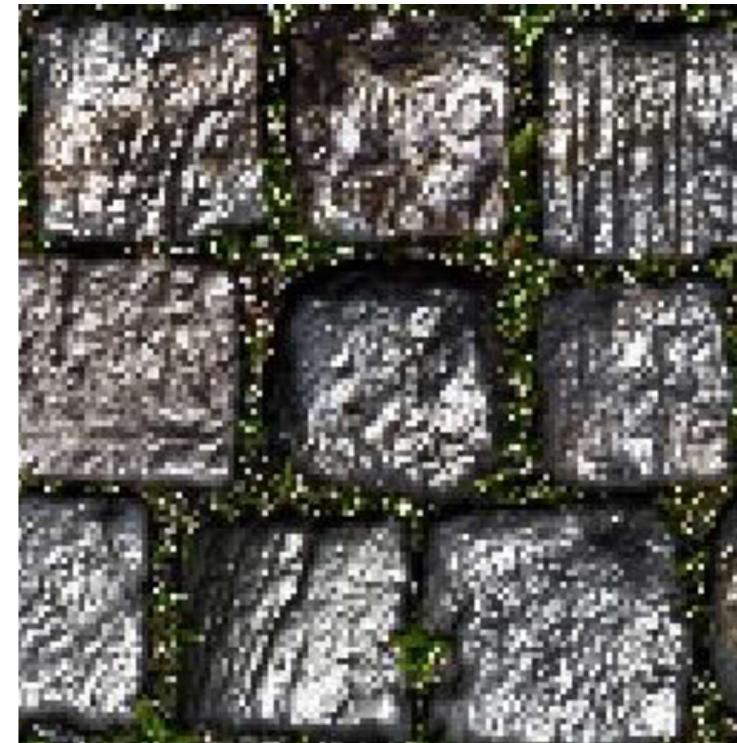
# ... blit the 'lr_output' to screen here
```

RESULT!

Downscaled inputs

- 2x downsample (Mip level 2)
- 1 sample per map
- PBR function run once per pixel
- Optimal but noisy

Poor normal map downsampling gives nasty surface artifacts.



DOWNSAMPLED OUTPUT



```
while app.process_events():
    # Allocate a tensor for output + call the render function
    output = spy.Tensor.empty_like(albedo_map)
    module.render(pixel = spy.call_id(),
                  material = {
                      "albedo": albedo_map,
                      "normal": normal_map,
                      "roughness": roughness_map
                  },
                  light_dir = spy.math.normalize(spy.float3(0.2, 0.2, 1.0)),
                  view_dir = spy.float3(0, 0, 1),
                  _result = output)

    # Downsample the output tensor to quarter res.
    output = downsample(output, 2)

    # Blit tensor to screen.
    app.blit(output, size=spy.int2(2048,2048))
    app.present()
```

RESULT!

Downscaled output

- 2x downsample (Mip level 2)
- 16 high res albedo, normal + roughness pixels sampled
- Full PBR function run 16 times
- Represents an **ideal** output

If our input mip maps were *perfect*, this is what we'd get



SIDE BY SIDE

The Ideal

Render at full res then downsample



The Reality

Downsample inputs and then render



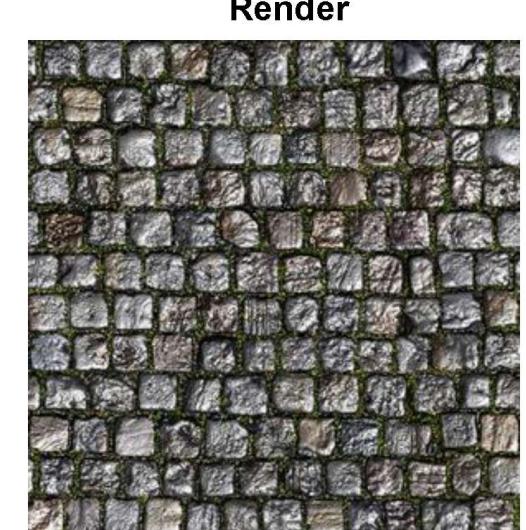
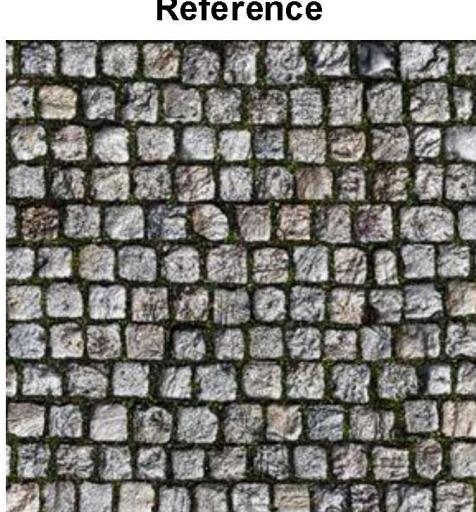
CALCULATING THE LOSS



```
// Render the difference between the rendered pixel  
// a reference pixel.  
  
float3 loss(  
    int2 pixel,  
    float3 reference,  
    MaterialParameters material,  
    float3 light_dir,  
    float3 view_dir)  
{  
    float3 color = render(pixel, material,  
                          light_dir, view_dir);  
    float3 error = color - reference;  
    return error * error; // Squared error  
}
```

```
# Loss between downsampled full res output (the  
# reference), and result from quarter res inputs.  
loss_output = spy.Tensor.empty_like(output)  
module.loss(pixel = spy.call_id(),  
           material = {  
               "albedo": downsample(albedo_map, 2),  
               "normal": downsample(normal_map, 2),  
               "roughness": downsample(roughness_map, 2),  
           },  
           reference = output,  
           light_dir = spy.math.normalize(  
               spy.float3(0.2, 0.2, 1.0)),  
           view_dir = spy.float3(0, 0, 1),  
           _result = loss_output)
```

CALCULATING THE LOSS

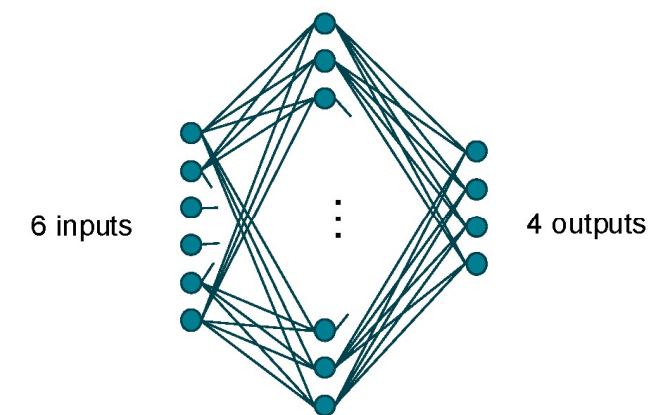


Make mipmap better == Make the loss smaller

IMPLEMENTING OUR TOY NEURAL TEXTURE EXAMPLE

```
struct NeuralTexture2D
{
    NetworkParameters<6, 32> layer1;
    NetworkParameters<32, 4> layer2;

    float4 sample(float2 uv, float2 du, float2 dv)
    {
        float input[6] = {uv.x, uv.y, du.x, du.y, dv.x, dv.y};
        let latentResult = layer1.eval(input);
        let output = layer2.eval(latentResult);
        return float4(output[0], output[1], output[2], output[3]);
    }
}
```

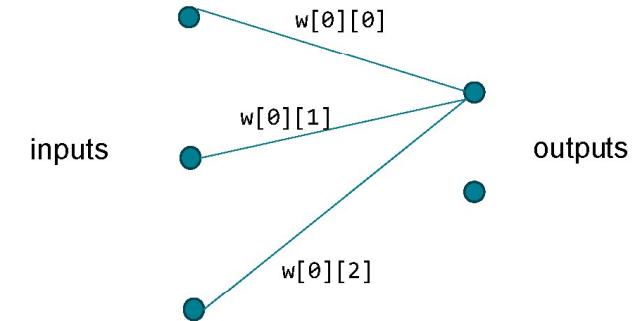


COMPUTATION OF A 3-INPUT, 2-OUTPUT LAYER

```
outputs[0] = biases[0] +  
           inputs[0] * w[0][0] +  
           inputs[1] * w[0][1] +  
           inputs[2] * w[0][2];  
  
outputs[0] = max(0.0, outputs[0]);
```

Weighted Sum
of Inputs

Activation



COMPUTATION OF A 3-INPUT, 2-OUTPUT LAYER

```
outputs[0] = biases[0] +  
    inputs[0] * w[0][0] +  
    inputs[1] * w[0][1] +  
    inputs[2] * w[0][2];
```

Weighted Sum
of Inputs

```
outputs[0] = max(0.0, outputs[0]);
```

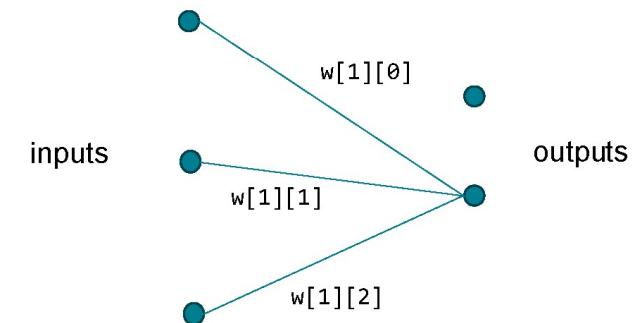
Activation

```
outputs[1] = biases[1] +  
    inputs[0] * w[1][0] +  
    inputs[1] * w[1][1] +  
    inputs[2] * w[1][2];
```

Weighted Sum
of Inputs

```
outputs[1] = max(0.0, outputs[1]);
```

Activation



COMPUTATION OF A 3-INPUT, 2-OUTPUT LAYER

`outputs[0] = biases[0] +`

`inputs[0] * w[0][0] +
 inputs[1] * w[0][1] +
 inputs[2] * w[0][2];`

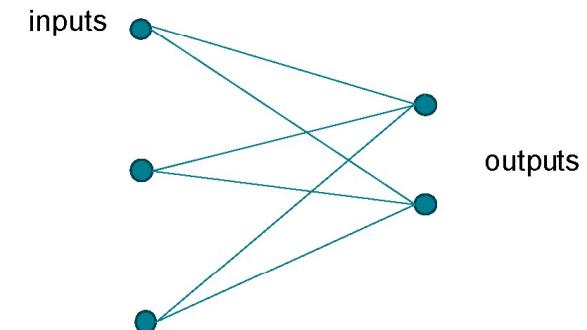
`outputs[0] = max(0.0, outputs[0]);`

`outputs[1] = biases[1] +`

`inputs[0] * w[1][0] +
 inputs[1] * w[1][1] +
 inputs[2] * w[1][2];`

`outputs[1] = max(0.0, outputs[1]);`

$$\left[\begin{array}{c} b_0 \\ b_1 \\ b_2 \end{array} \right] + \left[\begin{array}{ccc} w_{00} & w_{01} & w_{02} \\ w_{10} & w_{11} & w_{12} \end{array} \right] \left[\begin{array}{c} i_0 \\ i_1 \\ i_2 \end{array} \right]$$



SIMPLIFIED VIEW OF A FEED-FORWARD LAYER



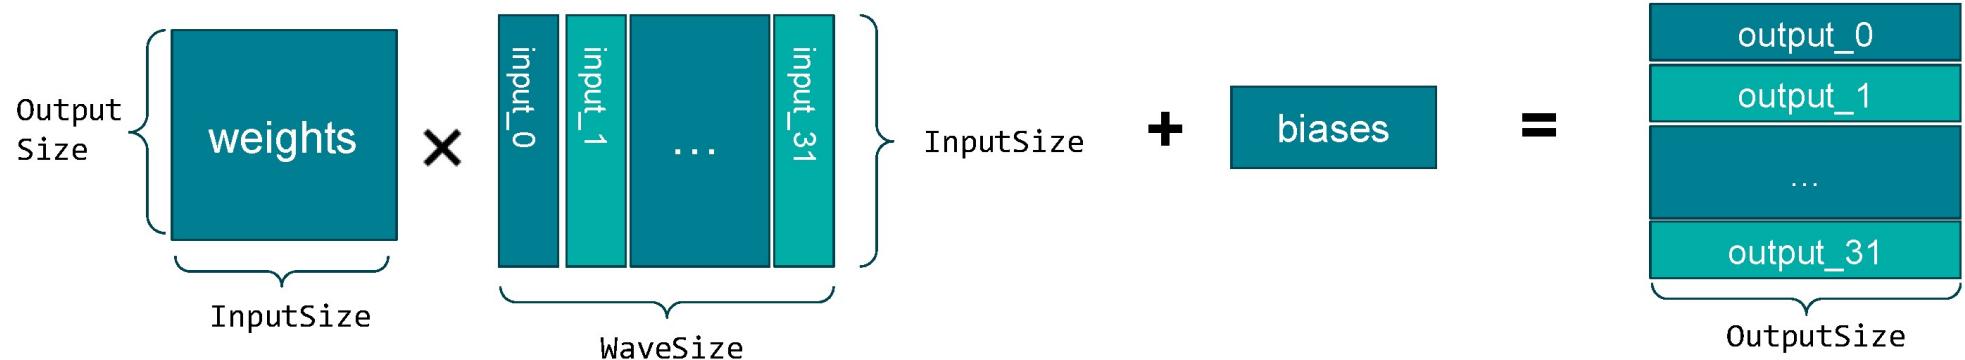
```
struct NetworkParameters<int InputSize, int OutputSize>
{
    float weights[InputSize][OutputSize];
    float biases[OutputSize];
    Array<float, OutputSize> eval(float input[InputSize])
    {
        float output[OutputSize] = biases;
        output += MatMulVec(weights, input);
        for (int i = 0; i < OutputSize; ++i)
            output[i] = max(output[i] * 0.01f, output[i]); // ReLU
        return output;
    }
}
```

EACH SUBGROUP IS COMPUTING A MATRIX-MATRIX MULTIPLICATION

```
Array<float, OutputSize> eval(float input[InputSize])  
{  
    float output[OutputSize] = biases;  
    output += MatMulVec(weights, input);  
    ...  
    return output;  
}
```

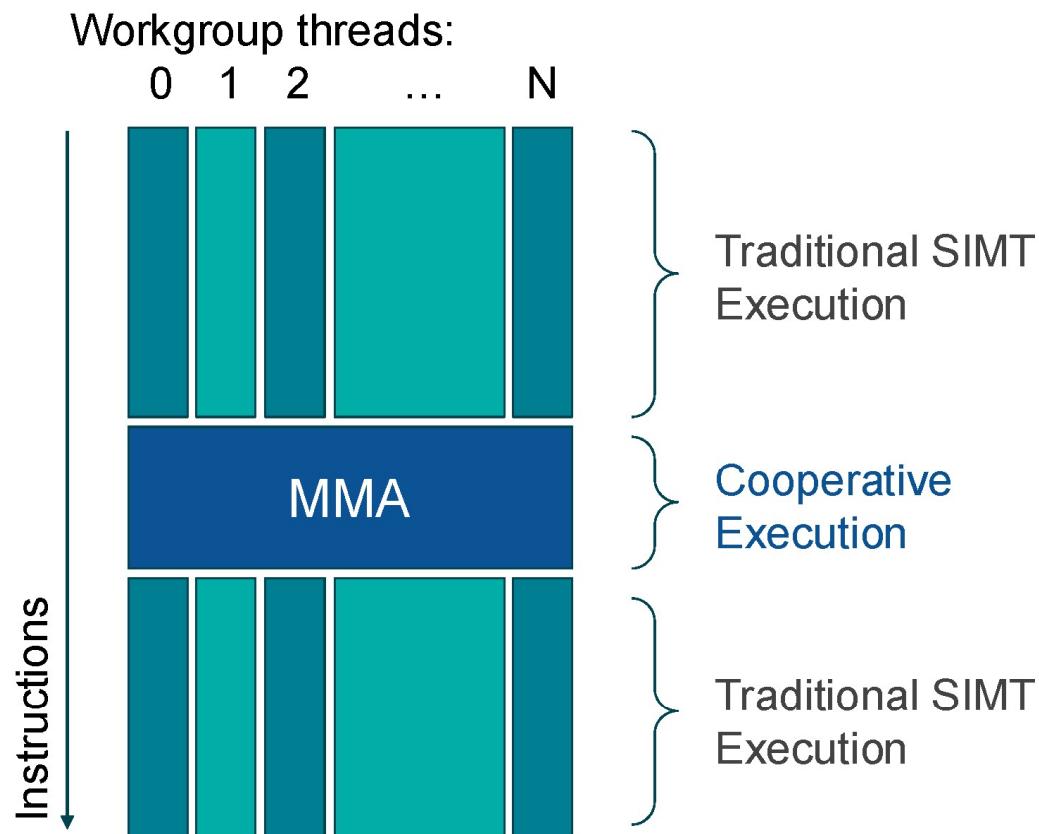
Thread	Computation
0	MatMulVec(weights, input_0)
1	MatMulVec(weights, input_1)
2	MatMulVec(weights, input_2)
3	MatMulVec(weights, input_3)
...	
31	MatMulVec(weights, input_31)

Computation done by the entire subgroup: **Matrix Multiply Accumulate** (a.k.a. **GEmM**)



- *Wave == Subgroup(Vulkan) == Warp (CUDA) == SIMD Group (Metal)

MODERN GPUS HAVE HARDWARE SUPPORT FOR MATRIX-MULTIPLY-ACCUMULATION



MMA EXPOSED AS INTRINSICS IN VARIOUS APIs



DirectX® 12

WaveMatrix



VK_KHR_cooperative_matrix

VK_NV_cooperative_matrix2



SIMD Group Matrix

MMA INTRINSICS RESTRICTIONS



- Matrix operands in a Matrix-Multiply-Accumulation operation must be the same across all threads in a subgroup.
- Input vectors must be explicitly packed into a matrix and loaded in one operation.
 - Often require explicit use of group shared memory, which isn't available outside compute shaders.

COOPERATIVE VECTORS: USING MMA HARDWARE FROM SIMT CODE



DirectX® 12

Shader Model 6.9



VK_NV_cooperative_vector



Metal 4
Machine Learning

SIMPLIFIED VIEW OF A FEED-FORWARD LAYER



```
struct NetworkParameters<int InputSize, int OutputSize>
{
    float weights[InputSize][OutputSize];
    float biases[OutputSize];
    Array<float, OutputSize> eval(float input[InputSize])
    {
        float output[OutputSize] = biases;
        output += MatMulVec(weights, input);
        for (int i = 0; i < OutputSize; ++i)
            output[i] = max(output[i] * 0.01f, output[i]); // ReLU
        return output;
    }
}
```

FEED-FORWARD LAYER USING COOPERATIVE VECTORS



```
struct NetworkParameters<int InputSize, int OutputSize>
{
    half* weights;
    half* biases;
    CoopVec<half, OutputSize> eval(CoopVec<half, InputSize> input)
    {
        let output = coopVecMatMulAdd<half, OutputSize>(
            input, CoopVecComponentType.Float16, // input and format
            weights, CoopVecComponentType.Float16, // weights and format
            biases, CoopVecComponentType.Float16, // biases and format
            CoopVecMatrixLayout.RowMajor, // matrix layout
            false, // transpose matrix before multiply?
            sizeof(half) * InputSize); // matrix stride
        return max(output * 0.01h, output); // Leaky ReLU activation
    }
}
```

Using 16-bit floats because 32-bit float is not supported by HW.

TRAINING WITH COOPERATIVE VECTORS



```
struct NetworkParameters<int InputSize, int OutputSize>
{
```

```
    half* weights;
    half* biases;
```

1. Where should we store the gradients of the weights?

```
CoopVec<half, OutputSize> eval(CoopVec<half, InputSize> input)
{
    let output = coopVecMatMulAdd<half, OutputSize>(
        input, CoopVecComponentType.Float16, // input and format
        weights, CoopVecComponentType.Float16, // weights and format
        biases, CoopVecComponentType.Float16, // biases and format
        CoopVecMatrixLayout.RowMajor, // matrix layout
        false, // transpose matrix before multiply?
        sizeof(half) * InputSize); // matrix stride
    return max(output * 0.01h, output); // Leaky ReLU activation
}
```

2. How to make this
differentiable?

STORING THE GRADIENTS OF THE WEIGHTS



```
struct NetworkParameters<int InputSize, int OutputSize>
{
    half* weights;
    half* biases;
    half* weightsGrad;
    half* biasesGrad;
}
```

1. Where should we store the gradients of the weights?

TRAINING WITH COOPERATIVE VECTORS



```
struct NetworkParameters<int InputSize, int OutputSize>
{
    ...
    CoopVec<half, OutputSize> eval(CoopVec<half, InputSize> input)
    {
        let output = coopVecMatMulAdd<half, OutputSize>(
            input, CoopVecComponentType.Float16, // input and format
            weights, CoopVecComponentType.Float16, // weights and format
            biases, CoopVecComponentType.Float16, // biases and format
            CoopVecMatrixLayout.RowMajor, // matrix layout
            false, // transpose matrix before multiply?
            sizeof(half) * InputSize); // matrix stride
        return max(output * 0.01h, output); // Leaky ReLU activation
    }
}
```

How to make `eval` differentiable?

- `CoopVec` type is non-differentiable.
- `coopVecMatMulAdd` intrinsic is non-differentiable.

WRAPPING COOPERATIVE VECTOR IN A DIFFERENTIABLE TYPE



```
struct MLVec<int N> : IDifferentiable
{
    CoopVec<half, N> data;
    typealias Differential = This;

    static Differential dadd(Differential d0, Differential d1)
    {
        return {d0.data + d1.data};
    }
    static Differential dmul<U:<__BuiltinRealType>>(U s, Differential d)
    {
        return {d.data * __realCast<half>(s)};
    }
    static Differential dzero()
    {
        return {};
    }

}
```

Wrap CoopVec in a MLVec type, that is declared to be Differentiable, with Differential type being itself.

UPDATE EVAL TO USE MLVEC TYPE



```
struct NetworkParameters<int InputSize, int OutputSize>
{
    ...
    half* weightsGrad;
    half* biasesGrad;

    MLVec<OutputSize> eval(MLVec<InputSize> input) { ... }

}
```

MAKE EVAL DIFFERENTIABLE WITH CUSTOM DERIVATIVE



```
struct NetworkParameters<int InputSize, int OutputSize>
{
    ...
    half* weightsGrad;
    half* biasesGrad;

    MLVec<OutputSize> eval(MLVec<InputSize> input) { ... }

    [BackwardDerivativeOf(eval)]
    void evalBwd(
        inout DifferentiablePair<MLVec<InputSize> input,
        MLVec<OutputSize> resultGrad)
    {
        // What goes here??
    }
}
```

Task: propagate gradients from each output element (resultGrad) to:

- weight and bias parameters
(to store in `weightsGrad` and `biasesGrad`).
- Each input element (return via `input.d`)

GRADIENT PROPAGATION FOR MATRIX-VECTOR MULTIPLICATION



Given matrix-vector multiplication

$$\mathbf{u} = \mathbf{M}\mathbf{v} + \mathbf{B}$$

We have:

$$d\mathbf{M} = d\mathbf{u} \otimes_{\text{outer}} \mathbf{v}$$

$$d\mathbf{B} = d\mathbf{u}$$

$$d\mathbf{v} = \mathbf{M}^T d\mathbf{u}$$

\otimes_{outer} : outer vector product

ACCUMULATING GRADIENTS ACROSS THREADS

Given matrix-vector multiplication

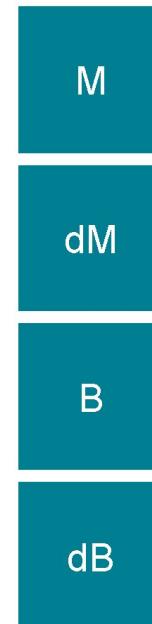
$$\mathbf{u} = \mathbf{M}\mathbf{v} + \mathbf{B}$$

We have:

$$d\mathbf{M} = d\mathbf{u} \otimes_{\text{outer}} \mathbf{v}$$

$$dB = d\mathbf{u}$$

$$d\mathbf{v} = \mathbf{M}^T d\mathbf{u}$$



Thread	Computation
0	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_0, \mathbf{B})$
1	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_1, \mathbf{B})$
2	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_2, \mathbf{B})$
...	
N	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_N, \mathbf{B})$

\otimes_{outer} : outer vector product

Each thread accesses different \mathbf{v} , but the same \mathbf{M} and \mathbf{B}

ACCUMULATING GRADIENTS ACROSS THREADS

Given matrix-vector multiplication

$$\mathbf{u} = \mathbf{M}\mathbf{v} + \mathbf{B}$$

We have:

$$d\mathbf{M} \doteq d\mathbf{u} \otimes_{\text{outer}} \mathbf{v}$$

$$d\mathbf{B} \doteq d\mathbf{u}$$

$$d\mathbf{v} = \mathbf{M}^T d\mathbf{u}$$



Thread	Computation
0	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_0, \mathbf{B})$ $d\mathbf{M} \doteq \text{outerProduct}(\dots)$ $d\mathbf{B} \doteq d\mathbf{U}$
1	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_1, \mathbf{B})$ $d\mathbf{M} \doteq \text{outerProduct}(\dots)$ $d\mathbf{B} \doteq d\mathbf{U}$
2	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_2, \mathbf{B})$ $d\mathbf{M} \doteq \text{outerProduct}(\dots)$ $d\mathbf{B} \doteq d\mathbf{U}$
...	
N	$\mathbf{u} = \text{matMulVecAdd}(\mathbf{M}, \mathbf{v}_N, \mathbf{B})$ $d\mathbf{M} \doteq \text{outerProduct}(\dots)$ $d\mathbf{B} \doteq d\mathbf{U}$

Needs to be atomic to prevent race conditions!

INTRINSICS TO SPEEDUP ATOMIC GRADIENT ACCUMULATION



Given matrix-vector multiplication

$$\mathbf{u} = \mathbf{M}\mathbf{v} + \mathbf{B}$$

We have:

$$d\mathbf{M} += d\mathbf{u} \otimes_{\text{outer}} \mathbf{v}$$

$$d\mathbf{B} += d\mathbf{u}$$

$$d\mathbf{v} = \mathbf{M}^T d\mathbf{u}$$

Each step has its corresponding coop-vector intrinsic:

`coopVecOuterProductAccumulate`

`coopVecReduceSumAccumulate`

`coopVecMatMul`

BACK-PROP GRADIENTS THROUGH EVAL



```
[BackwardDerivativeOf(eval)]
void evalBwd(
    inout DifferentialPair<MLVec<InputSize>> input,
    MLVec<OutputSize> resultGrad)
{
    let fwd = eval(input.p);
    // Back-prop resultGrad through activation.
    for (int i = 0; i < OutputSize; i++)
        if (fwd.data[i] < 0.0)
            resultGrad.data[i] *= 0.01h;

    // Back-prop gradients to the weights matrix.
    coopVecOuterProductAccumulate(
        resultGrad.data,
        input.v.data,
        weightsGrad, 0,
        CoopVecMatrixLayout.TrainingOptimal, CoopVecComponentType.Float16);

    // Back-prop gradients to the biases vector.
    coopVecReduceSumAccumulate(resultGrad.data, biasesGrad, biasesOffset);
```

coopVecOuterProductAccumulate
requires weightsGrad to be stored
in TrainingOptimal layout.

$$dM += d\mathbf{u} \otimes_{\text{outer}} \mathbf{v}$$

$$dB += d\mathbf{u}$$

BACK-PROP GRADIENTS THROUGH EVAL



```
[BackwardDerivativeOf(eval)]
void evalBwd(
    inout DifferentialPair<MLVec<InputSize>> input,
    MLVec<OutputSize> resultGrad)
{
    ...
    // Back-prop gradients to the input vector.
    let dInput = coopVecMatMul<half, InputSize>(
        resultGrad.data, CoopVecComponentType.Float16,
        weights, CoopVecComponentType.Float16,
        CoopVecMatrixLayout.ColumnMajor, false, sizeof(half)*InputSize);
    input = {input.p, {dInput}};
}
```

$$d\mathbf{v} = \mathbf{M}^T d\mathbf{u}$$

Compute
i.e. `transpose(weights) * resultGrad`
by setting the layout of weights to `ColumnMajor`.

CALLERS OF FEED-FORWARD LAYER CAN NOW BE AUTODIFF'ED



```
struct NeuralTexture2D
{
    NetworkParameters<6, 32> layer1;
    NetworkParameters<32, 4> layer2;

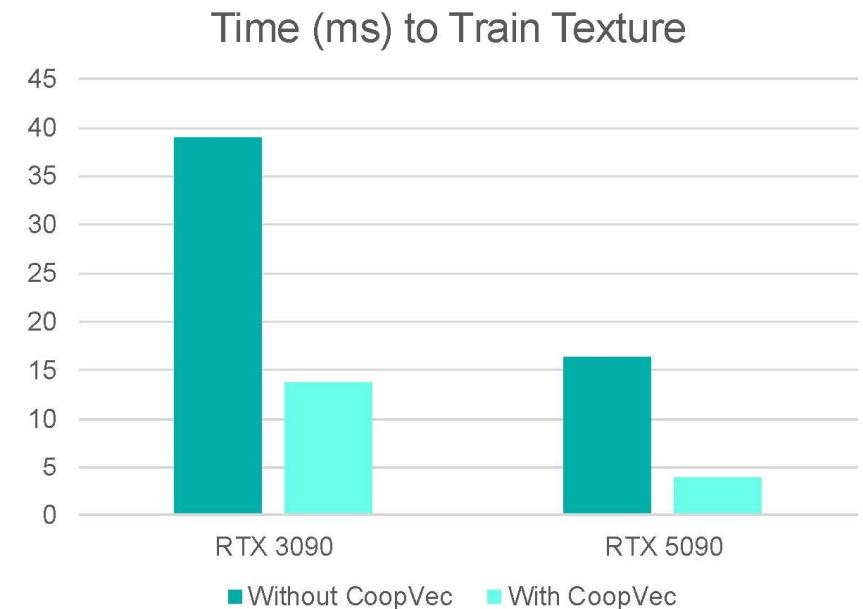
    [Differentiable]
    float4 sample(float2 uv, float2 du, float2 dv)
    {
        float input[6] = {uv.x, uv.y, du.x, du.y, dv.x, dv.y};
        let latentResult = layer1.eval(input);
        let output = layer2.eval(latentResult);
        return float4(output[0], output[1], output[2], output[3]);
    }
}
```

PERFORMANCE IMPROVEMENTS WITH COOP-VECTOR

- **Benchmark:** Neural Texture Training
- **Input:** 100MB PNG

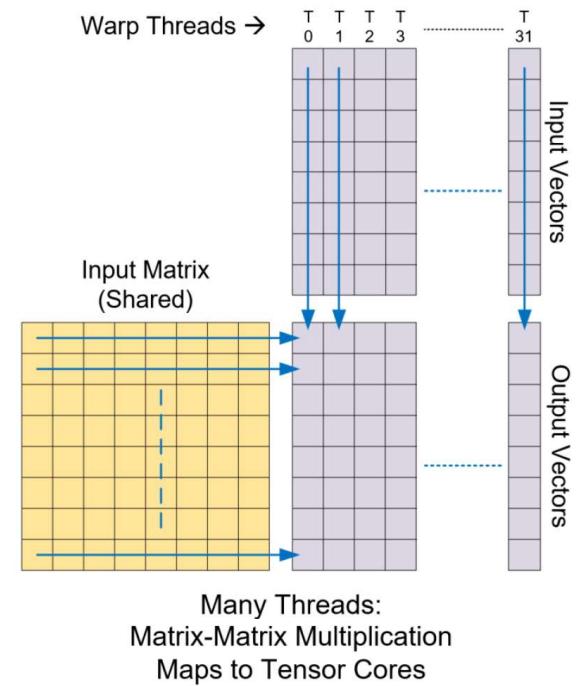
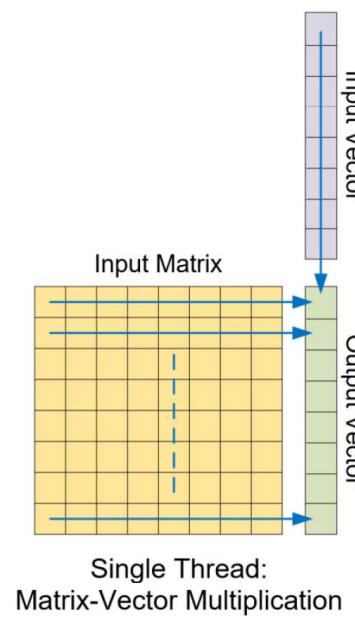


Over **3x** speedup compared to highly optimized code without coop-vec.



BATCHED MATRIX-VECTOR MULTIPLICATION

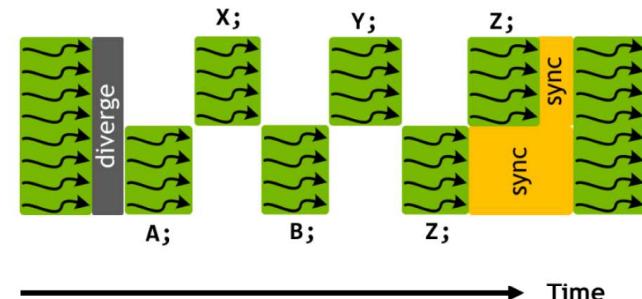
- Hardware Tensor Cores:
 - Matrix-Matrix multiplication using entire wave/warp
 - Low precision (FP16, FP8, INT8, even FP4)
- Cooperative Vector API:
 - Matrix-Vector multiplication in each thread
- Mapping CoopVec onto Tensor Cores:
 - Combine M-V multiplications from all threads in a wave
 - Divergence becomes a problem (more on that later)



REFRESHER ON GPU SIMT ARCHITECTURE

- NVIDIA GPUs group execution threads into 32-thread “warps”
 - Other vendors have similar concepts with different thread counts
 - We use the term “wave” throughout this course for consistency
- All threads in a wave execute one instruction on different data
 - Some may be inactive
 - Following different code branches in a wave causes execution divergence
 - Divergent code branches are serialized into synchronous groups
- Memory accesses work best when addresses are packed and aligned
 - Not following optimal address patterns causes data divergence
 - Divergent memory accesses are serialized into optimal groups

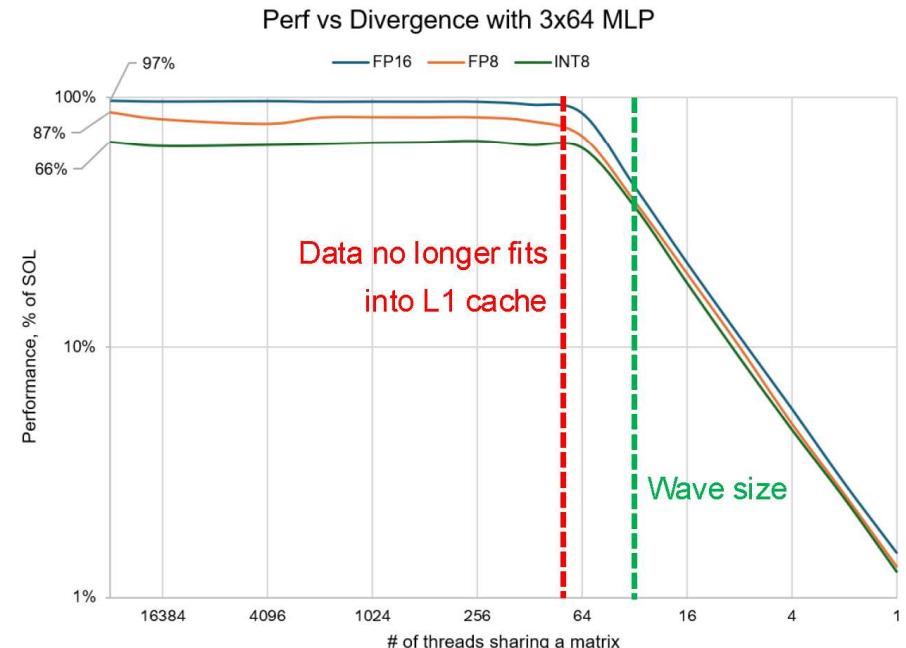
```
if (threadIdx.x < 4) {  
    A;  
    B;  
} else {  
    X;  
    Y;  
}  
Z;  
__syncwarp()
```



Source: Volta architecture whitepaper

TENSOR CORES AND DIVERGENCE

- Tensor Cores compute a single MMA using the entire wave
 - MMA = Matrix Multiply Accumulate
 - We can wake up inactive threads if needed
- Wave-wide `coopVecMatMul[Add]` might be incompatible:
 - Vector inputs are trivially combined into a matrix
 - Matrix inputs must be the same for all threads
 - Programming model allows matrix inputs to be different per-thread
- Solution: Serialize divergent matrix operations
 - Handled transparently by the driver
 - Has severe performance effects
- **Avoid CoopVec operations with divergent matrices**
 - Group draw calls by material
 - Sort threads manually or use Shader Execution Reordering (SER)



Measured on RTX 4090 using a directed test – Log/Log scale

https://github.com/jeffbolz/vk_cooperative_vector_perf

TENSOR CORE MATRIX LAYOUT PROBLEM



- Tensor cores use custom matrix layouts
 - Components of matrices are shuffled between threads
 - Specific layout depends on the GPU and data types
- Need to shuffle data before and after MMAs
 - Weight matrix must be pre-shuffled in memory
 - DX12 and Vulkan provide functions to shuffle the weights
 - Input and output shuffling is handled transparently by the driver
- Output of one MMA can be used as input of another MMA without shuffle
 - Redundant shuffles can be eliminated

Row\Col	0	1	2	3	4	5	6	7
0	T0 : { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
3	T3: { c0, c1, c2, c3, c4, c5, c6, c7 }							
4	T16: { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
7	T19: { c0, c1, c2, c3, c4, c5, c6, c7 }							

Row\Col	0	1	2	3	4	5	6	7
0	T8 : { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
3	T11: { c0, c1, c2, c3, c4, c5, c6, c7 }							
4	T24: { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
7	T27: { c0, c1, c2, c3, c4, c5, c6, c7 }							

Row\Col	0	1	2	3	4	5	6	7
0	T4 : { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
3	T7: { c0, c1, c2, c3, c4, c5, c6, c7 }							
4	T20: { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
7	T23: { c0, c1, c2, c3, c4, c5, c6, c7 }							

Row\Col	0	1	2	3	4	5	6	7
0	T12 : { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
3	T15: { c0, c1, c2, c3, c4, c5, c6, c7 }							
4	T28: { c0, c1, c2, c3, c4, c5, c6, c7 }							
..								
7	T31: { c0, c1, c2, c3, c4, c5, c6, c7 }							

Output matrix layout for an FP16 MMA operation (maps to 4 instructions)
<https://docs.nvidia.com/cuda/parallel-thread-execution/index.html>

WEIGHT MATRIX LAYOUT CONVERSIONS



- Query shuffled matrix size from VK/DX12
 - Could be larger than the optimally packed data
- Upload row- or column-major matrix to GPU
- Convert layout (perform the shuffle) on the GPU
 - Inference Optimal
 - Training Optimal
- Use converted matrix in CoopVec operations
 - Specify layout as coopVecMatMulAdd argument

FEED-FORWARD LAYER USING COOPERATIVE VECTORS

```
struct NetworkParameters<int InputSize, int OutputSize>
{
    half* weights;
    half* biases;
    CoopVec<half, OutputSize> eval(CoopVec<half, InputSize> input)
    {
        let output = coopVecMatMulAdd<half, OutputSize>(
            input, CoopVecComponentType.Float16,
            weights, CoopVecComponentType.Float16,
            biases, CoopVecComponentType.Float16,
            CoopVecMatrixLayout.InferencingOptimal,
            false, // is matrix transposed
            sizeof(half) * InputSize); // matrix stride
        return max(output * 0.01h, output);
    }
}
```

[PREVIOUS SECTION](#)

MATRIX LAYOUT CONVERSION APIs



Vulkan (VK_NV_cooperative_vector)

- `vkConvertCooperativeVectorMatrixNV`
 - CPU conversion / size query
- `vkCmdConvertCooperativeVectorMatrixNV`
 - GPU conversion

DX12 (Agility SDK 717 Preview)

- `Device::GetLinearAlgebraMatrixConversion... DestinationInfo`
 - Size query
- `CommandList::ConvertLinearAlgebraMatrix`
 - GPU conversion
 - Note: No CPU conversion function!

MAPPING AN MLP ONTO TENSOR CORES 1/3

ORIGINAL CODE

```
coopVecMatMul(layer0, input, weightBuffer, offset0);
layer0 = max(layer0, 0);

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
layer1 = max(layer1, 0);

coopVecMatMul(output, layer1, weightBuffer, offset2);
```

ADD SHUFFLES AROUND MMA OPS

```
shflinput = shuffle(input);
coopVecMatMul(temp0, shflinput, weightBuffer, offset0);
layer0 = unshuffle(temp0);

layer0 = max(layer0, 0);

shfllayer0 = shuffle(layer0);
coopVecMatMul(temp1, shfllayer0, weightBuffer, offset1);
layer1 = unshuffle(temp1);

layer1 = max(layer1, 0);

shfllayer1 = shuffle(layer1);
coopVecMatMul(temp2, shfllayer1, weightBuffer, offset2);
output = unshuffle(temp2);
```

Baseline implementation – functional but slow

MAPPING AN MLP ONTO TENSOR CORES 2/3



REMOVE UNSHUFFLE / SHUFFLE PAIRS

```
shflinput = shuffle(input);
coopVecMatMul(temp0, shflinput, weightBuffer, offset0);
layer0 = unshuffle(temp0);

layer0 = max(layer0, 0);

shfllayer0 = shuffle(layer0);
coopVecMatMul(temp1, shfllayer0, weightBuffer, offset1);
layer1 = unshuffle(temp1);

layer1 = max(layer1, 0);

shfllayer1 = shuffle(layer1);
coopVecMatMul(temp2, shfllayer1, weightBuffer, offset2);
output = unshuffle(temp2);
```

```
shflinput = shuffle(input);
coopVecMatMul(temp0, shflinput, weightBuff, offset0);
temp0_1 = max(temp0, 0);

coopVecMatMul(temp1, temp0_1, weightBuff, offset1);
temp1_1 = max(temp1, 0);

coopVecMatMul(temp2, temp1_1, weightBuff, offset1);
output = unshuffle(temp2);
```

The shuffle removal is called “Layer fusion”

Sometimes it fails

MAPPING AN MLP ONTO TENSOR CORES 3/3



PEEL THE DIVERGENT MATRICES

```
shflinput = shuffle(input);
coopVecMatMul(shflinput, weightBuff, offset0);
temp0_1 = max(temp0, 0);

coopVecMatMul(temp1, temp0_1, weightBuff, offset1);
temp1_1 = max(temp1, 0);

coopVecMatMul(temp2, temp1_1, weightBuff, offset1);
output = unshuffle(temp2);
```

```
shflinput = shuffle(input);

foreach (unique combination of offsets)
{
    coopVecMatMul(temp0, shflinput, weightBuff, offset0);
    temp0_1 = max(temp0, 0);

    coopVecMatMul(temp1, temp0_1, weightBuff, offset1);
    temp1_1 = max(temp1, 0);

    coopVecMatMul(temp2, temp1_1, weightBuff, offset2);

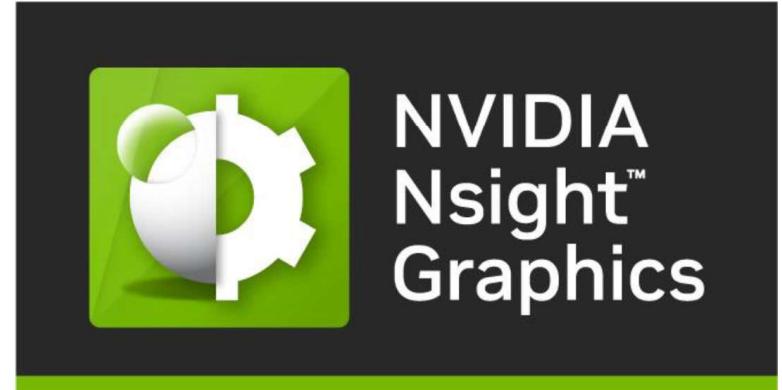
    mergeThreadResults(temp2);
}

output = unshuffle(temp2);
```

DIAGNOSING LAYER FUSION FAILURES



- No publicly available diagnostic tool at this time
- Experimental approach:
 - Measure performance of the original code
 - Remove all code between coopVecMatMul[Add] operations
 - See if performance increases dramatically
 - Re-introduce pieces of that code looking for perf cliffs



Nsight Graphics is a great tool,
but it doesn't fully support CoopVec yet

PREVENTING LAYER FUSION FAILURES 1/3



Avoid elementwise operations on CoopVec's between layers

BAD

```
coopVecMatMul(layer0, input, weightBuffer, offset0);

for (int i = 0; i < 64; ++i)
{
    // Leaky ReLU
    if (layer0[i] < 0)
        layer0[i] *= 0.01;
}

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
```

GOOD

```
coopVecMatMul(layer0, input, weightBuffer, offset0);

// Use vector operations to express the same math
layer0 = max(layer0, layer0 * 0.01);

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
```

Useful vector intrinsics: min, max, clamp, **step**

Look in “hlsl.meta.slang” for more

PREVENTING LAYER FUSION FAILURES 2/3



Use vector load and store operations instead of elementwise ones

BAD

```
coopVecMatMul(layer0, input, weightBuffer, offset0);

CoopVec<half, 64> bias;
for (int i = 0; i < 64; ++i)
{
    bias[i] = biasBuffer.Load<half>(biasOffset0
        + i * sizeof(half));
}

layer0 += bias;

layer0 = max(layer0, 0); // ReLU

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
```

BETTER

```
coopVecMatMul(layer0, input, weightBuffer, offset0);

let bias = CoopVec<half, 64>.load(biasBuffer, biasOffset0);
layer0 += bias;

layer0 = max(layer0, 0); // ReLU

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
```

PREVENTING LAYER FUSION FAILURES 3/3



Use vector load and store operations instead of elementwise ones

BETTER

```
coopVecMatMul(layer0, input, weightBuffer, offset0);

let bias = CoopVec<half, 64>.load(biasBuffer, biasOffset0);
layer0 += bias;

layer0 = max(layer0, 0); // ReLU

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
```

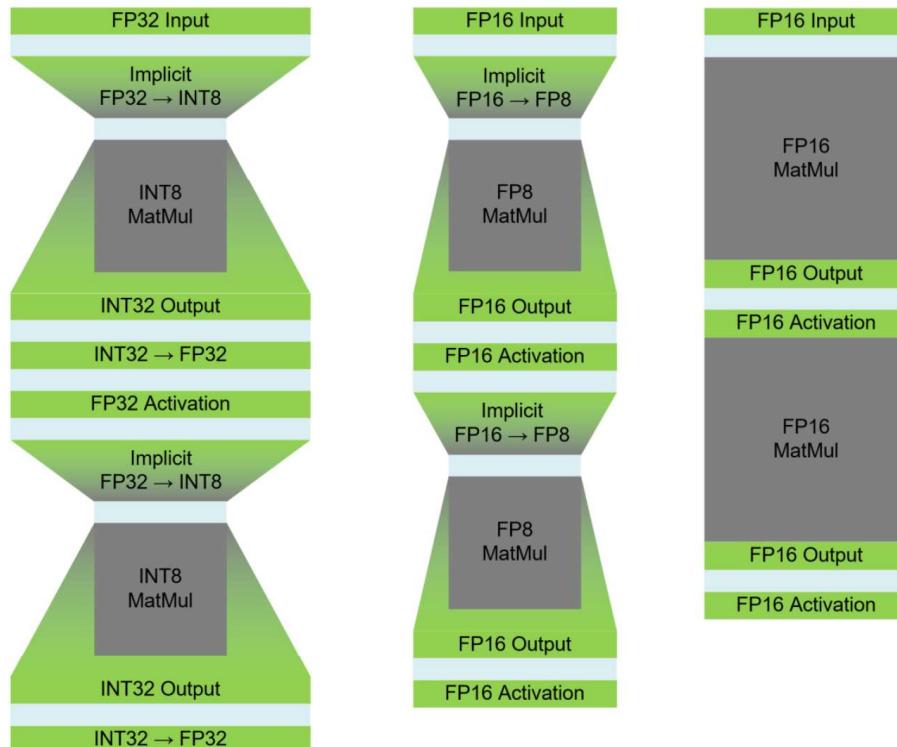
GOOD

```
coopVecMatMulAdd(layer0, input, weightBuffer, offset0,
                  biasBuffer, biasOffset0);

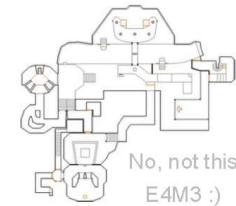
layer0 = max(layer0, 0); // ReLU

coopVecMatMul(layer1, layer0, weightBuffer, offset1);
```

MATH PRECISION CONSIDERATIONS

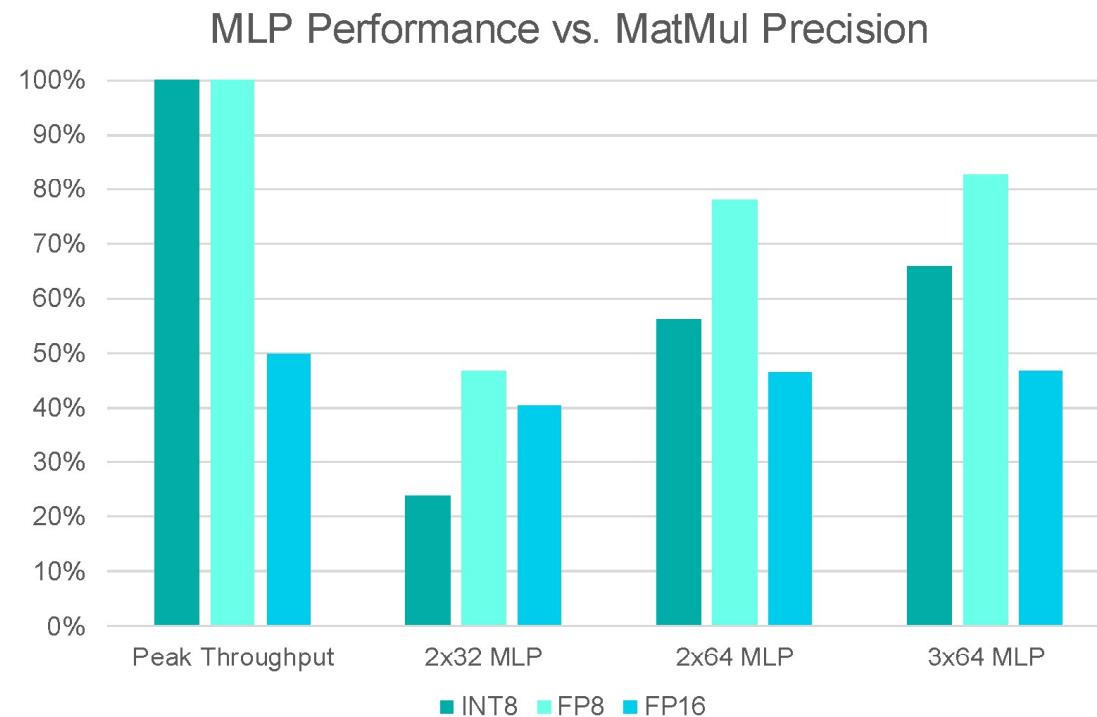


- **INT8 – Most difficult to use**
 - Extra scale operations needed for each layer
 - INT->FP conversions are not cheap
 - 32-bit per element vectors take many registers
- **FP8 (E4M3 and E5M2)**
 - No scaling or conversions needed, small register footprint
 - Very low precision, mostly suitable for hidden layers
- **FP16 – Easiest to use**
 - Lower peak throughput than FP8 or INT8
 - Enough precision for all practical inference needs



No, not this
E4M3 :)

MATH PRECISION COMPARISON



PERFORMANCE TIPS - SUMMARY



- Convert matrices to optimal layouts offline
- Minimize matrix divergence
- Avoid elementwise access to vectors between network layers
- Watch out for layer fusion failures

Math Precision Tradeoffs

	INT8	FP8	FP16
Precision	✓	✗	✓ ✓
Throughput	✓	✓ ✓	✗
Register Pressure	✗	✓	✓
Ease of Use	✗	✓	✓ ✓
GPU Compatibility	✓ ✓	✓	✓ ✓

THE TRANSITION STORY

Training + Inference



↑
↓
SlangPy



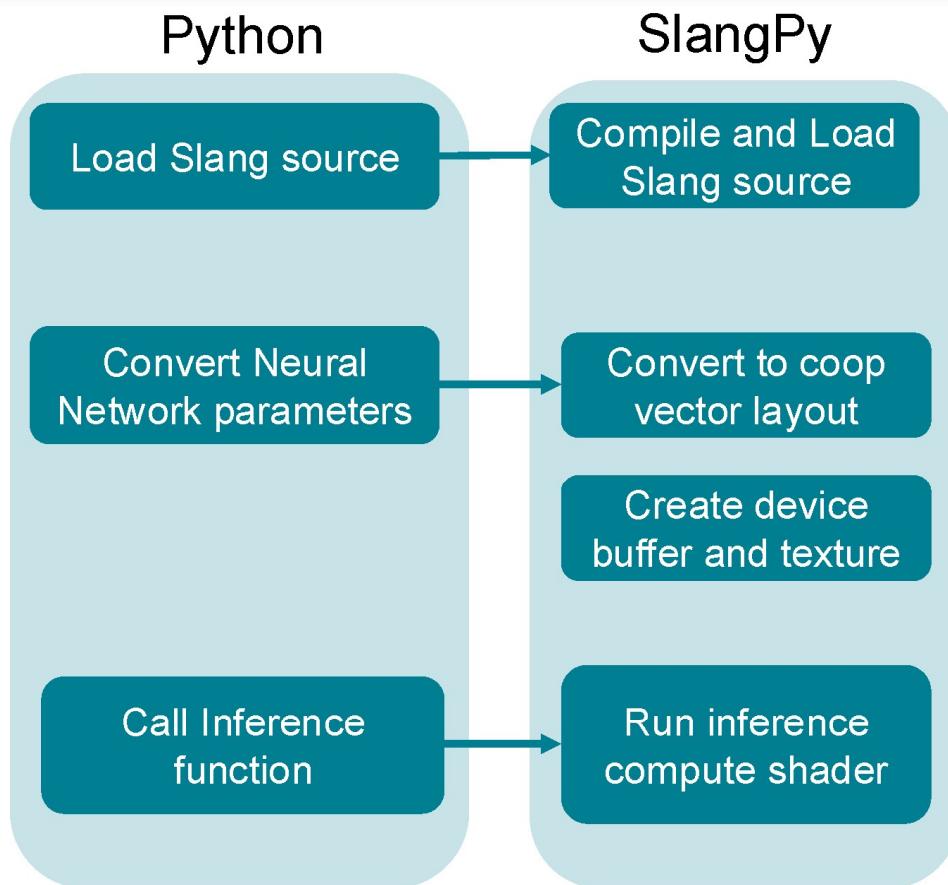
Inference



↑
↓



PYTHON IMPLEMENTATION



```
import slangpy as spy
module = spy.Module.load_from_file(app.device, "inference.slang")
app = App(width=512*3+10*2, height=512, title="Mipmap Example",
          device_type=spy.DeviceType.vulkan)

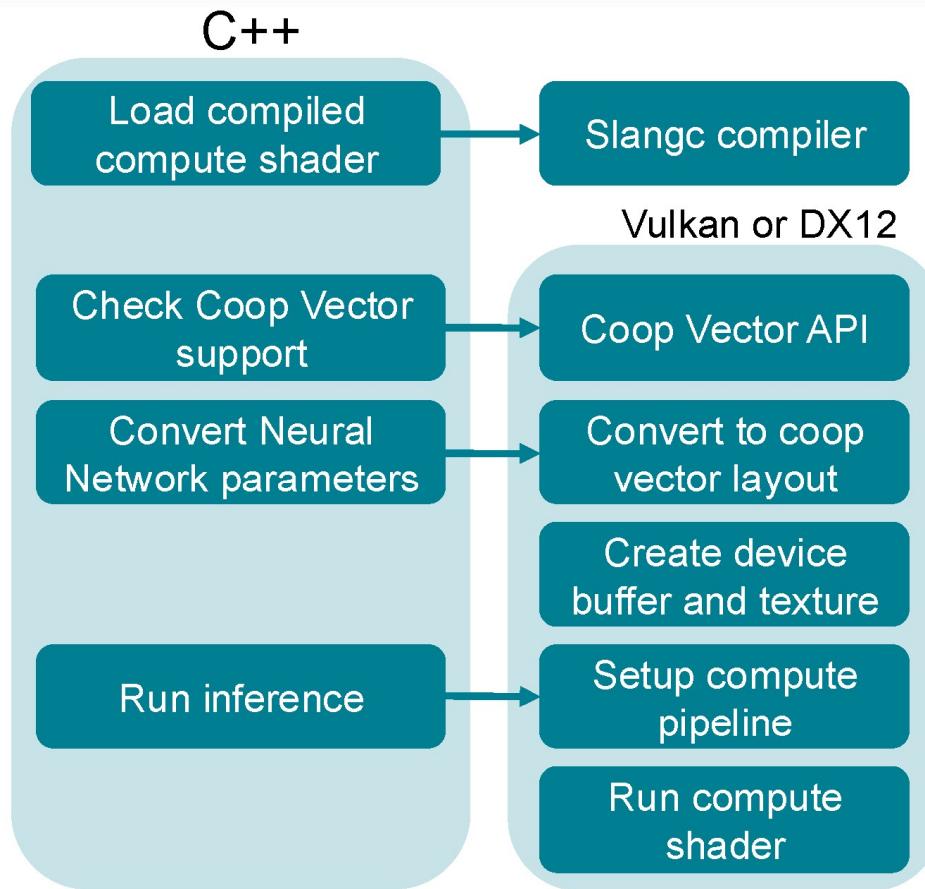
# Load values of biases and weights
weights_np = np.array(data['weights'], dtype=np.float16).reshape(outputs, inputs)
biases_np = np.array(data['biases'], dtype=np.float16)

# Convert weights into coopvec layout
desc = app.device.coopvec_create_matrix_desc(self.outputs, self.inputs,
                                              self.layout,
                                              spy.DataType.Float16, 0)
weight_count = desc.size // 2 # sizeof(half)
params_np = np.zeros((weight_count, ), dtype=np.float16)
app.device.coopvec_convert_matrix_host(weights_np, params_np,
                                         dst_layout=self.layout)

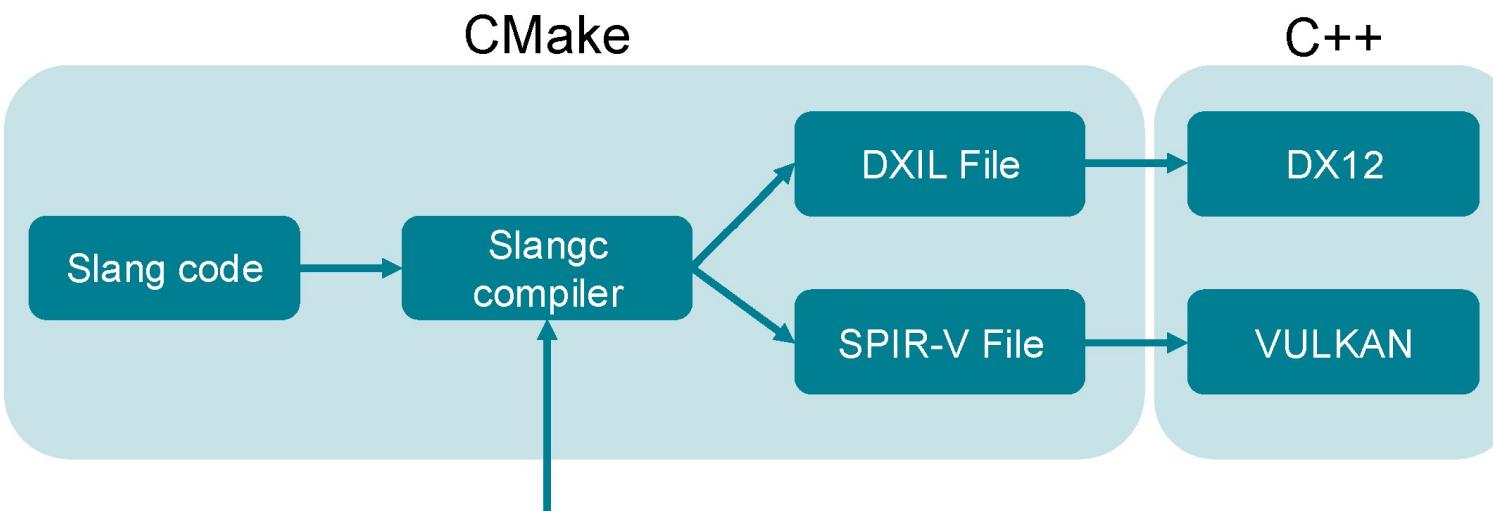
self.biases = app.device.create_buffer(struct_size=2, element_count=self.outputs,
                                       data=biases_np)
self.weights = app.device.create_buffer(struct_size=2, element_count=weight_count,
                                       data=params_np)

lr_output = spy.Tensor.empty_like(image)
module.inference(pixel = spy.call_id(),
                  resolution = res,
                  network = network,
                  _result = lr_output)
```

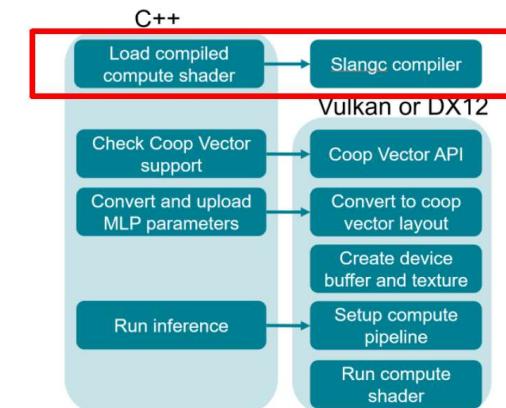
C++ IMPLEMENTATION



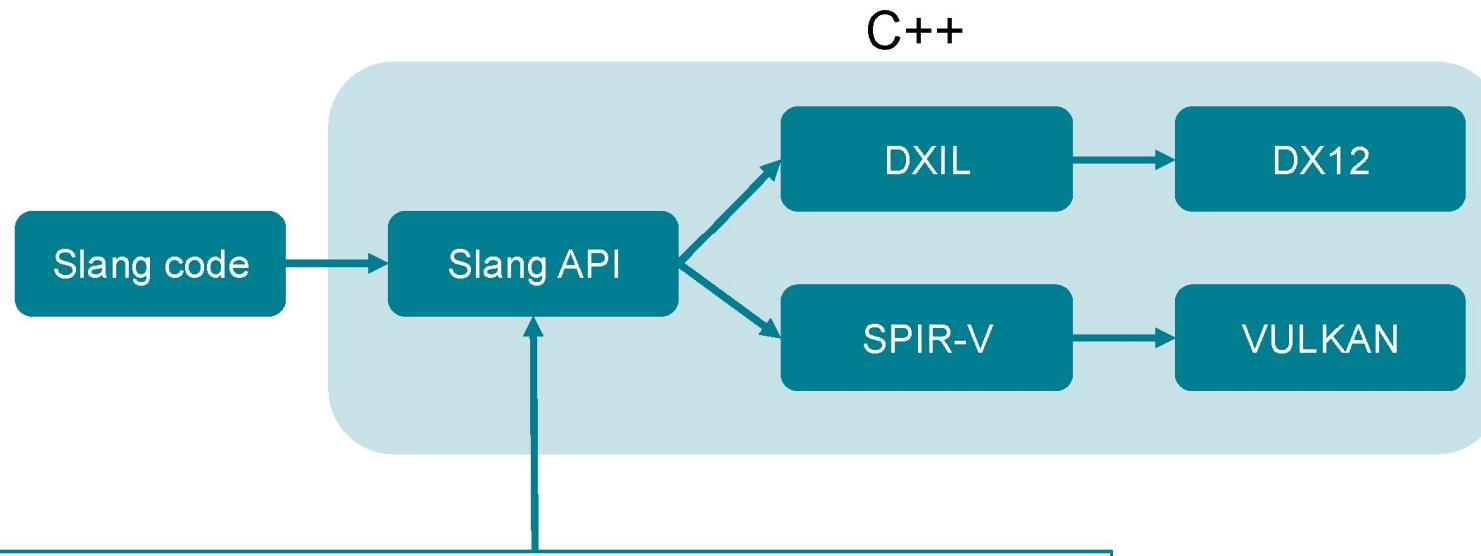
STEP 1 – SLANG AS SHADER MODULE



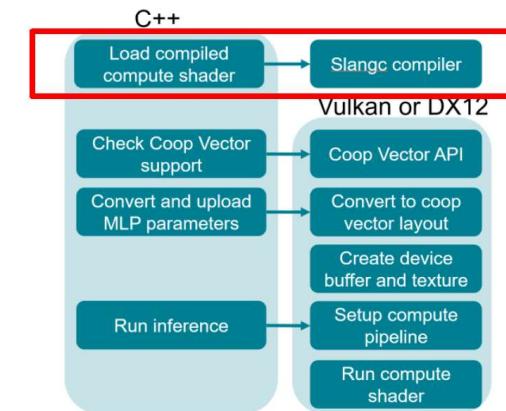
-capability spvCooperativeVectorNV
-capability pvCooperativeVectorTrainingNV



STEP 1 – SLANG AS SHADER MODULE



- Cooperative Vector parameters
 - -capability `spvCooperativeVectorNV`
 - -capability `pvCooperativeVectorTrainingNV`
- Compile time parameters
- Template parameters
- Reflection



STEP 2 - FEATURE DETECTION FOR COOPERATIVE VECTORS (VULKAN)



Required Vulkan device extensions:

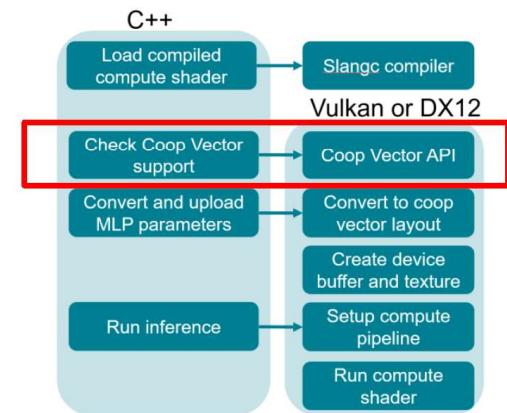
`VK_NV_cooperative_vector`

Required Vulkan physical device features:

`VK_STRUCTURE_TYPE_PHYSICAL_DEVICE_COOPERATIVE_VECTOR_FEATURES_NV`

`cooperativeVector = true`

`cooperativeVectorTraining = true`



STEP 2 - FEATURE DETECTION FOR COOPERATIVE VECTORS (DX12)

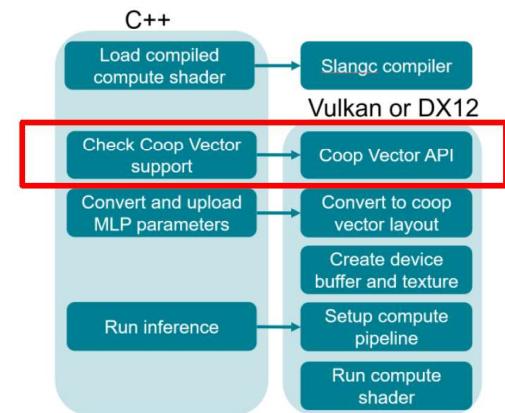


Call `D3D12EnableExperimentalFeatures` to enable following features:

- `D3D12ExperimentalShaderModels`
- `D3D12CooperativeVectorExperiment`

Call D3D12 device method `CheckFeatureSupport` for `D3D12_FEATURE_D3D12_OPTIONS_EXPERIMENTAL` and check:

- `D3D12_COOPERATIVE_VECTOR_TIER_1_0` for inference support
- `D3D12_COOPERATIVE_VECTOR_TIER_1_1` for training support



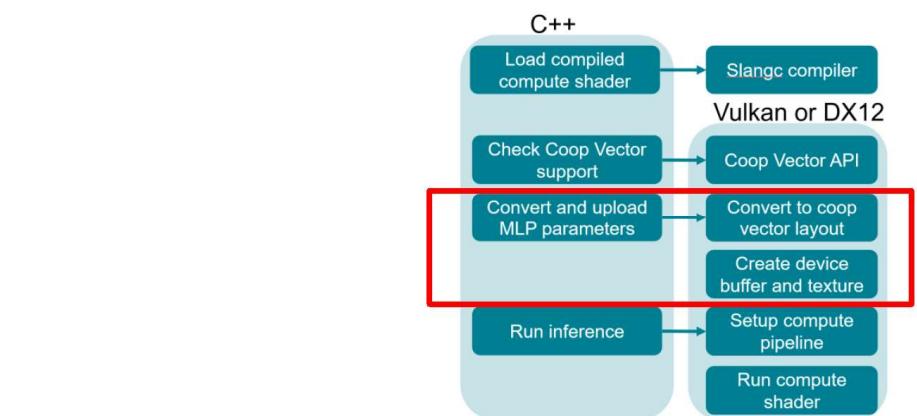
STEP 3 – MATRIX OPTIMAL LAYOUTS



FEED-FORWARD LAYER USING COOPERATIVE VECTORS

```
struct NetworkParameters<int Inputs, int Outputs>
{
    ByteAddressBuffer weights, biases;

    CoopVec<half, Outputs> forward(CoopVec<half, Inputs> x)
    {
        return coopVecMatMulAdd<half, Outputs>(
            x, CoopVecComponentType.Float16,
            weights, 0, CoopVecComponentType.Float16,
            biases, 0, CoopVecComponentType.Float16,
            CoopVecMatrixLayout.InferencingOptimal,
            false,
            0
        );
    }
}
```



- Tensor cores use custom matrix layouts
 - Components of matrices are shuffled between threads
 - Specific layout depends on the GPU and data types

STEP 3 – MATRIX LAYOUT CONVERSION

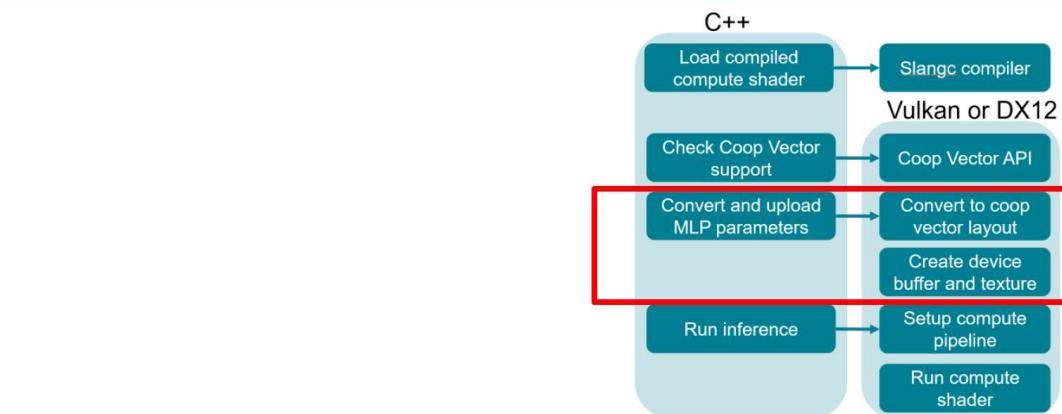


1. Create parameters buffer in optimal layout
 - Query shuffled matrix size from VK/DX12
 - Upload row- or column-major matrix to GPU
 - Convert layout (perform the shuffle) on the GPU
 - Inference Optimal
 - Training Optimal

Vulkan

- `vkConvertCooperativeVectorMatrixNV`
 - CPU conversion / size query
- `vkCmdConvertCooperativeVectorMatrixNV`
 - GPU conversion

2. Create constant buffer
3. Create output texture



DX12

- `GetLinearAlgebraMatrixConversionDestinationInfo`
 - Size query
- `ConvertLinearAlgebraMatrix`
 - GPU conversion

STEP 4 – CONVERT INFERENCE FUNCTION TO SHADER

SLANG

```
struct NetworkParameters<int Inputs, int Outputs> {
    ByteAddressBuffer<half> weights, biases;

    CoopVec<half, Outputs> forward(CoopVec<half, Inputs> x);
}
```

```
struct Network {
    NetworkParameters<16, 32> layer0, layer1, layer2;

    float3 eval(no_diff float2 uv);
}
```

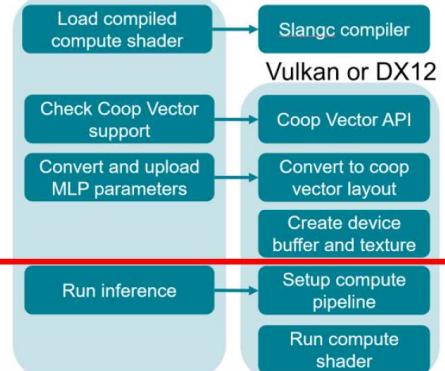
```
float3 inference(int2 pixel, int2 resolution, Network network)
{
    float2 uv = (float2(pixel) + 0.5f) / float2(resolution);
    return network.eval(uv);
}
```

C++

```
lr_output = spy.Tensor.empty_like(image)

module.inference(pixel = spy.call_id(), resolution = res, network = network, _result = lr_output)
```

C++



SLANGPY

Compute Shader

STEP 4 – CONVERT INFERENCE FUNCTION TO SHADER

SLANG

```
struct NetworkParameters<int Inputs, int Outputs> {
    ByteAddressBuffer <half> weights, biases;
    CoopVec<half, Outputs> forward(CoopVec<half, Inputs> x);
}

struct Network {
    NetworkParameters<16, 32> layer0;
    NetworkParameters<32, 32> layer1;
    NetworkParameters<32, 3> layer2;

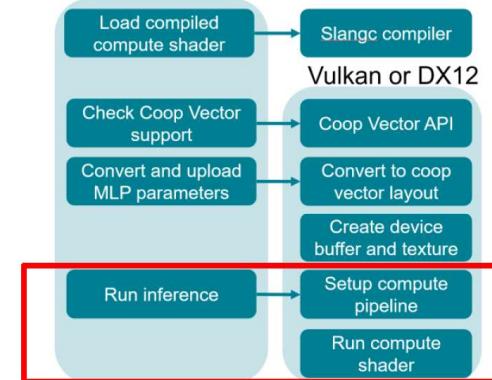
    float3 eval(no_diff float2 uv);
}

float3 inference(int2 pixel, int2 resolution, Network network)
{
    float2 uv = (float2(pixel) + 0.5f) / float2(resolution);
    return network.eval(uv);
}
```

```
Network network;
ConstantBuffer<NeuralConstants> gConst;
RWTexture2D<float4> outputTexture;

[shader("compute")]
[numthreads(8, 8, 1)]
void main_cs(uint3 pixel : SV_DispatchThreadID)
{
    outputTexture[pixel.xy] = inference(pixel.xy,
                                         gConst.resolution,
                                         network);
}
```

C++



STEP 4 – CONVERT INFERENCE FUNCTION TO SHADER

SLANG

```
struct NetworkParameters<int Inputs, int Outputs> {
    ByteAddressBuffer <half> weights, biases;
    CoopVec<half, Outputs> forward(CoopVec<half, Inputs> x);
}

struct Network {
    NetworkParameters<16, 32> layer0;
    NetworkParameters<32, 32> layer1;
    NetworkParameters<32, 3> layer2;

    float3 eval(no_diff float2 uv);
}

Network network;

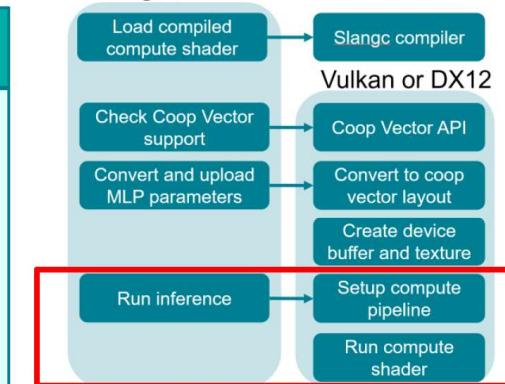
ConstantBuffer<NeuralConstants> gConst;
RWTexture2D<float4> outputTexture;

[shader("compute")]
[numthreads(8, 8, 1)]
void main_cs(uint3 pixel : SV_DispatchThreadID){
    // Compute shader
}
```

C++

```
// Initialize parameters buffers
// Setup up constant buffer
// Bind output texture
// Dispatch compute shader
nvrhi::ComputeState state;
m_CommandList->setComputeState(state);
m_CommandList->dispatch(textureWidth, textureHeight, 1);
```

C++



Run inference

Setup compute pipeline
Run compute shader

Load compiled compute shader
Slangc compiler
Vulkan or DX12

Check Coop Vector support
Coop Vector API

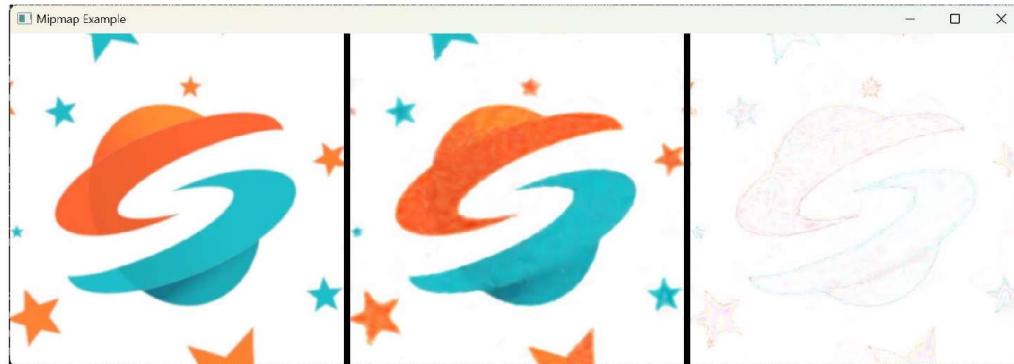
Convert and upload MLP parameters
Convert to coop vector layout
Create device buffer and texture

Run inference
Setup compute pipeline
Run compute shader

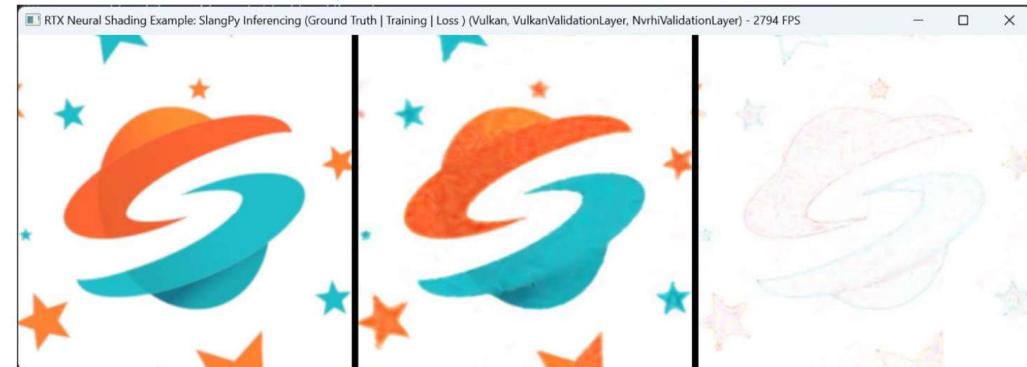
YOUR NEURAL SHADER IS NOW SHIPPED!



Python



C++

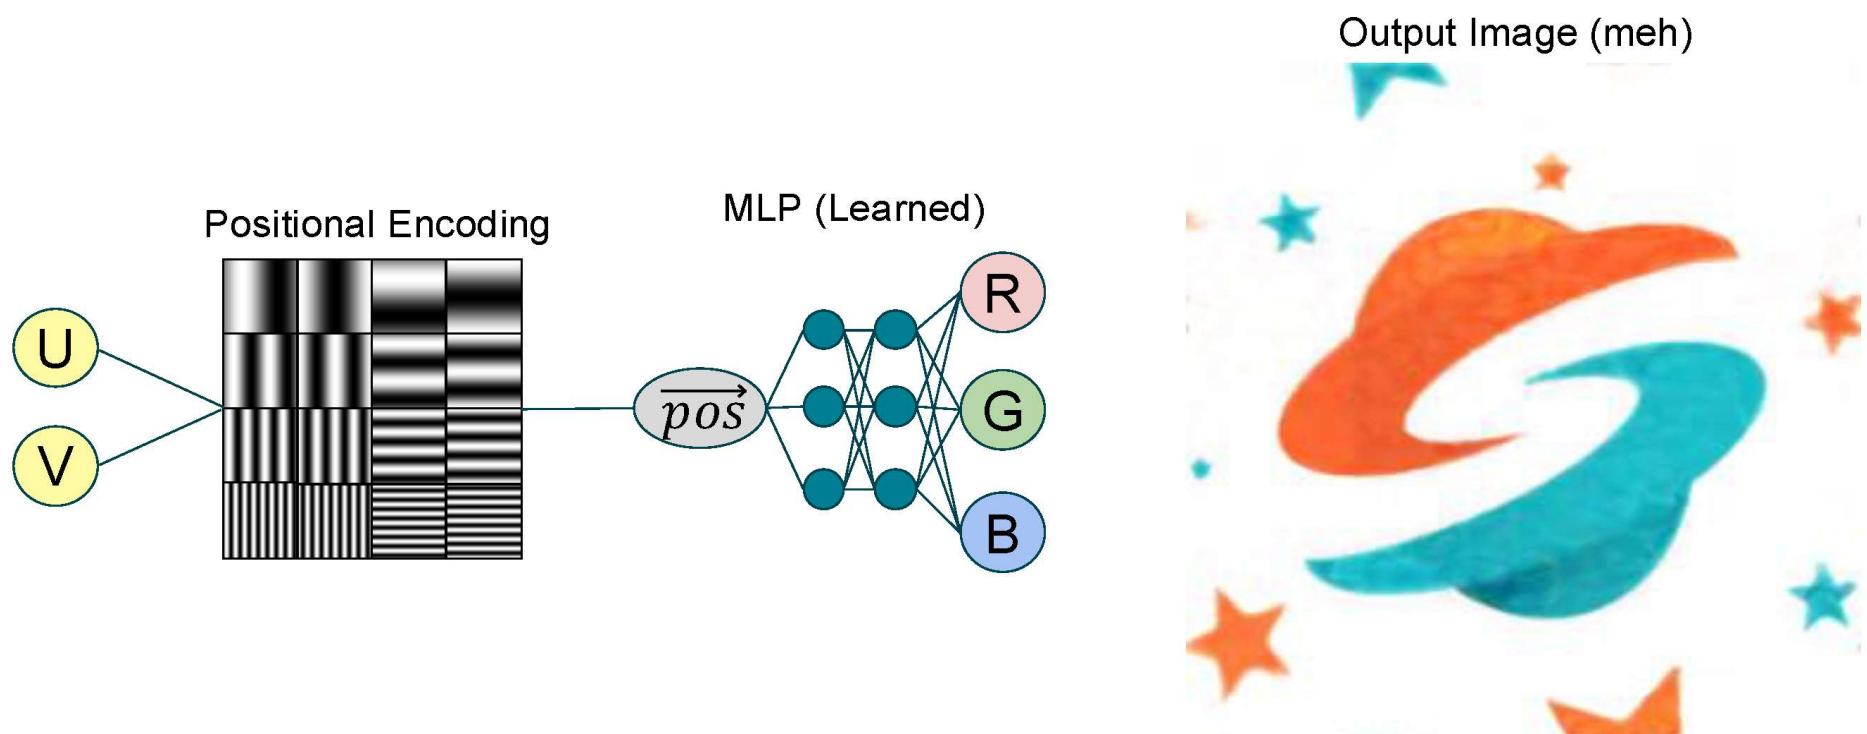




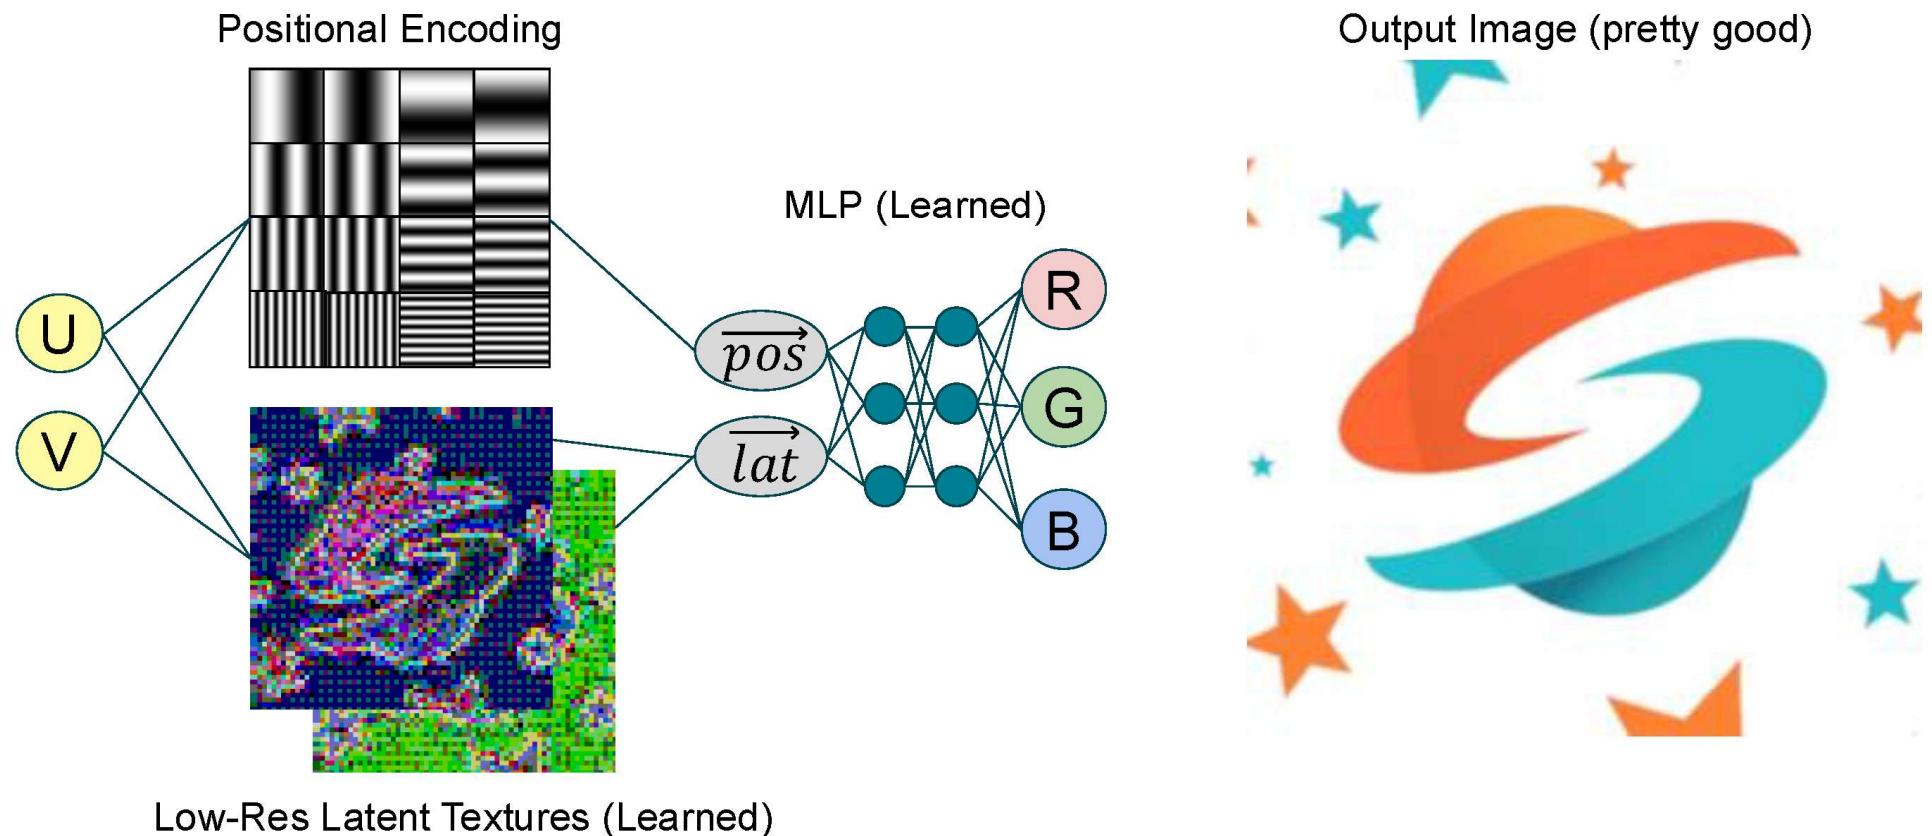
 SIGGRAPH 2025
Vancouver+ 10-14 August

NEURAL TEXTURE COMPRESSION

FLASHBACK: PREVIOUS BEST RESULT



ADDING LATENT TEXTURES



KEY FACTS ABOUT NTC

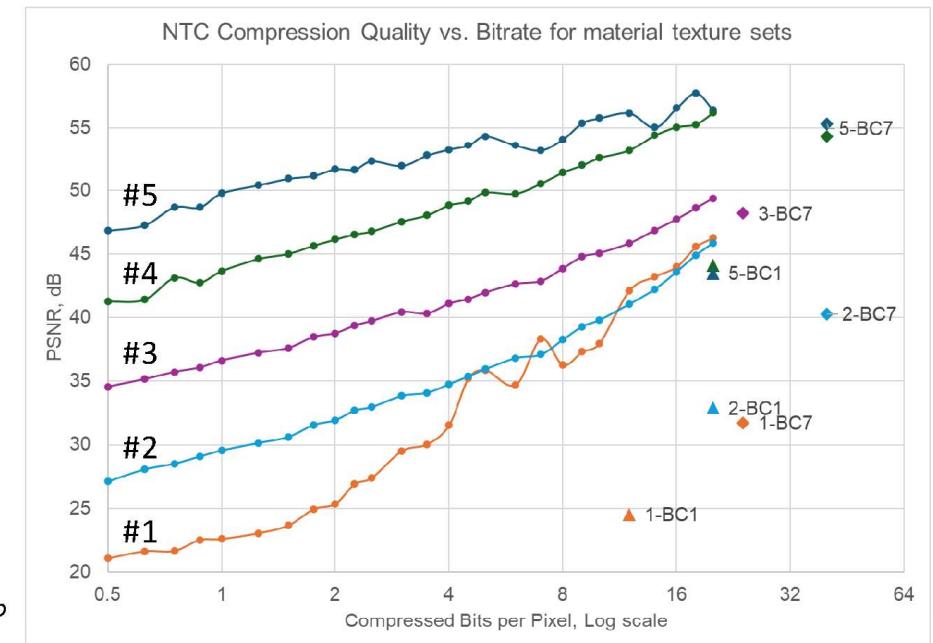
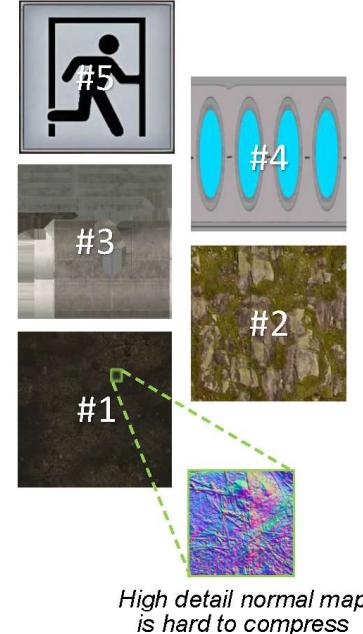
- Wide range of latent texture parameters
 - Independent resolution, channel count, bits per feature controls
 - Overall practical range of 0.5 — 20 bits per pixel storing up to 16 channels
 - MLP weights are tiny compared to the latents (~12 KB vs. many MB)
- Each texel can be decoded independently
 - Suitable for use as a replacement for regular material textures
- Neural compression without hallucinations
 - Compression is training of the network and latent textures
 - Network is small and overfitted to reproduce only one material
- Original paper describing NTC and STF:
 - K. Vaidyanathan et al. "Random-Access Neural Compression of Material Textures"



Crops from an NTC compressed texture at 0.5 and 20.0 bpp

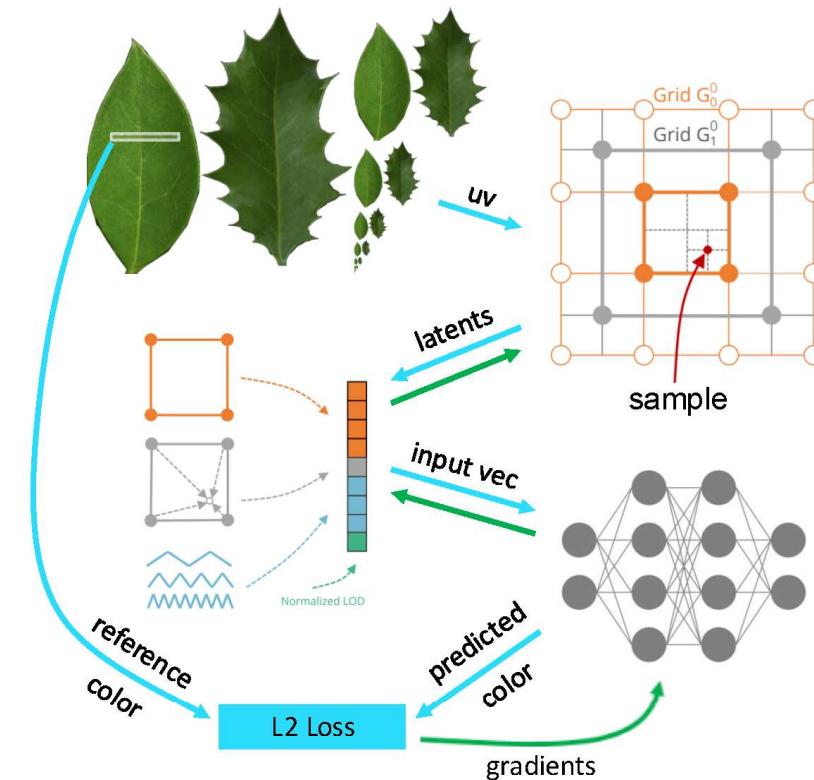
NEURAL TEXTURE COMPRESSION QUALITY

- Consistently better than BCn
 - More benefit for complex materials
 - Up to 8x better at similar visual quality
 - Introduces blur, not blockiness
- Mostly monotonic vs. BPP
 - More information means higher quality
 - Training finds effective ways to use the bits
- Benefits from correlated channels
 - All textures for a PBR material are better compressed together than separately



NTC TRAINING PROCESS

- Pick a set of pixels to process on this step
 - 64k pixels per step is usually good
- For each pixel in parallel:
 - Calculate positional encoding based on UV
 - Sample the latent textures at UV
 - Prepare MLP input vector
 - Evaluate MLP
 - Load reference texture channels
 - Compute gradient \approx (predicted – reference)
 - Backpropagate gradients
 - Accumulate gradients for MLP and features
- Run an optimizer step
 - Add the quantization noise
- Repeat 100-200k times



NTC MODES OF OPERATION 1/3

INFERENCE ON SAMPLE

- Replace material texture sampling with NTC decode
- Apply Stochastic Texture Filteringing (STF) instead of hardware filtering
 - NTC gives one pixel per decode, direct trilinear or aniso filtering would be too expensive
- Benefits:
 - Minimal VRAM footprint
- Drawbacks:
 - Makes shaders larger and slower (depends...)
 - STF is a requirement

Random-Access Neural Compression of Material Textures

KARTHIK VAIDYANATHAN*, NVIDIA, USA
MARCO SALVI*, NVIDIA, USA
BARTLOMIEJ WRONSKI*, NVIDIA, USA
TOMAS AKENINE-MÖLLER, NVIDIA, Sweden
PONTUS EBELIN, NVIDIA, Sweden
AARON LEFOHN, NVIDIA, USA

BC high: PSNR (↑) 194.4 dB, TLDP (↓) 0.224
1024 × 1024 at 5.3 MB

NTC: PSNR (↑) 222.0 dB, TLDP (↓) 0.177
4096 × 4096 at 3.8 MB

reference not compressed
4096 × 4096 at 256 MB

Filtering After Shading With Stochastic Texture Filteringing

MATT PHARR*, NVIDIA, USA
BARTLOMIEJ WRONSKI*, NVIDIA, USA
MARCO SALVI, NVIDIA, USA
MARCOS FAJARDO, Shiokara—Engawa Research, Spain

2D texture maps and 3D voxel arrays are widely used to add rich detail to the surfaces and volumes of rendered scenes, and filtered texture lookups are integral to producing high-quality images. We show that applying the texture textures before BSDF apparent in common through the use of

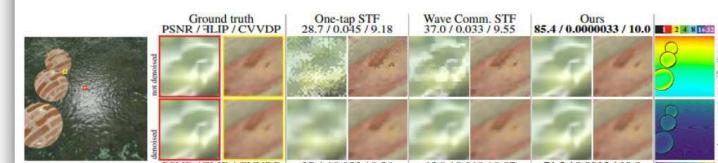
Collaborative Texture Filteringing

T. Akenine-Möller, P. Ebelin, M. Pharr, and B. Wronski

NVIDIA

Ground truth
PSNR / TLDP / CVVDP

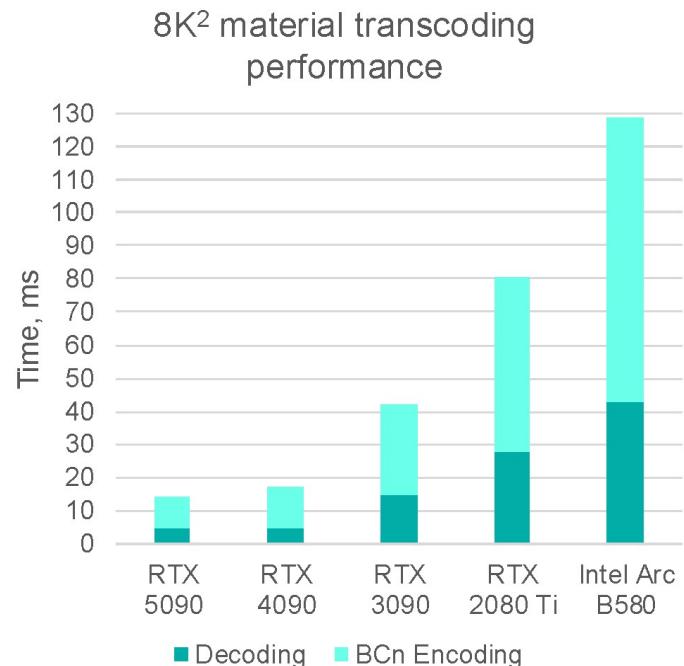
One-tap STF 28.7 / 0.045 / 9.18	Wave Comm. STF 37.0 / 0.033 / 9.55	Ours 85.4 / 0.0000033 / 10.0
PSNR / TLDP / CVVDP 37.4 / 0.033 / 9.36	45.8 / 0.019 / 9.87	71.2 / 0.0013 / 10.0



NTC MODES OF OPERATION 2/3

INFERENCE ON LOAD

- Transcode NTC materials into BCn at load time
- Sample BCn textures as usual
- **Benefits:**
 - Easy integration
 - Rendering performance is unaffected
- **Drawbacks:**
 - No VRAM savings, only disk space / network traffic



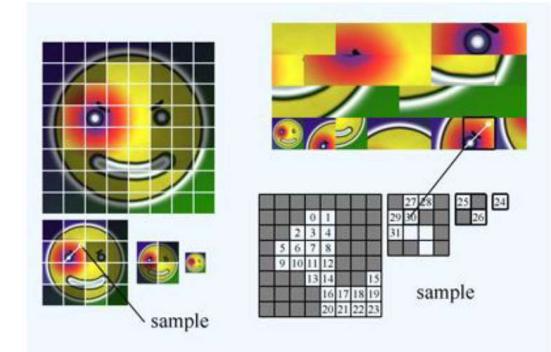
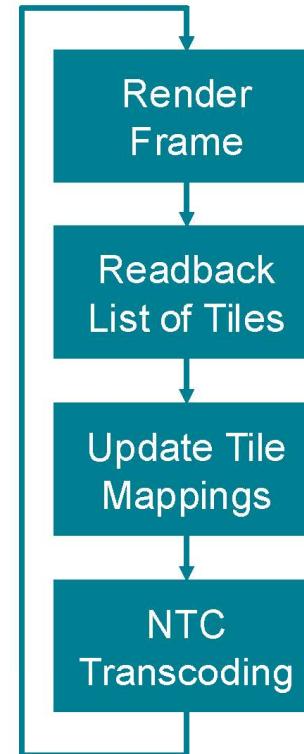
Using FP8 inference with CoopVec for NTC

Encoding 2x BC7 and 1x BC5

NTC MODES OF OPERATION 3/3

INFERENCE ON FEEDBACK

- Extension for a virtual texturing system
- Track which texture MIPs and tiles are needed
- Transcode the necessary tiles to BCn at runtime
- **Benefits:**
 - Low performance impact on render passes
- **Drawbacks:**
 - VRAM usage includes both NTC and partial BCn data
 - Potentially uneven frame times



Virtual texture

Image: Sean Barrett, GDC 2008

<https://silverspaceship.com/src/svt/>

NTC PERFORMANCE COMPARISON



Intel Sponza scene in the NTC SDK Renderer

Mode	Render Time	Texture Memory
On Load	0.39 ms	2041 MB
On Sample	0.82 ms	243 MB
On Feedback	0.65 ms	555 MB

Measured on an RTX PRO 6000 Blackwell WE

BENEFITS OF NEURAL TEXTURE COMPRESSION



PRACTICAL

- Saves disk space
- Saves download traffic
- Saves VRAM (maybe)
- Can be used on any platform
 - High-end PC – on-sample
 - Low-end and consoles – on-load or on-feedback
- Can be used now
 - SDK available: [github.com / NVIDIA-RTX / RTXNTC](https://github.com/NVIDIA-RTX/RTXNTC)

CONCEPTUAL

- Allows using higher detail materials
 - Fewer bits per pixel => more pixels with the same storage
- Can be extended with perceptual loss functions
 - Higher compression ratios with better visual detail
 - Ongoing research direction

MORE NEURAL COMPRESSION TECHNIQUES



Real-Time Neural Materials using Block-Compressed Features

C. Weinreich[†], L. De Oliveira[†], A. Houdard[†] and G. Nader[†]

Ubisoft La Forge

Neural Texture Block Compression

S. Fujieda and T. Harada

Advanced Micro Devices, Inc.

Uncompressed Diffuse

Ref. BC
48 MB, 42.53 dB



SIGGRAPH 2024
DENVER+ 28 JUL – 1 AUG

NEURAL LIGHT GRID

or adventures in trying to get something useful out of ML in production

Michał Iwanicki
Peter-Pike Sloan
Ari Silvennoinen
Peter Shirley

© 2024 SIGGRAPH ADVANCES IN REAL-TIME RENDERING IN GAMES course ALL RIGHTS RESERVED.



Global Illumination In "Once Human": A Hybrid Approach for 16km Open World

Maode Shi, Lele Pan



Decoding Light: Neural Compression of Global Illumination

Luyan Cao, LIGHTSPEED STUDIOS

INSPIRATION: REAL MATERIALS ARE COMPLEX



Real-Time Neural Appearance Models

Tizian Zeltner[†]
Benedikt Bitterli[†]

Fabrice Rousselle[†]
Alex Evans

Andrea Weidlich[†]
Tomáš Davidovič

Petriklarberg[†]
Simon Kallweit

Jan Novák[†]
Aaron Lefohn

NVIDIA

[†]equal contribution, order determined by a rock-paper-scissors tournament.



INSPIRATION: REAL MATERIALS ARE COMPLEX

SIGGRAPH 2025
Vancouver+ 10-14 August



WE CAN RENDER THESE, BUT NOT IN REAL-TIME



Metal teapot handle



Plastic slicer handle



Blue teapot ceramic



Metal slicer blade



Aged metal inkwell

THESE ARE COMPLEX MATERIAL GRAPHS...



Metal teapot handle



Plastic slicer handle



Aged metal inkwell

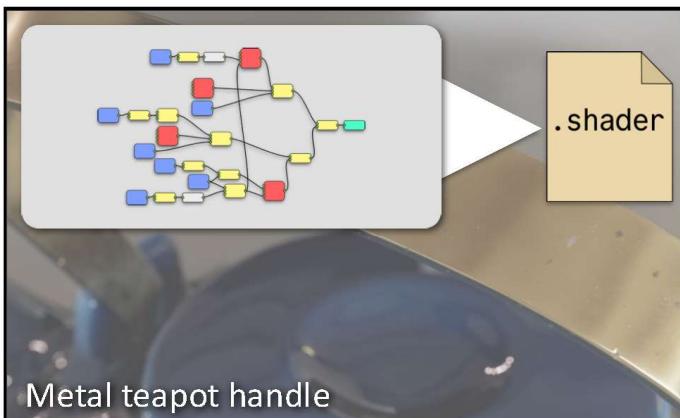


Blue teapot ceramic

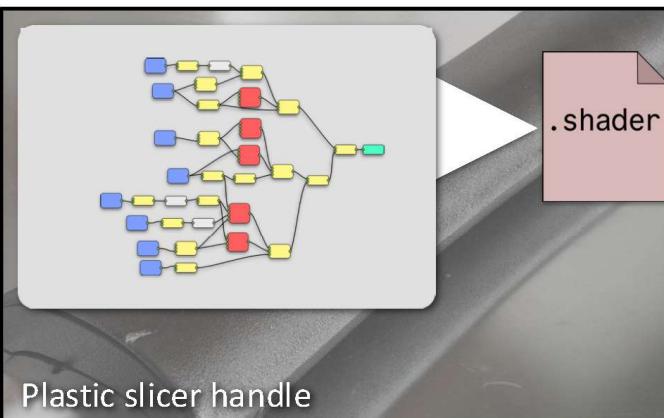


Metal slicer blade

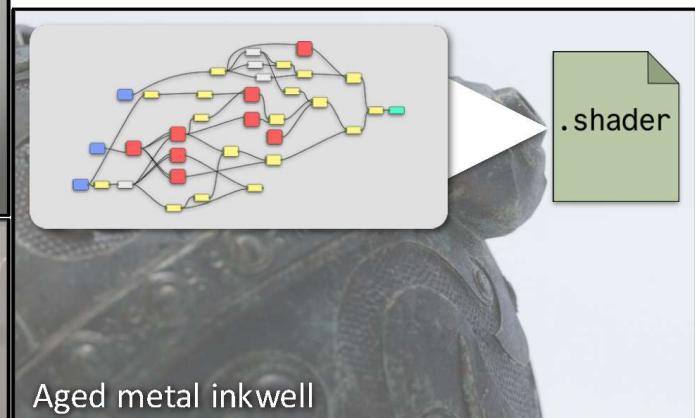
...AND WE DON'T KNOW HOW TO SIMPLIFY THEM WELL



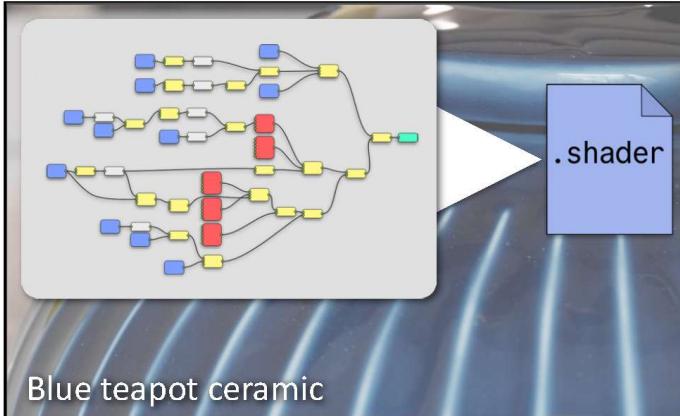
Metal teapot handle



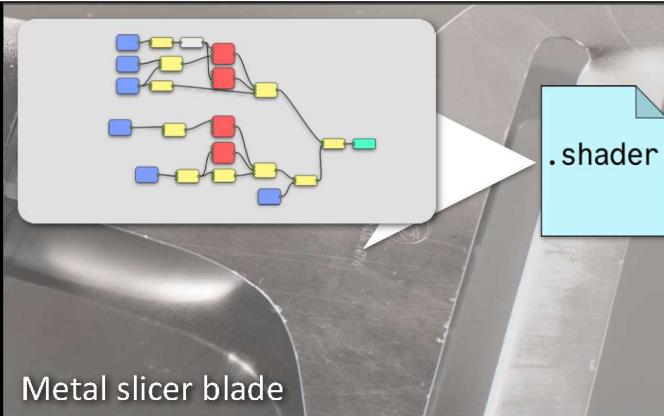
Plastic slicer handle



Aged metal inkwell



Blue teapot ceramic

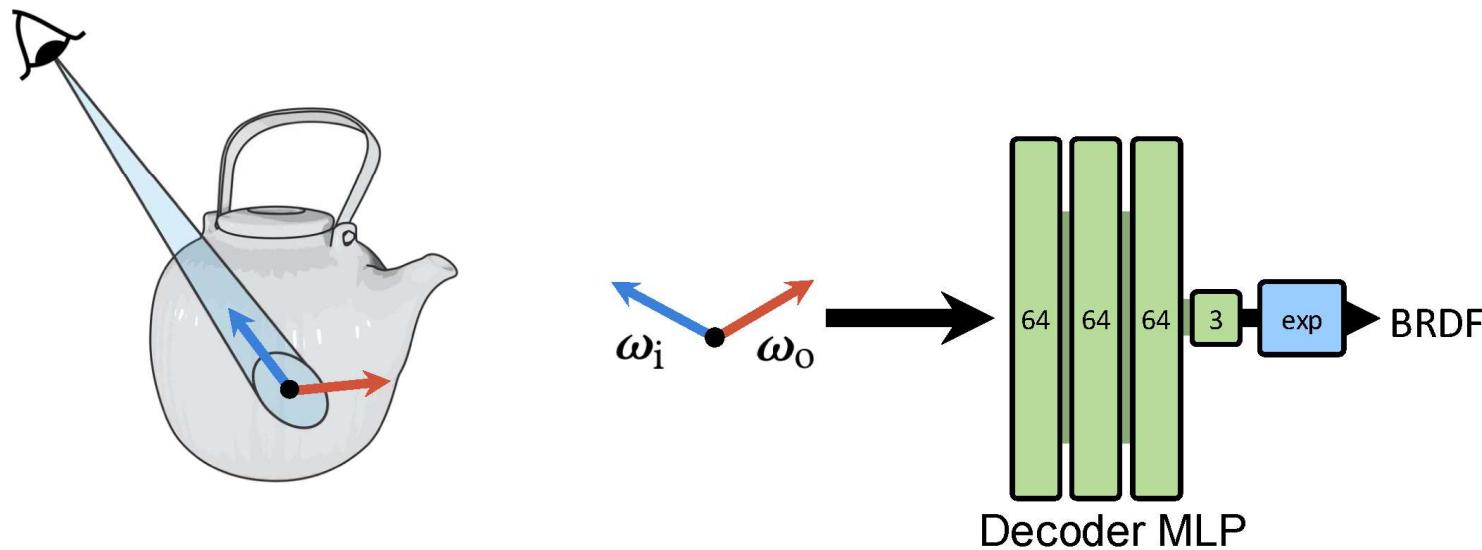


Metal slicer blade

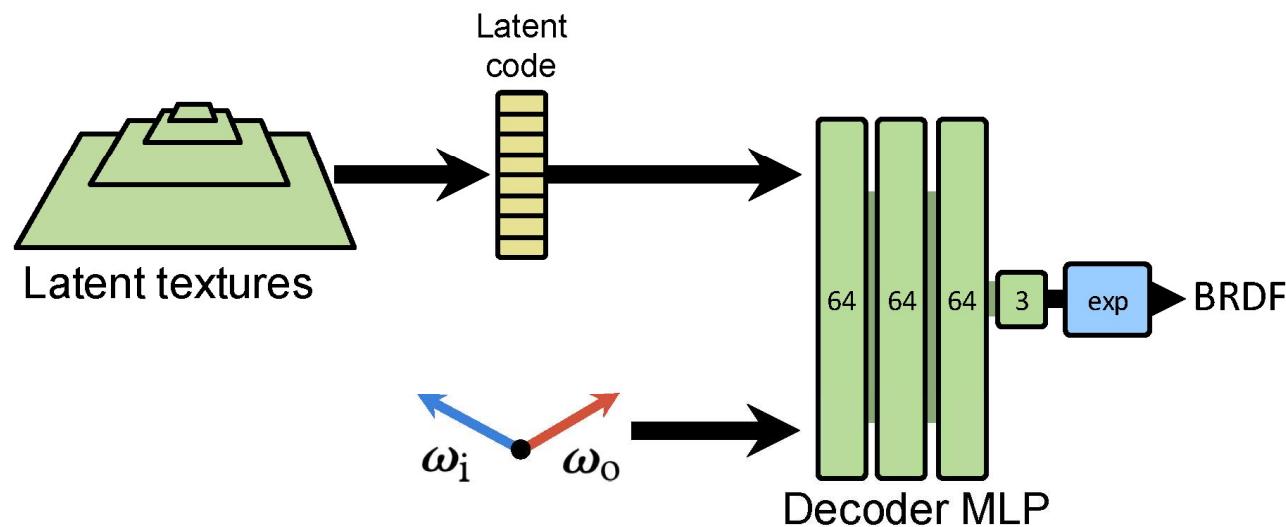
WHAT IF WE USED A NEURAL NETWORK FOR THE TASK?



SIMPLEST FIRST: PASS DIRECTIONS TO A NETWORK



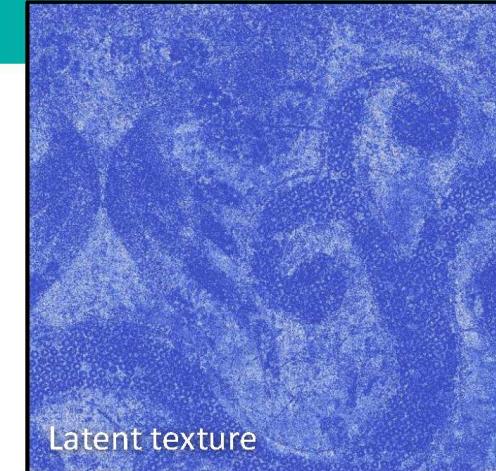
...AND THEN TEXTURE IT



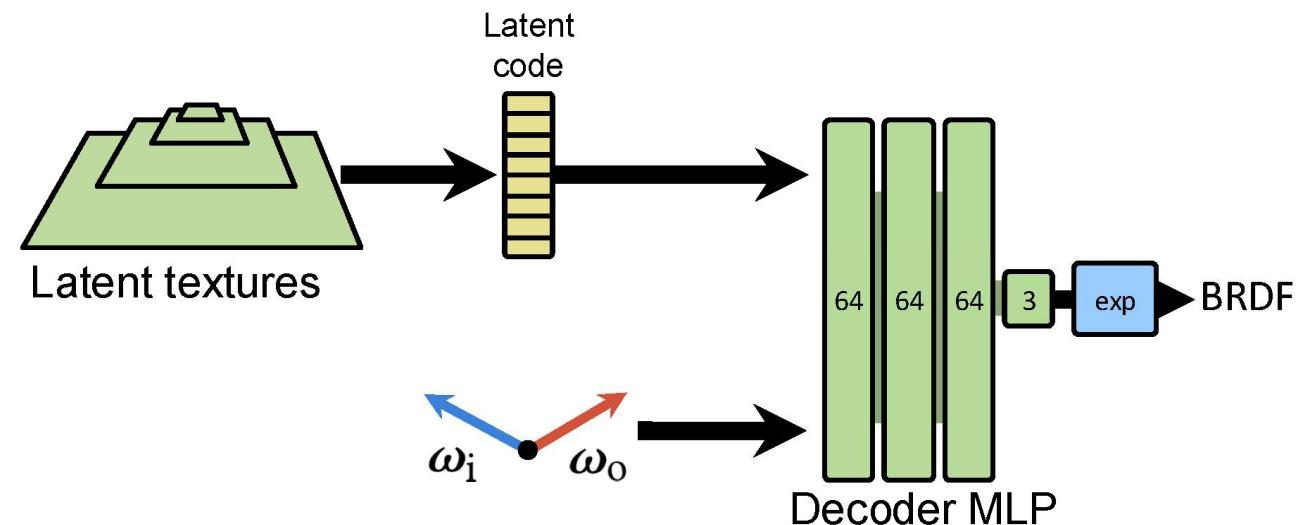
THIS WORKS!



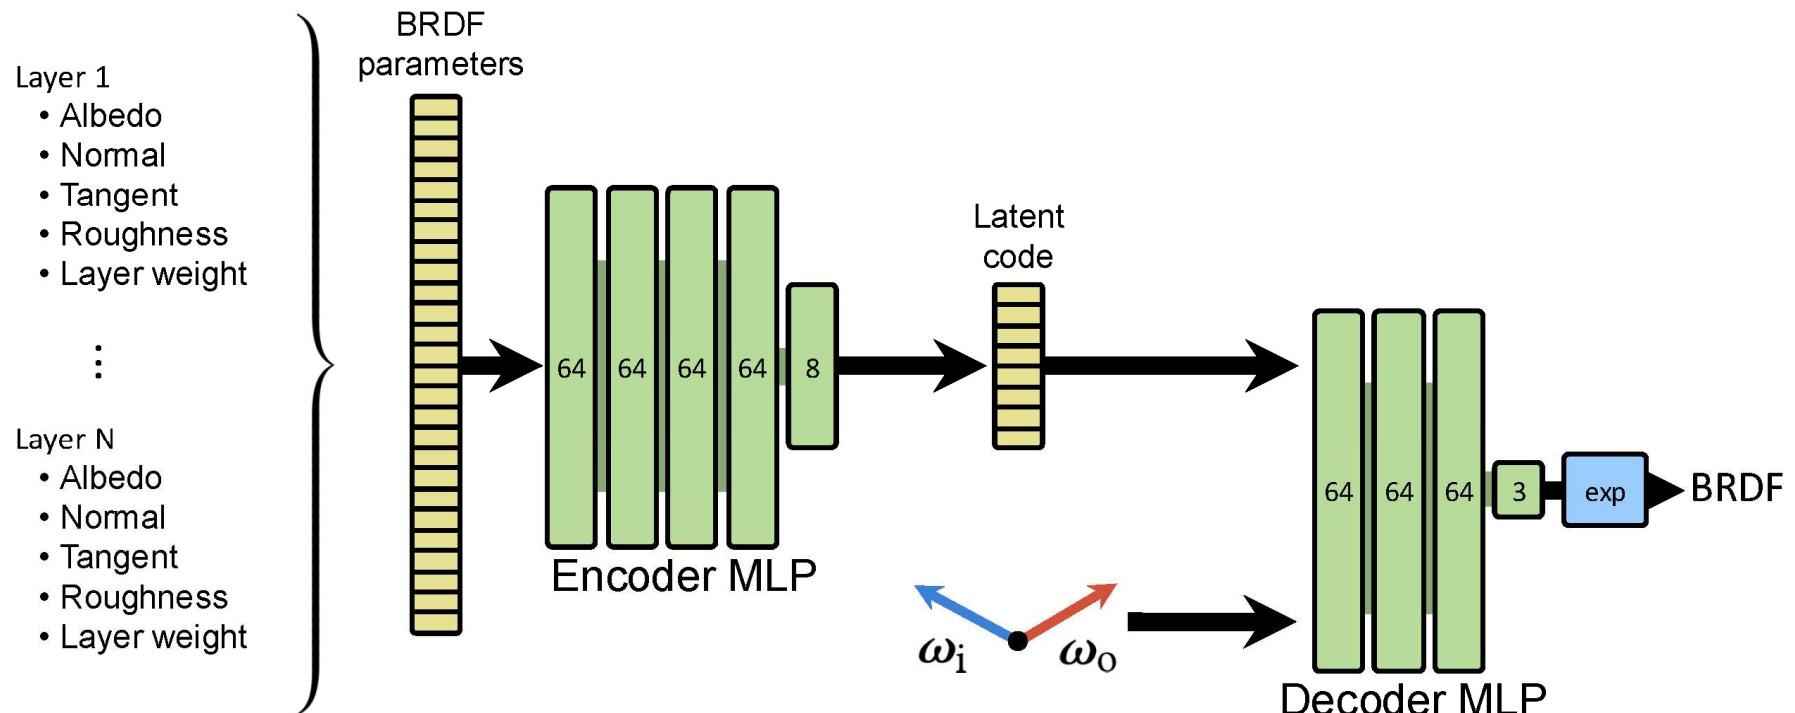
...UP TO A POINT



WE'RE LACKING A “GLOBAL” VIEW



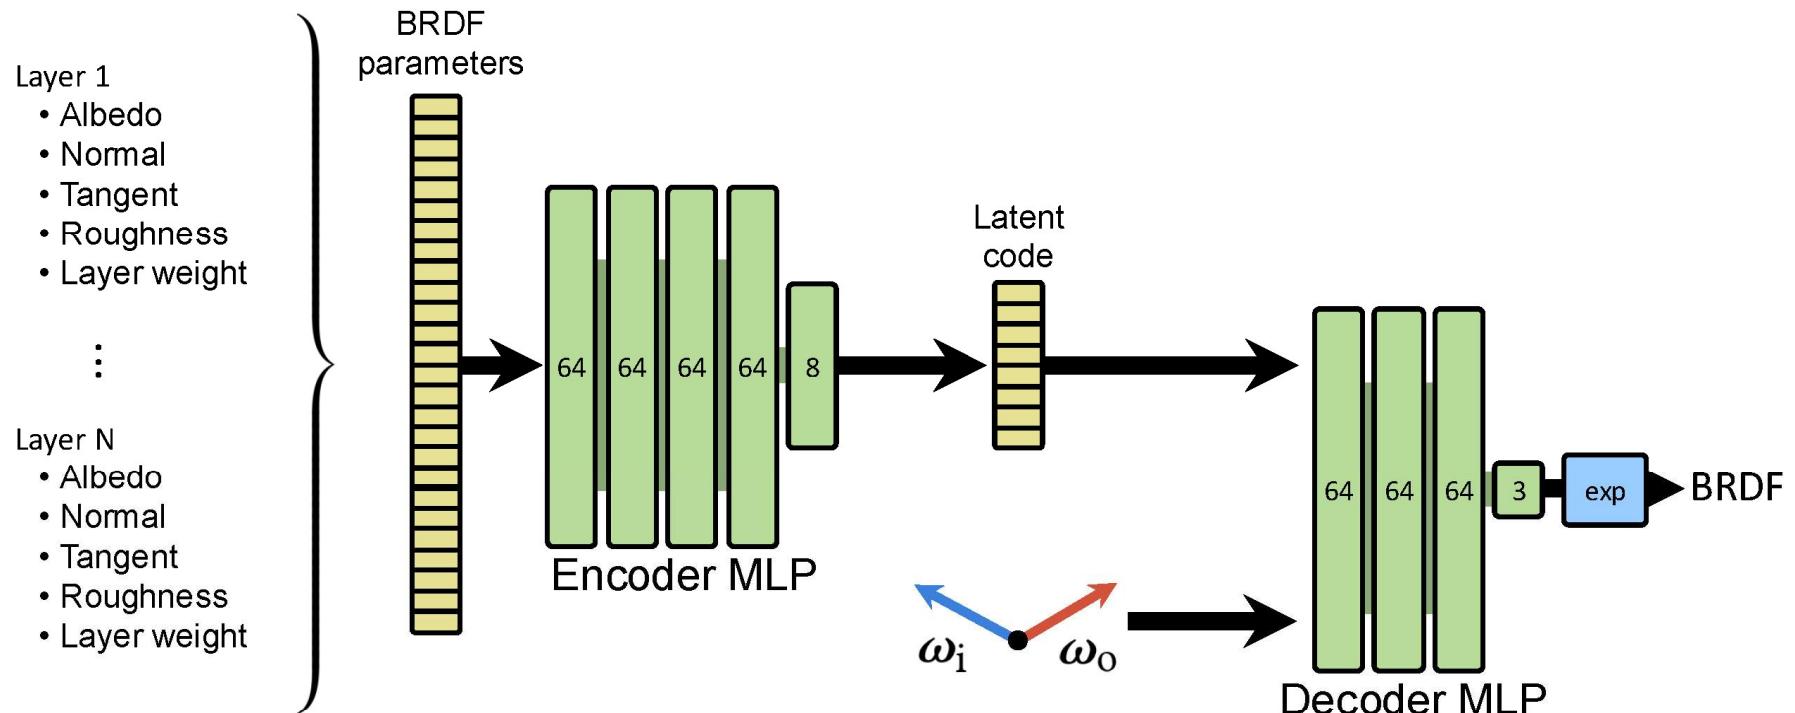
IDEA: LEARN TO TRANSLATE THE ORIGINAL TEXTURES



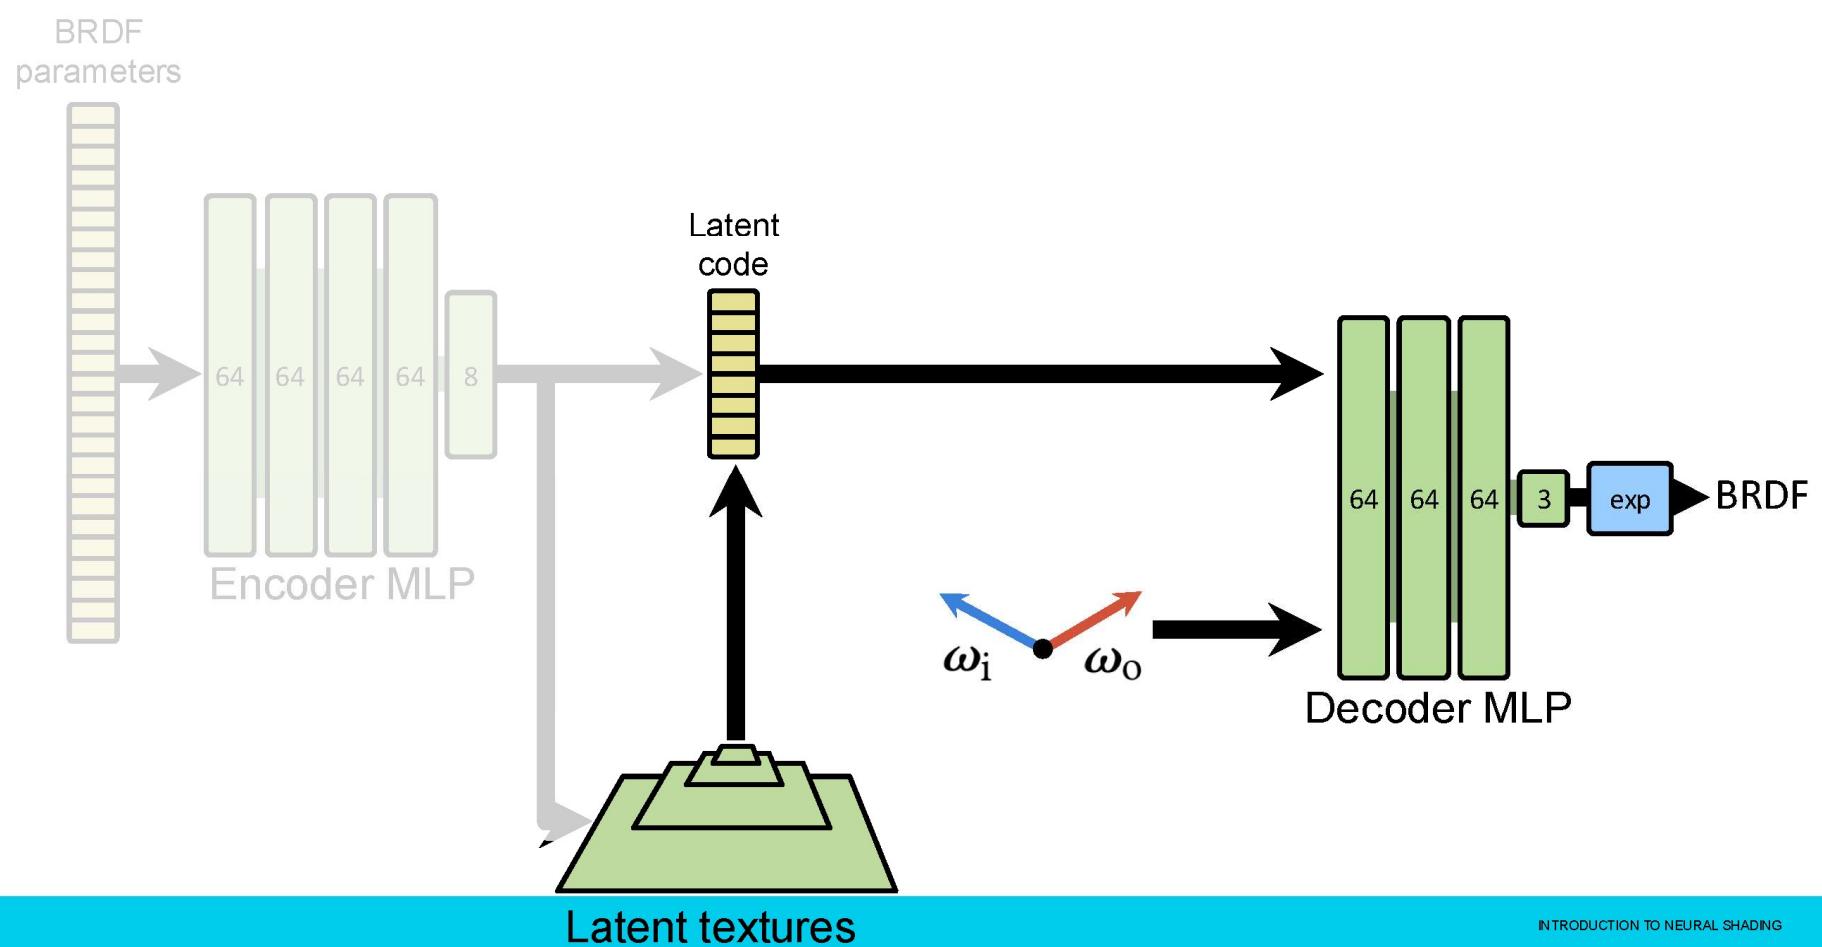
IDEA: LEARN TO TRANSLATE THE ORIGINAL TEXTURES

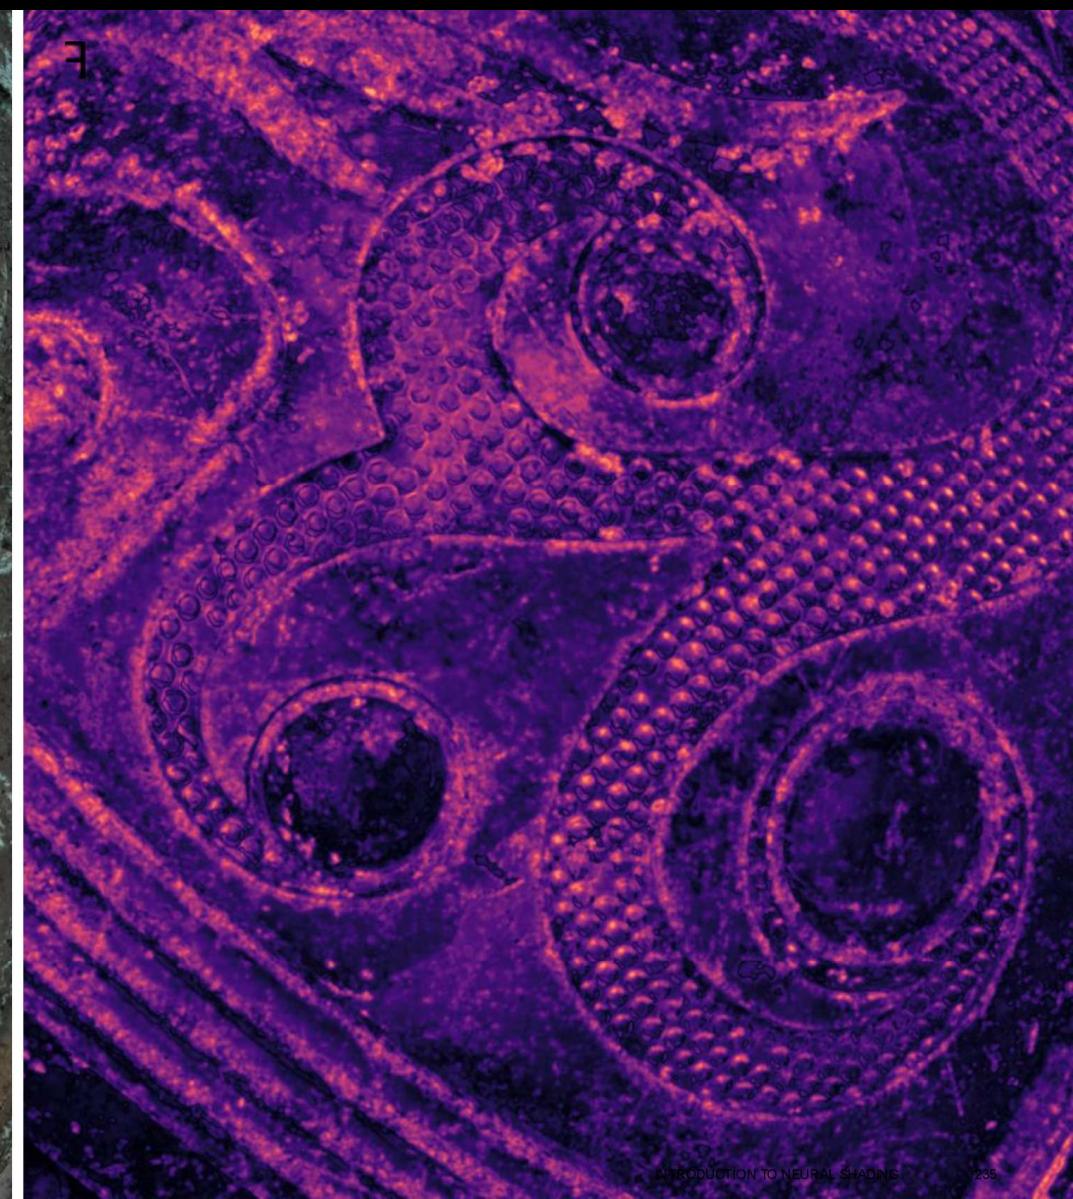


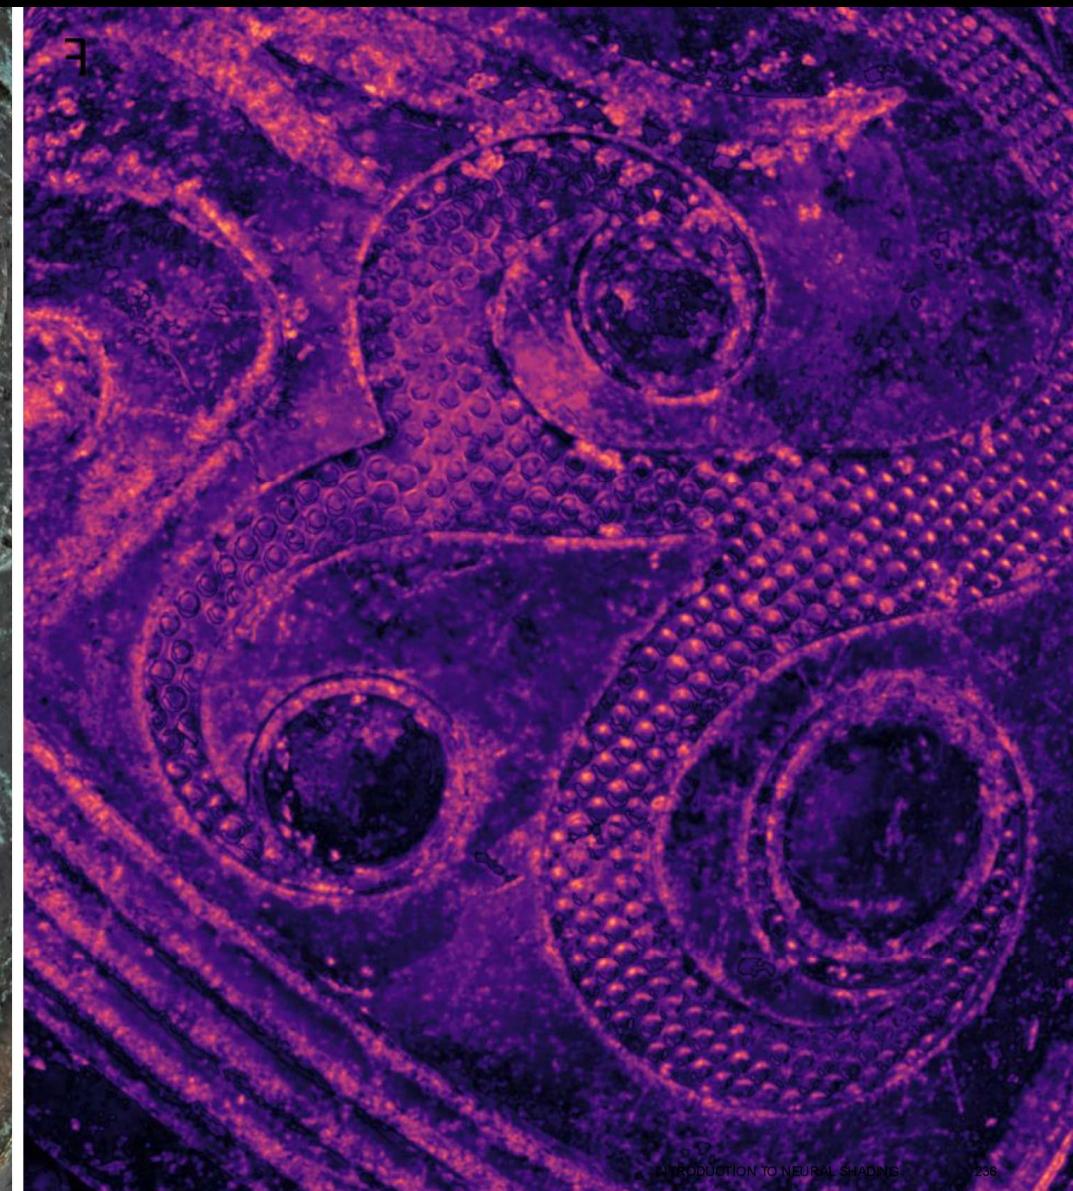
YOU ONLY NEED THIS TRICK DURING TRAINING!



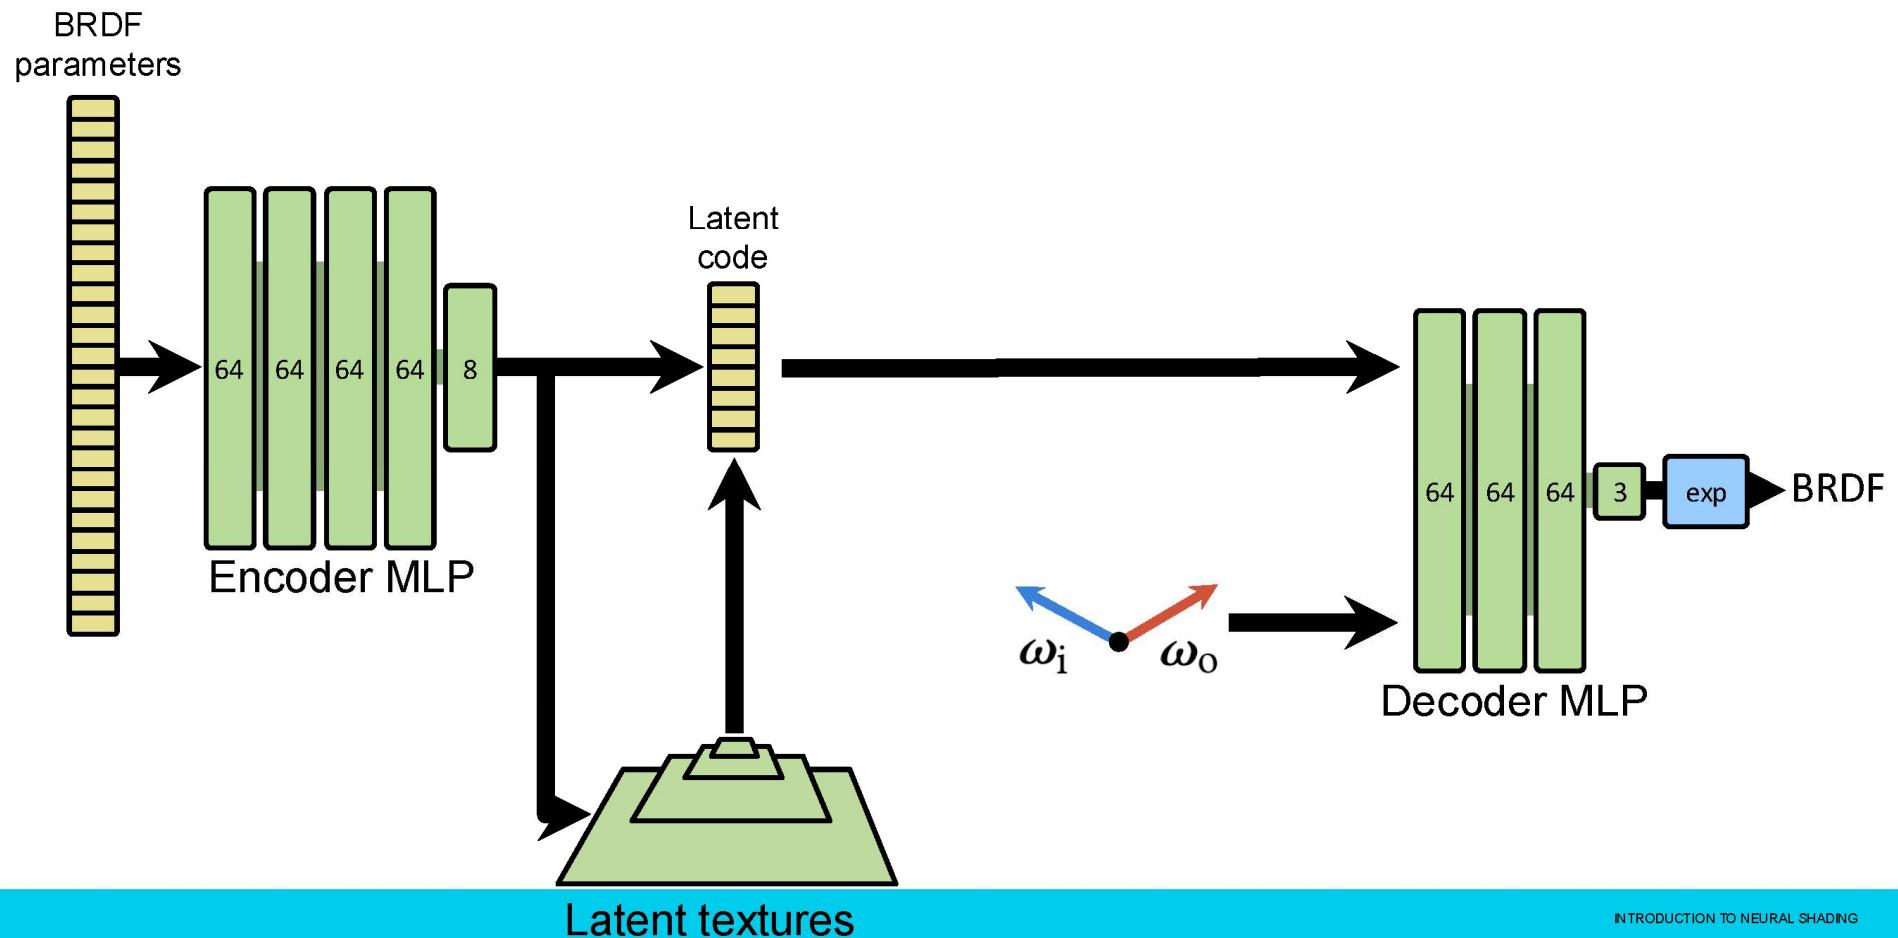
YOU ONLY NEED THIS TRICK DURING TRAINING!



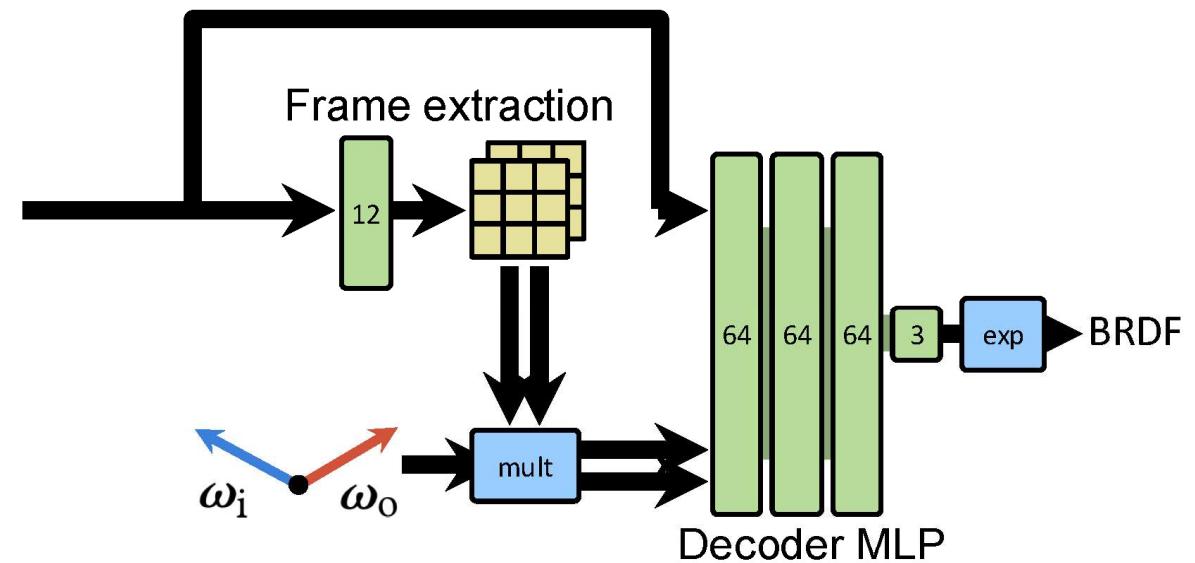




LET'S ADD A FIXED-FUNCTION FRAME TRANSFORM

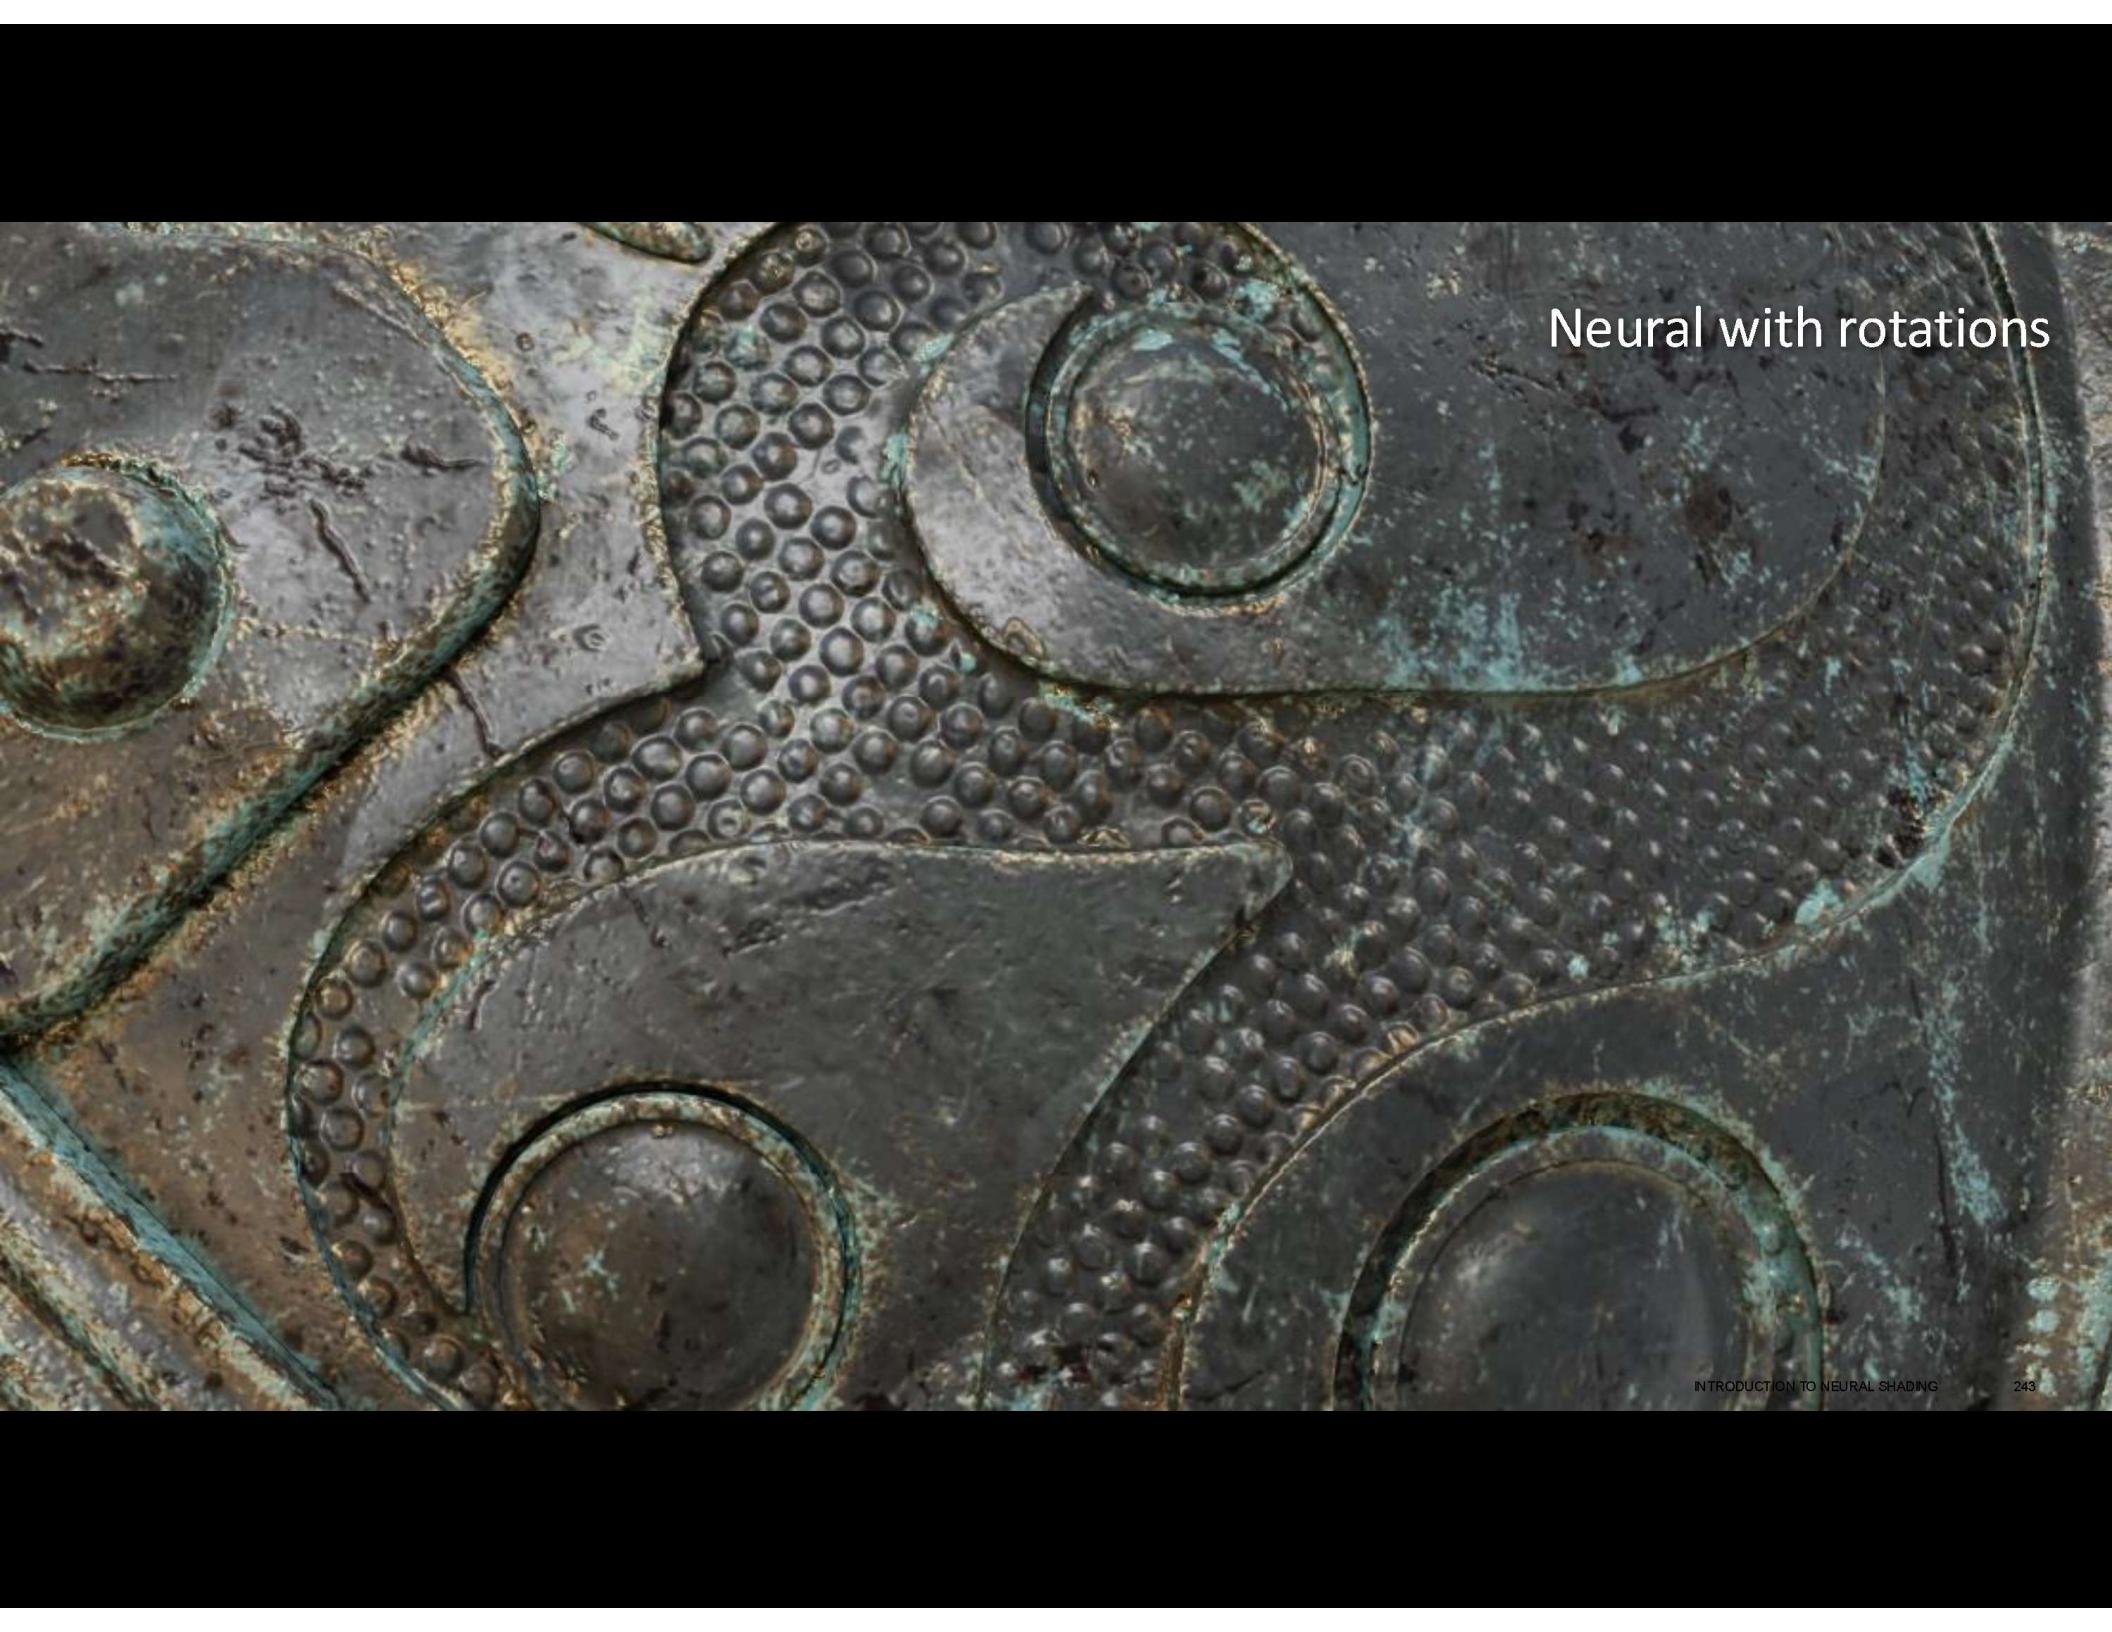


LET'S ADD A FIXED-FUNCTION FRAME TRANSFORM

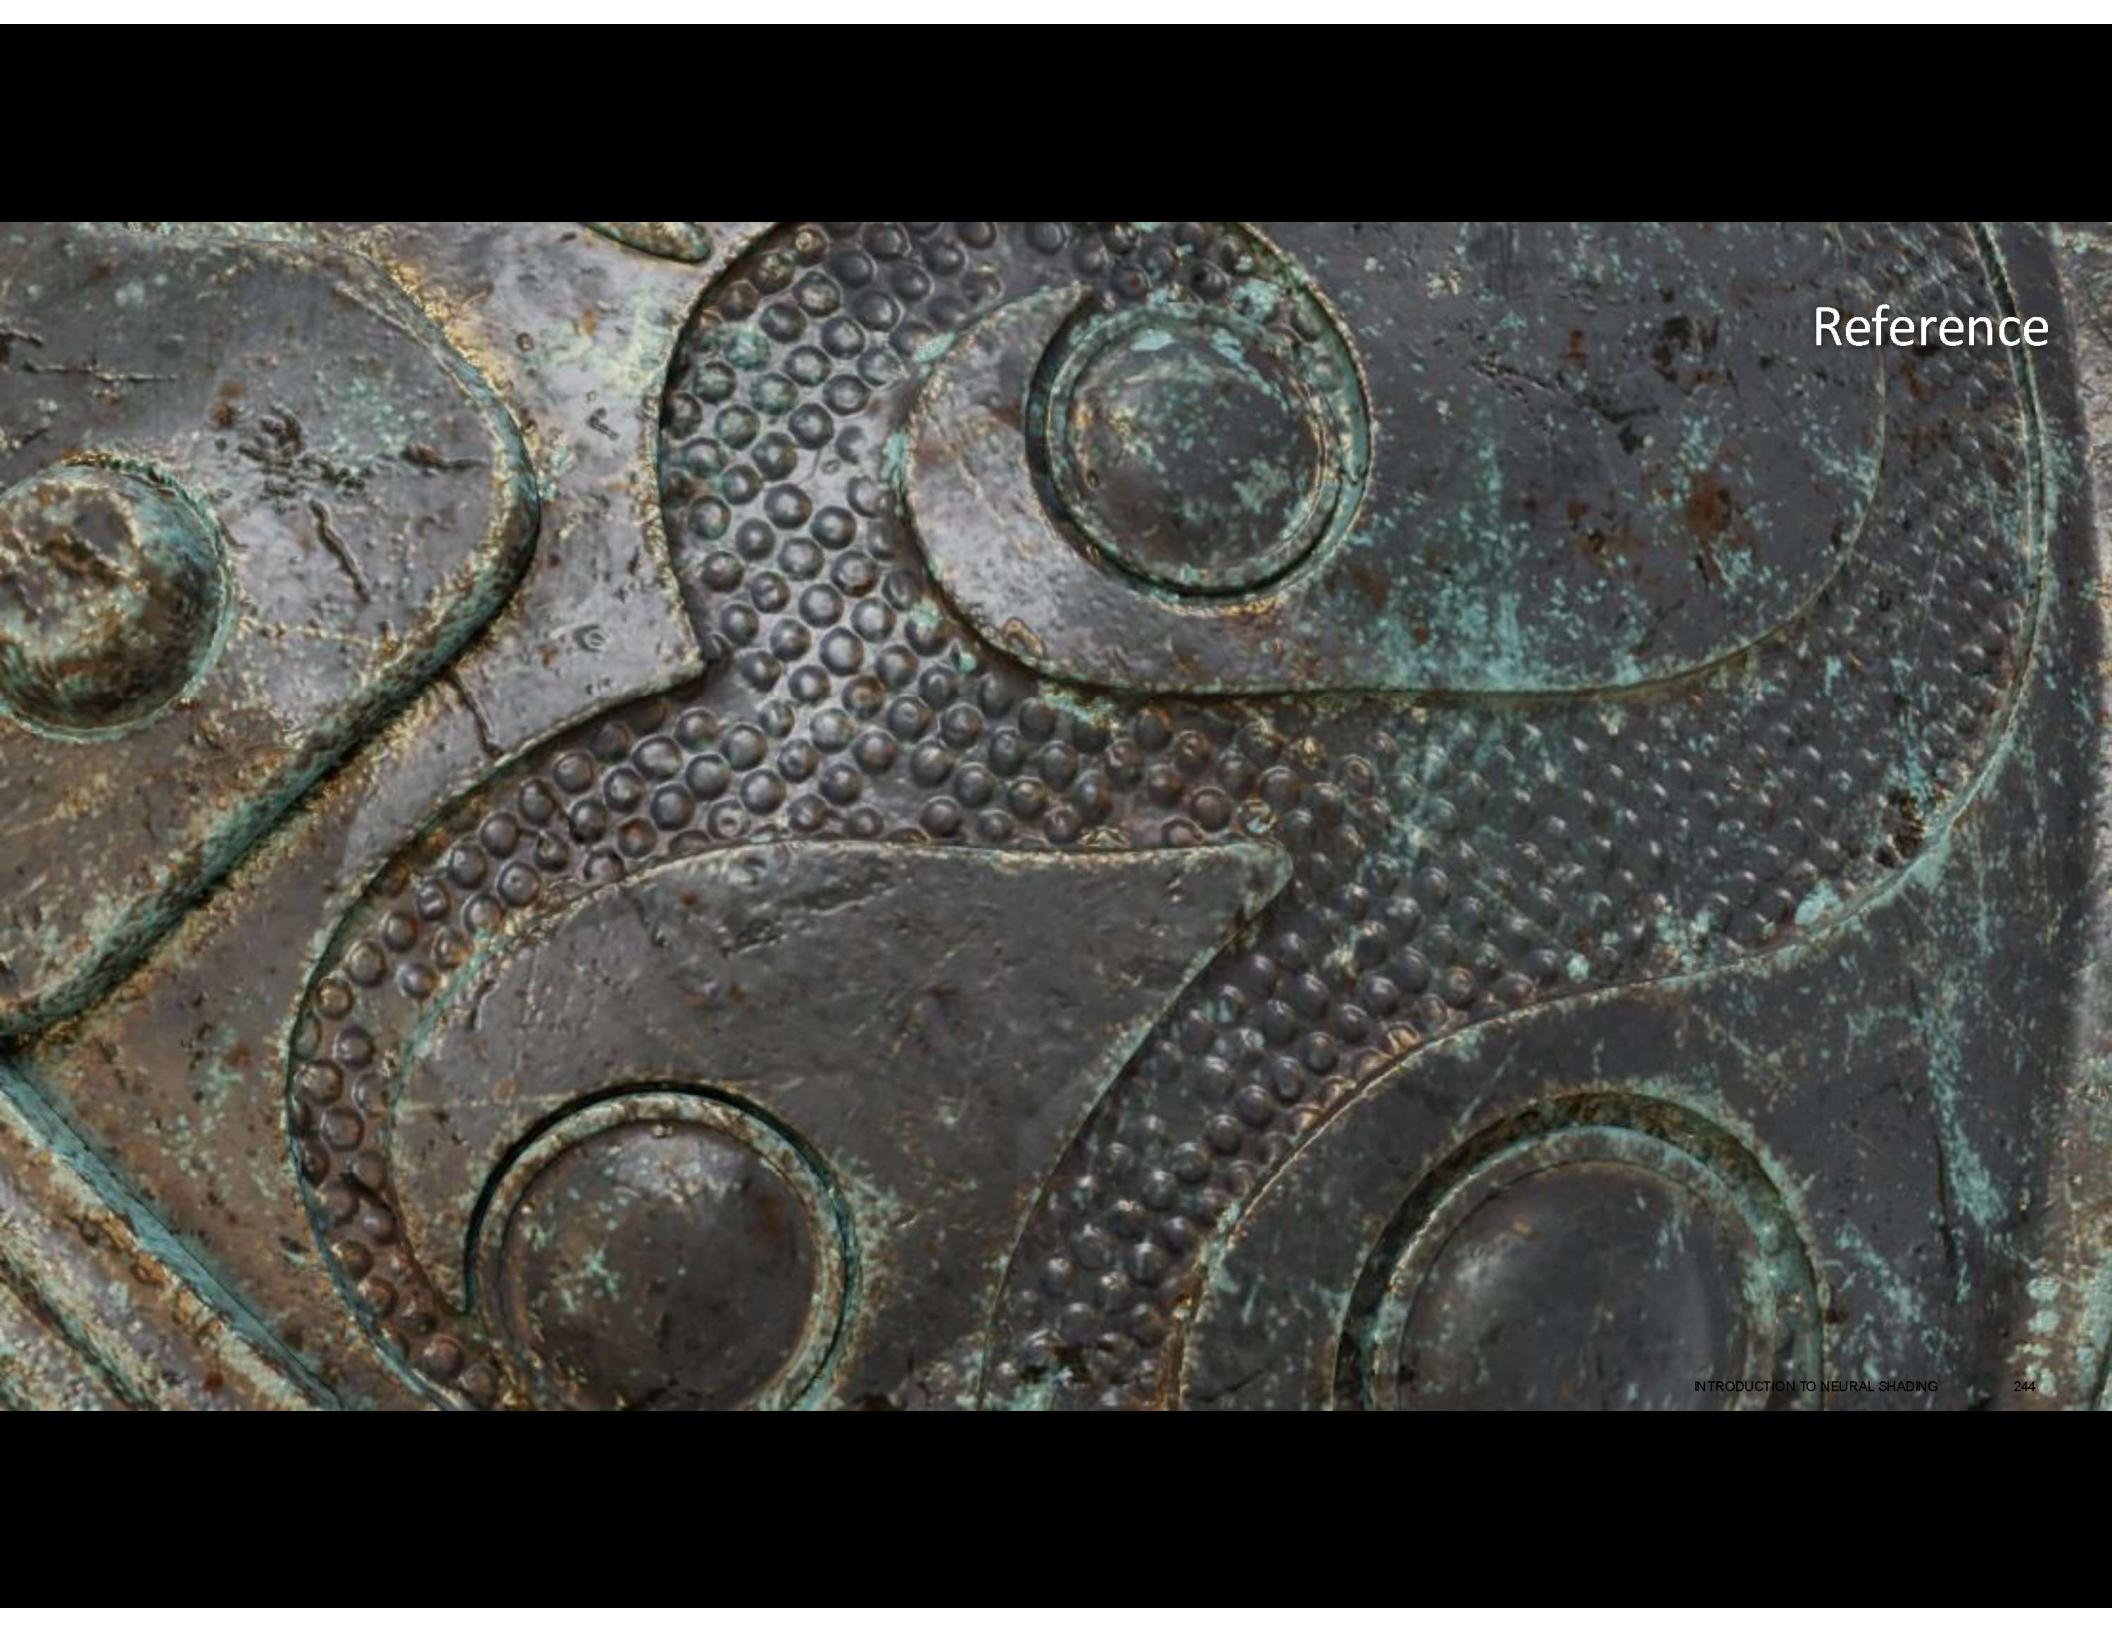




Neural without rotations

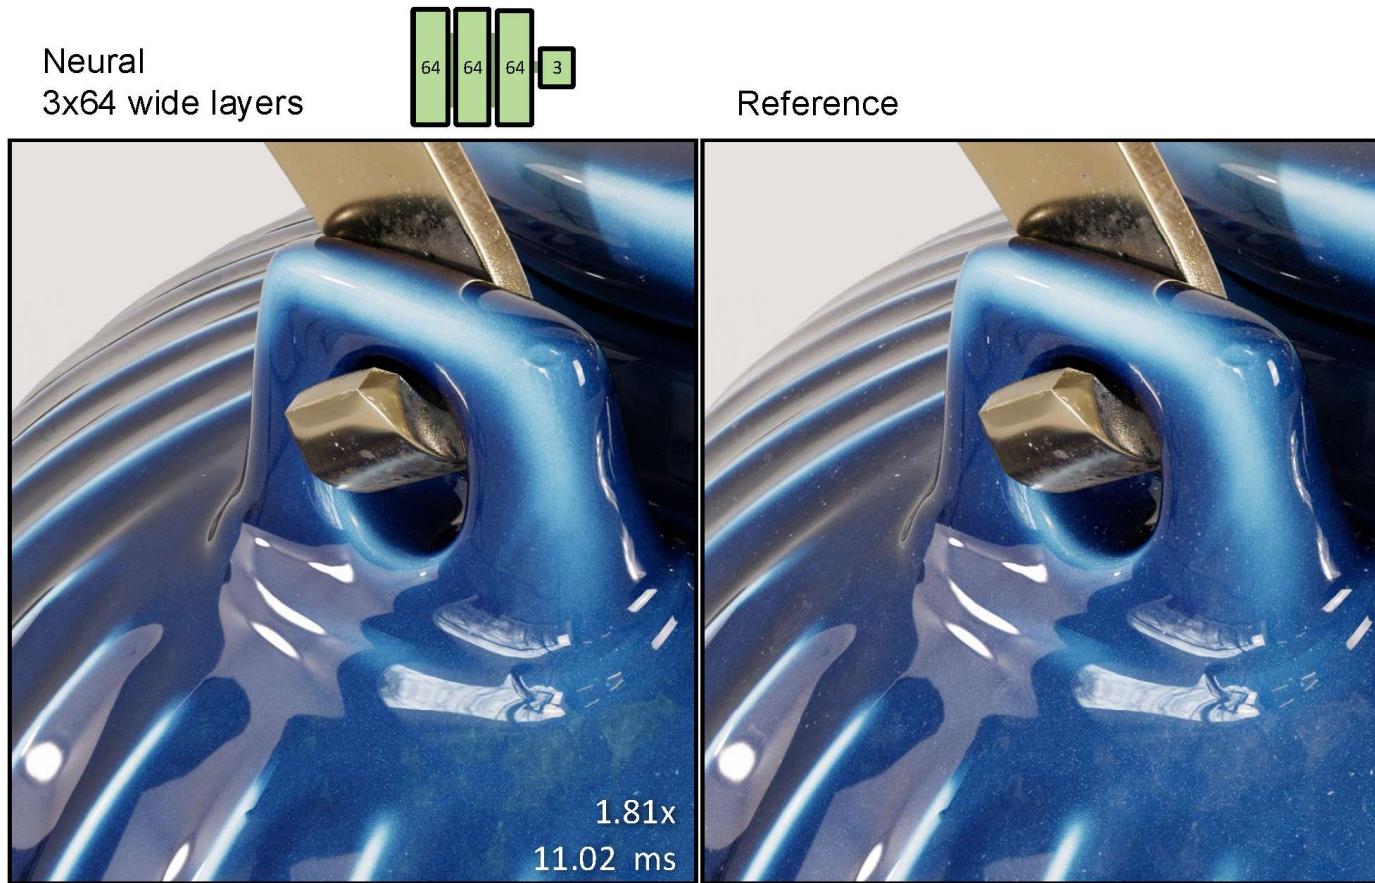


Neural with rotations



Reference

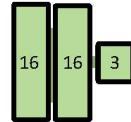
END RESULT



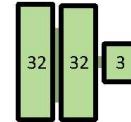
NEURAL SHADING GIVES YOU A QUALITY/COST DIAL



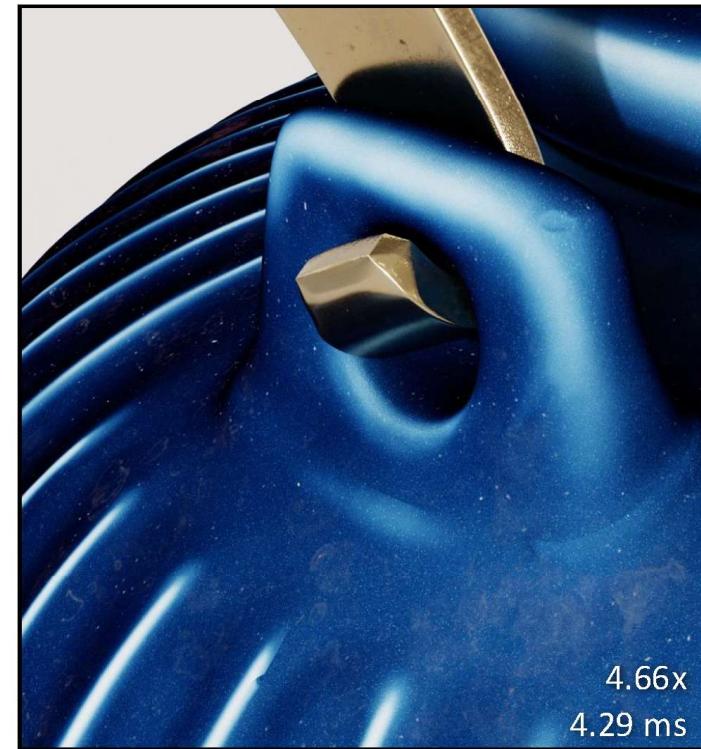
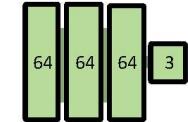
Neural
2x16 wide layers



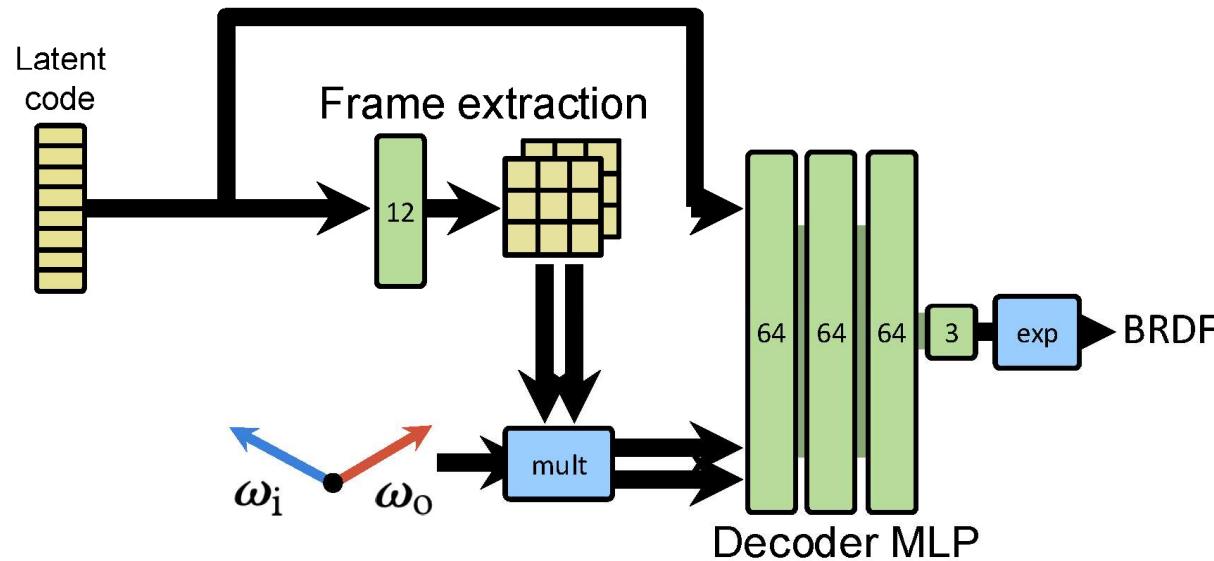
Neural
2x32 wide layers



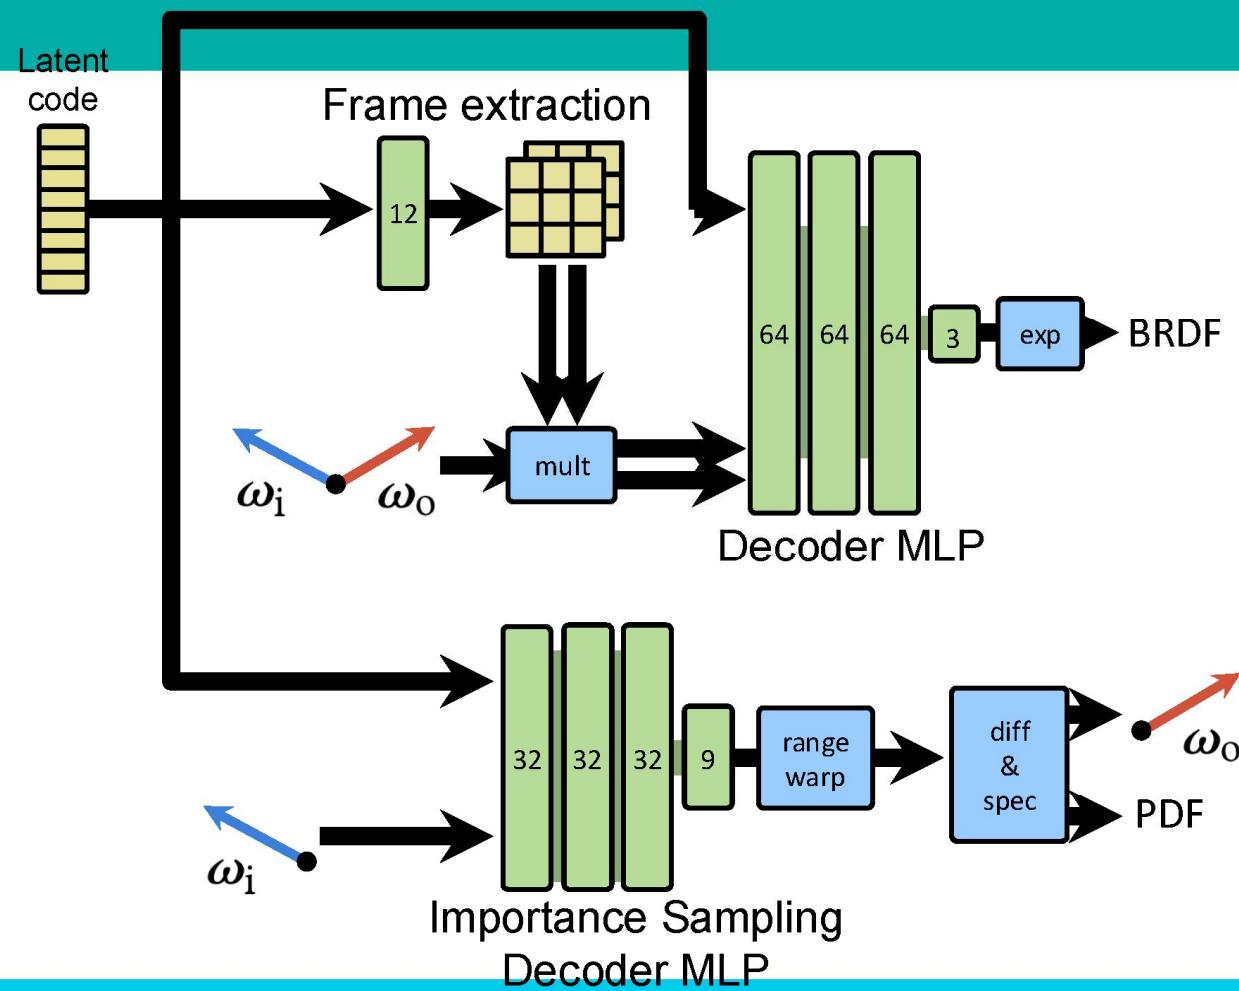
Neural
3x64 wide layers



ADDING ADDITIONAL CAPABILITIES



ADDING ADDITIONAL CAPABILITIES



Thank you.